

Quantitative Convergence Analysis of Stochastic
Maximum Principles
for Mean-Field Games with Common Noise

Nisong Monyimba

June 24, 2026

Abstract

This thesis develops a mathematically rigorous and computationally implementable framework for mean-field games with common noise, integrating three research streams — SPDE theory, multi-timescale reinforcement learning, and quantitative finance — into a single coherent pipeline. The central claim is that market inefficiency can be productively understood as distributional disequilibrium under shared shocks: a measurable distance in probability space between the observed distribution of institutional investor positions and the game-theoretic equilibrium.

The work provides six principal contributions.

1. **Quantitative well-posedness** (Theorem 3.1). Under standard regularity assumptions, the mean-field game with common noise admits a unique equilibrium, and the fixed-point map is a contraction in a weighted Wasserstein-2 metric. The stability constant is explicit in the model parameters ($C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$), directly informing numerical step-size selection and reinforcement-learning rate design. The contraction is established in expectation under synchronous coupling of the common noise.
2. **Numerical convergence** (Theorem 4.1). A combined Euler–Maruyama / particle / SPDE discretisation converges with weak error $O(\Delta t^{1/2} + \Delta x)$, and a CFL condition coupling parabolic, stochastic transport, and explicit feedback constraints is derived and validated on the linear-quadratic benchmark.
3. **Three-timescale reinforcement-learning convergence** (Theorem 5.1). A multi-timescale actor-critic algorithm converges $\mathbb{P}$ -almost surely along subsequences ($\liminf W_2 = 0$) to the equilibrium measure flow in continuous state-action spaces, under compact Lipschitz function approximators. An exponential convergence rate under entropy regularisation and a log-Sobolev inequality is also established (Theorem 5.2).
4. **SPDE-grounded hierarchical architecture** (Proposition 6.1). The master-equation decomposition into idiosyncratic, mean-field, and common-noise operators canonically motivates a three-tier architecture (SAC-PPO-TD3), which significantly outperforms monolithic baselines in sample efficiency and stability, with a control experiment isolating the SPDE-guided timescale contribution.
5. **Alpha signals from the equilibrium** (Proposition 8.1). Four trading signals are derived from the linear-quadratic MFG equilibrium. Three are adaptations of prior work (optimal spread, mean-reversion speed, factor exposure); the novel contribution is a crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ that quantifies distributional disequilibrium. Methods for estimating the unobservable equilibrium measure (GMM, Kalman filtering, reduced-form proxies) are provided, and a mean-variance portfolio construction with a crowding discount is developed.

6. Empirical validation on S&P 500 constituents (2000–2024) (Chapter 9).

A strategy based on the four signals, using point-in-time 13F holdings and CFTC COT data, achieves an out-of-sample Sharpe ratio of 1.91 (net of 5–10bps transaction costs), with a maximum drawdown of -18.3% and positive skewness (0.32). Statistical significance is $p = 0.002$ (block bootstrap, 10,000 samples). Performance is robust to alternative cost assumptions, calibration windows, discount functions, and sub-periods, and incremental over both a naïve 13F mean-reversion benchmark (Sharpe 0.72) and a 13F residual momentum benchmark (Sharpe 0.89). The crowding measure c_t is strongly monotonic in subsequent alpha: the top quintile yields average monthly returns of 2.8% versus 0.9% in the bottom quintile.

Selected core estimates — the Gronwall inequality, the Riccati system derivation, the CFL condition, and the optimality of the linear feedback control — have been formally verified in the Lean 4 interactive theorem prover; the code is publicly available. The thesis concludes with a candid discussion of limitations (regularity assumptions, subsequential convergence, model-risk dominance in the empirical Sharpe ratio) and a map of six open problems, including jump-diffusion extensions, online RL with streaming data, and full formal verification of the pipeline.

All theoretical results hold under the standing assumptions collected in Section 1.9; these are strong and their practical implications are discussed throughout. The reported Sharpe ratio is conditional on the LQ model specification and estimation procedures; model risk dominates statistical risk.

Keywords: Mean-field games, common noise, Wasserstein distance, contraction, multi-timescale reinforcement learning, actor-critic, hierarchical architecture, alpha generation, crowding measure.

Contents

1	Introduction	11
1.1	Overview	11
1.2	What Is a Mean-Field Game?	13
1.2.1	The N -Player Problem: A Financial Market	13
1.2.2	Why the N -Player Problem Is Intractable	14
1.2.3	The Mean-Field Limit: $N \rightarrow \infty$	14
1.2.4	Why Common Noise Is Non-Trivial	15
1.3	The FBSDE Characterisation	15
1.3.1	Overview of the System	15
1.3.2	Heuristic Derivation	15
1.3.3	The Forward Equation: Population Measure Flow	16
1.3.4	The Backward Equation: Adjoint Variables	16
1.3.5	The Optimality Condition	16
1.3.6	The Consistency Condition	16
1.3.7	The Complete FBSDE System	16
1.4	State of the Art	17
1.4.1	Mean-Field Game Theory	17
1.4.2	Reinforcement Learning for Mean-Field Games	18
1.4.3	Quantitative Finance	18
1.4.4	Wasserstein Distance in Financial Applications	19
1.4.5	Gaps in the Literature	19
1.5	The Three Critical Gaps	19
1.6	Principal Contributions	20
1.7	Connections Between Fields	22
1.8	Thesis Roadmap	23
1.8.1	Chapter Dependency Graph	23
1.8.2	Chapter-by-Chapter Summary	23
1.9	Key Notation and Standing Assumptions	23
1.9.1	Notation Reference	24
1.9.2	Standing Assumptions	24
1.10	Guide to Reading	24
1.11	Why This Problem, Why Now	25
2	Mathematical Preliminaries	26
2.1	Overview	26
2.2	Wasserstein-2 Space	27
2.2.1	Probability Measures with Finite Second Moment	27
2.2.2	Optimal Transport and the Wasserstein Distance	27

2.2.3	The Sinkhorn Algorithm	28
2.2.4	Metric Space Structure	28
2.2.5	Weak Convergence in $\mathcal{P}_2$	29
2.2.6	Geodesics and the Lions Derivative	29
2.2.7	Law of Large Numbers in Wasserstein Distance	30
2.2.8	Wasserstein Barycentre	30
2.3	Brownian Motion and Itô Calculus	31
2.3.1	Brownian Motion: Construction and Properties	31
2.3.2	Quadratic Variation	31
2.3.3	The Itô Integral	32
2.3.4	Itô's Formula	32
2.3.5	Itô's Product Rule and Integration by Parts	33
2.3.6	Stochastic Differential Equations	33
2.4	Fokker–Planck SPDE	34
2.4.1	The Forward Equation	34
2.4.2	Strong Form: The SPDE in Density Form	34
2.4.3	Sobolev Regularity	35
2.4.4	The Gelfand Triple	35
2.5	Martingale Theory	35
2.5.1	Martingales and Stopping Times	35
2.5.2	Martingale Representation Theorem	36
2.5.3	The Novikov Condition and Girsanov's Theorem	36
2.6	Sobolev Spaces and Functional Analysis	37
2.6.1	Sobolev Embedding Theorem	37
2.6.2	Hahn–Banach Extension Theorem	37
2.6.3	Lumer–Phillips Theorem and Contraction Semigroups	37
2.6.4	Spectral Gap and Log-Sobolev Inequality	37
2.7	Stochastic Maximum Principle	38
2.7.1	The Control Problem	38
2.7.2	The Hamiltonian	38
2.7.3	The Adjoint Process and SMP	39
2.7.4	Connection to Hamilton–Jacobi–Bellman	39
2.8	Summary: Toolkit Reference	40
3	Quantitative Well-Posedness of the Common-Noise MFG	42
3.1	Overview	42
3.2	Standing Assumptions and Notation	43
3.3	Statement of the Main Theorem	44
3.4	The Fixed-Point Map Φ	44
3.5	Stability Estimate for the Map Φ	45
3.5.1	Synchronous Coupling	45
3.5.2	Itô Expansion of the Squared Difference	46
3.5.3	Lipschitz Estimates and Gronwall Setup	46
3.5.4	Control Difference via SMP and Lions Derivatives	46
3.5.5	Adjoint Process Difference	47
3.5.6	Combining the Estimates	47
3.5.7	The Wasserstein Stability Estimate	48
3.6	Contraction in Weighted Wasserstein Metric	48

3.7	Existence and Uniqueness	48
3.8	The ε -Nash Property	49
3.9	Discussion of the Contraction Constant	49
3.10	Formal Verification in Lean 4	50
4	Numerical Analysis of the Common-Noise MFG System	52
4.1	Overview	52
4.2	Discretisation of the FBSDE System	53
4.2.1	Discrete Forward SDE	53
4.2.2	Discrete Adjoint BSDE	53
4.2.3	Discrete Optimality Condition	54
4.2.4	Discrete Consistency Condition	54
4.3	Particle Approximation of the Measure Flow	54
4.4	Error Analysis for the Euler–Maruyama Scheme	55
4.5	Finite-Difference Approximation of the Fokker–Planck SPDE	56
4.5.1	Spatial Discretisation	56
4.5.2	Stability and Convergence	56
4.6	Combined Error Estimate: Proof of the Main Theorem	57
4.7	CFL Condition for the Coupled System	58
4.8	Numerical Experiments: Linear-Quadratic Benchmark	59
4.8.1	Experimental Setup	59
4.8.2	Results	59
4.8.3	Discussion	60
4.9	Conclusion	61
5	Reinforcement Learning Convergence for Common-Noise MFGs	62
5.1	Overview	62
5.2	Reinforcement Learning Formulation of the MFG	64
5.3	Three-Timescale Actor-Critic Architecture	64
5.3.1	Discrete-Time Algorithm	65
5.3.2	Continuous-Time Limiting Dynamics	65
5.4	Statement of Main Theorems	66
5.5	Proof of Subsequential Almost-Sure Convergence	67
5.5.1	Step 1: Compactness and Tightness	68
5.5.2	Step 2: Critic Convergence (Fast Timescale)	68
5.5.3	Step 3: Actor Convergence (Intermediate Timescale)	68
5.5.4	Step 4: Measure Filter Convergence (Slowest Timescale)	69
5.5.5	Step 5: Coupling and Subsequential Convergence	69
5.6	Exponential Convergence under Entropy Regularisation	69
5.6.1	Entropy-Regularised Objective	70
5.6.2	Log-Sobolev Inequality and Exponential Convergence	70
5.6.3	Discussion of the Log-Sobolev Assumption	70
5.7	Practical Implementation Considerations and Limitations	71
5.7.1	Function Approximation	71
5.7.2	Computing the Wasserstein Distance	71
5.7.3	Parallelisation and Scalability	71
5.7.4	Learning Rate Schedules	71
5.7.5	Limitations of the Theoretical Guarantees	71

5.8	Conclusion and Connection to Chapter 6	72
6	Hierarchical Reinforcement Learning Architecture for Common-Noise MFGs	73
6.1	Overview	73
6.2	SPDE Decomposition of the Master Equation	74
6.3	Mapping Operators to Algorithmic Components	75
6.3.1	Fast Critic: Soft Actor-Critic (SAC) for Idiosyncratic Diffusion . .	76
6.3.2	Intermediate Actor: Proximal Policy Optimization (PPO) for Mean-Field Interaction	76
6.3.3	Slow Meta-Learner: Twin Delayed DDPG (TD3) for Common-Noise Adaptation	77
6.4	Formal Timescale Separation	77
6.5	Implementation of the Hierarchical Architecture	79
6.5.1	System Overview	79
6.5.2	Communication Protocol	80
6.5.3	Embedding the Population Measure	80
6.5.4	Algorithm Summary	81
6.6	Experimental Evaluation	81
6.6.1	Linear-Quadratic Benchmark	81
6.6.2	Multi-Asset Market Simulation	82
6.6.3	Ablation Study: Timescale Separation	83
6.7	Limitations and Future Directions	83
6.8	Conclusion	84
7	Case Study: Systemic Risk and Optimal Bailouts in a Common-Noise MFG	86
7.1	Overview	86
7.2	Economic Environment	87
7.2.1	The Banking System	87
7.2.2	Reserve Dynamics	87
7.2.3	Bank Default and Government Objective	88
7.2.4	Interaction with Bank Optimisation	89
7.3	Stackelberg Mean-Field Game Formulation	89
7.3.1	Representative Bank's Problem	89
7.3.2	Government's Problem	90
7.3.3	FBSDE Characterisation of the Equilibrium	90
7.4	Closed-Form Solution of the Linear-Quadratic MFG	91
7.4.1	The Riccati System	91
7.4.2	Equilibrium Measure Dynamics	91
7.4.3	Optimal Government Bailout Policy	92
7.4.4	Fixed Point and Equilibrium Consistency	92
7.5	$O(1/N)$ Approximation Error	93
7.6	Policy Implications	94
7.6.1	Bailout Generosity and Moral Hazard	94
7.6.2	Common Noise Amplification	94
7.6.3	Distributional Impact	94
7.7	Numerical Illustration	95

7.8	Conclusion and Bridge to Alpha Generation	96
8	Alpha Generation from the Mean-Field Game Equilibrium	97
8.1	Introduction	97
8.2	The Mean-Field Game Framework	98
8.2.1	The Linear-Quadratic Specification	98
8.2.2	The Riccati System	99
8.3	Derivation of the Four Alpha Signals	99
8.3.1	Optimal Spread	100
8.3.2	Mean-Reversion Speed	100
8.3.3	Factor Exposure	100
8.3.4	Crowding Measure	101
8.4	Estimation of the Equilibrium Measure μ_t^*	101
8.4.1	Structural Calibration via GMM	102
8.4.2	Filtering Approach via Kalman Filter	102
8.4.3	Reduced-Form Approximation	102
8.5	Portfolio Construction	103
8.5.1	Signal Aggregation	103
8.5.2	Mean-Variance Optimisation	103
8.5.3	Crowding Discount	104
8.5.4	Implementation Algorithm	104
8.6	Robustness and Model Risk	104
8.7	Formal Verification in Lean 4	106
8.8	Conclusion	107
9	Empirical Validation: S&P 500 Constituents, 2000–2024	108
9.1	Overview	108
9.2	Data and the Empirical Measure μ_t	109
9.2.1	Universe and Time Period	109
9.2.2	Price and Returns Data	109
9.2.3	Estimating the Empirical Measure μ_t	109
9.2.4	Descriptive Statistics	111
9.3	Calibration of the Equilibrium Measure μ_t^*	111
9.3.1	LQ-MFG Specification	111
9.3.2	Calibration Methodology	111
9.3.3	Calibrated Parameters	112
9.4	Construction of the Four Alpha Signals	112
9.5	Portfolio Construction	113
9.5.1	Signal Aggregation and Mean-Variance Weights	113
9.5.2	Crowding Discount	113
9.5.3	Transaction Costs and Turnover Constraint	113
9.5.4	Leverage and Short-Sale Constraints	113
9.6	Backtesting Framework	113
9.6.1	Walk-Forward Procedure	114
9.6.2	Purged Data	114
9.6.3	Benchmark Portfolios	114
9.7	Main Empirical Results	115
9.7.1	Performance Metrics	115

9.7.2	Cumulative Returns	115
9.7.3	Statistical Significance	115
9.7.4	The Role of the Crowding Measure	116
9.8	Robustness Checks	116
9.8.1	Alternative Transaction Cost Assumptions	117
9.8.2	Alternative Calibration Windows	117
9.8.3	Alternative Factor Exposure Construction	117
9.8.4	Alternative Crowding Discount Functions	118
9.8.5	Subperiod Analysis	118
9.8.6	Alternative Benchmark Universes	118
9.8.7	Unconstrained Turnover	118
9.8.8	Alternative Empirical Measure Constructions	119
9.8.9	Long-Only Version	119
9.8.10	13F Residual Momentum Benchmark	120
9.8.11	Sensitivity to Estimation Error in μ_t^*	120
9.8.12	Placebo Test: Randomised Signals	120
9.9	Discussion and Limitations	120
9.9.1	Interpretation of the Sharpe Ratio	120
9.9.2	Capacity Analysis (with Strong Caveats)	121
9.9.3	Comparison to Prior Literature	121
9.10	Conclusion	121
10	Synthesis and Outlook	123
10.1	Overview	123
10.2	Summary of Contributions	123
10.2.1	Theoretical Foundations (Chapters 1–3)	123
10.2.2	Numerical Methods (Chapter 4)	124
10.2.3	Reinforcement Learning (Chapters 5–6)	124
10.2.4	Financial Applications (Chapters 7–9)	124
10.3	Limitations	125
10.3.1	Mathematical Limitations	125
10.3.2	Algorithmic Limitations	126
10.3.3	Empirical Limitations	126
10.3.4	Formal Verification Limitations	127
10.4	Open Problems and Future Research	127
10.4.1	Extension to Jump-Diffusion Common Noise	127
10.4.2	Fully Online Reinforcement Learning with Streaming Data	128
10.4.3	Asymmetric Information and Heterogeneous Beliefs	128
10.4.4	Application to Other Asset Classes and Macroeconomic Settings	128
10.4.5	Relaxing the Log-Sobolev Assumption	129
10.4.6	Formal Verification of the Full Pipeline	129
10.5	Broader Impact	130
10.5.1	For Financial Economics	130
10.5.2	For Reinforcement Learning	130
10.5.3	For Quantitative Investment Practice	130
10.5.4	For the Reproducibility of Mathematical Finance	130
10.6	Conclusion	131

A	Proofs	135
A.1	Detailed Derivation of the Contraction Estimate (Chapter 3)	135
A.1.1	Derivation of C_α	135
A.1.2	Derivation of C_{BSDE}	135
A.1.3	Combining the Estimates	136
A.2	Discretisation Error Lemma (Chapter 4)	136
A.3	Derivation of the Riccati Systems (Chapters 7 and 8)	137
A.3.1	Setup	137
A.3.2	Quadratic Terms	137
A.3.3	Linear Terms	138
A.3.4	Proof that v_t is Independent of b_t	139
A.3.5	Government's Riccati Equation (Chapter 7)	139
A.3.6	Explicit Bound for the Propagation of Chaos Constant (Chapter 7)	139
B	Formal Verification in Lean 4	140
B.1	Gronwall's Inequality	140
B.2	Optimality of the Linear Feedback Control	142
B.3	Contraction Factor in the Weighted Wasserstein Metric	144
B.4	Wasserstein-2 Distance Between Gaussian Measures	145
B.5	Talagrand Transportation Inequality and Exponential Convergence	145
C	Additional Results	148
C.1	Index of Numbered Results	148
C.2	Expanded Notation Glossary	148

List of Figures

List of Tables

1.1	Summary of Principal Contributions and Their Novelty	22
2.1	Toolkit Reference	41
3.1	Components of the Contraction Constant	50
4.1	Summary of Particle Approximation Properties	55
4.2	CFL Condition Components	59
4.3	Convergence results for the LQ-MFG benchmark	60
4.4	Computational Cost Comparison (LQ-MFG, $T = 1$)	60
5.1	Comparison of RL for MFG Convergence Results	63
5.2	Summary of Assumptions and Their Implications	67
6.1	SPDE Operator Decomposition and Timescales	75
6.2	Timescale Separation Verification for LQ-MFG	79
6.3	LQ-MFG Benchmark Results	82
6.4	Multi-Asset Market Simulation Results	83
6.5	Ablation Study on Update Frequencies	83
6.6	Summary of Limitations and Mitigations	85
7.1	Summary of Model Parameters and Their Economic Interpretation	93
7.2	Comparative Statics of the Equilibrium	95
7.3	Baseline Parameters for Numerical Illustration	95
7.4	Riccati Coefficients at Selected Times	96
7.5	Wasserstein Error for Different N at $t = 0.5$	96
8.1	Summary of Alpha Signals	101
8.2	Comparison of Equilibrium Measure Estimation Methods	103
8.3	Model Risk Sources and Mitigations	105
9.1	Summary statistics for institutional ownership of S&P 500 constituents, 2000–2024	111
9.2	Time-series average calibrated LQ-MFG parameters, 2005–2024	112
9.3	Out-of-sample performance metrics, January 2005 – December 2024	115
9.4	Panel regression of monthly stock returns on lagged MFG signals	116
9.5	Fama–MacBeth regression of monthly stock returns on lagged MFG signals	116
9.6	Average monthly strategy return by crowding measure quintile	117
9.7	Sharpe ratio under alternative transaction cost assumptions	117
9.8	Sharpe ratio for alternative calibration windows	117
9.9	Sharpe ratio for alternative crowding discount functions	118

9.10 Sharpe ratio by subperiod	118
9.11 Sharpe ratio for alternative universes	119
9.12 Performance of unconstrained vs. turnover-constrained strategies	119
9.13 Sharpe ratio for alternative empirical measure constructions	119
10.1 Summary of Principal Contributions and Their Novelty	132
10.2 Summary of Limitations and Practical Implications	133
10.3 Open Problems and Key References	134
C.1 Index of numbered results: Chapter 2 (mathematical toolkit)	149
C.2 Index of numbered results: Chapter 2 (continued)	150
C.3 Index of numbered results: Chapters 3–8 (main theoretical and applied contributions)	151
C.4 Expanded notation glossary	152

Chapter 1

Introduction

1.1 Overview

This chapter establishes the intellectual foundation for the thesis. It begins with a concrete financial-market narrative that motivates mean-field games with common noise (Section 1.2), derives the governing forward-backward stochastic differential equation (FBSDE) system from first principles (Section 1.3), and surveys the literature across stochastic partial differential equation (SPDE) theory, reinforcement learning, and quantitative finance (Section 1.4). We identify three precise gaps that this work addresses, explicitly situating each within the existing body of knowledge. The six principal contributions are then presented together with a candid discussion of their novelty relative to prior work and their interconnections. We demonstrate the deep structural connections between the three fields, provide a ten-chapter roadmap, and collect all key notation and standing assumptions. Three tailored reading paths are offered, and the chapter closes with a discussion of the mathematical, computational, and market conditions that make this research both possible and timely.

By the end of this chapter the reader will be able to answer five questions:

1. How does common noise fundamentally alter the geometry and stability of mean-field equilibria?
2. What is the FBSDE system, and what role does each equation play in characterising the equilibrium?
3. What does the literature already contain, and what are the three specific gaps this thesis addresses — understood as incremental extensions rather than fundamental breakthroughs?
4. What exactly does this thesis prove, and how do the six principal results interconnect and relate to prior work?
5. How should one navigate the remaining nine chapters?

Executive summary. This thesis models financial markets as a large population of interacting institutional investors. Each investor adjusts its portfolio (defined as log-positions across assets) based on the aggregate behaviour of all others — the distribution of positions across the market — which evolves as a stochastic process taking values

in the space of probability measures equipped with the Wasserstein-2 geometry. The Wasserstein-2 geometry is the central analytical and economic object of the thesis, providing both the metric for stability analysis and the basis for measuring market inefficiency. An equilibrium is reached when the distribution generated by everyone acting optimally coincides with the distribution that each investor takes as given.

The central claim of this thesis is that **market inefficiency can be productively understood as distributional disequilibrium under shared shocks** — a measurable distance in probability space. This thesis reframes market inefficiency as a geometric misalignment in the space of distributions, measured by the Wasserstein-2 distance between the empirical market measure and the theoretical equilibrium measure.

A key structural result is that common noise preserves the contraction structure of the MFG fixed-point map in expectation under synchronous coupling and homogeneous common-noise exposure. The contraction is established at the level of expected Wasserstein dynamics; it does not extend to pathwise contraction or heterogeneous noise structures. The first-order stochastic terms cancel in expectation, and the contribution of the common noise enters only through quadratic variation terms in the Itô expansion of optimally coupled processes. This cancellation is a second-moment phenomenon formalised in Theorem 3.9.

We prove that this model is mathematically stable (in the sense of a contraction in a weighted Wasserstein metric) even under common shocks, design regularised reinforcement learning algorithms with convergence guarantees (established for regularised, compact function classes as a theoretical baseline; extending these results to overparameterised deep neural networks remains an open problem), and show that deviations from equilibrium yield economically meaningful trading signals whose directional behaviour aligns with the theoretical structure.

This thesis develops a mathematically rigorous and computationally implementable framework for mean-field games with common noise. It integrates three research streams that have not yet been combined into a single framework simultaneously providing (i) quantitative stability estimates, (ii) learning guarantees under common noise, and (iii) empirical validation on real financial data:

- **SPDE theory** — the analysis of forward-backward stochastic partial differential equations arising from the Stochastic Maximum Principle;
- **Reinforcement learning** — multi-timescale actor-critic algorithms with provable convergence guarantees for regularised function classes;
- **Quantitative finance** — deployable alpha generation from equilibrium theory, validated by rigorous statistical testing.

The framework yields six principal contributions spanning well-posedness, numerical analysis, learning theory, and empirical finance. Selected core estimates are supported by formal verification in Lean 4, and the empirical application provides evidence broadly consistent with the equilibrium predictions. **All theoretical results are subject to the standing assumptions collected in Section 1.9; these assumptions are strong and their implications for practical implementation are discussed throughout.**

1.2 What Is a Mean-Field Game?

1.2.1 The N -Player Problem: A Financial Market

Consider a financial market populated by N institutional investors — hedge funds, asset managers, pension funds, and proprietary trading desks. We work on a complete filtered probability space $(\Omega, \mathcal{F}, (\mathcal{F}_t)_{t \in [0, T]}, \mathbb{P})$ satisfying the usual conditions. The filtration $(\mathcal{F}_t)$ is the augmentation of that generated by a collection of independent Brownian motions: a d_0 -dimensional common noise $W^0 = (W_t^0)_{t \in [0, T]}$ and, for each agent $i \in \{1, \dots, N\}$, a d -dimensional idiosyncratic noise $W^i = (W_t^i)_{t \in [0, T]}$, with $W^i \perp W^j$ for $i \neq j$ and all W^i independent of W^0 . All stochastic processes are assumed to be progressively measurable with respect to this filtration. The set of admissible controls $\mathcal{A}_{\text{adm}}$ consists of $\mathcal{F}_t$ -adapted processes with values in a compact convex set $A \subset \mathbb{R}^k$ satisfying $\mathbb{E} \int_0^T |\alpha_t|^2 dt < \infty$.

Each investor i controls its own portfolio state $X_t^i \in \mathbb{R}^d$, representing the vector of log-positions across d assets (i.e., the natural logarithm of the fraction of wealth allocated to each asset, or equivalently the log-portfolio weights). The objective is to maximise a risk-adjusted return over the finite horizon $[0, T]$:

$$J_i(\alpha^i, \alpha^{-i}) = \mathbb{E} \left[\int_0^T f(t, X_t^i, \mu_t^N, \alpha_t^i) dt + g(X_T^i, \mu_T^N) \right], \quad (1.1)$$

where

- $\alpha_t^i \in A \subset \mathbb{R}^k$ is investor i 's control (trading rate), assumed to be $\mathcal{F}_t$ -adapted;
- α^{-i} denotes the strategies of all other investors;
- $f : [0, T] \times \mathbb{R}^d \times \mathcal{P}_2(\mathbb{R}^d) \times A \rightarrow \mathbb{R}$ is the running reward (return net of trading costs);
- $g : \mathbb{R}^d \times \mathcal{P}_2(\mathbb{R}^d) \rightarrow \mathbb{R}$ is the terminal reward (e.g., liquidation value).

The payoff depends on the empirical measure of the population,

$$\mu_t^N := \frac{1}{N} \sum_{j=1}^N \delta_{X_t^j} \in \mathcal{P}_2(\mathbb{R}^d), \quad (1.2)$$

which captures price impact, crowding effects, and information spillovers. Here $\mathcal{P}_2(\mathbb{R}^d)$ denotes the space of probability measures on $\mathbb{R}^d$ with finite second moment, equipped with the Wasserstein-2 distance W_2 . The dependence of J_i on α^{-i} enters solely through the empirical measure μ_t^N , reflecting the mean-field interaction.

Each portfolio evolves according to the McKean–Vlasov stochastic differential equation (SDE)

$$dX_t^i = b(t, X_t^i, \mu_t^N, \alpha_t^i) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^i \sim \mu_{\text{init}}, \quad (1.3)$$

where

- $b : [0, T] \times \mathbb{R}^d \times \mathcal{P}_2(\mathbb{R}^d) \times A \rightarrow \mathbb{R}^d$ is the drift (expected price movement net of trading impact);
- $\sigma > 0$ is the idiosyncratic volatility; for $d > 1$ assets we set σ to be a scalar (for notational simplicity);

- $\sigma_0 \in \mathbb{R}^{d \times d_0}$ is the common-noise coefficient, and $W_t^0 \in \mathbb{R}^{d_0}$ is a shared Brownian motion representing aggregate shocks;
- $\mu_{\text{init}} \in \mathcal{P}_2(\mathbb{R}^d)$ is the initial distribution of positions.

Example 1.1 (Linear-quadratic portfolio MFG). For a single asset ($d = 1$), set

$$b(t, x, \mu, \alpha) = \kappa(\bar{x} - x) + \alpha, \quad f(t, x, \mu, \alpha) = -\left(x^2 + \frac{\lambda}{2}\alpha^2\right), \quad g(x, \mu) = -(x - \bar{x}_T)^2,$$

where $\bar{x} = \int y \mu(dy)$ is the population mean. Here the investor seeks to minimise a quadratic cost functional. This canonical LQ-MFG admits an explicit solution via Riccati equations and serves as the computational benchmark throughout Chapters 4 and 9.

1.2.2 Why the N -Player Problem Is Intractable

The N -player game is a stochastic differential game whose Nash equilibrium requires solving N coupled Hamilton–Jacobi–Bellman equations on $\mathbb{R}^{dN}$. Fundamental obstacles include dimensionality (for $d = 50$, $N = 200$ the state space has dimension 10,000), broken symmetry (heterogeneity in realised paths destroys closed-form solutions except in LQ cases), and exponential complexity of dynamic programming. These difficulties motivate the mean-field limit.

1.2.3 The Mean-Field Limit: $N \rightarrow \infty$

The Key Observation. (Lasry and Lions, 2007) and independently (Huang et al., 2006) observed that as $N \rightarrow \infty$ the empirical measure μ_t^N concentrates on a (possibly random) measure flow μ_t . By the conditional law of large numbers given the common noise W^0 , and using conditional propagation of chaos,

$$\mu_t^N \xrightarrow[N \rightarrow \infty]{W_2} \mu_t,$$

with convergence rate $O(N^{-1/2})$ for sufficiently regular measures ((Fournier and Guillin, 2015); see Chapter 2). Each investor then solves a single-agent control problem against the aggregate measure flow $\mu = (\mu_t)_{t \in [0, T]}$.

Definition 1.2 (MFG Nash equilibrium). A pair (μ^*, α^*) is an MFG Nash equilibrium if:

- (1) *Optimality*: Given μ^* , the control α^* maximises the representative agent’s objective with the empirical measure replaced by μ^* ;
- (2) *Consistency*: For almost all $t \in [0, T]$, $\mu_t^* = \mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0})$, where $\mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0})$ is a random probability measure measurable with respect to $\mathcal{F}_t^{W^0}$, and X^* satisfies (1.3) under α^* with μ^N replaced by μ^* .

This is the fixed-point problem $\Phi(\mu) = \mu$ on the space of measure flows, where $\Phi(\mu) := \mathcal{L}(X^\mu \mid \mathcal{F}_T^{W^0})$. **All contraction statements in this thesis refer to the fixed-point map on measure flows, not to pathwise contraction of individual trajectories.** Theorem 3.9 establishes that, under monotonicity and sufficient regularity ensuring stability of optimal transport maps, Φ is a contraction locally in time in general, and globally under strengthened convexity conditions in the weighted Wasserstein-2 metric $\|\mu^1 - \mu^2\|_{\lambda, T} := \sup_{t \leq T} e^{-\lambda t} \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2]^{1/2}$, guaranteeing existence and uniqueness.

The ε -Nash Property. If all N investors adopt α^* , the resulting profile is an ε_N -Nash equilibrium with $\varepsilon_N = O(N^{-1/2})$ ((Carmona and Delarue, 2018). Chapter 4 provides a quantitative bound calibrating this error against realistic transaction costs.

1.2.4 Why Common Noise Is Non-Trivial

Deterministic Case ($\sigma_0 = 0$). The measure flow μ_t is deterministic and the MFG reduces to a coupled system of PDEs. Existence, uniqueness, and stability follow under monotonicity conditions ((Lasry and Lions, 2007); (Carmona and Delarue, 2018).

Common-Noise Case ($\sigma_0 > 0$): New Challenges. When $\sigma_0 > 0$, the measure flow μ_t becomes a random element of $\mathcal{P}_2(\mathbb{R}^d)$, evolving as a stochastic process on an infinite-dimensional space. Three fundamental difficulties appear: (i) the Fokker–Planck equation becomes an SPDE; (ii) the consistency condition is measure-valued and must hold conditional on $\mathcal{F}^{W^0}$; (iii) proving convergence of particle systems requires compactness criteria on the Skorokhod space $D([0, T]; \mathcal{P}_2(\mathbb{R}^d))$ of càdlàg measure-valued processes, endowed with the usual J_1 topology (Jakubowski’s criterion or Mitoma’s theorem; see Chapter 2).

Key Structural Observation. The key asymmetry is that individual states are adapted to the full filtration $\mathcal{F}_t$, while the measure flow is adapted only to the common-noise filtration $\mathcal{F}_t^{W^0}$. When two candidate MFG solutions (μ^1, μ^2) are coupled on the same probability space sharing the identical common noise W^0 , the common-noise terms partially cancel in expectation under synchronous coupling and homogeneous common-noise exposure. **The contraction is established at the level of expected Wasserstein dynamics; it does not imply pathwise cancellation.** The first-order stochastic differential cancels only after taking expectations, and the contribution of the common noise enters through Lipschitz estimates on the adjoint BSDE and the control difference, affecting the drift of the expected squared Wasserstein distance. This cancellation is a second-moment (in-expectation) phenomenon under synchronous coupling, formalised in Theorem 3.9.

1.3 The FBSDE Characterisation

1.3.1 Overview of the System

The MFG equilibrium of Section 1.2 can be characterised by a coupled forward-backward stochastic differential equation (FBSDE) system, coupled with a measure-valued SPDE — the central mathematical object of this thesis. Every subsequent chapter either analyses this system theoretically (Chapters 2–3), discretises it (Chapter 4), learns it via reinforcement learning (Chapters 5–6), or extracts tradable signals from it (Chapters 7–9).

1.3.2 Heuristic Derivation

The dynamic programming principle yields a backward SPDE for the value function; the stochastic maximum principle gives the adjoint BSDE; the forward population evolution is governed by the Fokker–Planck SPDE; consistency closes the loop. Full rigorous

derivation is provided in Chapter 2, following the probabilistic approach of (Carmona and Delarue, 2018).

1.3.3 The Forward Equation: Population Measure Flow

Definition 1.3 (Fokker–Planck SPDE). Under the MFG equilibrium, the population measure $\mu_t \in \mathcal{P}_2(\mathbb{R}^d)$ satisfies, for all test functions $\phi \in C_c^\infty(\mathbb{R}^d)$,

$$d\langle \mu_t, \phi \rangle = \langle \mu_t, \mathcal{L}_t^* \phi \rangle dt + \langle \mu_t, \nabla \phi \cdot \sigma_0 \rangle dW_t^0, \quad \mu_0 = \mu_{\text{init}}, \quad (1.4)$$

where $\langle \mu, \phi \rangle := \int \phi d\mu$ and $\mathcal{L}_t^*$ is the formal adjoint of the generator.

The rigorous interpretation uses a Gelfand triple $V \subset H \subset V^*$ with $H = L^2(\mathbb{R}^d)$, $V = H^1(\mathbb{R}^d)$; solutions live in the Skorokhod space $D([0, T]; \mathcal{P}_2(\mathbb{R}^d))$ (see Chapter 2).

1.3.4 The Backward Equation: Adjoint Variables

Definition 1.4 (Adjoint BSDE). The adjoint process (p_t, q_t, r_t) satisfies

$$-dp_t = \partial_x H(t, X_t, \mu_t, p_t, \alpha_t^*) dt - q_t dW_t^i - r_t dW_t^0, \quad p_T = \partial_x g(X_T, \mu_T), \quad (1.5)$$

with Hamiltonian $H = p \cdot b + f$.

All processes are adapted to $\mathcal{F}_t$, while μ_t is adapted to $\mathcal{F}_t^{W^0}$. Under sufficient smoothness (Assumption A3), formal identification with Lions derivatives holds; these identifications are not required for existence proofs.

1.3.5 The Optimality Condition

Proposition 1.5 (Stochastic Maximum Principle). *Under Assumption A5 (uniform concavity of $\alpha \mapsto H$), the optimal control satisfies $\alpha_t^* = \arg \max_{\alpha \in A} H(t, X_t, \mu_t, p_t, \alpha)$. For the LQ example, $\alpha_t^* = p_t/\lambda$.*

1.3.6 The Consistency Condition

$$\mu_t = \mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0}) \quad \text{for all } t \in [0, T], \quad (1.6)$$

where X^* follows (1.3) with μ^N replaced by μ and control α^* . This is proved simultaneously with existence in Theorem 3.9.

1.3.7 The Complete FBSDE System

The system is solved as a fixed point in the space of measure flows: for a given μ , solve the forward–backward system for (X, p) , update μ via consistency, and iterate. The five coupled components are:

$$\begin{aligned} (\text{Fwd}) \quad & dX_t^* = b(t, X_t^*, \mu_t, \alpha_t^*) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^* \sim \mu_{\text{init}}, \\ (\text{FP}) \quad & d\langle \mu_t, \phi \rangle = \langle \mu_t, \mathcal{L}_t^* \phi \rangle dt + \langle \mu_t, \nabla \phi \cdot \sigma_0 \rangle dW_t^0, \quad \forall \phi \in C_c^\infty(\mathbb{R}^d), \\ (\text{Bwd}) \quad & -dp_t = \partial_x H(t, X_t, \mu_t, p_t, \alpha_t^*) dt - q_t dW_t^i - r_t dW_t^0, \quad p_T = \partial_x g(X_T, \mu_T), \\ (\text{Opt}) \quad & \alpha_t^* = \arg \max_{\alpha \in A} H(t, X_t, \mu_t, p_t, \alpha), \\ (\text{Cons}) \quad & \mu_t = \mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0}). \end{aligned} \quad (1.7)$$

Theorem 3.9 establishes, under Assumptions A1–A7, existence, uniqueness, and a quantitative stability estimate. Each component corresponds to an observable or estimable market quantity, enabling mapping to tradable signals (Chapter 8).

1.4 State of the Art

This section surveys the literature across three domains: mean-field game theory, reinforcement learning for MFGs, and quantitative finance applications. We highlight both foundational results and recent advances, and we identify the specific gaps that this thesis addresses. Throughout, we emphasise that this thesis builds incrementally on recent work, particularly (Carmona and Delarue, 2018), (Ahuja et al., 2019), (Angiuli et al., 2024), and (Cardaliaguet and Lehalle, 2018).

1.4.1 Mean-Field Game Theory

Foundations. The mean-field game framework was introduced independently by (Lasry and Lions, 2007) and (Huang et al., 2006). Lasry and Lions established existence and uniqueness of solutions for deterministic MFGs under monotonicity conditions via PDE methods. The probabilistic formulation was developed by (Carmona and Delarue, 2018), who provided a comprehensive treatment using forward-backward SDEs and extended the theory to the common-noise setting. (Cardaliaguet et al., 2019) introduced the master equation, a single equation on the Wasserstein space that encodes the entire MFG equilibrium, and developed the Lions derivative calculus necessary for its analysis.

Common-noise MFGs. The presence of common noise introduces substantial mathematical challenges. (Ahuja et al., 2019) proved well-posedness for a class of common-noise MFGs with linear-convex structure and weak-monotone cost functions, using a continuation method for FBSDEs with monotone functionals. (Cardaliaguet et al., 2025) established SPDE regularity results for MFGs with common noise and degenerate idiosyncratic noise. **The contraction method used in this thesis builds directly on the probabilistic FBSDE framework of (Carmona and Delarue, 2018) and the monotone functional approach of (Ahuja et al., 2019). The incremental contribution is the derivation of explicit, parameter-dependent stability constants that are not readily available in the prior literature.**

Trade crowding. (Cardaliaguet and Lehalle, 2018) formulated the classical optimal liquidation problem as a mean-field game of controls, providing a closed-form solution and demonstrating that traders do not need instantaneous knowledge of the whole system but can learn by observing others’ behaviour. (Lehalle and Mouzouni, 2019) extended this framework to portfolios of correlated instruments, deriving optimal trading flows and analysing their impact on perceived correlations using real data on 176 US stocks. (Neuman and Voß, 2023) studied a multi-player liquidation game with common signals and proved convergence at rate $O(N^{-2})$. This thesis builds upon this line of work by incorporating common noise and deriving a richer set of trading signals, most notably the novel crowding measure $c_t = W_2(\mu_t, \mu_t^*)$.

1.4.2 Reinforcement Learning for Mean-Field Games

Two-timescale methods. The application of reinforcement learning to MFGs has seen rapid progress. (Angiuli et al., 2024) introduced a unified two-timescale RL algorithm that solves either MFG or mean-field control problems depending on the ratio of learning rates. (Hu and Laurière, 2023) developed a two-timescale actor-critic method specifically for common-noise MFGs. These works demonstrate convergence for finite state-action spaces under discretisation, but do not provide explicit Wasserstein control or extend to the continuous-time SPDE setting.

Three-timescale methods and common noise. Most relevant to this thesis, (Angiuli et al., 2024) proved the convergence of a three-timescale Q-learning algorithm for mean-field control games with common noise, establishing that the multi-scale dynamics correctly capture the representative agent’s viewpoint. Their proof generalises Borkar (1997, 2008)’s two-timescale approach to three timescales, but the setting is restricted to finite state and action spaces. **The present thesis extends this ODE-method approach to the continuous-time, continuous-space setting with explicit Wasserstein distance control. The proof structure follows that of (Angiuli et al., 2024) closely; the novel technical challenge is establishing tightness and convergence of the measure-valued component in the Wasserstein topology.**

Entropy regularisation and exponential convergence. Recent work has established exponential convergence rates for entropy-regularised RL. (Ma et al., 2024) proved that entropy-regularised policy gradient achieves exponential convergence under a log-Sobolev inequality. Theorem 5.2 of this thesis adapts this approach to the MFG setting, providing an exponential rate for the Wasserstein distance between learned and equilibrium measures. **The result holds under an idealised log-Sobolev assumption and serves as a theoretical benchmark rather than a guarantee for practical deep RL implementations.**

1.4.3 Quantitative Finance

Factor models and crowding. Classical factor models (Fama and French, 1993) lack equilibrium grounding. The crowding problem — the tendency of quantitative strategies to become overcrowded and subsequently underperform — has been documented by (Khandani and Lo, 2011), but lacked a theoretically derived measure of crowding. The optimal execution literature ((Almgren and Chriss, 2000) is predominantly single-agent.

MFG applications in finance. Beyond the trade crowding models of (Cardaliaguet and Lehalle, 2018) and (Lehalle and Mouzouni, 2019), several works have applied MFGs to financial problems. (Carmona et al., 2015) modelled systemic risk as a mean-field game of interbank borrowing and lending. (Casgrain and Jaimungal, 2020) studied algorithmic trading with price impact using MFG methods. This thesis contributes to this growing literature by providing a unified framework from MFG equilibrium to empirically validated trading signals, with a novel emphasis on distributional disequilibrium measured by the Wasserstein distance.

1.4.4 Wasserstein Distance in Financial Applications

The use of the Wasserstein metric in finance has gained traction. (Fournier and Guillin, 2015) established the rate of convergence of the empirical measure in Wasserstein distance, providing the theoretical foundation for particle approximations. (Esfahani and Kuhn, 2018) developed distributionally robust optimisation using Wasserstein balls, providing tractable reformulations and performance guarantees. (Blanchet et al., 2019) applied Wasserstein profile inference to machine learning problems. (Kim et al., 2015) showed that temporal differences in return distributions measured by Wasserstein distance predict cross-sectional returns. The present thesis uses the Wasserstein distance in a novel way: as a measure of deviation from the MFG equilibrium, capturing distributional disequilibrium anchored to an economic equilibrium concept.

1.4.5 Gaps in the Literature

The survey reveals three specific gaps, each framed as an incremental extension of existing work:

- **Gap 1 (Theoretical):** While existence and uniqueness of common-noise MFGs have been established ((Ahuja et al., 2019); (Carmona and Delarue, 2018), explicit parameter-dependent stability bounds suitable for numerical analysis and learning algorithms are not readily available. This thesis contributes Theorem 3.9, which provides such constants under additional regularity assumptions.
- **Gap 2 (Algorithmic):** Existing RL frameworks for MFGs either assume finite state-action spaces ((Angiuli et al., 2024) or lack explicit Wasserstein control. No approach simultaneously provides convergence in the common-noise, continuous-space setting, explicit Wasserstein control, and a multi-timescale architecture aligned with the SPDE structure. This thesis contributes Theorems 5.1–5.2 and the hierarchical architecture of Chapter 6, albeit under strong function-approximation assumptions.
- **Gap 3 (Applied):** A fully integrated pipeline combining MFG equilibrium solution, theoretically justified trading signals, and rigorous empirical validation on real equity data remains largely unexplored. The work of (Lehalle and Mouzouni, 2019) provides an empirical analysis of correlations but does not derive and test a full trading strategy. This thesis resolves this gap through the alpha signals of Chapter 8 and the empirical backtest of Chapter 9.

1.5 The Three Critical Gaps

Gap 1 (Theoretical). Explicit parameter-dependent stability bounds suitable for numerical and learning algorithms are not readily available. While (Ahuja et al., 2019) and (Carmona and Delarue, 2018) established well-posedness, they did not provide the quantitative contraction constants needed for algorithmic design. Addressed by Theorem 3.9.

Gap 2 (Algorithmic). No existing RL framework simultaneously provides provable convergence to the equilibrium measure flow in the common-noise, continuous-space setting, with explicit Wasserstein control and a multi-timescale architecture aligned with

the SPDE structure. (Angiuli et al., 2024) proved three-timescale convergence for finite spaces; this thesis extends the ideas to continuous spaces under compactness and Lipschitz assumptions. **The extension is not trivial, but the proof methodology follows the same ODE framework.** Addressed by Theorems 5.1–5.2 and Proposition 6.1.

Gap 3 (Applied). A fully integrated pipeline combining MFG equilibrium solution, theoretically justified trading signals, and empirical validation on real equity data remains largely unexplored. Resolved by the sequence: Chapter 3 (theory) → Chapter 4 (numerics) → Chapter 8 (signals) → Chapter 9 (validation).

1.6 Principal Contributions

All results are subject to the standing assumptions of Section 1.9 and additional structural conditions discussed in corresponding chapters. Each contribution is explicitly situated relative to prior work. **Theorem 3.9 is the central structural result; all subsequent results — numerical stability, RL convergence, and empirical signal construction — depend explicitly on its contraction estimate.**

Theory (Well-Posedness)

Theorem 1.6 (Quantitative well-posedness). *Under Assumptions A1–A7, the FBSDE system (1.7) admits a unique solution. Under monotonicity and sufficient regularity ensuring stability of optimal transport maps (e.g., absolute continuity with uniformly bounded densities), the solution map is a contraction (locally in time in general, globally under strengthened convexity) in the weighted Wasserstein-2 metric, yielding an explicit stability estimate*

$$\mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2] \leq C_0 \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds,$$

where $C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$, with C_{reg} depending on bounds of Lions derivatives of b, f, g . The ε -Nash property holds with $\varepsilon_N = O(N^{-1/2})$. (Chapter 3)

Relation to prior work: Extends (Carmona and Delarue, 2018) and (Ahuja et al., 2019) by providing explicit parameter dependence in the contraction constant. The contraction is in expectation, not pathwise.

Numerics

Theorem 1.7 (Numerical convergence). *For the Euler–Maruyama discretisation of the Fokker–Planck SPDE, under sufficient smoothness, the weak error is $O(\Delta t^{1/2} + \Delta x)$. A CFL condition is derived. (Chapter 4)*

Relation to prior work: Builds on (Chassagneux et al., 2019) and (Achdou and Porretta, 2016), extending to the common-noise SPDE case.

Learning

Theorem 1.8 (Convergence of regularised multi-timescale actor-critic). *A regularised multi-timescale actor-critic algorithm with learning rates satisfying Robbins–Monro and timescale separation, under bounded Lipschitz function approximators (e.g., regularised neural networks), converges $\mathbb{P}$ -almost surely along subsequences (in the sense of $\liminf W_2 = 0$) to a fixed point of the limiting mean-field operator. Under uniqueness (Theorem 3.9), the limit is the MFG equilibrium. (Chapter 5)*

Relation to prior work: Extends (Angiuli et al., 2024) from finite spaces to continuous spaces with Wasserstein control; the subsequential convergence is a weaker result than full almost-sure convergence, which remains open.

Theorem 1.9 (Exponential rate under entropy regularisation). *Under entropy regularisation and a log-Sobolev inequality for the policy class (a restrictive assumption illustrating achievable rates under idealised conditions), the policy iteration scheme satisfies*

$$W_2(\mu_t, \mu^*) \leq C_0 e^{-\lambda t} W_2(\mu_0, \mu^*), \quad \lambda = \frac{\beta}{2C_{\text{Lip}}}.$$

(Chapter 5)

Relation to prior work: Adapts (Ma et al., 2024) to the MFG setting; the LSI assumption limits practical applicability.

Proposition 1.10 (Hierarchical architecture). *The SPDE decomposition implies a natural timescale separation exploited by a hierarchical RL architecture (SAC/PPO/TD3). (Chapter 6)*

Relation to prior work: Provides a structural motivation for algorithm choice; the specific algorithm mapping is heuristic, not proven optimal.

Finance

Proposition 1.11 (Four alpha signals). *The equilibrium structure implies four candidate signals: optimal spread, mean-reversion speed, factor exposure, and crowding measure $c_t = W_2(\mu_t, \mu_t^*)$. The optimal spread, mean-reversion speed, and factor exposure are adaptations of signals found in (Cardaliaguet and Lehalle, 2018) and (Lehalle and Mouzouni, 2019); the crowding measure c_t is the novel contribution. Unlike prices, μ_t^* is not directly observable; estimation introduces model risk (addressed in Chapter 9). Mean-variance portfolio weights with crowding discount are provided. (Chapter 8)*

Empirical Finding

Theorem 1.12 (Empirical validation). *A strategy based on Proposition 8.1 signals, applied to S&P 500 constituents (2000–2024, survivorship-bias-free), yields an estimated annualised Sharpe ratio of 1.9 under the model and estimation procedure, after transaction costs (5–10 bps) and slippage. Statistical significance: $p = 0.002$ (10,000 block-bootstrap). Maximum drawdown -18% . Metrics use purged walk-forward cross-validation. Robustness checks support qualitative conclusions. **The reported performance is conditional on model specification and estimation of μ_t^* ; the magnitude of the***

Sharpe ratio is sensitive to these choices. Model risk dominates statistical risk, and the empirical results should be interpreted as validation of the structural predictions of the model rather than evidence of a directly deployable trading strategy. The strategy assumes a small-player limit; a capacity analysis is provided with strong caveats. (Chapter 9)

Formal verification. The Gronwall estimate of Theorem 3.9, the Riccati system derivation, the CFL condition of Theorem 1.7, and the optimality of the linear feedback control have been formally verified in Lean 4. The code is available in the supplementary repository and is excerpted in Appendix B. **The verification is partial; full verification of the entire pipeline remains future work.**

Table 1.1: Summary of Principal Contributions and Their Novelty

Contribution	Prior Work	Incremental Novelty
Theorem 3.1 (Quantitative well-posedness)	(Carmona and Delarue, 2018), (Ahuja et al., 2019)	Explicit parameter-dependent contraction constant; contraction in expectation
Theorem 4.1 (Numerical convergence)	(Chassagneux et al., 2019), (Achdou and Porretta, 2016)	Combined scheme with CFL condition for common-noise SPDE
Theorem 5.1 (RL convergence)	(Angiuli et al., 2024)	Continuous-space extension with Wasserstein control; subsequential convergence
Theorem 5.2 (Exponential rate)	(Ma et al., 2024)	Application to MFG; idealised LSI benchmark
Proposition 6.1 (Hierarchical architecture)	Hu–Laurière (2023)	SPDE-grounded timescale decomposition; heuristic SAC-PPO-TD3 mapping
Proposition 8.1 (Alpha signals)	(Cardaliaguet and Lehalle, 2018), (Lehalle and Mouzouni, 2019)	Four-signal framework; novel crowding measure c_t
Empirical Finding 9.1	(Lehalle and Mouzouni, 2019)	Full out-of-sample strategy on S&P 500; Sharpe 1.91 (conditional on model)

1.7 Connections Between Fields

Connection 1 (SPDE $\rightarrow$ RL architecture). The master equation’s decomposition into idiosyncratic, mean-field, and common-noise operators motivates a hierarchical RL architecture (SAC/PPO/TD3). Proposition E.2 provides timescale separation.

Connection 2 (Stability $\rightarrow$ Crisis robustness). The coupled Wasserstein dynamics show the contraction constant remains controlled under large common shocks in expectation. The crowding signal $c_t = W_2(\mu_t, \mu_t^*)$ provides a quantitative measure of market stress. (In extreme stress, orthogonality properties may degrade; this is acknowledged as a limitation.)

Connection 3 (Equilibrium $\rightarrow$ Alpha). Definition 7.1 formalises market efficiency as $\mu_t = \mu_t^*$. The Wasserstein distance measures deviation from Nash equilibrium. When c_t is large, profitable deviations are predicted; when small, the market is approximately efficient.

1.8 Thesis Roadmap

1.8.1 Chapter Dependency Graph

Theory stream: Chapter 2 $\rightarrow$ 3 $\rightarrow$ 4 $\rightarrow$ 5 $\rightarrow$ 6. Finance stream: Chapter 7 $\rightarrow$ 8 $\rightarrow$ 9. All synthesised in Chapter 10.

1.8.2 Chapter-by-Chapter Summary

- **Chapter 2:** Preliminaries (Wasserstein geometry, Itô calculus, SPDEs, Lions derivative).
- **Chapter 3:** Quantitative well-posedness (Theorem 3.9).
- **Chapter 4:** Numerical analysis (Theorem 1.7).
- **Chapter 5:** RL convergence (Theorems D.1, D.5).
- **Chapter 6:** Hierarchical architecture (Proposition E.2).
- **Chapter 7:** Case study — systemic risk, bridging theory to application.
- **Chapter 8:** Alpha generation (Proposition 1.11).
- **Chapter 9:** Empirical validation (Empirical Finding 1.12).
- **Chapter 10:** Synthesis, limitations, open problems.

1.9 Key Notation and Standing Assumptions

Reference section; may be skipped on first reading.

1.9.1 Notation Reference

Symbol	Meaning
$\mathcal{P}_2(\mathbb{R}^d)$	probability measures with finite second moment
$W_2(\mu, \nu)$	Wasserstein-2 distance
$\langle \mu, \phi \rangle$	$\int \phi d\mu$
$\mathcal{L}(X)$	law of X
W_t^i, W_t^0	idiosyncratic and common Brownian motions
$\mathcal{F}_t^{W^0}$	common-noise filtration
b, f, g	drift, running reward, terminal reward
$H = p \cdot b + f$	Hamiltonian
p_t, q_t, r_t	adjoint processes
$U(t, \mu)$	master-equation value function; $D_\mu U$: Lions derivative
L_b, L_f, L_g	Lipschitz constants
β	entropy regularisation coefficient
$c_t = W_2(\mu_t, \mu_t^*)$	crowding measure
$D([0, T]; E)$	Skorokhod space (càdlàg), J_1 topology
$C_c^\infty(\mathbb{R}^d)$	smooth compactly supported functions
$\ \sigma_0\ _F$	Frobenius norm

1.9.2 Standing Assumptions

Assumption 1.13 (Lipschitz drift). $\|b(t, x, \mu, \alpha) - b(t, x', \mu', \alpha')\| \leq L_b(\|x - x'\| + W_2(\mu, \mu') + \|\alpha - \alpha'\|)$.

Assumption 1.14 (Lipschitz rewards). f and g are Lipschitz.

Assumption 1.15 (Additional regularity). b, f, g are C^1 in x and Lions-differentiable in μ , with bounded first derivatives (and second derivatives when required).

Assumption 1.16 (Bounded common noise). $\|\sigma_0\|_F \leq M_\sigma < \infty$.

Assumption 1.17 (Uniform concavity of Hamiltonian). $H(t, x, \mu, p, \alpha') \leq H(t, x, \mu, p, \alpha) + \partial_\alpha H \cdot (\alpha' - \alpha) - \frac{\lambda_H}{2} \|\alpha' - \alpha\|^2$.

Assumption 1.18 (Regular terminal reward). g is C^1 in x and Lions-differentiable in μ , with bounded derivatives.

Assumption 1.19 (Initial condition). μ_{init} has finite fourth moment.

Remark 1.20. Assumptions A1–A7 are standard in the MFG literature but are restrictive. In particular, A3 (Lions differentiability) and the requirement of absolute continuity for optimal transport stability are strong and are not satisfied by empirical discrete measures without regularisation. The theoretical results should be viewed as establishing a well-posedness baseline under idealised conditions; practical implementation requires regularisation and model risk quantification.

1.10 Guide to Reading

- **Path A (Pure mathematics):** Chapters 2, 3, 5.

- **Path B (Algorithms/RL):** Chapters 2 (overview), 5, 6.
- **Path C (Quantitative finance):** Chapters 2 (Wasserstein), 7, 8, 9.

1.11 Why This Problem, Why Now

Three developments converged in 2023–2025: mathematical maturity (([Cardaliaguet et al., 2025](#)) resolved SPDE regularity for common-noise MFGs; ([Ahuja et al., 2019](#)) established well-posedness for monotone functionals), computational feasibility (GPU-accelerated optimal transport via the Sinkhorn algorithm of ([Cuturi, 2013](#)), mature RL libraries), and market structure (post-COVID concentrated participation, common shocks, crowding events).

Central message. This thesis argues that modern financial markets are naturally modelled as mean-field games: large populations of agents optimise against aggregate behaviour under common shocks. A key structural result is that common noise preserves the contraction structure of the MFG fixed-point map in expectation under synchronous coupling and homogeneous exposure. The Wasserstein distance $W_2(\mu_t, \mu_t^*)$ provides a mathematically coherent measure of crowding — sensitive to displacement and mass redistribution, not merely density differences. When this distance is large, theory predicts profitable deviations; when small, the market is approximately efficient. Proposition 1.11 turns this insight into tradable signals; Chapter 9 provides empirical evidence, with appropriate statistical caution, that the approach has out-of-sample merit. **In this sense, the thesis reframes market inefficiency not as a price deviation, but as a distributional disequilibrium under shared shocks.**

Chapter 2

Mathematical Preliminaries

2.1 Overview

This chapter develops the complete mathematical toolkit required for all subsequent chapters. No prior knowledge of optimal transport or stochastic partial differential equations is assumed; every result is stated with either a complete proof or a detailed proof sketch, with references to the authoritative sources for full details. The standing assumptions A1–A7 referenced throughout are collected in Section 1.9 of Chapter 1.

The chapter covers six interlocking topics:

1. **Wasserstein-2 space** ($\mathcal{P}_2(\mathbb{R}^d), W_2$): optimal transport, metric space structure, geodesics, weak convergence, rates of convergence for empirical measures, and the Lions derivative.
2. **Brownian motion and Itô calculus**: construction, quadratic variation, the Itô integral, Itô’s formula, and stochastic differential equations.
3. **Fokker–Planck SPDE**: derivation from Itô’s formula, weak formulation, strong form, Sobolev regularity, and the Gelfand triple framework. We clarify the distinction between the linear Fokker–Planck SPDE and the nonlinear Kushner–Zakai equation.
4. **Martingale theory**: martingales, stopping times, Doob’s inequality, the Burkholder–Davis–Gundy inequality, martingale representation, and Girsanov’s theorem.
5. **Sobolev spaces and functional analysis**: Sobolev embedding, Hahn–Banach extension, Lumer–Phillips theorem, spectral gap, and log-Sobolev inequalities.
6. **Stochastic Maximum Principle**: the control problem, Hamiltonian, adjoint BSDE, optimality conditions, and connections to Hamilton–Jacobi–Bellman equations.

A summary table at the end of the chapter (Section 2.8) maps each tool to its first use in the thesis, serving as a quick reference. For readers seeking a deeper treatment, we recommend (Villani, 2009) for optimal transport, (Carmona and Delarue, 2018), Vol. I, Chapters 1–2) for Wasserstein analysis in the MFG context, and Karatzas–Shreve (1991) for stochastic calculus.

2.2 Wasserstein-2 Space

2.2.1 Probability Measures with Finite Second Moment

Let $\mathcal{B}(\mathbb{R}^d)$ denote the Borel σ -algebra on $\mathbb{R}^d$. The space $\mathcal{P}_2(\mathbb{R}^d)$ consists of all Borel probability measures μ on $\mathbb{R}^d$ with finite second moment:

$$\mathcal{P}_2(\mathbb{R}^d) := \left\{ \mu \in \mathcal{P}(\mathbb{R}^d) \mid \int_{\mathbb{R}^d} |x|^2 \mu(dx) < \infty \right\}. \quad (2.1)$$

Elements of $\mathcal{P}_2(\mathbb{R}^d)$ arise naturally in our setting:

- As the law $\mu = \mathcal{L}(X)$ of an $\mathbb{R}^d$ -valued random variable with $\mathbb{E}[|X|^2] < \infty$.
- As the empirical measure $\mu^N = \frac{1}{N} \sum_{i=1}^N \delta_{x_i}$ when the particles satisfy $\frac{1}{N} \sum_i |x_i|^2 < \infty$.
- As the mean-field distribution μ_t in the MFG system, provided the agents' positions remain square-integrable in expectation.

2.2.2 Optimal Transport and the Wasserstein Distance

Definition 2.1 (Coupling and Wasserstein-2 Distance). A coupling of $\mu, \nu \in \mathcal{P}_2(\mathbb{R}^d)$ is a probability measure $\pi \in \mathcal{P}(\mathbb{R}^d \times \mathbb{R}^d)$ with marginals μ and ν : $\pi(\cdot \times \mathbb{R}^d) = \mu$, $\pi(\mathbb{R}^d \times \cdot) = \nu$. The set of all couplings is denoted $\Pi(\mu, \nu)$. The Wasserstein-2 distance between μ and ν is

$$W_2(\mu, \nu) := \left(\inf_{\pi \in \Pi(\mu, \nu)} \int_{\mathbb{R}^d \times \mathbb{R}^d} |x - y|^2 \pi(dx, dy) \right)^{1/2}. \quad (2.2)$$

The infimum in (2.2) is attained: there exists an optimal coupling $\pi^* \in \Pi(\mu, \nu)$ achieving the infimum ((Villani, 2009), Theorem 5.10). The optimal π^* is supported on the subdifferential of a convex function, the Brenier potential.

Theorem 2.2 (Brenier's Theorem, 1991). *Let $\mu, \nu \in \mathcal{P}_2(\mathbb{R}^d)$ with μ absolutely continuous with respect to Lebesgue measure. Then there exists a unique (up to μ -a.e. equality) convex function $\varphi : \mathbb{R}^d \rightarrow \mathbb{R}$ such that the map $T = \nabla \varphi$ satisfies $T_{\#}\mu = \nu$ (i.e., $\nu(A) = \mu(T^{-1}(A))$ for all A) and*

$$W_2(\mu, \nu)^2 = \int_{\mathbb{R}^d} |x - \nabla \varphi(x)|^2 \mu(dx).$$

The optimal coupling is $\pi^ = (\text{id} \times T)_{\#}\mu$, i.e., the law of $(X, T(X))$ when $X \sim \mu$.*

Proof sketch. The proof proceeds in three steps: (i) Kantorovich duality reduces the problem to finding a pair of conjugate convex functions (φ, φ^*) ; (ii) the optimality condition $y \in \partial \varphi(x)$ for π^* -a.e. (x, y) implies the transport is a gradient map; (iii) uniqueness follows from strict convexity of the cost $c(x, y) = |x - y|^2$. See (Villani, 2009), Ch. 9) for the complete argument. $\square$

Remark 2.3 (Discrete measures and regularisation). Brenier's Theorem requires absolute continuity of the source measure μ . In financial applications (Chapter 9), the empirical market measure μ_t is a sum of Dirac masses, and the equilibrium measure μ_t^* is estimated

from discrete data. The optimal transport between discrete measures is not uniquely given by a gradient of a convex function in the classical sense. The Sinkhorn algorithm (Section 2.2.3) bridges this gap: by adding an entropic penalty, the discrete measures are effectively “blurred” into continuous approximations, recovering a gradient-like structure and enabling stable computation.

Example 2.4 (One-dimensional case: sorted coupling). In $d = 1$, the optimal transport has a particularly elegant form. For $\mu, \nu \in \mathcal{P}_2(\mathbb{R})$ with distribution functions F_μ, F_ν and quantile functions F_μ^{-1}, F_ν^{-1} :

$$W_2(\mu, \nu)^2 = \int_0^1 |F_\mu^{-1}(u) - F_\nu^{-1}(u)|^2 du. \quad (2.3)$$

The optimal map is $T(x) = F_\nu^{-1}(F_\mu(x))$ — transport the u -quantile of μ to the u -quantile of ν (sorted coupling). This is computed exactly in $O(N \log N)$ by sorting the particles. **The synchronous coupling used in Chapter 3 for common-noise cancellation can be explicitly constructed in 1-D via this quantile coupling.**

2.2.3 The Sinkhorn Algorithm

For computational purposes, we regularise the Kantorovich problem with an entropic penalty:

$$W_{2,\epsilon}(\mu, \nu)^2 := \inf_{\pi \in \Pi(\mu, \nu)} \int |x - y|^2 \pi + \epsilon \text{KL}(\pi \| \mu \otimes \nu), \quad (2.4)$$

where $\text{KL}(\pi \| \mu \otimes \nu) = \int \log \frac{d\pi}{d(\mu \otimes \nu)} d\pi$. The solution has the form $\pi_\epsilon^*(dx, dy) = u(x)K(x, y)v(y) \mu(dx)\nu(dy)$ where $K(x, y) = e^{-|x-y|^2/\epsilon}$ is the Gaussian kernel, and (u, v) satisfies the Sinkhorn fixed-point equations:

$$u_i = a_i / \sum_j K_{ij} v_j, \quad v_j = b_j / \sum_i K_{ij} u_i, \quad (2.5)$$

where $a_i = \mu(\{x_i\})$, $b_j = \nu(\{y_j\})$. This converges in $O(N^2/\epsilon^2)$ operations ((Cuturi, 2013). In Chapter 4, this regularised distance will be used to compute Wasserstein gradients efficiently.

2.2.4 Metric Space Structure

Theorem 2.5 ($(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is a Complete Separable Metric Space). *The Wasserstein-2 distance satisfies:*

- (i) *Non-degeneracy:* $W_2(\mu, \nu) = 0 \iff \mu = \nu$;
- (ii) *Symmetry:* $W_2(\mu, \nu) = W_2(\nu, \mu)$;
- (iii) *Triangle inequality:* $W_2(\mu, \rho) \leq W_2(\mu, \nu) + W_2(\nu, \rho)$;
- (iv) *Completeness:* every Cauchy sequence in $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ converges to an element of $\mathcal{P}_2(\mathbb{R}^d)$;
- (v) *Separability:* $\mathcal{P}_2(\mathbb{R}^d)$ contains a countable dense subset.

Proof of (iii) via gluing lemma. Given μ, ν, ρ , let $\pi_{12} \in \Pi(\mu, \nu)$ and $\pi_{23} \in \Pi(\nu, \rho)$ be any couplings. By the gluing lemma ((Villani, 2009), Lemma 7.6), there exists $\pi_{123} \in \mathcal{P}(\mathbb{R}^d \times \mathbb{R}^d \times \mathbb{R}^d)$ with marginals (π_{12}, π_{23}) , i.e., the joint distribution of (X, Y, Z) where $(X, Y) \sim \pi_{12}$, $(Y, Z) \sim \pi_{23}$. The marginal $\pi_{13} = (\text{proj}_{13})_{\#} \pi_{123}$ is a coupling of μ and ρ . By Minkowski's inequality:

$$W_2(\mu, \rho) \leq \left(\int |x - z|^2 \pi_{13}(dx, dz) \right)^{1/2} = \mathbb{E}[|X - Z|^2]^{1/2} \leq \mathbb{E}[|X - Y|^2]^{1/2} + \mathbb{E}[|Y - Z|^2]^{1/2}.$$

Optimising over π_{12} and π_{23} gives the triangle inequality. The remaining properties are proved similarly. $\square$

2.2.5 Weak Convergence in $\mathcal{P}_2$

Theorem 2.6 (Characterisation of W_2 -convergence). *The following are equivalent for $\mu_n, \mu \in \mathcal{P}_2(\mathbb{R}^d)$:*

- (i) $W_2(\mu_n, \mu) \rightarrow 0$;
- (ii) $\mu_n \xrightarrow{w} \mu$ weakly and $\int |x|^2 \mu_n(dx) \rightarrow \int |x|^2 \mu(dx)$;
- (iii) For every continuous $f : \mathbb{R}^d \rightarrow \mathbb{R}$ with $|f(x)| \leq C(1 + |x|^2)$: $\int f d\mu_n \rightarrow \int f d\mu$.

Condition (ii) shows that W_2 -convergence is strictly stronger than weak convergence: it additionally requires convergence of the second moment. **This is precisely the correct topology for the MFG analysis, since $\mu_t = \mathcal{L}(X_t)$ and $\mathbb{E}[|X_t|^2]$ must remain finite and continuous in t .**

2.2.6 Geodesics and the Lions Derivative

The space $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ has a rich geometric structure that motivates the Lions derivative used in the master equation.

Definition 2.7 (Wasserstein Geodesic). The geodesic (shortest path) from μ_0 to μ_1 in $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is the curve

$$\mu_t := ((1 - t) \text{id} + tT)_{\#} \mu_0, \quad t \in [0, 1],$$

where $T = \nabla \varphi$ is the Brenier optimal transport map from μ_0 to μ_1 . This interpolates linearly in position space: the x -atom of μ_0 travels in a straight line to $T(x)$ at constant speed.

Proposition 2.8 (Geodesic is length-minimising). *The path $(\mu_t)_{t \in [0, 1]}$ satisfies $W_2(\mu_s, \mu_t) = |t - s| W_2(\mu_0, \mu_1)$ for all $s, t \in [0, 1]$.*

The Lions derivative (or L -derivative) formalises differentiation of functionals $F : \mathcal{P}_2(\mathbb{R}^d) \rightarrow \mathbb{R}$ along such geodesics. Rather than being a pointwise derivative, it is a linear functional on the tangent space. For a functional F that is Lions-differentiable, there exists a measurable function $D_{\mu}F(\mu) : \mathbb{R}^d \rightarrow \mathbb{R}^d$ (unique up to μ -a.e. equality) such that for any $\nu \in \mathcal{P}_2(\mathbb{R}^d)$ with optimal transport map $T = \nabla \varphi$ from μ to ν ,

$$\left. \frac{d}{dt} \right|_{t=0} F(\mu_t) = \int_{\mathbb{R}^d} D_{\mu}F(\mu)(x) \cdot (T(x) - x) \mu(dx). \quad (2.6)$$

In other words, the directional derivative of F along a geodesic with initial velocity $v = T - \text{id}$ is given by the $L^2(\mu)$ inner product of $D_{\mu}F(\mu)$ and v . For a comprehensive treatment, see (Carmona and Delarue, 2018), Vol. I, Chapter 5).

2.2.7 Law of Large Numbers in Wasserstein Distance

Theorem 2.9 (LLN Rate in W_2 , (Fournier and Guillin, 2015)). Let $\mu \in \mathcal{P}_2(\mathbb{R}^d)$ and $X_1, X_2, \dots$ be i.i.d. with law μ . Define $\mu^N = \frac{1}{N} \sum_{i=1}^N \delta_{X_i}$.

(i) Almost-sure convergence: $W_2(\mu^N, \mu) \rightarrow 0$ a.s. as $N \rightarrow \infty$.

(ii) Rate: If $\int_{\mathbb{R}^d} |x|^q \mu(dx) < \infty$ for some $q > 4$, then for $d = 1$:

$$\mathbb{E}[W_2(\mu^N, \mu)] \leq C_\mu \cdot N^{-1/2}.$$

For $d \geq 3$: $\mathbb{E}[W_2(\mu^N, \mu)] \leq C_{\mu,d} \cdot N^{-1/d}$, under appropriate moment conditions. (See (Fournier and Guillin, 2015), Theorem 1 and Corollary 2.)

Proof of (i) in $d = 1$. By the explicit formula (2.3):

$$W_2(\mu^N, \mu)^2 = \int_0^1 |F_N^{-1}(u) - F^{-1}(u)|^2 du.$$

For fixed u , $F_N^{-1}(u) \rightarrow F^{-1}(u)$ a.s. by the Glivenko–Cantelli theorem. Since $|F_N^{-1}(u) - F^{-1}(u)|^2 \leq 4 \max(|F^{-1}(u)|, |F_N^{-1}(u)|)^2$ and the right-hand side is uniformly integrable by the second moment assumption, the dominated convergence theorem gives $W_2(\mu^N, \mu)^2 \rightarrow 0$ a.s. $\square$

Remark 2.10 (Curse of dimensionality and mean-field mitigation). The rate $N^{-1/d}$ for $d \geq 3$ highlights the curse of dimensionality in Wasserstein space: for a portfolio of $d = 50$ assets, the empirical measure converges exceedingly slowly. **In practice, the mean-field assumption — coupled with the specific convex structure of the rewards and the contraction property established in Theorem 3.1 — allows the RL architecture (Chapters 5–6) to circumvent this theoretical bottleneck.** The Wasserstein distance is used primarily as an analytical tool for stability, not as the primary object of high-dimensional statistical estimation.

Remark 2.11 (Significance for the thesis). Theorem 2.9 justifies two central approximations: (a) replacing the N -player empirical measure μ_t^N by the mean-field distribution μ_t with error $O(N^{-1/2})$ (for $d = 1$ under sufficient moments); (b) validating the block-bootstrap consistency in Theorem 9.1, where $B = 10,000$ samples give $O(10^{-2.5})$ estimation error.

2.2.8 Wasserstein Barycentre

Definition 2.12 (Fréchet Mean / Barycentre). Given measures $\mu_1, \dots, \mu_K \in \mathcal{P}_2(\mathbb{R}^d)$ with weights $\lambda_k > 0$, $\sum_k \lambda_k = 1$, the Wasserstein barycentre is

$$\bar{\mu} = \arg \min_{\nu \in \mathcal{P}_2} \sum_{k=1}^K \lambda_k W_2(\nu, \mu_k)^2.$$

The barycentre exists and is unique when at least one μ_k is absolutely continuous ((Agueh and Carlier, 2011). It is computed by the free-support Sinkhorn barycentre algorithm ((Cuturi and Doucet, 2014). *Note: this concept is included for completeness but is not directly used in the subsequent chapters.*

2.3 Brownian Motion and Itô Calculus

2.3.1 Brownian Motion: Construction and Properties

Definition 2.13 (Standard Brownian Motion). A standard Brownian motion (Wiener process) $B = (B_t)_{t \geq 0}$ on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ is a continuous stochastic process satisfying:

- (i) $B_0 = 0$ almost surely;
- (ii) Independent increments: for $0 \leq s < t$, $B_t - B_s$ is independent of $\mathcal{F}_s$;
- (iii) Gaussian increments: $B_t - B_s \sim N(0, t - s)$ for all $0 \leq s < t$;
- (iv) Continuous paths: $t \mapsto B_t(\omega)$ is continuous for $\mathbb{P}$ -almost every ω .

Proposition 2.14 (Key path properties). *Brownian motion paths are almost surely:*

- (i) Hölder continuous of order $\alpha < \frac{1}{2}$ (but not $\frac{1}{2}$);
- (ii) Nowhere differentiable: $\lim_{h \rightarrow 0} \frac{B_{t+h} - B_t}{h}$ does not exist for any t ;
- (iii) Of unbounded variation: $\sup_{\Pi} \sum_k |B_{t_k} - B_{t_{k-1}}| = \infty$ over all partitions Π ;
- (iv) Of finite quadratic variation: $[B, B]_t = t$ (see below).

2.3.2 Quadratic Variation

Definition 2.15 (Quadratic Variation). The quadratic variation of a semimartingale M is the process $[M, M]_t$ defined as the limit in probability:

$$[M, M]_t = \lim_{|\Pi| \rightarrow 0} \sum_k (M_{t_k} - M_{t_{k-1}})^2,$$

where the limit is over partitions $\Pi = \{0 = t_0 < t_1 < \dots < t_n = t\}$ as the mesh $|\Pi| = \max_k (t_k - t_{k-1}) \rightarrow 0$.

Proposition 2.16 (Quadratic variation of Brownian motion). *For standard Brownian motion B : $[B, B]_t = t$ a.s.*

Proof sketch. Let $V_n = \sum_{k=1}^n (B_{kh} - B_{(k-1)h})^2$ with $h = t/n$. Then $\mathbb{E}[V_n] = n \cdot h = t$ and $\text{Var}(V_n) = n \cdot 2h^2 = 2t^2/n \rightarrow 0$. So $V_n \rightarrow t$ in L^2 , and by Borel–Cantelli, a.s. along a subsequence. $\square$

Remark 2.17 (Why quadratic variation matters: the Itô correction). The fact that $[B, B]_t = t \neq 0$ is the source of all the “extra terms” in stochastic calculus. For a smooth function f and a smooth path x , $f(x_t) = f(x_0) + \int f'(x_s) dx_s$ (ordinary chain rule). But for a Brownian path, the Taylor remainder $\frac{1}{2}f''(B_s)(dB_s)^2$ contributes $\frac{1}{2}f''(B_s) ds$ (not zero), because $(dB)^2 = dt$. This is Itô’s formula.

2.3.3 The Itô Integral

The classical Riemann–Stieltjes integral $\int_0^T H_t dB_t$ is not well-defined pathwise because B has unbounded variation. Itô's construction extends it to square-integrable adapted integrands.

Definition 2.18 (Itô Integral). Let $(\mathcal{F}_t)$ be the natural filtration of B , and $H = (H_t)_{t \in [0, T]}$ be $(\mathcal{F}_t)$ -adapted with $\mathbb{E}[\int_0^T H_t^2 dt] < \infty$. The Itô integral $\int_0^T H_t dB_t$ is the $L^2(\mathbb{P})$ -limit of Riemann sums with left endpoints:

$$\int_0^T H_t dB_t := L^2\text{-}\lim_{n \rightarrow \infty} \sum_{k=0}^{n-1} H_{t_k} (B_{t_{k+1}} - B_{t_k}).$$

Theorem 2.19 (Itô Isometry). *For adapted square-integrable H :*

$$\mathbb{E} \left[\left(\int_0^T H_t dB_t \right)^2 \right] = \mathbb{E} \left[\int_0^T H_t^2 dt \right]. \quad (2.7)$$

Proof sketch. For simple processes, the expectation of the square expands to $\sum_k \mathbb{E}[H_{t_k}^2] \Delta t_k$ by independence. The general case follows by approximation. $\square$

Remark 2.20 (Significance for Theorem 3.1). The Itô isometry is used in the proof of Theorem 3.1: after deriving the Wasserstein distance ODE $dW_2^2 \leq CW_2^2 dt + dM_t$, we take expectations and use $\mathbb{E}[M_t] = 0$ (since $M_t = \int_0^t Z_s dW_s^0$ is a martingale by Itô isometry, as $\mathbb{E}[\int Z_s^2 ds] < \infty$). This eliminates the stochastic term and allows the Gronwall inequality to close.

2.3.4 Itô's Formula

Theorem 2.21 (Itô's Formula, one-dimensional). *Let $f \in C^2(\mathbb{R})$ and $X_t = X_0 + \int_0^t b_s ds + \int_0^t \sigma_s dB_s$ be an Itô process. Then:*

$$f(X_t) = f(X_0) + \int_0^t f'(X_s) dX_s + \frac{1}{2} \int_0^t f''(X_s) d[X, X]_s.$$

Expanding $d[X, X]_s = \sigma_s^2 ds$:

$$f(X_t) = f(X_0) + \int_0^t f'(X_s) b_s ds + \int_0^t f'(X_s) \sigma_s dB_s + \frac{1}{2} \int_0^t f''(X_s) \sigma_s^2 ds. \quad (2.8)$$

Theorem 2.22 (Itô's Formula, multidimensional). *Let $f \in C^2(\mathbb{R}^d)$ and let $X_t \in \mathbb{R}^d$ satisfy $dX_t = b_t dt + \sigma_t dW_t + \sigma_0 dW_t^0$ (where $W_t \in \mathbb{R}^d$, $W_t^0 \in \mathbb{R}^{d_0}$, and $\sigma_t \in \mathbb{R}^{d \times d}$, $\sigma_0 \in \mathbb{R}^{d \times d_0}$). Then:*

$$\begin{aligned} f(X_t) = f(X_0) &+ \int_0^t \nabla f(X_s) \cdot b_s ds + \int_0^t \nabla f(X_s) \cdot \sigma_s dW_s + \int_0^t \nabla f(X_s) \cdot \sigma_0 dW_s^0 \\ &+ \frac{1}{2} \int_0^t \text{tr}(\sigma_s \sigma_s^\top D^2 f(X_s)) ds + \frac{1}{2} \int_0^t \text{tr}(\sigma_0 \sigma_0^\top D^2 f(X_s)) ds. \end{aligned}$$

Example 2.23 (Applying Itô to a test function). Let X_t^i be a single agent satisfying $dX_t^i = b_t dt + \sigma dW_t^i + \sigma_0 dW_t^0$. Apply Theorem 2.22 with $f = \phi$ (a smooth test function):

$$d\phi(X_t^i) = \nabla\phi(X_t^i) \cdot b_t dt + \sigma \nabla\phi(X_t^i) \cdot dW_t^i + \sigma_0 \nabla\phi(X_t^i) \cdot dW_t^0 + \frac{\sigma^2}{2} \Delta\phi(X_t^i) dt + \frac{1}{2} \text{tr}(\sigma_0 \sigma_0^\top D^2\phi(X_t^i)) dt.$$

Averaging over $i = 1, \dots, N$ and writing $\langle \mu_t^N, \phi \rangle = \frac{1}{N} \sum_i \phi(X_t^i)$:

$$d\langle \mu_t^N, \phi \rangle = \left\langle \mu_t^N, b_t \cdot \nabla\phi + \frac{\sigma^2}{2} \Delta\phi + \frac{1}{2} \text{tr}(\sigma_0 \sigma_0^\top D^2\phi) \right\rangle dt + \underbrace{\frac{\sigma}{N} \sum_i \nabla\phi(X_t^i) dW_t^i}_{\rightarrow 0 \text{ as } N \rightarrow \infty} + \sigma_0 \langle \mu_t^N, \nabla\phi \rangle dW_t^0.$$

The idiosyncratic term vanishes in L^2 as $N \rightarrow \infty$ (mutual independence of W^i). The common noise term survives because W_t^0 is the same for all agents. Taking $N \rightarrow \infty$ yields the Fokker–Planck SPDE of Section 2.4.

2.3.5 Itô's Product Rule and Integration by Parts

Proposition 2.24 (Itô Product Rule). *For Itô processes X_t and Y_t :*

$$d(X_t Y_t) = X_t dY_t + Y_t dX_t + d[X, Y]_t,$$

where $[X, Y]_t = \int_0^t \sigma_s^X \sigma_s^Y ds$ is the quadratic covariation (vanishes if driven by independent Brownians).

Example 2.25 (Integration by parts for stochastic integrals). $B_t^2 = 2 \int_0^t B_s dB_s + t$. (Verify: apply (2.8) with $f(x) = x^2$, $f'(x) = 2x$, $f''(x) = 2$.)

2.3.6 Stochastic Differential Equations

Theorem 2.26 (Existence and Uniqueness for SDEs). *Let $b : [0, T] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$ and $\sigma : [0, T] \times \mathbb{R}^d \rightarrow \mathbb{R}^{d \times m}$ satisfy for some $L > 0$:*

$$(i) \text{ Lipschitz: } |b(t, x) - b(t, y)| + \|\sigma(t, x) - \sigma(t, y)\| \leq L|x - y|;$$

$$(ii) \text{ Linear growth: } |b(t, x)| + \|\sigma(t, x)\| \leq L(1 + |x|).$$

Then for any $\mathcal{F}_0$ -measurable X_0 with $\mathbb{E}[|X_0|^2] < \infty$, the SDE $dX_t = b(t, X_t) dt + \sigma(t, X_t) dW_t$ has a unique strong solution $(X_t)_{t \in [0, T]}$ satisfying

$$\mathbb{E} \left[\sup_{t \in [0, T]} |X_t|^2 \right] \leq C(1 + \mathbb{E}[|X_0|^2])e^{CT}.$$

Proof sketch — Picard iteration. Define $X_t^0 \equiv X_0$ and $X_t^{n+1} = X_0 + \int_0^t b(s, X_s^n) ds + \int_0^t \sigma(s, X_s^n) dW_s$. By Itô isometry and the Lipschitz condition:

$$\mathbb{E} \left[\sup_{s \leq t} |X_s^{n+1} - X_s^n|^2 \right] \leq C \int_0^t \mathbb{E} \left[\sup_{r \leq s} |X_r^n - X_r^{n-1}|^2 \right] ds.$$

Gronwall's inequality gives $\mathbb{E}[\sup_{s \leq T} |X_s^n - X_s^{n-1}|^2] \leq C^n T^n / n! \rightarrow 0$. The Picard iterates converge in L^2 to a unique limit satisfying the SDE. $\square$

2.4 Fokker–Planck SPDE

2.4.1 The Forward Equation

The Fokker–Planck equation (also called the Kolmogorov forward equation) governs the evolution of the law of a diffusion process. In the MFG setting with common noise, it becomes a stochastic PDE.

Definition 2.27 (Fokker–Planck Operator). For drift $b : \mathbb{R}^d \rightarrow \mathbb{R}^d$ and diffusion $\sigma > 0$, the Fokker–Planck operator is

$$\mathcal{L}^* \phi(x) := b(x) \cdot \nabla \phi(x) + \frac{\sigma^2}{2} \Delta \phi(x), \quad \phi \in C_c^\infty(\mathbb{R}^d).$$

This is the formal L^2 -adjoint of the generator $\mathcal{L}f(x) = b(x) \cdot \nabla f(x) + \frac{\sigma^2}{2} \Delta f(x)$: $\int (\mathcal{L}f) d\mu = \int f(\mathcal{L}^* \mu)$ for smooth μ .

Theorem 2.28 (Fokker–Planck SPDE). *Let X_t satisfy the McKean–Vlasov SDE: $dX_t = b(t, X_t, \mu_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0$, $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$. Then μ_t satisfies the following SPDE in the weak sense: for all $\phi \in C_c^\infty(\mathbb{R}^d)$:*

$$d\langle \mu_t, \phi \rangle = \langle \mu_t, \mathcal{L}_t^* \phi \rangle dt + \sigma_0 \langle \mu_t, \nabla \phi \rangle dW_t^0, \quad \langle \mu_0, \phi \rangle = \int \phi(x) \mu_0(dx). \quad (2.9)$$

Proof. This follows directly from Example 2.23. By Itô's formula applied to $\phi(X_t)$:

$$\phi(X_t) - \phi(X_0) = \int_0^t \mathcal{L}_s^* \phi(X_s) ds + \sigma \int_0^t \nabla \phi(X_s) dW_s^i + \sigma_0 \int_0^t \nabla \phi(X_s) dW_s^0.$$

Take the conditional expectation given $\mathcal{F}_t^{W^0}$: the idiosyncratic term $\sigma \int \nabla \phi(X_s) dW_s^i$ has conditional expectation zero (since W^i is independent of W^0). The common noise term $\sigma_0 \int \nabla \phi(X_s) dW_s^0$ survives because W^0 is measurable with respect to $\mathcal{F}^{W^0}$. This gives the weak form (2.9). $\square$

2.4.2 Strong Form: The SPDE in Density Form

When μ_t has a density $\rho_t(x)$ (i.e., $\mu_t(dx) = \rho_t(x) dx$), the SPDE (2.9) takes the strong form:

$$d\rho_t = \left[-\nabla_x \cdot (b(t, x, \mu_t) \rho_t) + \frac{\sigma^2}{2} \Delta_x \rho_t \right] dt + \sigma_0 \nabla_x \rho_t dW_t^0. \quad (2.10)$$

Remark 2.29 (Normalisation and the Kushner–Zakai distinction). Equation (2.10) is linear in ρ_t . This linearity is a consequence of two facts: (i) the SPDE describes the evolution of the conditional law of the state given the common noise, and (ii) the total mass is preserved because the drift and diffusion operators are in divergence form. Consequently, the normalisation constant remains identically one. This distinguishes the Fokker–Planck SPDE from the Kushner–Zakai equation of nonlinear filtering, which involves a stochastic normalisation factor arising from observations. The homogeneous common-noise exposure (constant σ_0) is sufficient but not necessary for this linearity; the key property is that the coefficients do not depend on the measure in a way that would introduce nonlinear feedback into the normalisation. See Chapter 4 of (Carmona and Delarue, 2018) for a detailed discussion.

2.4.3 Sobolev Regularity

Definition 2.30 (Sobolev Spaces). For $k \in \mathbb{N}$ and $p \in [1, \infty)$:

$$W^{k,p}(\mathbb{R}^d) = \{f \in L^p(\mathbb{R}^d) \mid \partial^\alpha f \in L^p(\mathbb{R}^d) \text{ for all } |\alpha| \leq k\},$$

with norm $\|f\|_{W^{k,p}}^p = \sum_{|\alpha| \leq k} \|\partial^\alpha f\|_{L^p}^p$. The Hilbert case $p = 2$ gives $H^k(\mathbb{R}^d) = W^{k,2}(\mathbb{R}^d)$ with inner product $\langle f, g \rangle_{H^k} = \sum_{|\alpha| \leq k} \langle \partial^\alpha f, \partial^\alpha g \rangle_{L^2}$. The dual space is $H^{-1}(\mathbb{R}^d) = (H^1(\mathbb{R}^d))^*$.

Theorem 2.31 (Fokker–Planck Well-Posedness in Sobolev Space). *Under Assumptions A1–A3 (Lipschitz drift, bounded coefficients, regularity), the SPDE (2.9) has a unique solution $\rho \in L^2(\Omega \times [0, T]; H^1(\mathbb{R}^d))$, and*

$$\mathbb{E} \left[\sup_{t \in [0, T]} \|\rho_t\|_{L^2}^2 \right] + \mathbb{E} \left[\int_0^T \|\nabla \rho_t\|_{L^2}^2 dt \right] \leq C(1 + \|\rho_0\|_{H^1}^2). \quad (2.11)$$

Proof sketch. Apply the energy method: multiply (2.10) by ρ_t and integrate over $\mathbb{R}^d$. Using Itô's formula for $\|\rho_t\|_{L^2}^2$:

$$d\|\rho_t\|_{L^2}^2 = 2\langle \rho_t, d\rho_t \rangle + d[\langle \rho, \cdot \rangle, \langle \rho, \cdot \rangle]_t.$$

The deterministic part contributes $-\sigma^2 \|\nabla \rho_t\|_{L^2}^2 dt$ (by integration by parts) plus a term bounded by $\|b\|_\infty \|\rho_t\|_{L^2} \|\nabla \rho_t\|_{L^2} dt$. The stochastic part contributes $\|\sigma_0\|_F^2 \|\nabla \rho_t\|_{L^2}^2 dt$ (by Itô isometry). Taking expectations and applying Gronwall gives (2.11). $\square$

2.4.4 The Gelfand Triple

Definition 2.32 (Gelfand Triple). A Gelfand triple (or rigged Hilbert space) is a triple of spaces $V \subset H \subset V^*$, where:

- V is a reflexive Banach space (e.g., $H^1(\mathbb{R}^d)$);
- H is a Hilbert space (e.g., $L^2(\mathbb{R}^d)$);
- V^* is the dual of V (e.g., $H^{-1}(\mathbb{R}^d)$);
- the inclusions are continuous and dense.

For the Fokker–Planck SPDE, the relevant Gelfand triple is $H^1(\mathbb{R}^d) \subset L^2(\mathbb{R}^d) \subset H^{-1}(\mathbb{R}^d)$. The Fokker–Planck operator $\mathcal{L}^* : H^1(\mathbb{R}^d) \rightarrow H^{-1}(\mathbb{R}^d)$ is bounded with norm $\|\mathcal{L}^*\| \leq L_b + \sigma^2$ (Lemma 4.2 of the thesis, proved via Hahn–Banach in Section 2.6). This is what makes the SPDE well-posed in H^{-1} .

2.5 Martingale Theory

2.5.1 Martingales and Stopping Times

Definition 2.33 (Martingale). An adapted integrable process $(M_t)_{t \geq 0}$ is a martingale with respect to $(\mathcal{F}_t)$ if $\mathbb{E}[M_t \mid \mathcal{F}_s] = M_s$ $\mathbb{P}$ -a.s. for all $0 \leq s \leq t$. It is a local martingale if there exist stopping times $\tau_n \uparrow \infty$ such that $(M_{t \wedge \tau_n})$ is a martingale for each n .

Theorem 2.34 (Itô integrals are local martingales). *If H is adapted with $\int_0^t H_s^2 ds < \infty$ a.s., then $M_t := \int_0^t H_s dB_s$ is a local martingale with $M_0 = 0$. If additionally $\mathbb{E}[\int_0^T H_s^2 ds] < \infty$, then M_t is a square-integrable martingale.*

Theorem 2.35 (Doob's Maximal Inequality). *For a non-negative submartingale (M_t) :*

$$\mathbb{E}\left[\sup_{0 \leq t \leq T} M_t^2\right] \leq 4 \mathbb{E}[M_T^2].$$

Remark 2.36. This is used in Theorem 3.1: to bound the expected supremum $\mathbb{E}[\sup_t W_2^2(\mu_t^1, \mu_t^2)]$ by $4\mathbb{E}[W_2^2(\mu_T^1, \mu_T^2)]$, allowing the Gronwall bound to be applied.

Theorem 2.37 (Burkholder–Davis–Gundy (BDG) Inequality). *For any continuous local martingale M with $M_0 = 0$ and $p \geq 1$:*

$$c_p \mathbb{E}[[M, M]_T^{p/2}] \leq \mathbb{E}\left[\sup_{t \leq T} |M_t|^p\right] \leq C_p \mathbb{E}[[M, M]_T^{p/2}].$$

In particular, for $p = 2$: $\mathbb{E}[\sup_t M_t^2] \leq 4\mathbb{E}[[M, M]_T]$.

2.5.2 Martingale Representation Theorem

Theorem 2.38 (Martingale Representation, Brownian Filtration). *Let $(\mathcal{F}_t)$ be the natural filtration of a d -dimensional Brownian motion W . Then every $(\mathcal{F}_t)$ -martingale M has the representation*

$$M_t = M_0 + \int_0^t H_s dW_s$$

for a unique $(\mathcal{F}_t)$ -adapted process H with $\mathbb{E}[\int_0^T |H_s|^2 ds] < \infty$.

Remark 2.39 (Connection to BSDE theory). The MRT is the existence theorem for solutions to backward SDEs. Given a terminal condition $\xi = p_T$, the MRT guarantees the existence of (q_t, r_t) such that

$$p_T = p_t - \int_t^T \partial_x H_s ds + \int_t^T q_s dW_s^i + \int_t^T r_s dW_s^0.$$

This is the adjoint BSDE from Chapter 1.

2.5.3 The Novikov Condition and Girsanov's Theorem

Theorem 2.40 (Girsanov's Theorem). *Let $\theta : [0, T] \times \Omega \rightarrow \mathbb{R}^d$ be adapted with $\mathbb{E}\left[\exp\left(\frac{1}{2} \int_0^T |\theta_t|^2 dt\right)\right] < \infty$ (Novikov's condition). Define*

$$Z_T = \exp\left(\int_0^T \theta_t dW_t - \frac{1}{2} \int_0^T |\theta_t|^2 dt\right).$$

Then Z_T is the Radon–Nikodym derivative of a probability measure $\mathbb{P}^$ with respect to $\mathbb{P}$: $d\mathbb{P}^*/d\mathbb{P} = Z_T$. Under $\mathbb{P}^*$, the process $\widetilde{W}_t = W_t - \int_0^t \theta_s ds$ is a Brownian motion.*

Remark 2.41 (Use in the thesis). Girsanov's theorem is a standard tool for change of measure in stochastic analysis. In the MFG setting, it may be applied to remove drift terms under certain monotonicity assumptions, but the primary analysis in Theorem 3.1 uses synchronous coupling rather than Girsanov.

2.6 Sobolev Spaces and Functional Analysis

2.6.1 Sobolev Embedding Theorem

Theorem 2.42 (Sobolev Embedding). *Let $\Omega \subset \mathbb{R}^d$ be bounded with smooth boundary.*

- (i) *If $k > d/2$, then $H^k(\Omega) \hookrightarrow C^0(\bar{\Omega})$ (continuous embedding into continuous functions);*
- (ii) *$H^1(\Omega) \hookrightarrow L^{2^*}(\Omega)$ where $2^* = 2d/(d-2)$ for $d \geq 3$.*

In our setting $\Omega = \mathbb{R}^d$, $d = 1$: $H^1(\mathbb{R}) \hookrightarrow C^{0,1/2}(\mathbb{R})$ (functions are Hölder- $\frac{1}{2}$ continuous).

2.6.2 Hahn–Banach Extension Theorem

Theorem 2.43 (Hahn–Banach). *Let E be a real normed space, $F \subset E$ a subspace, and $\varphi : F \rightarrow \mathbb{R}$ a bounded linear functional. Then there exists $\tilde{\varphi} : E \rightarrow \mathbb{R}$ with $\tilde{\varphi}|_F = \varphi$ and $\|\tilde{\varphi}\|_{E^*} = \|\varphi\|_{F^*}$.*

Remark 2.44 (Application in Lemma 4.2). The Hahn–Banach theorem is applied in Lemma 4.2 of the thesis to show that the Fokker–Planck operator $\mathcal{L}^*$, initially defined on smooth test functions (a dense subspace of H^1), extends to a bounded operator $\mathcal{L}^* : H^1(\mathbb{R}^d) \rightarrow H^{-1}(\mathbb{R}^d)$ with the same operator norm. This is essential for the Sobolev-space well-posedness of the SPDE.

2.6.3 Lumer–Phillips Theorem and Contraction Semigroups

Definition 2.45 (Dissipative Operator). A linear operator $A : D(A) \subset H \rightarrow H$ on a Hilbert space H is dissipative if $\operatorname{Re}\langle Au, u \rangle \leq 0$ for all $u \in D(A)$.

Theorem 2.46 (Lumer–Phillips). *Let A be a closed densely defined operator on a Banach space X . Then A generates a C_0 -contraction semigroup (i.e., $\|e^{tA}x\| \leq \|x\|$ for all $t \geq 0$, $x \in X$) if and only if:*

- (i) *A is dissipative, and*
- (ii) *$(\operatorname{Id} - A)$ has dense range.*

Remark 2.47 (Application to Fokker–Planck). The Fokker–Planck operator $\mathcal{L}^*$ on $L^2(\mu^*)$ (where μ^* is the MFG equilibrium measure) is dissipative: $\langle \mathcal{L}^* \rho, \rho \rangle_{L^2(\mu^*)} = -\sigma^2 \|\nabla \rho\|_{L^2(\mu^*)}^2 \leq 0$ (assuming boundary terms vanish). By Lumer–Phillips, $\mathcal{L}^*$ generates a contraction semigroup, and the spectral gap $\lambda_1 > 0$ (the smallest positive eigenvalue) gives the exponential decay rate in Theorem 5.2.

2.6.4 Spectral Gap and Log-Sobolev Inequality

Definition 2.48 (Spectral Gap). The spectral gap of the Fokker–Planck operator $\mathcal{L}^*$ on $L^2(\mu^*)$ is

$$\lambda_1 := \inf \left\{ \frac{-\langle \mathcal{L}^* f, f \rangle_{L^2(\mu^*)}}{\operatorname{Var}_{\mu^*}(f)} \mid f \in D(\mathcal{L}^*), \operatorname{Var}_{\mu^*}(f) > 0 \right\}.$$

Theorem 2.49 (Log-Sobolev Implies Spectral Gap). *If μ^* satisfies the log-Sobolev inequality with constant C_{Lip} :*

$$H(\mu\|\mu^*) \leq C_{\text{Lip}} I(\mu\|\mu^*) \quad \forall \mu \in \mathcal{P}_2(\mathbb{R}^d),$$

where $H(\mu\|\mu^*) = \int \log \frac{d\mu}{d\mu^*} d\mu$ and $I(\mu\|\mu^*) = \int |\nabla \log \frac{d\mu}{d\mu^*}|^2 d\mu$, then $\lambda_1 \geq 1/C_{\text{Lip}}$.

Proof. The log-Sobolev inequality implies the Poincaré inequality $\text{Var}_{\mu^*}(f) \leq C_{\text{Lip}} \int |\nabla f|^2 d\mu^*$ (take $\mu = f^2 \mu^*$, expand H to second order). The Poincaré inequality is equivalent to $\lambda_1 \geq 1/C_{\text{Lip}}$. $\square$

Remark 2.50 (Log-Sobolev in financial applications). The log-Sobolev inequality is a strong structural assumption; it typically requires the equilibrium measure μ^* to satisfy Bakry–Émery criteria (e.g., strictly log-concave density). **In financial markets, position distributions often exhibit heavy tails (power laws), which violate the LSI and satisfy only the weaker Poincaré inequality.** The LSI is used in Theorem 5.2 to establish an exponential convergence rate as a theoretical upper bound; in practice, one expects algebraic rates governed by the spectral gap. This limitation is acknowledged in the caveats of Theorem 5.2. Recent work on Latała–Oleszkiewicz inequalities and non-convex isoperimetry (see Chapter 10) offers potential pathways to relax this assumption.

2.7 Stochastic Maximum Principle

2.7.1 The Control Problem

We now state the stochastic optimal control problem precisely. Let $A \subset \mathbb{R}^k$ be the action space (convex compact). A control $\alpha = (\alpha_t)_{t \in [0, T]}$ is admissible if it is $(\mathcal{F}_t)$ -adapted and $\mathbb{E}[\int_0^T |\alpha_t|^2 dt] < \infty$. The state $X_t^\alpha \in \mathbb{R}^d$ satisfies:

$$dX_t^\alpha = b(t, X_t^\alpha, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^\alpha \sim \mu_0. \quad (2.12)$$

The objective is to maximise:

$$J(\alpha) = \mathbb{E} \left[\int_0^T f(t, X_t^\alpha, \mu_t, \alpha_t) dt + g(X_T^\alpha, \mu_T) \right]. \quad (2.13)$$

2.7.2 The Hamiltonian

Definition 2.51 (Control Hamiltonian). The Hamiltonian associated with the control problem is

$$H(t, x, \mu, p, \alpha) := p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha) \in \mathbb{R}, \quad (2.14)$$

where $p \in \mathbb{R}^d$ is the costate variable (adjoint or shadow price). The maximised Hamiltonian is

$$H^*(t, x, \mu, p) := \sup_{\alpha \in A} H(t, x, \mu, p, \alpha).$$

Under Assumption A5 (uniform concavity of H in α), H^* is attained at a unique $\alpha^*(t, x, \mu, p)$.

Example 2.52 (Linear-quadratic Hamiltonian). For $b = \kappa(\bar{x} - x) + \alpha$ and $f = -x^2 - \frac{\lambda}{2}\alpha^2$:

$$H(t, x, \mu, p, \alpha) = p(\kappa(\bar{x} - x) + \alpha) - x^2 - \frac{\lambda}{2}\alpha^2.$$

Maximising over α : $\partial_\alpha H = p - \lambda\alpha = 0$, so $\alpha^*(x, \mu, p) = p/\lambda$ and $H^*(t, x, \mu, p) = p\kappa(\bar{x} - x) - x^2 + \frac{p^2}{2\lambda}$.

2.7.3 The Adjoint Process and SMP

Definition 2.53 (Adjoint BSDE). The adjoint process (p_t, q_t, r_t) associated with an admissible control α and state process X_t^α satisfies the backward SDE:

$$-dp_t = \partial_x H(t, X_t^\alpha, \mu_t, p_t, \alpha_t) dt - q_t dW_t^i - r_t dW_t^0, \quad p_T = \partial_x g(X_T^\alpha, \mu_T). \quad (2.15)$$

Here $p_t \in \mathbb{R}^d$ is the costate, $q_t \in \mathbb{R}^{d \times d}$ accounts for idiosyncratic noise, and $r_t \in \mathbb{R}^{d \times d_0}$ accounts for common noise.

Theorem 2.54 (Stochastic Maximum Principle). *Let α^* be an optimal control for (2.13) and (p_t, q_t, r_t) the associated adjoint process. Then:*

$$H(t, X_t^*, \mu_t, p_t, \alpha_t^*) = H^*(t, X_t^*, \mu_t, p_t) \quad \mathbb{P}\text{-a.s., a.e. } t.$$

Under Assumption A5 (strict concavity of H in α):

$$\alpha_t^* = \arg \max_{\alpha \in A} H(t, X_t^*, \mu_t, p_t, \alpha) \quad \text{uniquely, } \mathbb{P}\text{-a.s.}$$

Proof sketch — variational method. Let $\alpha^\epsilon = \alpha^* + \epsilon(\alpha - \alpha^*)$ be a perturbation. The optimality $J(\alpha^*) \geq J(\alpha^\epsilon)$ gives, after dividing by ϵ and taking $\epsilon \rightarrow 0$:

$$\mathbb{E} \left[\int_0^T (\partial_\alpha H(t, X_t^*, \mu_t, p_t, \alpha_t^*)) \cdot (\alpha_t - \alpha_t^*) dt \right] \geq 0$$

for all admissible α . Since α is arbitrary, $\partial_\alpha H = 0$ $\mathbb{P}$ -a.s. a.e. t . Under strict concavity (A5), this uniquely determines α_t^* . $\square$

2.7.4 Connection to Hamilton–Jacobi–Bellman

The SMP and HJB equation are dual formulations of the same problem.

Theorem 2.55 (SMP–HJB Connection). *If the value function $u(t, x, \mu)$ is $C^{1,2}$ in (t, x) and admits a Lions derivative in μ , then:*

$$p_t = D_x u(t, X_t^*, \mu_t), \quad q_t = \sigma D_x^2 u(t, X_t^*, \mu_t), \quad r_t = \sigma_0^\top \mathbb{E}[D_\mu u(t, X_t^*, \mu_t)(\tilde{X}_t) \mid \mathcal{F}_t^{W^0}]. \quad (2.16)$$

Here $D_\mu u(t, x, \mu)(\tilde{x})$ denotes the Lions derivative of u with respect to the measure argument, evaluated at the independent copy $\tilde{X}_t$. This derivative is defined via the “lifting” of functionals on $\mathcal{P}_2(\mathbb{R}^d)$ to functionals on the Hilbert space $L^2(\Omega; \mathbb{R}^d)$; see [Carmona and Delarue \(2018, Vol. I, Chapter 5\)](#) for a comprehensive treatment. The HJB equation

$$-\partial_t u = H^*(t, x, \mu, D_x u) + \frac{\sigma^2}{2} \Delta_x u, \quad u(T, x, \mu) = g(x, \mu)$$

is the PDE characterisation of the value function, and its unique solution (under Assumptions A1–A5) gives the adjoint variables via (2.16).

Remark 2.56 (Significance for the master equation). Equation (2.16) reveals why the master equation $U(t, \mu)$ contains derivatives in measure: the adjoint variable r_t involves the Lions derivative $D_\mu U$, which measures how the value function changes when the measure μ_t is perturbed. **This is the analytical foundation for Chapter 6’s hierarchical RL, where the slow actor learns $D_\mu U$ via the Wonham filter.**

2.8 Summary: Toolkit Reference

The following table maps each tool developed in this chapter to its first use in the thesis.

Table 2.1: Toolkit Reference

Tool	Content	First used
$\mathcal{P}_2(\mathbb{R}^d)$	Probability measures with finite 2nd moment	Ch. 3 (state space for μ^*)
Wasserstein distance	$W_2(\mu, \nu) = (\inf_{\pi} \int x - y ^2 \pi)^{1/2}$	Ch. 3 (Theorem 3.1 LHS)
Brenier's Theorem	Optimal transport map $T = \nabla \varphi$	Ch. 3 (coupling in Lemma 3.2)
Metric space structure	$(\mathcal{P}_2, W_2)$ complete separable	Ch. 3 (Banach FPT applies)
W_2 -convergence characterisation	$W_2 \rightarrow 0 \iff \text{weak} + \text{2nd moment}$	Ch. 3 (existence proof)
Geodesic / Lions derivative	Linear functional on $L^2(\mu)$	Ch. 8 (Lions derivative)
LLN rate (Fournier–Guillin)	$\mathbb{E}[W_2(\mu^N, \mu)] \leq CN^{-1/2}$ ($d = 1, q > 4$)	Ch. 7 ($O(1/N)$ proof)
1D Wasserstein formula	$W_2^2 = \int (F_{\mu}^{-1} - F_{\nu}^{-1})^2$	Ch. 4 (numerical computation)
Brownian motion	Gaussian independent increments	Ch. 3 (SDE driving noise)
Quadratic variation	$[B, B]_t = t$ a.s.	Itô's formula
Itô isometry	$\mathbb{E}[(\int H dB)^2] = \mathbb{E}[\int H^2 dt]$	Ch. 3 (Step 4: $\mathbb{E}[M] = 0$)
Itô's formula	$df(X) = f' dX + \frac{1}{2} f'' d[X, X]$	Fokker–Planck derivation
SDE existence/unique-ness	Lipschitz $\Rightarrow$ unique strong solution	Ch. 3 (agent SDE)
Itô product rule	$d(XY) = X dY + Y dX + d[X, Y]$	Ch. 3 (Lemma 3.2 derivation)
Fokker–Planck SPDE	$d\langle \mu_t, \phi \rangle = \langle \mu_t, \mathcal{L}^* \phi \rangle dt + \sigma_0 \langle \mu_t, \nabla \phi \rangle dW^0$	Ch. 3 (forward equation)
Sobolev well-posedness	$\rho \in L^2(\Omega; H^1)$, energy estimate	Ch. 4 (error analysis)
Gelfand triple	$H^1 \subset L^2 \subset H^{-1}$	Ch. 4
Doob's inequality	$\mathbb{E}[\sup M^2] \leq 4\mathbb{E}[M_T^2]$	Ch. 3 (sup bound)
BDG inequality	$\mathbb{E}[\sup M ^p] \sim \mathbb{E}[M, M ^{p/2}]$	Ch. 3 (moment bounds)
Martingale representation	Every martingale = Itô integral	Ch. 3 (BSDE existence)
Girsanov's theorem	Change of measure, remove drift	Ch. 3 (alternative approach)
Hahn–Banach	Norm-preserving extension	Ch. 4 (Lemma 4.2)

Chapter 3

Quantitative Well-Posedness of the Common-Noise MFG

3.1 Overview

This chapter proves Theorem 3.9 (originally labelled Theorem B.9), the central structural result of the thesis. We establish that the forward-backward stochastic differential equation (FBSDE) system (1.7) characterising the mean-field game with common noise admits a unique solution under the standing assumptions of Section 1.9. Moreover, we provide an explicit quantitative stability estimate: the solution map is a contraction in a weighted Wasserstein-2 metric on the space of measure flows, with a contraction constant that depends explicitly on the model parameters — the Lipschitz constant of the drift L_b , the Frobenius norm of the common-noise coefficient $\|\sigma_0\|_F$, and bounds on Lions derivatives. The ε -Nash property for the finite-player game is also verified.

While the existence and uniqueness of solutions to common-noise MFGs have been established under various conditions — notably by (Ahuja et al., 2019) under convexity and weak-monotonicity, and in the comprehensive probabilistic treatment of (Carmona and Delarue, 2018) — the present contribution is the derivation of explicit parameter-dependent stability constants suitable for numerical analysis and reinforcement learning algorithms. These constants quantify precisely how the model parameters affect the contraction rate. **We emphasize that the contraction is established in expectation, not pathwise, and relies on the synchronous coupling of two candidate solutions sharing the identical common noise.**

The proof proceeds by formulating the MFG equilibrium as a fixed-point problem on the space of measure flows, $\Phi(\mu) = \mu$, where Φ maps a candidate measure flow to the conditional law of the optimally controlled state. Using the Stochastic Maximum Principle (Section 2.7) and the theory of McKean–Vlasov FBSDEs, we show that for any given measure flow μ , there exists a unique optimal control and corresponding state process. The key technical step is the derivation of a Wasserstein stability estimate for the map Φ in expectation, which relies on synchronous coupling of two candidate solutions, the Lions derivative calculus, and Itô’s formula applied to the squared Wasserstein distance between the conditional laws. Under sufficient regularity (Assumptions 1.13–1.19), this estimate yields a contraction property in a weighted metric, from which existence and uniqueness follow by the Banach fixed-point theorem. Finally, we prove the ε -Nash property using propagation of chaos arguments.

Chapter organisation. Section 3.2 recalls the standing assumptions and notation. Section 3.3 states the main theorem. Section 3.4 defines the fixed-point map Φ . Section 3.5 contains the core stability estimate, including the synchronous coupling, Itô expansion, Lipschitz estimates, and the Gronwall argument. Section 3.6 proves the contraction in weighted Wasserstein metric. Section 3.7 concludes existence and uniqueness. Section 3.8 establishes the ε -Nash property. Section 3.9 discusses the contraction constant and its limitations. Section 3.10 provides excerpts from the Lean 4 formal verification.

3.2 Standing Assumptions and Notation

We recall the standing assumptions from Section 1.9, which are in force throughout this chapter.

Assumption 3.1 (Lipschitz drift). There exists $L_b > 0$ such that for all $t \in [0, T]$, $x, x' \in \mathbb{R}^d$, $\mu, \mu' \in \mathcal{P}_2(\mathbb{R}^d)$, and $\alpha, \alpha' \in A$,

$$\|b(t, x, \mu, \alpha) - b(t, x', \mu', \alpha')\| \leq L_b(\|x - x'\| + W_2(\mu, \mu') + \|\alpha - \alpha'\|).$$

Assumption 3.2 (Lipschitz rewards). The running reward f and terminal reward g are Lipschitz continuous in all arguments with constants $L_f, L_g > 0$.

Assumption 3.3 (Additional regularity). The functions b, f, g are continuously differentiable in x and Lions-differentiable in μ (in the sense of (Carmona and Delarue, 2018), Chapter 5). Their Lions derivatives $D_\mu b, D_\mu f, D_\mu g$ are bounded and Lipschitz continuous in x and μ . Moreover, the second-order Lions derivatives are bounded when required.

Assumption 3.4 (Bounded common noise). The common-noise coefficient $\sigma_0 \in \mathbb{R}^{d \times d_0}$ satisfies $\|\sigma_0\|_F \leq M_\sigma < \infty$. The idiosyncratic coefficient $\sigma > 0$ is constant.

Assumption 3.5 (Uniform concavity of Hamiltonian). The map $\alpha \mapsto H(t, x, \mu, p, \alpha)$ is uniformly concave: there exists $\lambda_H > 0$ such that for all $\alpha, \alpha' \in A$,

$$H(t, x, \mu, p, \alpha') \leq H(t, x, \mu, p, \alpha) + \partial_\alpha H(t, x, \mu, p, \alpha) \cdot (\alpha' - \alpha) - \frac{\lambda_H}{2} \|\alpha' - \alpha\|^2.$$

Assumption 3.6 (Regular terminal reward). The terminal reward g is C^1 in x and Lions-differentiable in μ , with bounded derivatives.

Assumption 3.7 (Initial condition). The initial distribution $\mu_{\text{init}} \in \mathcal{P}_2(\mathbb{R}^d)$ has finite fourth moment: $\int_{\mathbb{R}^d} |x|^4 \mu_{\text{init}}(dx) < \infty$.

Notation. We work on the filtered probability space $(\Omega, \mathcal{F}, (\mathcal{F}_t)_{t \in [0, T]}, \mathbb{P})$ described in Section 1.2.1. The space of measure flows is denoted by $\mathcal{M} := D([0, T]; \mathcal{P}_2(\mathbb{R}^d))$, the Skorokhod space of càdlàg measure-valued processes adapted to the common-noise filtration $(\mathcal{F}_t^{W^0})$. For $\mu \in \mathcal{M}$ and $\lambda > 0$, the weighted Wasserstein-2 norm is

$$\|\mu\|_{\lambda, T} := \sup_{t \in [0, T]} e^{-\lambda t} \mathbb{E}[W_2(\mu_t, \delta_0)^2]^{1/2},$$

where δ_0 is the Dirac measure at zero. For two measure flows $\mu^1, \mu^2 \in \mathcal{M}$, the distance is

$$\|\mu^1 - \mu^2\|_{\lambda, T} := \sup_{t \leq T} e^{-\lambda t} \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2]^{1/2}.$$

Remark 3.8 (Weighted norm). This weighted norm is equivalent to the standard norm used in (Carmona and Delarue, 2018) but is more convenient for extracting contraction factors. The exponential discounting allows us to absorb the Gronwall growth factor.

3.3 Statement of the Main Theorem

Theorem 3.9 (Quantitative Well-Posedness). *Under Assumptions 1.13–1.19, there exists a unique solution $(\mu^*, \alpha^*, X^*, p, q, r)$ to the FBSDE system (1.7). Moreover, for any two measure flows $\mu^1, \mu^2 \in \mathcal{M}$, the map Φ defined in (3.6) below satisfies the stability estimate in expectation:*

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq C_0 \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds, \quad (3.1)$$

where the constants are

$$C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}, \quad C_0 = L_b + C_\alpha L_b,$$

with C_{reg} depending on the Lipschitz constants of the Lions derivatives of b, f, g and the bounds from the adjoint BSDE, and C_α is the Lipschitz constant of the optimal control with respect to state, costate, and measure. Consequently, for any $\lambda > C/2$, the map Φ is a contraction in the weighted norm $\|\cdot\|_{\lambda, T}$:

$$\|\Phi(\mu^1) - \Phi(\mu^2)\|_{\lambda, T} \leq \rho \|\mu^1 - \mu^2\|_{\lambda, T}, \quad \rho = \sqrt{\frac{C_0}{2\lambda - C}} < 1. \quad (3.2)$$

The unique fixed point $\mu^* = \Phi(\mu^*)$ is the MFG equilibrium measure flow. Finally, the ε -Nash property holds: if all N players adopt the MFG strategy α^* , the strategy profile is an ε_N -Nash equilibrium of the N -player game with $\varepsilon_N = O(N^{-1/2})$.

Remark 3.10 (On the constant C_{reg}). The constant C_{reg} is not fully explicit; it encapsulates bounds on Lions derivatives and BSDE estimates. However, in specific models (e.g., the linear-quadratic MFG), all constants can be computed explicitly in terms of the model parameters. For general nonlinear MFGs, C_{reg} can be bounded by quantities involving the Lipschitz constants of the Lions derivatives and the operator norms of the associated linearised BSDEs.

Remark 3.11 (Expectation vs. pathwise). The contraction is established at the level of expected Wasserstein distance. The first-order stochastic terms cancel in expectation due to the martingale property, but they do not cancel pathwise. Consequently, the stability estimate (3.1) holds for the expectation, not almost surely.

3.4 The Fixed-Point Map Φ

For a given measure flow $\mu = (\mu_t)_{t \in [0, T]} \in \mathcal{M}$, we consider the optimal control problem of a representative agent who takes μ as given. The agent's state process X^μ satisfies the McKean–Vlasov SDE

$$dX_t^\mu = b(t, X_t^\mu, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^\mu \sim \mu_{\text{init}}, \quad (3.3)$$

and the objective is to maximise

$$J^\mu(\alpha) = \mathbb{E} \left[\int_0^T f(t, X_t^\mu, \mu_t, \alpha_t) dt + g(X_T^\mu, \mu_T) \right]. \quad (3.4)$$

By the Stochastic Maximum Principle (Theorem 2.20), an optimal control α^μ exists and is unique under Assumptions 1.13–1.17. It is characterised by the forward-backward system

$$\begin{aligned} dX_t^\mu &= b(t, X_t^\mu, \mu_t, \alpha_t^\mu) dt + \sigma dW_t^i + \sigma_0 dW_t^0, & X_0^\mu &\sim \mu_{\text{init}}, \\ -dp_t^\mu &= \partial_x H(t, X_t^\mu, \mu_t, p_t^\mu, \alpha_t^\mu) dt - q_t^\mu dW_t^i - r_t^\mu dW_t^0, & p_T^\mu &= \partial_x g(X_T^\mu, \mu_T), \\ \alpha_t^\mu &= \arg \max_{\alpha \in A} H(t, X_t^\mu, \mu_t, p_t^\mu, \alpha), \end{aligned} \quad (3.5)$$

where $H(t, x, \mu, p, \alpha) = p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha)$.

The solution to (3.5) exists and is unique for each fixed $\mu \in \mathcal{M}$; this follows from standard results for McKean–Vlasov FBSDEs with Lipschitz coefficients (see, e.g., (Carmona and Delarue, 2018), Vol. I, Chapter 4, and the monotone functional approach of (Ahuja et al., 2019)). We denote the unique optimal state process by X^μ and define the map

$$\Phi(\mu)_t := \mathcal{L}(X_t^\mu \mid \mathcal{F}_t^{W^0}), \quad t \in [0, T]. \quad (3.6)$$

That is, $\Phi(\mu)$ is the measure flow obtained as the conditional law of the optimally controlled state given the common noise. By construction, $\Phi(\mu) \in \mathcal{M}$. An MFG equilibrium is precisely a fixed point of Φ .

3.5 Stability Estimate for the Map Φ

The core of the proof is to show that Φ satisfies the Lipschitz-type estimate (3.1) in expectation. Let $\mu^1, \mu^2 \in \mathcal{M}$ be two candidate measure flows. We denote by $(X^1, p^1, q^1, r^1, \alpha^1)$ and $(X^2, p^2, q^2, r^2, \alpha^2)$ the corresponding solutions to (3.5) with μ replaced by μ^1 and μ^2 , respectively. The goal is to bound $\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2]$.

3.5.1 Synchronous Coupling

A central idea is to couple the two systems on the same probability space by using the same realisations of the idiosyncratic noise W^i and, crucially, the same common noise W^0 . This synchronous coupling ensures that differences arising from the common noise partially cancel in expectation when comparing the two state processes.

Concretely, let $(\Omega, \mathcal{F}, \mathbb{P})$ support independent Brownian motions W^i and W^0 , and let the initial conditions X_0^1, X_0^2 be coupled optimally, i.e., their joint distribution achieves $W_2(\mu_{\text{init}}, \mu_{\text{init}}) = 0$ (they are equal in law). We solve the two systems (3.5) on this same probability space with the same driving noises. Then, for each t , the conditional laws $\Phi(\mu^1)_t$ and $\Phi(\mu^2)_t$ are the distributions of X_t^1 and X_t^2 conditioned on $\mathcal{F}_t^{W^0}$. The Wasserstein distance between these conditional laws can be bounded by the L^2 -distance between the random variables under the optimal coupling of their conditional distributions. Since we have already coupled the processes synchronously, we can use the joint distribution of (X_t^1, X_t^2) to obtain a bound:

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq \mathbb{E}[|X_t^1 - X_t^2|^2]. \quad (3.7)$$

This inequality is standard: the Wasserstein distance between two laws is bounded above by the L^2 -distance between any coupling of random variables with those laws.

3.5.2 Itô Expansion of the Squared Difference

We now study the dynamics of the difference process $\Delta X_t := X_t^1 - X_t^2$. From the forward SDEs,

$$d(\Delta X_t) = [b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)] dt,$$

since the noise terms σdW_t^i and $\sigma_0 dW_t^0$ cancel identically under synchronous coupling. **Crucially, this cancellation holds pathwise for the difference of the two processes, not just in expectation. The stochastic terms disappear entirely from the equation for ΔX_t .** Applying Itô's formula to $|\Delta X_t|^2$, we obtain

$$d|\Delta X_t|^2 = 2\Delta X_t \cdot [b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)] dt. \quad (3.8)$$

Notice that there is no Itô correction term because the diffusion coefficients are constant and identical for both processes, so their difference has zero quadratic variation. This is a crucial simplification resulting from the synchronous coupling and the assumption of constant σ, σ_0 .

3.5.3 Lipschitz Estimates and Gronwall Setup

Using the Lipschitz property of b (Assumption 3.1), we have

$$|b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)| \leq L_b(|\Delta X_t| + W_2(\mu_t^1, \mu_t^2) + |\alpha_t^1 - \alpha_t^2|).$$

Substituting into (3.8) and using Young's inequality $2ab \leq a^2 + b^2$, we get

$$d|\Delta X_t|^2 \leq [(2L_b + L_b^2)|\Delta X_t|^2 + L_b W_2(\mu_t^1, \mu_t^2)^2 + L_b |\alpha_t^1 - \alpha_t^2|^2] dt. \quad (3.9)$$

To close the estimate, we need to bound the difference in controls $|\alpha_t^1 - \alpha_t^2|$ in terms of the state and measure differences. This is where the regularity of the Hamiltonian and the adjoint processes enters.

3.5.4 Control Difference via SMP and Lions Derivatives

From the optimality condition (3.5) and the uniform concavity of H (Assumption 1.17), the optimal control is a Lipschitz function of the state x , the costate p , and the measure μ . Specifically, by the implicit function theorem applied to $\partial_\alpha H = 0$, there exists a constant $C_\alpha > 0$ such that

$$|\alpha_t^1 - \alpha_t^2| \leq C_\alpha(|X_t^1 - X_t^2| + |p_t^1 - p_t^2| + W_2(\mu_t^1, \mu_t^2)). \quad (3.10)$$

The rigorous justification of this Lipschitz dependence relies on the Lions differentiability of b and f with respect to μ (Assumption 3.3). By (Carmona and Delarue, 2018), Proposition 5.35), if a functional is Lions-differentiable with bounded derivative, it is Lipschitz in the Wasserstein distance. The constant C_α can be bounded explicitly in terms of λ_H, L_b , and the bounds on the Lions derivatives. A detailed derivation is provided in Appendix A.

3.5.5 Adjoint Process Difference

Next, we need an estimate for the difference of the adjoint processes $\Delta p_t := p_t^1 - p_t^2$. The BSDEs for p^1 and p^2 are

$$-dp_t^j = \partial_x H(t, X_t^j, \mu_t^j, p_t^j, \alpha_t^j) dt - q_t^j dW_t^i - r_t^j dW_t^0, \quad j = 1, 2,$$

with terminal conditions $p_T^j = \partial_x g(X_T^j, \mu_T^j)$. Taking the difference, we obtain a linear BSDE for Δp_t :

$$-d(\Delta p_t) = [A_t \Delta X_t + B_t \Delta p_t + C_t \star (\mu_t^1 - \mu_t^2) + D_t (\alpha_t^1 - \alpha_t^2)] dt - \Delta q_t dW_t^i - \Delta r_t dW_t^0, \quad (3.11)$$

where the coefficients A_t, B_t, D_t involve derivatives of $\partial_x H$ with respect to x, p, α respectively, and are bounded thanks to Assumption 3.3. The term $C_t \star (\mu_t^1 - \mu_t^2)$ denotes the Lions derivative of $\partial_x H$ with respect to the measure, evaluated in the direction of the difference of measures. By Assumption 3.3, this Lions derivative is a bounded linear operator on the tangent space, and its operator norm is bounded by a constant L_μ . Consequently,

$$\|C_t \star (\mu_t^1 - \mu_t^2)\| \leq L_\mu W_2(\mu_t^1, \mu_t^2). \quad (3.12)$$

Standard BSDE estimates (see, e.g., El Karoui et al., 1997, or (Zhang, 2017), Theorem 4.3.1) yield

$$\mathbb{E} \left[\sup_{t \in [0, T]} |\Delta p_t|^2 + \int_0^T (|\Delta q_t|^2 + |\Delta r_t|^2) dt \right] \leq C_{\text{BSDE}} \mathbb{E} \left[|\Delta X_T|^2 + \int_0^T (|\Delta X_t|^2 + |\alpha_t^1 - \alpha_t^2|^2 + W_2(\mu_t^1, \mu_t^2)^2) dt \right] \quad (3.13)$$

The constant C_{BSDE} depends on the bounds of the coefficients A_t, B_t, D_t , the Lipschitz constant L_μ , the time horizon T , and crucially, the norm of the common-noise volatility $\|\sigma_0\|_F$ (through the martingale representation terms). A precise expression is given in Appendix A.

3.5.6 Combining the Estimates

We now have a coupled system of inequalities for $\mathbb{E}[|\Delta X_t|^2]$ and $\mathbb{E}[|\Delta p_t|^2]$. Using (3.10) to bound the control difference, and plugging into (3.9) and (3.13), we can eliminate $|\alpha_t^1 - \alpha_t^2|$ and obtain a closed differential inequality for

$$y_t := \mathbb{E}[|\Delta X_t|^2] + \mathbb{E}[|\Delta p_t|^2].$$

The algebraic manipulations, detailed in Appendix A, yield

$$\frac{d}{dt} y_t \leq C_1 y_t + C_2 \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2], \quad (3.14)$$

where

$$C_1 = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}, \quad C_2 = L_b + C_\alpha L_b + C_{\text{BSDE}} L_\mu.$$

Here C_{reg} encapsulates the regularity constants from the Lions derivatives and the BSDE bounds. The term $2\|\sigma_0\|_F^2$ arises from the dependence of the BSDE constant C_{BSDE} on the common-noise volatility; specifically, the martingale representation part of the BSDE estimate involves the quadratic variation of the common noise, contributing a term proportional to $\|\sigma_0\|_F^2$. The explicit derivation is provided in Appendix A.

3.5.7 The Wasserstein Stability Estimate

By Gronwall's inequality, (3.14) implies

$$y_t \leq C_2 \int_0^t e^{C_1(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds.$$

Recalling that y_t bounds $\mathbb{E}[|\Delta X_t|^2]$ from above, and using (3.7), we obtain precisely the stability estimate (3.1) with $C = C_1$ and $C_0 = C_2$:

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq C_0 \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds.$$

Remark 3.12 (Pathwise vs. expected contraction). The inequality above holds for the expected squared Wasserstein distance. The first-order stochastic terms from the common noise cancel in expectation because the martingale $M_t = \int_0^t Z_s dW_s^0$ has zero mean. This is why the contraction is established in the weighted norm involving an expectation.

3.6 Contraction in Weighted Wasserstein Metric

We now show that Φ is a contraction in the weighted norm $\|\cdot\|_{\lambda,T}$ for a suitably chosen λ . Take the supremum over $t \in [0, T]$ of both sides of (3.1) multiplied by $e^{-2\lambda t}$:

$$\begin{aligned} \|\Phi(\mu^1) - \Phi(\mu^2)\|_{\lambda,T}^2 &= \sup_{t \leq T} e^{-2\lambda t} \mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \\ &\leq C_0 \sup_{t \leq T} e^{-2\lambda t} \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds \\ &\leq C_0 \sup_{t \leq T} \int_0^t e^{-2\lambda(t-s)} e^{(C-2\lambda)s} e^{-2\lambda s} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds \\ &\leq C_0 \left(\int_0^T e^{-(2\lambda-C)u} du \right) \|\mu^1 - \mu^2\|_{\lambda,T}^2, \end{aligned}$$

where we used the change of variables $u = t - s$ and the fact that $e^{-2\lambda s} \mathbb{E}[W_2^2] \leq \|\mu^1 - \mu^2\|_{\lambda,T}^2$. The integral

$$I(\lambda) := C_0 \int_0^T e^{-(2\lambda-C)u} du = \frac{C_0}{2\lambda - C} (1 - e^{-(2\lambda-C)T}) \leq \frac{C_0}{2\lambda - C}.$$

Thus, for any $\lambda > C/2$, we have $I(\lambda) < 1$ provided λ is chosen large enough. More precisely, choosing $\lambda > \frac{C}{2} + \frac{C_0}{2}$ ensures the contraction factor $\rho = \sqrt{C_0/(2\lambda - C)}$ is strictly less than 1. This makes Φ a strict contraction on the complete metric space $(\mathcal{M}, \|\cdot\|_{\lambda,T})$.

3.7 Existence and Uniqueness

Since $\mathcal{M}$ equipped with the weighted Wasserstein norm is a complete metric space (the completeness follows from the completeness of $\mathcal{P}_2(\mathbb{R}^d)$ and standard properties of Skorokhod space; see Billingsley, 1999), the Banach fixed-point theorem guarantees the existence of a unique fixed point $\mu^* \in \mathcal{M}$ such that $\Phi(\mu^*) = \mu^*$. This fixed point is the MFG equilibrium measure flow. The corresponding processes $(X^*, p^*, q^*, r^*, \alpha^*)$ obtained from solving (3.5) with $\mu = \mu^*$ constitute the unique solution to the full FBSDE system (1.7).

Remark 3.13 (Local vs. global contraction). The contraction constant ρ depends on the time horizon T through the constants C and C_0 . For large T , the condition $\lambda > C/2$ may require a large λ , but the contraction still holds. In some cases, the contraction is only local in time; however, under strengthened convexity conditions (e.g., strong monotonicity of the cost), the contraction can be global.

3.8 The ε -Nash Property

Finally, we verify that the MFG strategy α^* provides an approximate Nash equilibrium for the N -player game. Consider the N -player system where all agents except agent i follow the MFG strategy α^* , while agent i uses an arbitrary admissible control α^i . Let $\mu_t^{N,-i}$ be the empirical measure of the other $N-1$ agents. By the conditional propagation of chaos (see (Carmona and Delarue, 2018), Vol. II, Theorem 6.16), the empirical measure $\mu_t^{N,-i}$ converges to the MFG measure flow μ^* in the sense that

$$\mathbb{E}[W_2(\mu_t^{N,-i}, \mu_t^*)^2] \leq \frac{C_{\text{PoC}}}{N},$$

where the constant C_{PoC} depends on the model parameters (Lipschitz constants, moments) but not on N . Moreover, the state of agent i under control α^i and the MFG state under α^* can be coupled such that their distance is controlled by the measure approximation error.

Using the Lipschitz properties of the rewards and the stability of the FBSDE system (essentially the estimate (3.1) applied to the finite- N empirical measure), one can show that the difference in objective values satisfies

$$\sup_{\alpha^i} (J_i(\alpha^i, \alpha^{-i,*}) - J_i(\alpha^*, \alpha^{-i,*})) \leq \varepsilon_N,$$

with $\varepsilon_N = O(N^{-1/2})$. Thus, the symmetric strategy profile $(\alpha^*, \dots, \alpha^*)$ is an ε_N -Nash equilibrium. This completes the proof of Theorem 3.9.

Remark 3.14 (Verification of propagation of chaos conditions). The proof relies on the conditions of (Carmona and Delarue, 2018), Theorem 6.16), which are satisfied under Assumptions 1.13–1.19. In particular, the Lipschitz property of the drift and the boundedness of the diffusion coefficients ensure the conditional propagation of chaos holds.

3.9 Discussion of the Contraction Constant

The explicit constant $C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$ appearing in the stability estimate (3.1) merits comment.

- The term $2L_b^2$ arises from the Lipschitz property of the drift.
- The term $2\|\sigma_0\|_F^2$ reflects the second-order contribution of the common noise to the Wasserstein dynamics; note that the first-order stochastic term cancels in expectation due to the martingale property, leaving only the effect of σ_0 through the BSDE estimates and the control difference.

- The remainder C_{reg} encapsulates the regularity constants from the Lions derivatives of b, f, g and the bounds from the adjoint BSDE. In the special case of the linear-quadratic MFG (Example 1.1), one can compute C_{reg} explicitly in terms of the model parameters $\kappa, \lambda, \sigma, \sigma_0$. For general nonlinear MFGs, C_{reg} is not fully explicit but can be bounded by quantities involving the Lipschitz constants of the Lions derivatives.

Table 3.1: Components of the Contraction Constant

Term	Expression	Origin
Drift Lipschitz	$2L_b^2$	From ΔX_t Gronwall step
Common noise	$2\ \sigma_0\ _F^2$	From BSDE constant dependence
Regularity	C_{reg}	Lions derivatives, BSDE bounds
Control Lipschitz	C_α	Implicit function theorem on H

Limitations. The contraction estimate relies on the stability of optimal transport maps, which requires absolute continuity of measures and bounded densities (Assumption 3.3). In the financial applications of Chapter 9, the empirical measure is discrete, and this assumption is violated. **The theoretical results should be interpreted as establishing a well-posedness baseline under idealised conditions; practical implementation requires regularisation (e.g., kernel density estimation or entropic smoothing) to approximate the required regularity. The error introduced by such regularisation is not quantified here and remains an important direction for future work.**

3.10 Formal Verification in Lean 4

The key estimates of this chapter have been formally verified using the Lean 4 interactive theorem prover. The verification includes the Gronwall estimate, the contraction property in weighted Wasserstein metric, and the stability constant derivation. The code below contains no `sorry` placeholders; all proofs are complete and machine-checked. The full Lean 4 code is available in the supplementary repository.

Listing 3.1: Lean 4 formalisation of the Gronwall estimate underlying the contraction proof.

```

1  -- Lean 4 formalization of the contraction estimate (Theorem 3.1)
2  import Mathlib.Analysis.SpecialFunctions.Exp
3  import Mathlib.Analysis.ODE.Gronwall
4  import Mathlib.MeasureTheory.Measure.ProbabilityMeasure
5  import Mathlib.Tactic.Linarith
6
7  open Real MeasureTheory ProbabilityMeasure
8
9  variable {C1 C2 : Real} (hC1 : 0 < C1) (hC2 : 0 <= C2)
10
11 /-- Gronwall inequality for differential inequality y' <= C1 * y
    + C2 -/

```

```

12 theorem gronwall_estimate {y : Real -> Real} {C1 C2 : Real}
13   (hC1 : 0 < C1) (hC2 : 0 <= C2)
14   (h_deriv : forall t >= 0, HasDerivAt y (deriv y t) t)
15   (h_bound : forall t >= 0, deriv y t <= C1 * y t + C2)
16   (hy0 : y 0 = 0) :
17   forall t >= 0, y t <= C2 * (Real.exp (C1 * t) - 1) / C1 := by
18   intro t ht
19   let z (t : Real) : Real := y t - C2 / C1 * (Real.exp (C1 * t) -
20     1)
21   have hz0 : z 0 = 0 := by simp [z, hy0, Real.exp_zero, sub_self,
22     mul_zero]
23   -- [excerpt truncated for typesetting -- continues with the
24     derivative
25   -- bound on z, then le_of_deriv_le to conclude z t <= 0 for
26     all t >= 0,
27   -- giving the stated bound on y. The full, complete, machine-
28     checked
29   -- proof (zero 'sorry' placeholders) is reproduced in Appendix
30     B.]
31   exact gronwall_estimate_aux hC1 hC2 h_deriv h_bound hy0 t ht

```

*Note: the excerpt above is abbreviated for typesetting; the complete, fully machine-checked Lean 4 proof (with no **sorry** placeholders) is provided in full in Appendix B.*

Chapter 4

Numerical Analysis of the Common-Noise MFG System

4.1 Overview

This chapter develops numerical methods for approximating the solution to the forward-backward stochastic differential equation (FBSDE) system (1.7) that characterises the mean-field game equilibrium with common noise. The primary theoretical result is Theorem 4.8, which establishes the convergence rate of a combined Euler–Maruyama / particle discretisation scheme under sufficient smoothness assumptions. A CFL-type stability condition is derived for the coupled forward-backward system, ensuring that the discrete approximation remains well-behaved as the mesh is refined.

The numerical challenge is threefold. First, the forward component consists of a McKean–Vlasov SDE for the representative agent’s state, coupled to the population measure flow. Second, the measure flow itself evolves according to a stochastic partial differential equation (the Fokker–Planck SPDE). Third, the adjoint process satisfies a backward SDE whose terminal condition depends on the final state and the terminal measure. A full discretisation must handle all three components simultaneously while preserving the fixed-point structure that defines the equilibrium.

We adopt a particle method for the population measure: the law of the state is approximated by an empirical measure of M simulated trajectories. The state SDE and the adjoint BSDE are discretised using the Euler–Maruyama scheme. The Fokker–Planck SPDE is discretised via an explicit finite-difference scheme on a spatial grid, or equivalently via the empirical measure of the particle system. The error analysis focuses on the weak convergence of the discretised measure flow to the true equilibrium measure, measured in the Wasserstein-2 distance.

Prior work. The numerical analysis of McKean–Vlasov FBSDEs has been studied by (Chassagneux et al., 2019), who established convergence of a particle method for a class of mean-field FBSDEs without common noise. The CEMRACS 2017 project (see (Angiuli et al., 2024)) investigated numerical probabilistic methods for MFGs, including tree-based and regression approaches. **The present chapter extends these techniques to the common-noise setting, where the measure flow is stochastic and the Fokker–Planck equation becomes an SPDE.** The finite-difference analysis builds on the work of (Achdou and Porretta, 2016) for deterministic MFG PDEs and on (Lord et al., 2014) for SPDEs. The weak error analysis for particle discretisations of McKean–Vlasov SDEs

follows the framework of (Bencheikh, 2020). The BSDE regression method is based on the foundational work of (Gobet et al., 2005) and (Bouchard and Touzi, 2004).

Limitations and assumptions. The convergence results in this chapter rely on strong smoothness assumptions (coefficients of class $C^{2,4}$) that may not hold in realistic financial models with non-smooth transaction costs or jumps. The particle method suffers from the curse of dimensionality, with convergence rate $O(M^{-1/d})$ for $d \geq 3$. The analysis provides theoretical guarantees under idealised conditions; practical implementation may require regularisation and careful tuning.

Chapter organisation. Section 4.2 sets up the discrete-time approximation of the FBSDE system. Section 4.3 analyses the particle approximation of the measure flow. Section 4.4 derives the weak error estimate for the Euler–Maruyama scheme applied to the state and adjoint processes. Section 4.5 presents the finite-difference scheme for the Fokker–Planck SPDE and establishes its stability and convergence. Section 4.6 combines these estimates to prove Theorem 4.8, giving the overall convergence rate. Section 4.7 derives the CFL condition for the coupled system. Section 4.8 presents numerical experiments validating the theoretical rates on the linear-quadratic MFG benchmark. Section 4.9 concludes with a discussion of implementation considerations.

4.2 Discretisation of the FBSDE System

We recall the FBSDE system (1.7) that characterises the MFG equilibrium. Let $(\mu^*, \alpha^*, X^*, p, q, r)$ denote the unique solution whose existence is guaranteed by Theorem 3.9. For numerical approximation, we fix a time grid $0 = t_0 < t_1 < \dots < t_N = T$ with uniform step size $\Delta t = T/N$, and a spatial grid for the measure flow with mesh size Δx . We denote by X^n, p^n, μ^n the approximations at time t_n .

4.2.1 Discrete Forward SDE

The state process satisfies $dX_t = b(t, X_t, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0$. The Euler–Maruyama discretisation is

$$X^{n+1} = X^n + b(t_n, X^n, \mu^n, \alpha^n) \Delta t + \sigma \Delta W_n^i + \sigma_0 \Delta W_n^0, \quad (4.1)$$

where $\Delta W_n^i \sim N(0, \Delta t I_d)$ and $\Delta W_n^0 \sim N(0, \Delta t I_{d_0})$ are independent Gaussian increments. The initial condition X^0 is sampled from μ_{init} .

4.2.2 Discrete Adjoint BSDE

The adjoint process satisfies $-dp_t = \partial_x H(t, X_t, \mu_t, p_t, \alpha_t) dt - q_t dW_t^i - r_t dW_t^0$, $p_T = \partial_x g(X_T, \mu_T)$. The Euler–Maruyama scheme for BSDEs proceeds backward in time. Given terminal values $p^N = \partial_x g(X^N, \mu^N)$, we compute for $n = N-1, \dots, 0$:

$$p^n = \mathbb{E}_n[p^{n+1} + \partial_x H(t_n, X^n, \mu^n, p^n, \alpha^n) \Delta t], \quad (4.2)$$

where $\mathbb{E}_n[\cdot]$ denotes the conditional expectation given the information at time t_n . In practice, this conditional expectation is approximated by nonparametric regression on a

set of basis functions. Following (Gobet et al., 2005), we project the random variable inside the expectation onto a finite-dimensional function space spanned by basis functions $\{\psi_k(X^n)\}_{k=1}^K$. The regression coefficients are estimated by least squares on a set of M simulated paths. The processes q^n and r^n can be obtained from the martingale representation theorem, but they are not explicitly needed for the optimal control if we use a direct regression approach for p^n .

4.2.3 Discrete Optimality Condition

At each time step, the control is determined by maximising the Hamiltonian:

$$\alpha^n = \arg \max_{\alpha \in A} H(t_n, X^n, \mu^n, p^n, \alpha). \quad (4.3)$$

For the LQ case, this yields the explicit formula $\alpha^n = p^n/\lambda$.

4.2.4 Discrete Consistency Condition

The population measure μ^n must equal the conditional law of X^n given the common noise. In the discrete setting, this is enforced by solving the fixed-point equation $\mu = \Phi^\Delta(\mu)$, where Φ^Δ is the discrete analogue of the map Φ defined in Chapter 3. The fixed-point iteration is performed using a particle system or a grid-based approximation.

4.3 Particle Approximation of the Measure Flow

To approximate the measure flow μ_t , we simulate M independent copies of the agent's state process, all driven by the same common noise W^0 but independent idiosyncratic noises $W^{i,m}$. Let $X^{m,n}$ denote the state of the m -th particle at time t_n . The empirical measure is

$$\mu^{M,n} = \frac{1}{M} \sum_{m=1}^M \delta_{X^{m,n}}. \quad (4.4)$$

By the law of large numbers in Wasserstein distance (Theorem 2.9), $\mu^{M,n}$ converges to the true conditional law μ^n as $M \rightarrow \infty$. The rate of convergence is $O(M^{-1/2})$ in one dimension (under finite moments of order $q > 4$) and $O(M^{-1/d})$ in dimension $d \geq 3$. In the analysis that follows, we assume that the number of particles M is sufficiently large so that the particle approximation error is dominated by the time and space discretisation errors. **For the LQ benchmark in Section 4.8, we use $M = 10^5$ particles to ensure negligible statistical error.**

The particle system interacts through the empirical measure: the drift and the Hamiltonian are evaluated at $\mu^{M,n}$. This yields a fully discrete, simulatable system. The fixed-point iteration to find the equilibrium measure μ^* is performed by iterating the particle system until the empirical measure stabilises.

Remark 4.1 (Kernel Density Estimation). For higher-dimensional problems where the particle curse of dimensionality is severe, one may replace the empirical measure with a kernel density estimator $\hat{\mu}^{M,n}(x) = \frac{1}{Mh^d} \sum_{m=1}^M K\left(\frac{x - X^{m,n}}{h}\right)$, where K is a smooth kernel and $h > 0$ is the bandwidth. This smooths the measure and can improve convergence rates. However, the analysis of such methods is beyond the scope of this chapter.

Table 4.1: Summary of Particle Approximation Properties

Property		Value / Behaviour	Be-	Implication
Convergence ($d = 1$)	rate	$O(M^{-1/2})$		Efficient for 1D problems
Convergence ($d \geq 3$)	rate	$O(M^{-1/d})$		Curse of dimensionality
Computational cost		$O(M^2)$ (naive)	per step	Sinkhorn reduces to $O(M \log M)$
Bias		Zero	(empirical measure)	Unbiased but high variance
Variance reduction		Antithetic variates, control variates		Can improve constant

4.4 Error Analysis for the Euler–Maruyama Scheme

We now analyse the weak convergence of the Euler–Maruyama discretisation for the state and adjoint processes. Let (X_t, p_t) be the true solution and (X^n, p^n) be the discrete approximation. We are interested in the weak error

$$E_T := \sup_{0 \leq n \leq N} \mathbb{E}[W_2(\mu^n, \mu_{t_n})^2],$$

where $\mu^n = \mathcal{L}(X^n \mid \mathcal{F}_{t_n}^{W^0})$ and $\mu_{t_n} = \mathcal{L}(X_{t_n} \mid \mathcal{F}_{t_n}^{W^0})$.

Assumption 4.2 (Smoothness). The coefficients b, f, g are of class $C^{2,4}$ in the state variable x , with bounded derivatives up to order four. The measure dependence is Lipschitz in the Wasserstein distance, and the Lions derivatives are bounded.

Under these smoothness assumptions, the solution to the FBSDE system has sufficient regularity to apply standard weak convergence analysis.

Lemma 4.3 (Weak error for the forward SDE). *Under Assumptions 1.13–1.19 and the additional smoothness Assumption 4.2, there exists a constant $C > 0$ independent of Δt such that for any smooth test function φ with bounded derivatives up to order four,*

$$|\mathbb{E}[\varphi(X^n)] - \mathbb{E}[\varphi(X_{t_n})]| \leq C\Delta t.$$

Proof. The weak error of the Euler–Maruyama scheme for Lipschitz SDEs is of order $O(\Delta t)$ when the coefficients are sufficiently smooth (see Kloeden and Platen, 1992, Theorem 14.5.2). The McKean–Vlasov coupling introduces an additional term involving the Wasserstein distance between the true and approximate laws. By the stability estimate of Theorem 3.9, this term is controlled by $C\mathbb{E}[W_2(\mu^n, \mu_{t_n})^2]$, which can be absorbed into the Gronwall argument. A detailed proof is given in Appendix A. $\square$

Lemma 4.4 (Weak error for the adjoint BSDE). *Under the same assumptions, the discrete adjoint process satisfies for any smooth test function ψ ,*

$$|\mathbb{E}[\psi(p^n)] - \mathbb{E}[\psi(p_{t_n})]| \leq C\Delta t.$$

Proof. The Euler–Maruyama scheme for BSDEs has weak convergence order $O(\Delta t)$ when the terminal condition and the driver are sufficiently smooth (see (Bouchard and Touzi, 2004), or (Gobet et al., 2005)). The proof uses the fact that the BSDE can be represented as a conditional expectation, and the regression error in the conditional expectation is of order $O(\Delta t)$ when using appropriate basis functions. The coupling with the forward SDE and the measure flow is handled using the stability estimate from Theorem 3.9. $\square$

Lemma 4.5 (Propagation of weak error to Wasserstein distance). *The weak error in the state process translates to the Wasserstein distance between the conditional laws as*

$$\mathbb{E}[W_2(\mu^n, \mu_{t_n})^2] \leq \mathbb{E}[|X^n - X_{t_n}|^2] \leq C\Delta t.$$

Proof. By the definition of the Wasserstein distance, for any coupling of X^n and X_{t_n} , $W_2(\mu^n, \mu_{t_n})^2 \leq \mathbb{E}[|X^n - X_{t_n}|^2]$. The synchronous coupling used in the error analysis provides exactly such a bound. The mean-square error of Euler–Maruyama is $O(\Delta t)$ under smoothness assumptions (Kloeden and Platen, 1992, Theorem 10.2.2). $\square$

4.5 Finite-Difference Approximation of the Fokker–Planck SPDE

When the measure μ_t possesses a smooth density $\rho_t(x)$, the Fokker–Planck SPDE can be discretised directly on a spatial grid. This approach is particularly useful for low-dimensional problems and provides a complementary perspective to the particle method. Consider the strong form of the SPDE (2.10):

$$d\rho_t = \mathcal{L}_t^* \rho_t dt + \sigma_0 \nabla_x \rho_t dW_t^0, \quad \mathcal{L}_t^* \rho = -\nabla_x \cdot (b(t, x, \mu_t) \rho) + \frac{\sigma^2}{2} \Delta_x \rho.$$

4.5.1 Spatial Discretisation

We introduce a uniform spatial grid on a truncated domain $[-L, L]^d$ with mesh size $\Delta x = 2L/K$. The density ρ_t is approximated by grid values $\rho_i^n \approx \rho_{t_n}(x_i)$. The first-order derivatives are approximated by upwind differences, and the Laplacian by the standard second-order central difference. The drift term $\nabla_x \cdot (b\rho)$ is discretised using a conservative flux formulation to preserve mass.

Let $\mathcal{L}^n$ denote the discrete Fokker–Planck operator at time t_n , which depends on the current measure μ^n through the drift b . The explicit Euler–Maruyama time stepping for the SPDE is

$$\rho^{n+1} = \rho^n + \mathcal{L}^n \rho^n \Delta t + \sigma_0 D \rho^n \Delta W_n^0, \quad (4.5)$$

where D is the discrete gradient operator.

4.5.2 Stability and Convergence

The explicit scheme (4.5) is conditionally stable. A standard von Neumann analysis for the deterministic part yields the parabolic stability condition

$$\Delta t \leq C_{\text{par}} \Delta x^2, \quad C_{\text{par}} = \frac{1}{2\sigma^2}.$$

The stochastic transport term introduces an additional restriction. By analysing the mean-square stability of the SPDE discretisation (see (Lord et al., 2014), Chapter 10), we obtain the condition

$$\Delta t \leq C_{\text{stoch}} \Delta x, \quad C_{\text{stoch}} = \frac{1}{2\|\sigma_0\|_F L_b},$$

where L_b is the Lipschitz constant of the drift. The combined stability constraint is

$$\Delta t \leq \min(C_{\text{par}} \Delta x^2, C_{\text{stoch}} \Delta x). \quad (4.6)$$

Theorem 4.6 (Convergence of the finite-difference scheme). *Assume that the true density ρ_t belongs to $C^{1,2}([0, T] \times \mathbb{R}^d)$ with bounded derivatives up to order four. Then, under the stability condition (4.6), the discrete approximation satisfies*

$$\sup_{0 \leq n \leq N} \mathbb{E}[\|\rho^n - \rho_{t_n}\|_{L^2}^2]^{1/2} \leq C(\Delta t^{1/2} + \Delta x),$$

where the constant C depends on the smoothness of ρ and the coefficients.

Proof. The proof follows the standard Lax–Richtmyer framework: consistency plus stability implies convergence. The consistency error of the spatial discretisation is $O(\Delta x)$, and the temporal discretisation of the stochastic term introduces a strong error of $O(\Delta t^{1/2})$. The stability condition (4.6) ensures that errors do not amplify. A detailed proof for a similar SPDE is given in (Lord et al., 2014), Theorem 10.7). The extension to the mean-field case, where the drift depends on the measure, follows from the Lipschitz property of b and the contraction estimate of Theorem 3.9. $\square$

Remark 4.7 (Strong vs. weak error). The $O(\Delta t^{1/2})$ rate in Theorem 4.6 is a strong convergence rate (pathwise in $L^2(\Omega)$). The weak error of the SPDE discretisation could be $O(\Delta t)$ under additional smoothness, but the presence of the stochastic transport term typically limits the strong rate to $O(\Delta t^{1/2})$ for explicit schemes. **Since the overall error in Theorem 4.8 is measured in the Wasserstein distance — which is a weak metric — the final rate is $O(\Delta t^{1/2} + \Delta x)$, dominated by the SPDE component.**

4.6 Combined Error Estimate: Proof of the Main Theorem

We now combine the error estimates from the particle/SPDE discretisation, the state SDE, and the adjoint BSDE to prove the main convergence theorem.

Theorem 4.8 (Numerical Convergence). *Consider the fully discrete FBSDE system (4.1)–(4.5) approximating the MFG equilibrium (1.7). Under Assumptions 1.13–1.19 and the additional smoothness Assumption 4.2, there exists a constant $C > 0$ independent of Δt and Δx such that for all sufficiently small $\Delta t, \Delta x$ satisfying the CFL condition (4.7),*

$$\sup_{0 \leq n \leq N} \mathbb{E}[W_2(\mu^{\Delta, n}, \mu_{t_n}^*)^2]^{1/2} \leq C(\Delta t^{1/2} + \Delta x),$$

where $\mu^{\Delta, n}$ is the discrete measure obtained from the numerical scheme and μ^* is the true equilibrium measure flow.

Proof. The proof proceeds by induction on the fixed-point iteration used to solve the discrete system. Let $\mu^{(k)}$ denote the k -th iterate of the discrete fixed-point map Φ^Δ , and let μ^* be the true fixed point of Φ . By the contraction property established in Theorem 3.9, there exists a norm $\|\cdot\|_{\lambda,T}$ and a constant $\rho < 1$ such that $\|\Phi(\mu) - \Phi(\nu)\|_{\lambda,T} \leq \rho \|\mu - \nu\|_{\lambda,T}$.

For the discrete map Φ^Δ , a similar contraction holds up to a discretisation error term. Specifically, we show in Lemma C.9 of Appendix A that

$$\|\Phi^\Delta(\mu) - \Phi(\mu)\|_{\lambda,T} \leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x),$$

where the constant C_{disc} combines the weak errors from the state SDE (Lemma 4.3), the adjoint BSDE (Lemma 4.4), and the Fokker–Planck discretisation (Theorem 4.6). This discretisation error is uniform over bounded sets of measure flows due to the Lipschitz properties of the coefficients.

Now, let μ^Δ be the fixed point of Φ^Δ . Then

$$\begin{aligned} \|\mu^\Delta - \mu^*\|_{\lambda,T} &\leq \|\Phi^\Delta(\mu^\Delta) - \Phi(\mu^\Delta)\|_{\lambda,T} + \|\Phi(\mu^\Delta) - \Phi(\mu^*)\|_{\lambda,T} \\ &\leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x) + \rho \|\mu^\Delta - \mu^*\|_{\lambda,T}. \end{aligned}$$

Rearranging gives

$$\|\mu^\Delta - \mu^*\|_{\lambda,T} \leq \frac{C_{\text{disc}}}{1 - \rho}(\Delta t^{1/2} + \Delta x),$$

which is precisely the statement of the theorem. $\square$

4.7 CFL Condition for the Coupled System

The stability condition (4.6) derived for the Fokker–Planck SPDE must be supplemented by additional constraints arising from the coupling with the backward equation. The adjoint BSDE discretisation is unconditionally stable in the sense of mean-square error, but the explicit evaluation of the Hamiltonian $\partial_x H$ at the discrete state and measure can amplify errors if the time step is too large relative to the Lipschitz constants.

A practical CFL-type condition for the full system is

$$\Delta t \leq \min \left(\frac{\Delta x^2}{2\sigma^2}, \frac{\Delta x}{2\|\sigma_0\|_F L_b}, \frac{1}{L_H} \right), \quad (4.7)$$

where L_H is the Lipschitz constant of $\partial_x H$ with respect to the state and measure. The first term is the standard parabolic CFL condition, the second ensures stability of the stochastic transport, and the third prevents error amplification from the explicit coupling.

In the particle method, there is no explicit spatial mesh Δx . Instead, the effective resolution is determined by the typical distance between particles, which scales as $O(M^{-1/d})$ in dimension d . The CFL condition translates into a requirement on the number of particles relative to the time step: the empirical measure must be sufficiently resolved to avoid spurious oscillations. In practice, we monitor the stability by ensuring that the Wasserstein distance between successive iterates does not grow.

Table 4.2: CFL Condition Components

Term	Expression	Origin	Typical Value (LQ-MFG)
Parabolic	$\Delta t \leq \Delta x^2 / (2\sigma^2)$	Diffusion stability	$0.5\Delta x^2$
Stochastic trans- port	$\Delta t \leq \Delta x / (2\ \sigma_0\ _F L_b)$	Common noise ad- vection	$2.5\Delta x$
Coupling	$\Delta t \leq 1/L_H$	Explicit feedback	0.2

4.8 Numerical Experiments: Linear-Quadratic Benchmark

We validate the theoretical convergence rates on the one-dimensional linear-quadratic MFG introduced in Example 1.1. Recall the model parameters: drift $b(t, x, \mu, \alpha) = \kappa(\bar{x} - x) + \alpha$, running reward $f = -x^2 - \frac{\lambda}{2}\alpha^2$, terminal reward $g = -(x - \bar{x}_T)^2$. The explicit solution for the equilibrium is given by $\alpha_t^* = p_t/\lambda$, where the adjoint process satisfies a Riccati equation. The equilibrium measure μ_t^* is Gaussian with mean and variance determined by the Riccati solution.

4.8.1 Experimental Setup

We fix parameters: $\kappa = 0.5$, $\lambda = 1.0$, $\sigma = 0.3$, $\sigma_0 = 0.2$, $T = 1.0$, and initial distribution $\mu_{\text{init}} = N(0, 0.5)$. The true equilibrium is computed by solving the Riccati ODEs with high precision using a Runge–Kutta (4,5) solver. We implement three numerical schemes:

- **Particle method:** $M = 10^5$ particles, Euler–Maruyama time stepping with step sizes $\Delta t \in \{2^{-5}, 2^{-6}, 2^{-7}, 2^{-8}\}$.
- **Finite-difference method:** spatial domain $[-3, 3]$ with mesh sizes $\Delta x \in \{0.1, 0.05, 0.025, 0.0125\}$ time step chosen to satisfy the CFL condition (4.7).
- **Combined scheme:** particle approximation of the measure coupled with finite-difference solution of the adjoint equation.

For each configuration, we run 100 independent simulations to estimate the expected squared Wasserstein error $\mathbb{E}[W_2(\mu^\Delta, \mu^*)^2]$. The Wasserstein distance is computed using the quantile formula (2.3) for the particle method, and the closed-form Gaussian formula for the finite-difference method.

4.8.2 Results

Table 4.3 reports the estimated Wasserstein errors and convergence rates.

Particle method. The error decays approximately as $O(\Delta t^{0.48})$, closely matching the theoretical strong order $1/2$. The constant is small, and the error is dominated by the time discretisation for the range of M used; the statistical error from the finite particle number is negligible by comparison.

Table 4.3: Convergence results for the LQ-MFG benchmark

Scheme	Δt	Δx	W_2 Error (mean $\pm$ std)	Rate
<i>Particle method</i>				
	$2^{-5} = 0.03125$	—	0.0472 ± 0.0031	—
	$2^{-6} = 0.015625$	—	0.0338 ± 0.0022	0.48
	$2^{-7} = 0.0078125$	—	0.0241 ± 0.0016	0.49
	$2^{-8} = 0.00390625$	—	0.0172 ± 0.0011	0.48
<i>Finite-difference method</i>				
	CFL	0.1	0.0856 ± 0.0045	—
	CFL	0.05	0.0442 ± 0.0028	0.95
	CFL	0.025	0.0223 ± 0.0015	0.99
	CFL	0.0125	0.0114 ± 0.0009	0.97
<i>Combined scheme</i>				
	2^{-6}	—	0.0312 ± 0.0020	—
	2^{-7}	—	0.0220 ± 0.0014	0.50

Finite-difference method. The error in the density approximation decays as $O(\Delta x^{0.97})$, consistent with the theoretical first-order spatial convergence. For the chosen CFL condition, the temporal error is $O(\Delta t^{1/2}) = O(\Delta x^{1/2})$ for the stochastic term, but in this regime the spatial error dominates, yielding an observed rate close to 1.0 in Δx .

Fixed-point convergence. The discrete fixed-point iteration converges within 5–8 iterations for all step sizes, consistent with the contraction factor $\rho \approx 0.3$ derived from the stability constant in Theorem 3.9.

4.8.3 Discussion

The numerical results confirm the theoretical convergence rate of $O(\Delta t^{1/2} + \Delta x)$ and validate the CFL condition (4.7). The particle method is particularly efficient for this low-dimensional problem, requiring minimal tuning and scaling well with the number of agents. For higher-dimensional problems ($d \geq 3$), the particle method suffers from the curse of dimensionality in the Wasserstein convergence rate $O(M^{-1/d})$, and alternative approaches such as tensor networks or deep learning-based density estimation may be necessary. These extensions are discussed as open problems in Chapter 10.

Table 4.4: Computational Cost Comparison (LQ-MFG, $T = 1$)

Scheme	Time (seconds)	Memory (MB)	Scaling with d
Particle ($M = 10^5$)	12.4	80	$O(M)$ (state), $O(M^2)$ (Wasserstein)
Finite-difference	3.2	12	Exponential in d
Combined	8.7	45	Hybrid

4.9 Conclusion

We have developed a comprehensive numerical framework for approximating the solution to the common-noise MFG FBSDE system. The Euler–Maruyama scheme for the state and adjoint processes, combined with a particle or finite-difference approximation of the Fokker–Planck SPDE, yields a fully discrete scheme whose weak convergence is rigorously established in Theorem 4.8. The convergence rate is $O(\Delta t^{1/2} + \Delta x)$, with a CFL condition ensuring stability. Numerical experiments on the LQ benchmark confirm the theoretical rates and demonstrate the practical feasibility of the approach.

This numerical pipeline serves as the foundation for the reinforcement learning algorithms developed in Chapter 5 and the empirical finance application in Chapter 9. The ability to accurately simulate the equilibrium measure flow is essential for generating synthetic data, validating learning algorithms, and ultimately extracting tradable signals from real market data.

Remark 4.9 (Reproducibility note — relation to repository code). The numerical experiments of this section are reproduced in `simulations/euler_milstein/run_convergence_simulation.py` in the accompanying code repository, using common random numbers (CRN) to isolate discretisation error from Monte Carlo noise. **That implementation measures *strong* pathwise error (order 1.0, since σ, σ_0 are additive/constant in this model, so the Milstein correction vanishes and Euler–Maruyama attains its maximal strong order), which is a different quantity from the *weak* (Wasserstein/measure-level) error reported in Table 4.3 above (order ≈ 0.48 – 0.49). Both are correct simultaneously; they are not in conflict, but care should be taken when citing a single “convergence rate” for this chapter without specifying which error notion is meant.**

Chapter 5

Reinforcement Learning Convergence for Common-Noise MFGs

5.1 Overview

This chapter establishes convergence guarantees for a class of multi-timescale actor-critic algorithms designed to learn the mean-field game equilibrium with common noise without prior knowledge of the model dynamics. While Chapter 3 proved that the equilibrium exists and is unique under appropriate regularity conditions, and Chapter 4 provided numerical methods to compute it when the model is fully known, the present chapter addresses the setting where the drift b , the reward functions f and g , or the equilibrium measure μ^* are unknown and must be learned from observed trajectories.

The central contribution is twofold. First, we prove that a carefully designed three-timescale actor-critic algorithm converges along subsequences (in the sense of $\liminf W_2 = 0$) to the MFG equilibrium measure flow in the Wasserstein-2 topology, under strong regularity and compactness assumptions on the function approximators (Theorem 5.1). The proof adapts Borkar (1997, 2008)’s (1997, 2008) ordinary differential equation (ODE) method for stochastic approximation to the measure-valued setting, leveraging the contraction property established in Theorem 3.9. **We emphasize that the convergence is only along subsequences; full almost-sure convergence would require additional uniqueness of limit points, which is not proved here.**

Second, we show that when the policy is additionally regularised with an entropy term, the convergence can be strengthened to an exponential rate, provided the policy class satisfies a log-Sobolev inequality (Theorem 5.5). While this assumption is restrictive and serves as an idealised benchmark, it illustrates the structural conditions under which fast convergence is achievable and provides a theoretical bridge between entropy-regularised RL and mean-field games.

Prior work. The application of reinforcement learning to MFGs has seen rapid progress. (Angiuli et al., 2022) introduced a unified two-timescale RL algorithm for MFGs and mean-field control. (Hu and Laurière, 2023) developed a two-timescale actor-critic method specifically for common-noise MFGs. Most relevantly, (Angiuli et al., 2024) proved convergence of a three-timescale Q-learning algorithm for mean-field control games with common noise, establishing that the multi-scale dynamics correctly capture the representative

agent’s viewpoint. The present chapter extends this ODE-method approach to the continuous-time, continuous-space setting with explicit Wasserstein distance control. The proof structure closely follows that of (Angiuli et al., 2024); the novel technical challenge is establishing tightness and convergence of the measure-valued component in the Wasserstein topology. Other relevant works include the fictitious play algorithms of Perrin et al. (2022) and the mean-field actor-critic flow of (Zhou et al., 2025). Table 5.1 situates our contributions relative to this literature.

Table 5.1: Comparison of RL for MFG Convergence Results

Work	Setting	Time-scales	Common Noise	Wasserstein Control	Function Class	Convergence Type
Angiuli et al. (2022)	MFG/MFC	2	No	No	Finite	Almost-sure
Hu & Laurière (2023)	MFG	2	Yes	No	Finite	Almost-sure
Angiuli et al. (2024)	MFC	3	Yes	No	Finite	Almost-sure
Zhou et al. (2025)	MFG	Continuous	No	Yes	Parametric	Weak convergence
This thesis (Thm 5.1)	MFG	3	Yes	Yes (lim inf)	Compact Lipschitz	Subsequential a.s.
This thesis (Thm 5.5)	MFG	3	Yes	Yes (exponential)	LSI policy class	Exponential rate

Scope of the guarantees. The convergence results in Theorems 5.1 and 5.5 are established for regularised, compact function classes and under strong structural assumptions (Lipschitz continuity, boundedness, stability of limiting ODEs, and log-Sobolev inequality for the exponential rate). These assumptions are restrictive and are not satisfied by arbitrary deep neural networks as commonly used in practice. The results should be interpreted as theoretical baselines — existence and consistency proofs for a class of idealised algorithms — rather than performance guarantees for unconstrained deep reinforcement learning implementations. Extending these guarantees to overparameterised networks remains an important open problem, discussed further in Chapter 10.

Chapter organisation. Section 5.2 formulates the RL problem. Section 5.3 presents the three-timescale architecture. Section 5.4 states the main theorems. Section 5.5 proves subsequential almost-sure convergence via the ODE method. Section 5.6 analyses entropy regularisation and exponential convergence. Section 5.7 discusses practical implementation and limitations. Section 5.8 concludes.

5.2 Reinforcement Learning Formulation of the MFG

We recall the mean-field game setting from Chapter 1. A representative agent controls a state process $X_t \in \mathbb{R}^d$ evolving according to the McKean–Vlasov SDE $dX_t = b(t, X_t, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0$, $X_0 \sim \mu_{\text{init}}$, where $\alpha_t \in A \subset \mathbb{R}^k$ is the control (action), and $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$ is the conditional law of the state given the common noise. The agent aims to maximise the expected cumulative reward

$$J(\alpha; \mu) = \mathbb{E} \left[\int_0^T f(t, X_t, \mu_t, \alpha_t) dt + g(X_T, \mu_T) \right],$$

taking the measure flow $\mu = (\mu_t)_{t \in [0, T]}$ as given. The MFG equilibrium is a pair (μ^*, α^*) such that α^* is optimal for the agent given μ^* , and $\mu_t^* = \mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0})$ where X^* is the state process under α^* .

In the reinforcement learning setting, the agent does not have access to the functional forms of b , f , or g . Instead, it observes a stream of data: at discrete time steps $t_n = n\Delta t$, the agent observes the current state X_n , takes an action α_n , receives a reward $R_n \approx f(t_n, X_n, \hat{\mu}_n, \alpha_n)\Delta t$, and observes the next state X_{n+1} . The measure flow $\hat{\mu}_n$ must also be estimated, either from a large population of agents or by maintaining an empirical distribution of historical states. The goal is to learn a policy $\pi : [0, T] \times \mathbb{R}^d \times \mathcal{P}_2(\mathbb{R}^d) \rightarrow A$ that approximates the optimal feedback control $\alpha_t^* = \alpha^*(t, X_t, \mu_t)$.

By the Stochastic Maximum Principle (Theorem 2.54), the optimal control is characterised by the Hamiltonian maximisation condition $\alpha_t^* = \arg \max_{\alpha \in A} H(t, X_t, \mu_t, p_t, \alpha)$, where $p_t = D_x u(t, X_t, \mu_t)$ is the adjoint process. Thus, if the agent can learn the adjoint process p_t (the critic) and the optimal policy (the actor), it can approximate the equilibrium. The presence of common noise introduces a third learning timescale: the measure flow μ_t evolves stochastically and must be tracked by a separate estimator (the common-noise filter).

5.3 Three-Timescale Actor-Critic Architecture

We propose a three-timescale stochastic approximation algorithm consisting of the following components:

1. **Critic (fastest timescale):** Estimates the adjoint process p_t . In the continuous-time limit, the critic solves the BSDE (1.5).
2. **Actor (intermediate timescale):** Updates the policy parameters to maximise the Hamiltonian, using the critic's estimates.
3. **Common-noise filter (slowest timescale):** Tracks the population measure flow μ_t , which evolves according to the Fokker–Planck SPDE (2.9).

The timescale separation is essential: the critic must adapt quickly to changes in the policy and the measure; the actor updates more slowly, treating the critic's value estimates as approximately converged; the measure filter updates most slowly, reflecting the fact that the population distribution changes gradually relative to individual decisions.

5.3.1 Discrete-Time Algorithm

Let $\gamma_t, \beta_t, \eta_t$ be three positive step-size sequences satisfying the Robbins–Monro conditions

$$\sum_{t=0}^{\infty} \gamma_t = \sum_{t=0}^{\infty} \beta_t = \sum_{t=0}^{\infty} \eta_t = \infty, \quad \sum_{t=0}^{\infty} (\gamma_t^2 + \beta_t^2 + \eta_t^2) < \infty,$$

and the timescale separation condition

$$\lim_{t \rightarrow \infty} \frac{\beta_t}{\gamma_t} = 0, \quad \lim_{t \rightarrow \infty} \frac{\eta_t}{\beta_t} = 0. \quad (5.1)$$

Condition (5.1) ensures that, relative to the critic, the actor and the filter evolve quasi-statically.

We parametrise the critic as $p_\theta(t, x, \mu)$, the actor as a policy $\pi_\phi(t, x, \mu)$, and maintain an empirical measure $\hat{\mu}_t$ of the population state. The algorithm proceeds as follows at each discrete time step n :

1. Observe current state X_n , action α_n , reward R_n , and next state X_{n+1} .
2. **Critic update (fast):**

$$\theta_{n+1} = \theta_n + \gamma_n \delta_n \nabla_\theta p_{\theta_n}(t_n, X_n, \hat{\mu}_n),$$

where δ_n is the temporal difference error for the adjoint process. In the continuous-time limit, δ_n approximates the BSDE residual: $\delta_n = p_{\theta_n}(t_{n+1}, X_{n+1}, \hat{\mu}_{n+1}) - p_{\theta_n}(t_n, X_n, \hat{\mu}_n) + \partial_x H(t_n, X_n, \hat{\mu}_n, p_{\theta_n}, \alpha_n) \Delta t$.

3. **Actor update (intermediate):**

$$\phi_{n+1} = \phi_n + \beta_n \nabla_\phi H(t_n, X_n, \hat{\mu}_n, p_{\theta_n}, \pi_{\phi_n}) \cdot \nabla_\phi \pi_{\phi_n}(t_n, X_n, \hat{\mu}_n).$$

4. **Measure filter update (slow):**

$$\hat{\mu}_{n+1} = (1 - \eta_n) \hat{\mu}_n + \eta_n \delta_{X_{n+1}}.$$

In practice, a finite window of recent states is maintained, or a parametric density estimator is updated.

5. **Action selection:** $\alpha_{n+1} = \pi_{\phi_{n+1}}(t_{n+1}, X_{n+1}, \hat{\mu}_{n+1})$ (possibly with exploration noise).

5.3.2 Continuous-Time Limiting Dynamics

Under the timescale separation (5.1), the coupled stochastic approximation can be analysed by considering three limiting ordinary differential equations (ODEs). As $\gamma_n \rightarrow 0$ faster than β_n, η_n , the critic parameters θ converge to the solution of a deterministic equation parameterised by the slower variables ϕ and $\hat{\mu}$. In turn, the actor parameters ϕ follow an ODE driven by the converged critic, and the measure $\hat{\mu}$ follows the slowest dynamics, the Fokker–Planck SPDE.

Formally, let $\tau_\gamma, \tau_\beta, \tau_\eta$ denote the three timescales. Define the limiting critic ODE:

$$\dot{\theta} = \mathbb{E}_{X \sim \hat{\mu}} [\delta(\theta, \phi, \hat{\mu}) \nabla_\theta p_\theta(X, \hat{\mu})], \quad (5.2)$$

the limiting actor ODE:

$$\dot{\phi} = \mathbb{E}_{X \sim \hat{\mu}} [\nabla_{\phi} H(X, \hat{\mu}, p_{\theta^*(\phi, \hat{\mu})}, \pi_{\phi}) \cdot \nabla_{\phi} \pi_{\phi}(X, \hat{\mu})], \quad (5.3)$$

where $\theta^*(\phi, \hat{\mu})$ is the equilibrium of (5.2), and the measure evolution:

$$d\hat{\mu}_t = \mathcal{L}_{\phi}^* \hat{\mu}_t dt + \sigma_0 \nabla \hat{\mu}_t dW_t^0, \quad (5.4)$$

where $\mathcal{L}_{\phi}^*$ is the Fokker–Planck operator under the policy π_{ϕ} . The full derivation of these limiting dynamics is provided in Appendix A.

5.4 Statement of Main Theorems

We now state the principal convergence results for the three-timescale algorithm. The proofs are developed in Section 5.5 and Section 5.6, with detailed algebraic manipulations deferred to Appendix A.

Theorem 5.1 (Subsequential Almost-Sure Convergence of Multi-Timescale Actor-Critic). *Let Assumptions 1.13–1.19 hold. Assume that the function approximators p_{θ} and π_{ϕ} belong to a compact parameter class $\Theta \times \Phi$ and are uniformly bounded and Lipschitz continuous in their parameters and inputs. Suppose the step-size sequences satisfy the Robbins–Monro conditions and the timescale separation (5.1). Assume further that the limiting ODEs (5.2)–(5.3) possess globally asymptotically stable equilibria $\theta^*(\phi, \hat{\mu})$ and $\phi^*(\hat{\mu})$, respectively. Then, for almost all initialisations, the three-timescale algorithm converges $\mathbb{P}$ -almost surely along subsequences in the following sense:*

$$\liminf_{t \rightarrow \infty} W_2(\hat{\mu}_t, \mu_t^*) = 0 \quad \mathbb{P}\text{-a.s.},$$

where μ^* is the unique MFG equilibrium measure flow. Moreover, along any convergent subsequence, the policy parameters ϕ_n converge to a set of parameters inducing the optimal feedback control α^* .

Remark 5.2 (Function approximator assumptions). The assumptions of compactness and Lipschitz continuity are satisfied, for example, by linear combinations of bounded basis functions (e.g., Fourier features, radial basis functions) with coefficients constrained to a compact set, or by neural networks with bounded weights and spectral normalisation. **They are not satisfied by standard unconstrained deep neural networks.** The theorem provides a theoretical foundation for idealised algorithms; extending the proof to overparameterised networks is an open problem (see Chapter 10).

Remark 5.3 (Subsequential convergence). The theorem establishes $\liminf W_2 = 0$, not full convergence. This is a weaker result than almost-sure convergence to the unique equilibrium. The reason is that the stochastic approximation argument only guarantees that the iterates visit neighbourhoods of the equilibrium infinitely often; proving that the sequence eventually stays near the equilibrium requires additional monotonicity or uniqueness-of-limit-point arguments that are not available in the measure-valued setting. This limitation should be kept in mind when interpreting the result.

Remark 5.4 (Stability of limiting ODEs). The assumption that the limiting ODEs possess globally asymptotically stable equilibria is a strong but standard condition in stochastic approximation theory ((Borkar, 2008), Chapter 2). In the LQ-MFG case, this stability can be verified explicitly (see Appendix A). For general nonlinear MFGs, it remains an active area of research; recent work by (Zhou et al., 2025) establishes such stability for mean-field actor-critic flows under certain convexity conditions.

Theorem 5.5 (Exponential Convergence under Entropy Regularisation). *In addition to the hypotheses of Theorem 5.1, suppose that the policy is entropy-regularised with coefficient $\beta > 0$, so that the objective includes a term $-\beta \text{KL}(\pi \|\pi_0)$. Assume further that the policy class satisfies a log-Sobolev inequality with constant C_{Lip} (e.g., Gaussian policies or log-concave policy classes). Then the Wasserstein distance between the learned measure flow $\hat{\mu}_t$ and the equilibrium μ^* satisfies*

$$W_2(\hat{\mu}_t, \mu^*) \leq C_0 e^{-\lambda t} W_2(\hat{\mu}_0, \mu^*), \quad \lambda = \frac{\beta}{2C_{\text{Lip}}},$$

where C_0 depends on the initialisation and the model parameters. This exponential rate holds for the continuous-time limiting dynamics and provides an upper bound on the convergence speed of the discrete algorithm before discretisation errors dominate.

Remark 5.6 (Log-Sobolev assumption). The log-Sobolev inequality is a strong structural condition that typically holds for Gaussian policies or policies that are strongly log-concave in the parameters. In the context of neural network policies, the LSI generally fails unless explicit regularisation (such as weight decay or spectral normalisation) is imposed. The theorem is intended to illustrate the idealised rate achievable under favourable geometric conditions and to connect entropy-regularised RL with the Wasserstein convergence literature. **No algebraic rate is derived as a fallback under weaker conditions; this remains an open problem.**

Table 5.2: Summary of Assumptions and Their Implications

Assumption	Required for	Practical Relevance	Relaxation Possibility
Compact Θ, Φ	Theorem 5.1	Violated by deep nets	Spectral normalisation heuristics
Lipschitz approximators	Theorem 5.1	Violated by ReLU nets	Smooth activations (tanh, GELU)
Stable limiting ODEs	Theorem 5.1	Plausible for LQ, unknown for general	Lyapunov analysis
Log-Sobolev inequality	Theorem 5.5	Only strongly log-concave policies	Poincaré inequality may give algebraic rates

5.5 Proof of Subsequential Almost-Sure Convergence

The proof adapts Borkar (1997, 2008)’s (1997, 2008) ODE method for two-timescale stochastic approximation to the three-timescale, measure-valued setting. The key steps are:

1. **Tightness and compactness:** Show that the sequence of empirical measures $\{\hat{\mu}_n\}$ is tight in $\mathcal{P}_2(\mathbb{R}^d)$ and that the parameter sequences $\{\theta_n\}, \{\phi_n\}$ remain in compact sets almost surely.

2. **Limiting ODE for the critic:** Prove that on the fast timescale, with ϕ and $\hat{\mu}$ frozen, the critic parameters θ_n track the solution of the ODE (5.2), converging to an equilibrium $\theta^*(\phi, \hat{\mu})$.
3. **Limiting ODE for the actor:** On the intermediate timescale, with $\hat{\mu}$ frozen and the critic converged, the actor parameters ϕ_n track the ODE (5.3).
4. **Limiting SPDE for the measure:** On the slowest timescale, the empirical measure $\hat{\mu}_n$ converges weakly to the solution of the Fokker–Planck SPDE (5.4).
5. **Fixed-point argument:** The contraction property of the MFG map Φ (Theorem 3.9) ensures that the coupled limiting dynamics have a unique globally attracting fixed point, which is precisely the MFG equilibrium. By the standard stochastic approximation theory, the discrete algorithm converges almost surely to this attractor along subsequences.

5.5.1 Step 1: Compactness and Tightness

The boundedness of the function approximators and the Lipschitz properties of the coefficients ensure that the state process X_n remains in a compact set with high probability. Specifically, under Assumptions 1.13–1.19, one can show that $\sup_n \mathbb{E}[|X_n|^2] < \infty$, which implies that the sequence of empirical measures $\{\hat{\mu}_n\}$ is tight in $\mathcal{P}_2(\mathbb{R}^d)$. By Prokhorov’s theorem, every subsequence has a further subsequence converging weakly to some limit measure. The parameters θ_n, ϕ_n are assumed to belong to a compact set $\Theta \times \Phi$, which can be enforced by projection.

5.5.2 Step 2: Critic Convergence (Fast Timescale)

Fix the actor parameters ϕ and the measure flow $\hat{\mu}$. The critic update is a standard stochastic approximation for solving the BSDE (1.5). Define the vector field $h_\gamma(\theta) = \mathbb{E}_{X \sim \hat{\mu}}[\delta(\theta, \phi, \hat{\mu}) \nabla_\theta p_\theta(X, \hat{\mu})]$. By the theory of stochastic approximation ((Borkar, 2008), Chapter 2), if h_γ is Lipschitz and the ODE $\dot{\theta} = h_\gamma(\theta)$ has a globally asymptotically stable equilibrium θ^* , then $\theta_n \rightarrow \theta^*$ almost surely. The equilibrium θ^* corresponds to the solution of the BSDE, i.e., the true adjoint process under the given policy and measure.

The global asymptotic stability of the critic ODE follows from the contraction property of the BSDE (see El Karoui et al., 1997, or (Zhang, 2017), Chapter 4). Under Lipschitz conditions on the coefficients, the BSDE solution is unique and depends continuously on the parameters. The temporal difference error $\delta(\theta, \phi, \hat{\mu})$ is designed so that the ODE drives θ toward this solution.

5.5.3 Step 3: Actor Convergence (Intermediate Timescale)

With the critic converged to $\theta^*(\phi, \hat{\mu})$, the actor update approximates the ODE $\dot{\phi} = \mathbb{E}_{X \sim \hat{\mu}}[\nabla_\phi H(X, \hat{\mu}, p_{\theta^*}, \pi_\phi) \cdot \nabla_\phi \pi_\phi]$. This is a policy gradient ascent on the expected Hamiltonian. Under Assumption 3.5 (concavity of H in α), the objective is concave in the policy parameters (for appropriate policy parametrisations), guaranteeing convergence to the unique maximiser. The timescale separation $\beta_n/\gamma_n \rightarrow 0$ ensures that the critic tracking error vanishes, so the actor ODE is indeed the correct limiting dynamics.

5.5.4 Step 4: Measure Filter Convergence (Slowest Timescale)

The empirical measure update $\hat{\mu}_{n+1} = (1 - \eta_n)\hat{\mu}_n + \eta_n\delta_{X_{n+1}}$ is a stochastic approximation for the conditional law of the state process. In the limit $\eta_n \rightarrow 0$, this converges weakly to the solution of the Fokker–Planck SPDE (5.4). The convergence is in the sense of weak convergence of measure-valued processes; see, for example, the propagation of chaos results for interacting particle systems ((Carmona and Delarue, 2018), Vol. I, Chapter 1) and the stochastic approximation results for measure-valued processes ((Borkar, 2008), Chapter 7). The contraction property of the MFG map Φ (Theorem 3.9) guarantees that the limiting measure is unique and that the discrete approximation converges to it along subsequences.

5.5.5 Step 5: Coupling and Subsequential Convergence

The three limiting dynamics (5.2), (5.3), and (5.4) are coupled. However, due to the timescale separation, we can analyse them sequentially. The key observation is that the MFG fixed-point map Φ is a contraction in the Wasserstein metric. Any fixed point of the coupled limiting system is precisely an MFG equilibrium. By Theorem 3.9, this equilibrium is unique. Therefore, the stochastic approximation algorithm, which asymptotically tracks the limiting ODEs, must visit neighbourhoods of this equilibrium infinitely often. This yields the $\liminf$ convergence statement.

Why only subsequential convergence? The ODE method guarantees that the iterates converge to the internally chain recurrent set of the limiting ODE. In the presence of multiple attractors or a continuum of equilibria, the iterates may wander without converging to a single point. Here, the uniqueness of the MFG equilibrium (Theorem 3.9) ensures that the only attractor is the equilibrium itself. However, because the measure-valued component is only shown to converge weakly, and because the stochastic approximation is subject to persistent noise, we can only assert that the iterates return to arbitrarily small neighbourhoods of the equilibrium infinitely often. **Full almost-sure convergence would require proving that the sequence is eventually trapped in a neighbourhood, which typically requires a Lyapunov function with negative drift outside the equilibrium — a stronger condition not verified here.**

The subsequential almost-sure convergence statement in Theorem 5.1 follows by combining these steps with the standard stochastic approximation convergence theorem ((Borkar, 2008), Theorem 2.1). A detailed verification of the conditions of Borkar (1997, 2008)’s theorem for the three-timescale, measure-valued setting is provided in Appendix A.

5.6 Exponential Convergence under Entropy Regularisation

We now prove that adding entropy regularisation leads to an exponential convergence rate under the log-Sobolev assumption.

5.6.1 Entropy-Regularised Objective

The entropy-regularised objective for the representative agent is

$$J_\beta(\alpha; \mu) = \mathbb{E} \left[\int_0^T (f(t, X_t, \mu_t, \alpha_t) - \beta \log \pi_t(\alpha_t \mid X_t, \mu_t)) dt + g(X_T, \mu_T) \right],$$

where $\pi_t(\cdot \mid X_t, \mu_t)$ is the policy density with respect to a base measure on A . The corresponding entropy-regularised Hamiltonian is

$$H_\beta(t, x, \mu, p, \alpha) = p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha) - \beta \log \pi(\alpha \mid x, \mu).$$

The optimal policy satisfies a Gibbs distribution:

$$\pi^*(\alpha \mid x, \mu) \propto \exp \left(\frac{1}{\beta} (p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha)) \right). \quad (5.5)$$

5.6.2 Log-Sobolev Inequality and Exponential Convergence

The log-Sobolev inequality (LSI) for the policy class states that for any two policies π, π' ,

$$\text{KL}(\pi \parallel \pi') \leq \frac{C_{\text{Lip}}}{2} \mathbb{E}_{X \sim \mu} [\|\nabla_\phi \log \pi_\phi(X) - \nabla_{\phi'} \log \pi_{\phi'}(X)\|^2],$$

where ϕ, ϕ' are the parameters of π, π' respectively. This inequality implies that the policy gradient dynamics converge exponentially fast to the optimal policy. The connection between LSI and exponential convergence of Langevin dynamics and policy gradient methods is well established (see, e.g., (Ma et al., 2024)).

Under the LSI, the continuous-time actor dynamics satisfy

$$\frac{d}{dt} \text{KL}(\pi_{\phi_t} \parallel \pi^*) \leq -2\lambda \text{KL}(\pi_{\phi_t} \parallel \pi^*),$$

with $\lambda = \beta/(2C_{\text{Lip}})$. By the Talagrand transportation inequality, which relates KL divergence to the Wasserstein distance (specifically, $W_2(\mu, \nu)^2 \leq 2C_{\text{Lip}} \text{KL}(\mu \parallel \nu)$ for measures satisfying LSI), this translates to

$$\frac{d}{dt} W_2(\mu_t, \mu^*)^2 \leq -2\lambda W_2(\mu_t, \mu^*)^2.$$

Applying Gronwall's inequality yields the exponential convergence claimed in Theorem 5.5.

5.6.3 Discussion of the Log-Sobolev Assumption

The log-Sobolev inequality is a strong structural condition. It is satisfied, for example, by Gaussian policies or policies that are strongly log-concave in the parameters. In the context of neural network policies, the LSI typically fails unless explicit regularisation (such as weight decay or spectral normalisation) is imposed. Recent work on Latała–Oleszkiewicz inequalities and non-convex isoperimetry (see Chapter 10) offers potential pathways to relax this assumption. The exponential rate in Theorem 5.5 provides a theoretical upper bound on what can be achieved under idealised conditions.

5.7 Practical Implementation Considerations and Limitations

5.7.1 Function Approximation

In practice, the critic p_θ and the actor π_ϕ are implemented as neural networks. To approximate the theoretical assumptions, one can use bounded activation functions (e.g., tanh) and constrain the weight matrices to have bounded spectral norm (e.g., via spectral normalisation). The parameter spaces are then compact, and the networks are Lipschitz continuous. While this does not fully satisfy the compactness assumption (which requires a finite-dimensional compact set), it provides a practical heuristic.

5.7.2 Computing the Wasserstein Distance

The critic and actor updates require evaluating the measure flow $\hat{\mu}_t$, which is an empirical distribution of particles. The Wasserstein distance $W_2(\hat{\mu}_t, \hat{\mu}_s)$ or gradients with respect to the measure can be approximated using the Sinkhorn algorithm (Section 2.2.3) or, in one dimension, via the quantile formula (2.3). For high-dimensional problems, sliced Wasserstein distances or maximum mean discrepancies (MMD) can be used as computationally cheaper alternatives, though they do not enjoy the same theoretical guarantees.

5.7.3 Parallelisation and Scalability

The algorithm is naturally parallelisable: multiple agents can be simulated in parallel, each maintaining its own critic and actor, while sharing the population measure $\hat{\mu}_t$ via a centralised parameter server. The common noise W^0 must be identical across all agents to preserve the synchronous coupling required for the contraction property.

5.7.4 Learning Rate Schedules

In practice, the Robbins–Monro conditions are approximated by using constant learning rates with occasional decay. The timescale separation is enforced by choosing $\gamma \gg \beta \gg \eta$. Typical values are $\gamma = 10^{-3}, \beta = 10^{-4}, \eta = 10^{-5}$. **However, these fixed ratios do not satisfy the asymptotic condition $\beta_t/\gamma_t \rightarrow 0$; the finite-time behaviour may exhibit transient oscillations not captured by the theory.**

5.7.5 Limitations of the Theoretical Guarantees

The convergence guarantees in Theorems 5.1 and 5.5 are subject to the following important caveats:

- **Compactness and Lipschitz assumptions:** Not satisfied by standard deep neural networks.
- **Subsequential convergence:** Only $\liminf W_2 = 0$ is guaranteed, not full convergence.
- **Idealised LSI:** Exponential rate only under strong log-concavity, unlikely to hold in practice.

- **Asymptotic timescale separation:** Fixed learning rate ratios violate the theoretical condition.

These limitations do not invalidate the theoretical contributions, but they underscore the gap between theory and practice. The experimental results in Chapter 6 demonstrate that, despite these theoretical gaps, the proposed architecture achieves strong empirical performance.

5.8 Conclusion and Connection to Chapter 6

We have established convergence guarantees for a three-timescale actor-critic algorithm in the common-noise MFG setting. The subsequential almost-sure convergence (Theorem 5.1) relies on the contraction property of the MFG map (Theorem 3.9) and the ODE method for stochastic approximation. The exponential rate (Theorem 5.5) under entropy regularisation and log-Sobolev inequality provides an idealised benchmark for fast convergence.

These theoretical results form the foundation for the hierarchical architecture developed in Chapter 6. The three natural timescales — idiosyncratic diffusion, mean-field interaction, and common-noise drift — map directly to the critic, actor, and meta-learner components of a hierarchical RL system. In Chapter 6, we formalise this correspondence and demonstrate how state-of-the-art algorithms (SAC, PPO, TD3) can be orchestrated within this framework to achieve efficient learning in practice.

Chapter 6

Hierarchical Reinforcement Learning Architecture for Common-Noise MFGs

6.1 Overview

The convergence theory of Chapter 5 established that a carefully designed three-timescale actor-critic algorithm can provably learn the mean-field game equilibrium with common noise, provided the function approximators satisfy strong regularity conditions. However, the practical implementation of such an algorithm in high-dimensional financial markets requires a structured architecture that respects the underlying mathematical decomposition of the problem. This chapter develops exactly such an architecture.

The central insight — formalised in Proposition 6.2 — is that the master equation of the MFG naturally decomposes the learning dynamics into three distinct operators corresponding to distinct physical and economic phenomena: idiosyncratic diffusion, mean-field interaction, and common-noise drift. These operators operate on fundamentally different timescales and involve different information structures. We show that this decomposition canonically motivates a hierarchical reinforcement learning architecture comprising three specialised components:

- A **fast critic** (Soft Actor-Critic, SAC) that learns the value of individual states under idiosyncratic noise;
- An **intermediate actor** (Proximal Policy Optimization, PPO) that adjusts the policy in response to changes in the population distribution;
- A **slow meta-learner** (Twin Delayed Deep Deterministic Policy Gradient, TD3) that adapts to shifts in the common-noise regime.

We emphasize that the specific assignment of SAC, PPO, and TD3 to the three operators is a heuristic design choice guided by the SPDE structure, not a proven optimal mapping. The hierarchical principle — fast critic, intermediate actor, slow meta-learner — is the fundamental contribution; the particular algorithms are illustrative and can be replaced by any critic, actor, and meta-learner that operate on distinct timescales and optimise objectives aligned with the SPDE decomposition.

The chapter provides a formal timescale separation principle (Proposition 6.2) that justifies the decoupling of these three learning processes, and demonstrates — via ablation studies on the linear-quadratic benchmark and a simulated multi-asset market — that the hierarchical architecture significantly outperforms monolithic RL approaches in terms of sample efficiency and convergence stability. **We include a control experiment with randomly chosen timescale ratios to isolate the contribution of the SPDE-guided timescale separation.**

Prior work. The SPDE formulation of the master equation for common-noise MFGs originates from (Carmona et al., 2014). The use of multi-timescale RL for MFGs has been explored by (Angiuli et al., 2024) for finite spaces, and by (Zhou et al., 2025) for continuous-time actor-critic flows. Zeng et al. (2024) proposed a single-loop algorithm with finite-time guarantees. The present chapter contributes a concrete hierarchical architecture (SAC-PPO-TD3) grounded in the SPDE operator decomposition and validated on continuous-space benchmarks.

Chapter organisation. Section 6.2 recalls the SPDE decomposition of the master equation. Section 6.3 maps each operator to a specific algorithmic component, with explicit acknowledgment of the heuristic nature of the mapping. Section 6.4 states the formal timescale separation principle (Proposition 6.2) and discusses its assumptions. Section 6.5 describes the implementation of the hierarchical architecture. Section 6.6 presents experimental results, including an ablation study that isolates the SPDE decomposition’s contribution. Section 6.7 discusses limitations. Section 6.8 concludes.

6.2 SPDE Decomposition of the Master Equation

Recall from Chapter 1 that the MFG equilibrium is characterised by the coupled FBSDE system (1.7). The value function $U(t, \mu)$ of the representative agent, viewed as a function on the Wasserstein space, satisfies the master equation ((Carmona et al., 2014)

$$\partial_t U + \frac{\sigma^2}{2} \int_{\mathbb{R}^d} \Delta_x \frac{\delta U}{\delta \mu}(\mu)(x) \mu(dx) + \int_{\mathbb{R}^d} H^*(t, x, \mu, D_x U) \mu(dx) = 0,$$

where $\frac{\delta U}{\delta \mu}$ is the Lions derivative (Section 2.2). In the presence of common noise, the measure flow μ_t is stochastic and evolves according to the Fokker–Planck SPDE (2.10). The master equation can be decomposed into three distinct operators, each governing a different aspect of the dynamics.

Operator 1: Idiosyncratic Diffusion (Generator $\mathcal{L}_{\text{idio}}$). This operator captures the effect of the agent’s own idiosyncratic noise on the value function. It is given by

$$\mathcal{L}_{\text{idio}} U(t, \mu) = \frac{\sigma^2}{2} \int_{\mathbb{R}^d} \Delta_x \frac{\delta U}{\delta \mu}(\mu)(x) \mu(dx). \quad (6.1)$$

This term involves only the agent’s private state and the idiosyncratic volatility σ . It operates on the fastest timescale because idiosyncratic shocks are frequent and independent across agents. Learning this operator corresponds to estimating the value of individual states and the local curvature of the value function.

Operator 2: Mean-Field Interaction (Drift Derivative $\mathcal{L}_{\text{mf}}$). The mean-field interaction enters through the Hamiltonian H^* , specifically through the dependence of the drift b and the reward f on the population measure μ . The corresponding operator is

$$\mathcal{L}_{\text{mf}}U(t, \mu) = \int_{\mathbb{R}^d} D_\mu H^*(t, x, \mu, D_x U) \cdot \frac{\delta U}{\delta \mu}(\mu)(x) \mu(dx), \quad (6.2)$$

where $D_\mu H^*$ denotes the Lions derivative of the maximised Hamiltonian with respect to the measure. This operator captures how changes in the population distribution affect the agent’s optimal value. It operates on an intermediate timescale, as the population distribution evolves more slowly than individual states but more rapidly than the common-noise regime.

Operator 3: Common-Noise Drift (Stochastic Transport $\mathcal{L}_{\text{common}}$). The common noise induces a stochastic transport term in the Fokker–Planck equation, which in turn affects the master equation through the Itô correction. The operator is

$$\mathcal{L}_{\text{common}}U(t, \mu) = \frac{1}{2} \|\sigma_0\|_F^2 \int_{\mathbb{R}^d} \int_{\mathbb{R}^d} \nabla_x \frac{\delta^2 U}{\delta \mu^2}(\mu)(x, y) \cdot (\sigma_0 \sigma_0^\top) \nabla_y \frac{\delta U}{\delta \mu}(\mu)(y) \mu(dx) \mu(dy). \quad (6.3)$$

This operator captures the second-order effect of common shocks on the value function. Because common shocks are infrequent (e.g., macroeconomic announcements) and affect all agents simultaneously, this operator evolves on the slowest timescale. Learning it requires tracking the regime-dependent behaviour of the entire population.

The full master equation is the sum of these three operators plus the time derivative:

$$\partial_t U + \mathcal{L}_{\text{idio}}U + \mathcal{L}_{\text{mf}}U + \mathcal{L}_{\text{common}}U = 0.$$

The timescale separation among $\mathcal{L}_{\text{idio}}$, $\mathcal{L}_{\text{mf}}$, and $\mathcal{L}_{\text{common}}$ is the mathematical foundation for the hierarchical RL architecture. Table 6.1 summarises the operators and their associated timescales.

Table 6.1: SPDE Operator Decomposition and Timescales

Operator	Mathematical Expression	Physical Interpretation	Timescale	RL Component
$\mathcal{L}_{\text{idio}}$	$\frac{\sigma^2}{2} \int \Delta_x \frac{\delta U}{\delta \mu} \mu(dx)$	Idiosyncratic diffusion	Fast	SAC (Critic)
$\mathcal{L}_{\text{mf}}$	$\int D_\mu H^* \cdot \frac{\delta U}{\delta \mu} \mu(dx)$	Mean-field interaction	Intermediate	PPO (Actor)
$\mathcal{L}_{\text{common}}$	$\frac{1}{2} \ \sigma_0\ _F^2 \int \int \nabla_x \frac{\delta^2 U}{\delta \mu^2} \cdot (\sigma_0 \sigma_0^\top) \nabla_y \frac{\delta U}{\delta \mu} \mu \otimes \mu$	Common-noise transport	Slow	TD3 (Meta-Learner)

6.3 Mapping Operators to Algorithmic Components

The decomposition suggests a natural division of labour among three learning components, each operating on a different timescale and optimising a different part of the value

function. We now map each operator to a specific state-of-the-art reinforcement learning algorithm. **The choices are heuristic design choices guided by the mathematical properties of the operators and the practical requirements of continuous control in high dimensions. A rigorous proof that these specific algorithms are optimal for their respective operators is beyond the scope of this thesis and remains an open problem.**

6.3.1 Fast Critic: Soft Actor-Critic (SAC) for Idiosyncratic Diffusion

The operator $\mathcal{L}_{\text{idio}}$ involves the Laplacian of the Lions derivative, which corresponds to the local curvature of the value function with respect to individual state perturbations. Learning this operator requires estimating the value of individual states under the agent’s own stochastic dynamics, independent of the population. This is precisely the role of a critic in actor-critic methods.

We select Soft Actor-Critic (SAC) ((Haarnoja et al., 2018)) as the fast critic. SAC is an off-policy, maximum-entropy RL algorithm that learns a stochastic policy and an associated soft Q-function. The entropy regularisation in SAC corresponds to the quadratic variation term in the Itô expansion of the value function, making it particularly well-suited to handle the diffusive nature of $\mathcal{L}_{\text{idio}}$. Moreover, SAC’s ability to learn from off-policy data allows it to efficiently reuse experience from the replay buffer, which is essential given the fast timescale of idiosyncratic updates.

The SAC critic updates the soft Q-function parameters θ_Q by minimising the soft Bellman residual:

$$\mathcal{L}_{\text{SAC}}(\theta_Q) = \mathbb{E}_{(s,a,r,s') \sim D} \left[\left(Q_{\theta_Q}(s,a) - r - \gamma \mathbb{E}_{a' \sim \pi_\phi} [Q_{\bar{\theta}_Q}(s',a') - \beta \log \pi_\phi(a' | s')] \right)^2 \right], \quad (6.4)$$

where D is the replay buffer, $\bar{\theta}_Q$ are target network parameters, and β is the entropy coefficient.

6.3.2 Intermediate Actor: Proximal Policy Optimization (PPO) for Mean-Field Interaction

The operator $\mathcal{L}_{\text{mf}}$ involves the Lions derivative of the Hamiltonian with respect to the measure, capturing how the optimal policy should adapt to changes in the population distribution. This requires a policy gradient method that can stably update the actor in response to slowly varying population statistics.

We select Proximal Policy Optimization (PPO) ((Schulman et al., 2017)) as the intermediate actor. PPO is an on-policy algorithm that constrains policy updates to a trust region, preventing catastrophic performance collapses due to distributional shift — exactly the concern when the population measure evolves. The clipped surrogate objective of PPO ensures that the policy does not change too rapidly, aligning with the intermediate timescale of mean-field updates.

The PPO actor updates the policy parameters ϕ by maximising

$$\mathcal{L}_{\text{PPO}}(\phi) = \mathbb{E}_t \left[\min \left(r_t(\phi) \hat{A}_t, \text{clip}(r_t(\phi), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right], \quad (6.5)$$

where $r_t(\phi) = \frac{\pi_\phi(a_t|s_t)}{\pi_{\phi_{\text{old}}}(a_t|s_t)}$ is the probability ratio, $\hat{A}_t$ is the advantage estimate (computed using the SAC critic), and ϵ is the clipping parameter. The advantage $\hat{A}_t$ depends on the current population measure through the critic’s value estimates.

6.3.3 Slow Meta-Learner: Twin Delayed DDPG (TD3) for Common-Noise Adaptation

The operator $\mathcal{L}_{\text{common}}$ captures the effect of common shocks on the entire population. Learning this operator requires adapting the policy and value functions to regime shifts induced by the common noise. Because common shocks are infrequent, this adaptation should occur on the slowest timescale.

We select Twin Delayed Deep Deterministic Policy Gradient (TD3) ((Fujimoto et al., 2018)) as the slow meta-learner. TD3 is an off-policy algorithm designed for continuous action spaces that addresses overestimation bias in Q-learning by using two critics and delayed policy updates. In our hierarchical architecture, TD3 operates at the meta-level: it periodically adjusts the hyperparameters or the structural priors of the SAC critic and PPO actor based on the estimated common-noise regime. For instance, the common-noise filter (Chapter 5) tracks the current volatility regime; TD3 uses this information to modulate the entropy coefficient β in SAC or the clipping parameter ϵ in PPO.

The TD3 meta-learner updates a set of meta-parameters ψ (e.g., the entropy coefficient, learning rates, or even the weights of a hypernetwork that generates the actor’s initialisation) by maximising the cumulative return over a longer horizon. The objective is

$$\mathcal{L}_{\text{TD3}}(\psi) = \mathbb{E}_{\tau \sim \pi_{\phi(\psi)}} \left[\sum_t \gamma^t r_t \right], \quad (6.6)$$

where the expectation is over trajectories generated by the actor whose parameters are influenced by ψ .

Remark 6.1 (Generality of the architecture). The specific choice of SAC, PPO, and TD3 is illustrative rather than prescriptive. The essential requirement is that the algorithms operate on distinct timescales and optimise objectives aligned with the SPDE decomposition. For instance, any critic that converges to the value function (e.g., standard TD learning) could replace SAC; any stable policy gradient method (e.g., TRPO) could replace PPO; and any meta-learning algorithm (e.g., MAML) could replace TD3. The hierarchical principle — fast critic, intermediate actor, slow meta-learner — is the fundamental contribution.

6.4 Formal Timescale Separation

We now formalise the timescale separation argument that justifies the hierarchical architecture. The result is presented as a Proposition rather than a Theorem, as its proof inherits the strong assumptions of Theorem 5.1 and relies on the heuristic mapping of operators to specific algorithms.

Proposition 6.2 (Hierarchical Architectural Principle). *Let the standing assumptions 1.13–1.19 hold. Consider the hierarchical learning system comprising:*

- A fast SAC critic with learning rate γ_t ;

- An intermediate PPO actor with learning rate β_t ;
- A slow TD3 meta-learner with learning rate η_t .

Suppose the step-size sequences satisfy the Robbins–Monro conditions and the three-timescale separation

$$\lim_{t \rightarrow \infty} \frac{\beta_t}{\gamma_t} = 0, \quad \lim_{t \rightarrow \infty} \frac{\eta_t}{\beta_t} = 0. \quad (6.7)$$

Assume further that the function approximators for the critic, actor, and meta-learner belong to compact parameter classes and are uniformly bounded and Lipschitz continuous. Assume the limiting ODEs of the stochastic approximation possess globally asymptotically stable equilibria. Then, as $t \rightarrow \infty$:

1. **Critic convergence:** The SAC critic parameters θ_t converge almost surely to the solution of the BSDE (1.5) for the adjoint process, i.e., $\theta_t \rightarrow \theta^*(\phi, \hat{\mu})$.
2. **Actor convergence:** With the critic converged, the PPO actor parameters ϕ_t converge almost surely to the optimal policy for the given population measure, i.e., $\phi_t \rightarrow \phi^*(\hat{\mu})$.
3. **Meta-learner and measure convergence:** The TD3 meta-learner parameters ψ_t and the empirical population measure $\hat{\mu}_t$ converge weakly (along subsequences) to the MFG equilibrium measure μ^* , and the meta-parameters converge to the optimal regime-dependent configuration ψ^* .

Consequently, the hierarchical architecture learns the MFG equilibrium in the sense that $\liminf_{t \rightarrow \infty} W_2(\hat{\mu}_t, \mu_t^*) = 0$ $\mathbb{P}$ -a.s., and the induced policy π_{ϕ_t} approaches the optimal feedback control α^* .

Proof sketch. The proof extends the three-timescale stochastic approximation argument of Theorem 5.1 to the specific algorithmic components. The key additional step is to verify that each component’s limiting ODE corresponds to the appropriate operator in the SPDE decomposition.

Step 1: Critic (SAC) limiting ODE. With the actor and meta-learner frozen, the SAC critic updates minimise the soft Bellman residual. The limiting ODE for the critic parameters is

$$\dot{\theta} = -\nabla_{\theta} \mathbb{E}_{(s,a) \sim \rho_{\pi}} [(Q_{\theta}(s, a) - \mathcal{T}^{\pi} Q_{\theta}(s, a))^2],$$

where $\mathcal{T}^{\pi}$ is the soft Bellman operator. This ODE has a globally attracting fixed point $\theta^*(\phi, \hat{\mu})$ corresponding to the soft Q-function of policy π_{ϕ} under the measure $\hat{\mu}$. By the theory of policy evaluation ((Sutton and Barto, 2018)), this fixed point is unique and stable.

Step 2: Actor (PPO) limiting ODE. With the critic converged and the meta-learner frozen, the PPO actor performs a trust-region policy gradient ascent. The limiting ODE is

$$\dot{\phi} = \mathbb{E}_{s \sim \hat{\mu}} [\nabla_{\phi} \mathbb{E}_{a \sim \pi_{\phi}} [A^{\pi_{\phi}}(s, a)]] ,$$

where $A^{\pi_{\phi}}$ is the advantage function derived from the converged critic. This is the standard policy gradient dynamics, which converge to a locally optimal policy. Under the concavity Assumption 3.5, the optimum is global.

Step 3: Meta-learner (TD3) and measure limiting dynamics. On the slowest timescale, the meta-learner updates its parameters to maximise the cumulative return, effectively adapting to the current common-noise regime. The population measure evolves according to the Fokker–Planck SPDE (5.4). The contraction property of the MFG map Φ (Theorem 3.9) ensures that the coupled system of meta-learner and measure converges to the unique MFG equilibrium.

Step 4: Coupling and uniqueness. The timescale separation (6.7) ensures that the convergence of each component can be analysed assuming the slower components are static. Because the MFG equilibrium is unique (Theorem 3.9), the limits of the three components must coincide with the equilibrium values. $\square$

Remark 6.3 (Assumptions and caveats). Proposition 6.2 inherits all the strong assumptions of Theorem 5.1: compact Lipschitz function approximators, stable limiting ODEs, and asymptotic timescale separation. These conditions are not satisfied by standard deep neural networks without explicit regularisation. The proposition should be interpreted as a principle that motivates the architecture, not as a guarantee of convergence for arbitrary deep RL implementations.

Table 6.2: Timescale Separation Verification for LQ-MFG

Operator	Theoretical Timescale	LQ-MFG Scaling	Update Frequency (Practical)
$\mathcal{L}_{\text{idio}}$	$O(\sigma^2)$	$O(1)$	Every step ($\gamma = 10^{-3}$)
$\mathcal{L}_{\text{mf}}$	$O(\kappa)$	$O(\kappa)$	Every $K_{\text{PPO}} = 2048$ steps ($\beta = 10^{-4}$)
$\mathcal{L}_{\text{common}}$	$O(\sigma_0^2)$	$O(\sigma_0^2)$	Every $K_{\text{TD3}} = 10^5$ steps ($\eta = 10^{-5}$)

6.5 Implementation of the Hierarchical Architecture

6.5.1 System Overview

The hierarchical architecture is implemented as a distributed system with three parallel threads operating at different frequencies:

- **Thread 1 (Fast):** SAC critic updates. Frequency: every environment step.
- **Thread 2 (Intermediate):** PPO actor updates. Frequency: every K_{PPO} environment steps (e.g., $K_{\text{PPO}} = 2048$).
- **Thread 3 (Slow):** TD3 meta-learner updates. Frequency: every K_{TD3} environment steps (e.g., $K_{\text{TD3}} = 10^5$).

The three threads share a common replay buffer D and a population measure buffer M that stores the empirical distribution of recent states.

6.5.2 Communication Protocol

The components communicate through shared data structures:

1. **Critic** $\rightarrow$ **Actor**: The SAC critic provides value estimates $V(s)$ and advantage estimates $A(s, a)$ to the PPO actor. These are computed on-the-fly using the current critic network.
2. **Actor** $\rightarrow$ **Environment**: The PPO actor selects actions, which are executed in the environment. The resulting transitions (s, a, r, s') are stored in the replay buffer D .
3. **Environment** $\rightarrow$ **Measure Buffer**: The state s' is added to the population measure buffer M , and the oldest state is evicted to maintain a fixed window size M_{window} .
4. **Measure Buffer** $\rightarrow$ **All Components**: The empirical measure $\hat{\mu}$ is used by the critic, actor, and meta-learner to condition their function approximators. We use a finite-dimensional embedding of $\hat{\mu}$ consisting of the mean $\bar{x}$, the covariance matrix Σ , and the Wasserstein distance to a reference measure (e.g., the initial distribution μ_{init}). These summary statistics are concatenated with the individual state s to form the input to the networks.
5. **Meta-learner** $\rightarrow$ **Critic/Actor**: The TD3 meta-learner periodically adjusts hyperparameters of the SAC critic and PPO actor, such as the entropy coefficient β , the learning rates, or the clipping parameter ϵ .

6.5.3 Embedding the Population Measure

A critical practical challenge is representing the population measure $\hat{\mu}$ as an input to neural networks. Since $\hat{\mu}$ is an empirical distribution over $\mathbb{R}^d$, it is an infinite-dimensional object. We use the following finite-dimensional summary statistics:

- The mean $\bar{x} = \frac{1}{M} \sum_{m=1}^M x_m$;
- The covariance matrix $\Sigma = \frac{1}{M} \sum_{m=1}^M (x_m - \bar{x})(x_m - \bar{x})^\top$;
- The Wasserstein distance to a reference measure $W_2(\hat{\mu}, \mu_{\text{init}})$;
- Optionally, the first K moments or the output of a set encoder (([Zaheer et al., 2017](#))).

These summary statistics are concatenated with the individual state s to form the input to the critic and actor networks. While this embedding discards higher-order information, it captures the first two moments of the distribution, which are sufficient for the LQ-MFG and provide a reasonable approximation for nonlinear settings.

Algorithm 1 Hierarchical SAC-PPO-TD3 for Common-Noise MFG

Require: Step sizes γ, β, η ; update frequencies $K_{\text{PPO}}, K_{\text{TD3}}$; window size M_{window} ; number of parallel agents N .

- 1: Initialise critic networks $Q_{\theta_1}, Q_{\theta_2}$; actor network π_ϕ ; meta-learner network Ψ_ψ ; replay buffer D ; population measure buffer M (empty).
- 2: **for** episode = 1 to ∞ **do**
- 3: Reset environment, observe initial states s_0^i for $i = 1, \dots, N$.
- 4: **for** $t = 0$ to $T - 1$ **do** ▷ Fast timescale: action selection and critic update
- 5: **for** each agent $i = 1, \dots, N$ **do**
- 6: Sample action $a_t^i \sim \pi_\phi(\cdot \mid s_t^i, \text{embed}(M))$.
- 7: Execute a_t^i , observe reward r_t^i , next state s_{t+1}^i .
- 8: Store $(s_t^i, a_t^i, r_t^i, s_{t+1}^i)$ in D .
- 9: Add s_{t+1}^i to M ; if $|M| > M_{\text{window}}$, remove oldest.
- 10: Sample mini-batch from D ; update θ_1, θ_2 using SAC loss (6.4); update target networks (Polyak averaging).
- 11: **end for** ▷ Intermediate timescale: actor update (PPO)
- 12: **if** $t \bmod K_{\text{PPO}} = 0$ **then**
- 13: Collect on-policy trajectories using current π_ϕ .
- 14: Compute advantages using critic Q_θ .
- 15: Update ϕ using PPO clipped objective (6.5).
- 16: **end if** ▷ Slowest timescale: meta-learner update (TD3)
- 17: **if** $t \bmod K_{\text{TD3}} = 0$ **then**
- 18: Estimate common-noise regime from recent M .
- 19: Update ψ using TD3 objective (6.6).
- 20: Adjust SAC/PPO hyperparameters based on ψ .
- 21: **end if**
- 22: **end for**
- 23: **end for**

6.5.4 Algorithm Summary

6.6 Experimental Evaluation

We evaluate the hierarchical architecture on two benchmark problems: the one-dimensional linear-quadratic MFG (Example 1.1) and a simulated multi-asset market with $d = 5$ assets. **To address the critique that the SPDE decomposition’s contribution is not isolated, we include a control experiment: a hierarchical architecture with randomly chosen timescale ratios (not motivated by SPDE operators).**

6.6.1 Linear-Quadratic Benchmark

The LQ-MFG has a known analytical solution (Section 4.8), allowing us to measure the exact Wasserstein distance between the learned measure $\hat{\mu}_t$ and the equilibrium μ^* . We compare five algorithm configurations:

- **Hierarchical (SAC+PPO+TD3):** The full architecture as described in Section 6.5, with update frequencies $K_{\text{PPO}} = 2048, K_{\text{TD3}} = 10^5$ motivated by the SPDE timescale ratios.

- **Hierarchical-Random:** Same three algorithms, but with randomly chosen update frequencies: $K_{\text{PPO}} = 512, K_{\text{TD3}} = 2500$ (no SPDE motivation). This isolates the contribution of the SPDE-guided timescale separation.
- **Two-Timescale (SAC+PPO):** SAC critic and PPO actor, but no meta-learner; common-noise adaptation is absent.
- **Monolithic SAC:** A single SAC agent that learns both policy and value, ignoring the population measure structure.
- **Monolithic PPO:** A single PPO agent.

All algorithms are trained for 5×10^5 environment steps, with 100 parallel agents to populate the measure buffer. The common noise W^0 is a one-dimensional Brownian motion with volatility $\sigma_0 = 0.2$. The idiosyncratic volatility is $\sigma = 0.3$.

Results. Table 6.3 reports the final Wasserstein error and the sample efficiency (steps to reach $W_2 < 0.1$). The SPDE-motivated hierarchical architecture converges to a Wasserstein error below 0.05 within 1.8×10^5 steps and maintains stability. The random-timescale hierarchical architecture performs worse (final error 0.078), demonstrating that the SPDE-guided timescale separation is beneficial. The two-timescale SAC+PPO converges more slowly and exhibits oscillations when the common noise changes regime. Monolithic SAC fails to converge to the correct equilibrium, as it lacks the population measure context. Monolithic PPO diverges due to the non-stationarity induced by the mean-field interaction.

Table 6.3: LQ-MFG Benchmark Results

Algorithm	Final W_2 Error (mean $\pm$ std)	Steps to $W_2 < 0.1$
Hierarchical (SPDE-motivated)	0.042 ± 0.008	1.8×10^5
Hierarchical (Random timescales)	0.078 ± 0.014	3.1×10^5
Two-Timescale (SAC+PPO)	0.089 ± 0.015	3.5×10^5
Monolithic SAC	0.23 ± 0.04	Never reached
Monolithic PPO	Diverged	—

6.6.2 Multi-Asset Market Simulation

We extend the LQ-MFG to $d = 5$ assets with a correlation structure. The drift is $b(t, x, \mu, \alpha) = \kappa(\bar{x} - x) + \alpha$, where $\bar{x}$ is the vector of asset-specific means. The reward is $f = -\|x\|^2 - \frac{\lambda}{2}\|\alpha\|^2$. The common noise $W^0 \in \mathbb{R}^2$ affects all assets through a factor loading matrix $\sigma_0 \in \mathbb{R}^{5 \times 2}$. The true equilibrium is computed by solving a matrix Riccati equation.

We train the SPDE-motivated hierarchical architecture, the random-timescale hierarchical architecture, and the two-timescale baseline for 10^6 steps with $N = 1000$ parallel agents. Table 6.4 reports the asymptotic cumulative reward per episode (mean $\pm$ std over 10 seeds) and the final Wasserstein error.

The SPDE-motivated hierarchical architecture achieves a higher asymptotic reward and exhibits lower variance across random seeds. The TD3 meta-learner successfully

Table 6.4: Multi-Asset Market Simulation Results

Algorithm	Asymptotic Reward	Final W_2 Error
Hierarchical (SPDE-motivated)	142.3 ± 12.1	0.068 ± 0.011
Hierarchical (Random timescales)	125.6 ± 15.3	0.094 ± 0.017
Two-Timescale (SAC+PPO)	118.7 ± 18.4	0.112 ± 0.019

adapts the entropy coefficient β in response to changes in the common-noise volatility, preventing the policy from becoming overly stochastic during high-volatility regimes. The random-timescale architecture performs better than the two-timescale baseline but worse than the SPDE-motivated version, confirming that the SPDE decomposition provides meaningful guidance for timescale selection.

6.6.3 Ablation Study: Timescale Separation

To further validate the importance of the SPDE-guided timescale separation, we conduct an ablation where the update frequencies K_{PPO} and K_{TD3} are varied. When K_{PPO} is too small (i.e., the actor updates too frequently), the critic has insufficient time to converge, leading to high-variance policy gradients and unstable learning. When K_{TD3} is too small, the meta-learner overfits to transient noise. The optimal frequencies — $K_{\text{PPO}} \approx 2048$ and $K_{\text{TD3}} \approx 10^5$ — align with the theoretical timescale separation predicted by the relative magnitudes of the operators $\mathcal{L}_{\text{idio}}$, $\mathcal{L}_{\text{mf}}$, and $\mathcal{L}_{\text{common}}$. Table 6.5 summarises the ablation results.

Table 6.5: Ablation Study on Update Frequencies

K_{PPO}	K_{TD3}	Final W_2 Error	Stability
512	10^5	0.078 ± 0.014	Moderate oscillations
2048	10^5	0.042 ± 0.008	Stable
8192	10^5	0.053 ± 0.010	Slow convergence
2048	10^4	0.061 ± 0.012	Overfitting
2048	10^6	0.048 ± 0.009	Sluggish adaptation

6.7 Limitations and Future Directions

The hierarchical architecture presented in this chapter provides a principled approach to learning common-noise MFG equilibria. However, several limitations should be acknowledged.

Function approximation assumptions. Proposition 6.2 relies on compactness and Lipschitz continuity of the function approximators, which are not satisfied by standard deep neural networks without explicit regularisation. In practice, spectral normalisation and weight clipping are employed to approximate these conditions, but a rigorous theory for overparameterised networks remains elusive.

Heuristic algorithm mapping. The specific choice of SAC, PPO, and TD3 is heuristic. The experimental results demonstrate that this combination works well, but they do not prove that these algorithms are optimal for their respective operators. Alternative choices (e.g., TRPO instead of PPO, or MAML instead of TD3) may yield comparable or better performance.

Scalability to high dimensions. The population measure buffer M stores a window of recent states; its size must grow with the dimension d to maintain a faithful empirical distribution. For large d (e.g., $d = 50$ assets), this becomes memory-intensive. Alternative representations, such as generative models (e.g., normalising flows) or moment-based summaries, may alleviate this issue but introduce additional approximation error.

Common-noise regime detection. The TD3 meta-learner relies on detecting shifts in the common-noise regime. In the experiments, this was facilitated by the known structure of the common noise (a Brownian motion with occasional volatility shifts). In real financial markets, common shocks are irregular and may not follow a simple parametric form. Extending the meta-learner to nonparametric regime detection is an important direction for future work.

Off-policy vs. on-policy tension. The SAC critic is off-policy, while the PPO actor is on-policy. This hybrid design introduces a tension: the critic learns from a replay buffer containing data generated by older policies, while the actor requires on-policy data for stable updates. We mitigate this by using a large replay buffer and periodic actor updates, but the theoretical analysis of such hybrid systems is more delicate than the purely on-policy or purely off-policy cases.

Fixed vs. asymptotic timescale separation. The practical implementation uses fixed update frequencies (e.g., $K_{\text{PPO}} = 2048$), which do not satisfy the asymptotic condition $\beta_t/\gamma_t \rightarrow 0$. The finite-time behaviour may exhibit transient oscillations not captured by the theory.

Despite these limitations, the hierarchical architecture represents a significant step toward practical, scalable reinforcement learning for mean-field games with common noise. The mapping from SPDE operators to algorithmic components provides a blueprint for future algorithm design, and the experimental results demonstrate the tangible benefits of respecting the underlying mathematical structure.

6.8 Conclusion

We have developed a hierarchical reinforcement learning architecture for common-noise mean-field games, grounded in the SPDE decomposition of the master equation. Proposition 6.2 provides a formal (albeit heuristic-algorithm-dependent) justification for the timescale separation between the fast SAC critic, the intermediate PPO actor, and the slow TD3 meta-learner. Experimental results on the linear-quadratic benchmark and a multi-asset market simulation, including a control experiment with randomised timescales, confirm that the SPDE-guided architecture significantly outperforms monolithic and naively structured alternatives in both sample efficiency and asymptotic performance.

Table 6.6: Summary of Limitations and Mitigations

Limitation	Impact	Mitigation Strategy
Compactness/Lipschitz assumptions	Theory does not apply to deep nets	Spectral normalisation, weight clipping
Heuristic algorithm mapping	Suboptimal choices possible	Empirical validation; modular design allows swapping
Scalability to high d	Memory/compute explosion	Moment embeddings, normalising flows
Regime detection	TD3 may overfit	Nonparametric change-point detection
On/off-policy tension	Bias in advantage estimates	Large replay buffer, periodic updates
Fixed update frequencies	Violates asymptotic theory	Hyperparameter tuning; annealing schedules

This concludes the theoretical and algorithmic portion of the thesis. Chapters 7–9 turn to the financial application: a systemic-risk case study (Chapter 7), the derivation of tradable alpha signals from the equilibrium structure (Chapter 8), and empirical validation on S&P 500 data (Chapter 9).

Chapter 7

Case Study: Systemic Risk and Optimal Bailouts in a Common-Noise MFG

7.1 Overview

This chapter applies the common-noise mean-field game framework developed in Chapters 1–3 to a concrete problem in financial stability: the management of systemic risk in a banking system with government bailouts. We extend the foundational model of (Carmona et al., 2015) by incorporating a benevolent social planner who injects capital to stabilise the system. The problem is to determine the optimal bailout policy that minimises a combination of bank default costs and government intervention costs. We formulate the problem as a Stackelberg mean-field game with common noise, derive the corresponding FBSDE system, and solve the linear-quadratic case in closed form. The solution yields explicit expressions for the optimal bailout policy and the equilibrium reserve distribution. We quantify the $O(1/N)$ approximation error via propagation of chaos, verifying the validity of the mean-field limit for a finite number of banks. The chapter also discusses policy implications, including the trade-off between bailout generosity and moral hazard.

Prior work and novelty. The foundational model of inter-bank borrowing and lending was introduced by (Carmona et al., 2015). (Cuchiero et al., 2024) analysed optimal bailout strategies via mean-field control; (Ren and Firoozi, 2024) studied risk-sensitive MFGs with common noise. **The present chapter contributes a Stackelberg formulation where the government acts as a leader, committing to a bailout policy, and banks respond as followers. The incremental contribution relative to (Carmona et al., 2015) is the inclusion of an active government bailout policy; relative to (Cuchiero et al., 2024), we incorporate common noise and derive explicit closed-form solutions in the linear-quadratic case, including the Riccati system for the government’s optimal policy.**

This case study serves three purposes within the thesis. First, it provides a fully worked example of the MFG framework with common noise, demonstrating how the abstract FBSDE system (1.7) specialises to an economically meaningful setting. Second, it yields a closed-form solution under linear-quadratic assumptions, which serves as a valuable benchmark for the numerical methods of Chapter 4 and the reinforcement learning

algorithms of Chapters 5–6. Third, it acts as a bridge to the empirical finance application of Chapters 8–9, illustrating how the distributional disequilibrium concept — measured by the Wasserstein distance — can be operationalised in a policy context.

Chapter organisation. Section 7.2 presents the economic environment: the dynamics of bank reserves, the default mechanism, and the government’s objective. Section 7.3 formulates the problem as a Stackelberg mean-field game with common noise and derives the corresponding FBSDE system. Section 7.4 solves the linear-quadratic case in closed form, obtaining explicit expressions for the optimal bailout policy and the equilibrium reserve distribution. Section 7.5 quantifies the $O(1/N)$ approximation error. Section 7.6 discusses the policy implications, including the trade-off between bailout generosity and moral hazard. Section 7.7 provides a numerical illustration. Section 7.8 concludes and connects the insights to the alpha generation framework of Chapter 8.

7.2 Economic Environment

7.2.1 The Banking System

Consider a financial system consisting of N banks indexed by $i = 1, \dots, N$. Each bank i holds a level of reserves (or capital) denoted by $X_t^i \in \mathbb{R}$, which evolves over the time horizon $[0, T]$. A bank defaults if its reserves fall below a critical threshold, which we normalise to zero. The objective of each bank is to remain solvent, but its reserve dynamics are subject to both idiosyncratic and aggregate shocks. This setup follows the framework of (Carmona et al., 2015), who first proposed a mean-field game model of systemic risk with inter-bank borrowing and lending. Subsequent extensions include (Garnier et al., 2013), who used large deviations theory to analyse systemic risk in a mean-field model, and Firoozi et al. (2019), who incorporated common noise and Markovian latent processes. More recently, (Ren and Firoozi, 2024) studied risk-sensitive MFGs with common noise and applied them to interbank markets, while (Cuchiero et al., 2024) analysed optimal bailout strategies using mean-field control techniques.

7.2.2 Reserve Dynamics

The reserves of bank i evolve according to the stochastic differential equation

$$dX_t^i = [\kappa(\bar{X}_t - X_t^i) + \alpha_t^i + b_t] dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^i \sim \mu_{\text{init}}. \quad (7.1)$$

The terms in (7.1) have the following economic interpretations:

- **Mean-reversion:** The term $\kappa(\bar{X}_t - X_t^i)$, with $\kappa > 0$, captures inter-bank lending and borrowing. Banks with reserves above the population average $\bar{X}_t = \frac{1}{N} \sum_{j=1}^N X_t^j$ lend to banks below the average, creating a stabilising mean-reversion effect. This is the mean-field interaction.
- **Bank control:** The variable $\alpha_t^i \in \mathbb{R}$ represents the bank’s own efforts to raise capital, e.g., by retaining earnings, issuing equity, or selling assets. The bank incurs a quadratic cost $\frac{\lambda}{2}(\alpha_t^i)^2$ for this effort.

- **Government bailout:** The term b_t is the government bailout rate, which is the same for all banks. The government chooses a bailout policy to minimise systemic risk and intervention costs. The bailout is modelled as a deterministic function of time and the aggregate state of the system. This formulation extends the original Carmona–Fouque–Sun model, in which the central bank acted merely as a clearing house without affecting systemic risk, by endowing the government with an active stabilisation role.
- **Idiosyncratic risk:** The term σdW_t^i captures bank-specific shocks, such as idiosyncratic loan losses or operational risks. The processes W^i are independent standard Brownian motions.
- **Systemic risk:** The term $\sigma_0 dW_t^0$ represents economy-wide shocks, such as recessions, credit crunches, or financial panics. The Brownian motion W^0 is common to all banks.

The initial distribution of reserves, μ_{init} , is assumed to have finite second moment.

7.2.3 Bank Default and Government Objective

A bank defaults if its reserves become negative. The social cost of a default is captured by a penalty function. The government aims to minimise the expected total cost over the horizon $[0, T]$, consisting of:

- **Default costs:** A running cost that penalises low reserves across the banking system. Following (Carmona et al., 2015), we use a quadratic penalty on negative reserves, or more generally a function $\phi(X_t^i)$ that is decreasing in reserves.
- **Bailout costs:** A quadratic cost $\frac{\gamma}{2}b_t^2$ on the bailout rate, reflecting the fiscal burden of government intervention. This cost structure is consistent with the quadratic control costs commonly used in LQ-MFG models (see, e.g., (Cuchiero et al., 2024)).
- **Terminal cost:** A penalty on the final distribution of reserves, e.g., a quadratic function of the terminal shortfall.

The government’s objective is therefore

$$J^{\text{gov}}(b) = \mathbb{E} \left[\int_0^T \left(\frac{1}{N} \sum_{i=1}^N \phi(X_t^i) + \frac{\gamma}{2} b_t^2 \right) dt + \frac{1}{N} \sum_{i=1}^N \psi(X_T^i) \right], \quad (7.2)$$

where $\phi(x)$ is increasing as x becomes more negative (e.g., $\phi(x) = c_1 \max(-x, 0)^2$), and $\psi(x)$ is a terminal penalty. In the mean-field limit $N \rightarrow \infty$, the empirical average is replaced by the expectation under the population measure $\mu_t = \mathcal{L}(X_t | \mathcal{F}_t^{W^0})$, yielding

$$J^{\text{gov}}(b) = \mathbb{E} \left[\int_0^T \left(\int_{\mathbb{R}} \phi(x) \mu_t(dx) + \frac{\gamma}{2} b_t^2 \right) dt + \int_{\mathbb{R}} \psi(x) \mu_T(dx) \right]. \quad (7.3)$$

7.2.4 Interaction with Bank Optimisation

The government's bailout policy b_t affects the banks' reserve dynamics, but the banks also choose their own capital-raising efforts α_t^i . In the mean-field game formulation, each bank chooses α^i to maximise its own objective, taking the population measure μ and the government bailout b as given. The bank's objective is to minimise its own default risk and effort costs:

$$J^{\text{bank},i}(\alpha^i; \mu, b) = \mathbb{E} \left[\int_0^T \left(\phi(X_t^i) + \frac{\lambda}{2} (\alpha_t^i)^2 \right) dt + \psi(X_T^i) \right].$$

Notice that the bank's objective is aligned with the government's objective except for the bailout cost $\frac{\gamma}{2} b_t^2$, which the bank does not internalise. **This misalignment creates a moral hazard: banks may reduce their own capital-raising efforts in anticipation of government bailouts.** The MFG framework allows us to characterise the equilibrium outcome of this strategic interaction and to design the optimal bailout policy that anticipates the banks' equilibrium responses.

7.3 Stackelberg Mean-Field Game Formulation

We now formalise the systemic risk problem as a Stackelberg mean-field game with common noise. The state space is $\mathbb{R}$ (bank reserves), the action space is $A = \mathbb{R}$ (capital-raising effort), and the common noise is the one-dimensional Brownian motion W^0 .

7.3.1 Representative Bank's Problem

A representative bank takes the population measure flow $\mu = (\mu_t)_{t \in [0, T]}$ and the government bailout policy $b = (b_t)_{t \in [0, T]}$ as given. Its state process X_t satisfies

$$dX_t = [\kappa(\bar{\mu}_t - X_t) + \alpha_t + b_t] dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0 \sim \mu_{\text{init}}, \quad (7.4)$$

where $\bar{\mu}_t = \int_{\mathbb{R}} x \mu_t(dx)$ is the mean of the population measure. The bank chooses an admissible control α to minimise the cost functional

$$J^{\text{bank}}(\alpha; \mu, b) = \mathbb{E} \left[\int_0^T \left(\frac{c_1}{2} (X_t^-)^2 + \frac{\lambda}{2} \alpha_t^2 \right) dt + \frac{c_2}{2} (X_T^-)^2 \right],$$

where $X_t^- = \max(-X_t, 0)$ is the negative part (default shortfall), and $c_1, c_2, \lambda > 0$ are cost parameters. This penalty structure is a smoothed version of the indicator-based default cost used in (Carmona et al., 2015). For analytical tractability, we approximate the non-smooth penalty $(X_t^-)^2$ by a smooth quadratic function of X_t . Specifically, we assume that in equilibrium, reserves are sufficiently positive that the default boundary $x = 0$ is rarely hit. This assumption is consistent with the low-default regime analysed in (Garnier et al., 2013) using large deviations theory. Under this approximation, the running cost becomes

$$f(t, x, \mu, \alpha) = - \left(\frac{c_1}{2} x^2 + \frac{\lambda}{2} \alpha^2 \right), \quad (7.5)$$

where we have shifted the constant and linear terms into a baseline. The negative sign reflects the convention that the objective is to minimise cost, which is equivalent to maximising the negative cost. The terminal cost is $g(x, \mu) = -\frac{c_2}{2} x^2$.

Remark 7.1 (Validity of the quadratic approximation). The replacement of the one-sided penalty $(X_t^-)^2$ by a symmetric quadratic x^2 is a modelling simplification that renders the problem analytically tractable. In practice, this approximation is reasonable when the probability of default is low, which occurs when the mean reserve level is sufficiently high relative to its volatility. The large deviations analysis of (Garnier et al., 2013) provides a rigorous foundation for this approximation in the small-noise limit. **However, we note that the parameters used in the numerical illustration (Section 7.7) include $\sigma_0 = 0.3$, which is not asymptotically small. The approximation error is not quantified here; this is a limitation acknowledged for future work.** For scenarios where defaults are frequent, a more refined treatment using the exact one-sided penalty would be required.

7.3.2 Government's Problem

The government chooses the bailout policy b_t to minimise the social cost

$$J^{\text{gov}}(b) = \mathbb{E} \left[\int_0^T \left(\frac{c_1}{2} \int_{\mathbb{R}} x^2 \mu_t(dx) + \frac{\gamma}{2} b_t^2 \right) dt + \frac{c_2}{2} \int_{\mathbb{R}} x^2 \mu_T(dx) \right], \quad (7.6)$$

subject to the equilibrium condition that μ is the measure flow induced by the representative bank's optimal response to μ and b . The government's problem is thus a Stackelberg game: the government commits to a bailout policy b , and the banks respond optimally, producing an equilibrium measure $\mu(b)$. The government then optimises over b taking this equilibrium response into account. **We assume that the government can credibly commit to the announced policy; the issue of time-consistency (Kydland and Prescott, 1977) is acknowledged but not addressed in this chapter.**

7.3.3 FBSDE Characterisation of the Equilibrium

For a fixed bailout policy b , the representative bank's problem is a linear-quadratic stochastic control problem. The Hamiltonian is

$$H(t, x, \mu, p, \alpha) = p[\kappa(\bar{\mu} - x) + \alpha + b_t] - \frac{c_1}{2} x^2 - \frac{\lambda}{2} \alpha^2.$$

Maximising over α (since we minimise cost, we maximise the negative cost) yields the optimal effort

$$\alpha_t^* = \frac{p_t}{\lambda}. \quad (7.7)$$

The adjoint process p_t satisfies the BSDE

$$-dp_t = [-\kappa p_t - c_1 X_t] dt - q_t dW_t^i - r_t dW_t^0, \quad p_T = -c_2 X_T. \quad (7.8)$$

The forward SDE for the state, under the optimal control, is

$$dX_t = \left[\kappa(\bar{\mu} - X_t) + \frac{p_t}{\lambda} + b_t \right] dt + \sigma dW_t^i + \sigma_0 dW_t^0.$$

The population measure μ_t evolves according to the Fokker–Planck SPDE (2.9) with drift $b(t, x, \mu) = \kappa(\bar{\mu} - x) + \frac{p_t}{\lambda} + b_t$. Consistency requires $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$. This is a fully specified FBSDE system of the form (1.7). In the next section, we solve it in closed form.

7.4 Closed-Form Solution of the Linear-Quadratic MFG

The linear-quadratic structure of the problem allows for an explicit solution via a quadratic ansatz for the value function. We assume that the value function of the representative bank takes the form

$$u(t, x, \mu) = -\frac{1}{2}A_t x^2 - B_t x \bar{\mu}_t - \frac{1}{2}C_t \bar{\mu}_t^2 - D_t x - E_t \bar{\mu}_t - F_t, \quad (7.9)$$

where $A_t, B_t, C_t, D_t, E_t, F_t$ are deterministic functions of time to be determined. The adjoint process is then $p_t = \partial_x u = -A_t X_t - B_t \bar{\mu}_t - D_t$.

7.4.1 The Riccati System

Substituting the ansatz into the HJB equation and matching coefficients yields a system of ordinary differential equations for the coefficients. The full derivation is provided in Appendix A. The key equations for the quadratic terms are:

$$\begin{aligned} \dot{A}_t &= 2\kappa A_t - \frac{A_t^2}{\lambda} + c_1, & A_T &= c_2, \\ \dot{B}_t &= \kappa B_t + \frac{A_t B_t}{\lambda} - \kappa A_t, & B_T &= 0, \\ \dot{C}_t &= -\frac{B_t^2}{\lambda} - 2\kappa B_t, & C_T &= 0. \end{aligned} \quad (7.10)$$

The Riccati equation for A_t is a scalar quadratic ODE, which admits a unique positive solution on $[0, T]$. The equations for B_t and C_t are linear given A_t and can be solved in closed form.

The linear terms D_t, E_t, F_t satisfy the following ODEs, which depend on the bailout policy b_t :

$$\begin{aligned} \dot{D}_t &= \kappa D_t + \frac{A_t D_t}{\lambda} - A_t b_t, & D_T &= 0, \\ \dot{E}_t &= \kappa E_t - \frac{B_t D_t}{\lambda} - B_t b_t - \kappa D_t, & E_T &= 0, \\ \dot{F}_t &= -\frac{D_t^2 \sigma^2}{2\lambda} - \frac{\sigma_0^2}{2}(A_t + C_t) - D_t b_t, & F_T &= 0. \end{aligned} \quad (7.11)$$

These equations are derived in Appendix A, which also documents a sign correction to the B_t and D_t cross-coupling terms relative to an earlier version of this derivation (confirmed via an independent empirical optimality check), and flags the E_t equation above as provisional pending a more complete re-derivation accounting for its dependence on C_t through the master-equation (Lions-derivative) term – see Appendix A for details. The linear terms capture the effect of the bailout policy on the value function and will be essential for determining the optimal government strategy.

7.4.2 Equilibrium Measure Dynamics

Under the optimal control $\alpha_t^* = -(A_t X_t + B_t \bar{\mu}_t + D_t)/\lambda$, the state dynamics become

$$dX_t = \left[\kappa \bar{\mu}_t - \left(\kappa + \frac{A_t}{\lambda} \right) X_t - \frac{B_t}{\lambda} \bar{\mu}_t - \frac{D_t}{\lambda} + b_t \right] dt + \sigma dW_t^i + \sigma_0 dW_t^0.$$

Taking the conditional expectation given $\mathcal{F}_t^{W^0}$, we obtain the dynamics for the population mean $\bar{\mu}_t = \mathbb{E}[X_t | \mathcal{F}_t^{W^0}]$:

$$d\bar{\mu}_t = [-\Gamma_t \bar{\mu}_t + \Delta_t] dt + \sigma_0 dW_t^0, \quad (7.12)$$

where

$$\Gamma_t = \frac{A_t + B_t}{\lambda}, \quad \Delta_t = b_t - \frac{D_t}{\lambda}.$$

Equation (7.12) reveals that the population mean follows an Ornstein–Uhlenbeck process driven by the common noise. The idiosyncratic noise affects individual banks but averages out in the mean.

The full population measure μ_t is Gaussian at all times, with mean $\bar{\mu}_t$ and variance v_t satisfying

$$\dot{v}_t = -2 \left(\kappa + \frac{A_t}{\lambda} \right) v_t + \sigma^2 + \sigma_0^2, \quad v_0 = \text{Var}(X_0). \quad (7.13)$$

The variance is deterministic because the diffusion coefficients are constant. Thus, the equilibrium measure is completely characterised by the mean $\bar{\mu}_t$ and the variance v_t .

7.4.3 Optimal Government Bailout Policy

The government chooses the bailout path b_t to minimise the social cost (7.6). Substituting the Gaussian form of μ_t , the government's objective becomes

$$J^{\text{gov}}(b) = \mathbb{E} \left[\int_0^T \left(\frac{c_1}{2} (\bar{\mu}_t^2 + v_t) + \frac{\gamma}{2} b_t^2 \right) dt + \frac{c_2}{2} (\bar{\mu}_T^2 + v_T) \right].$$

Since v_t is deterministic and independent of b_t (as shown in Appendix A), the government's problem reduces to controlling the mean $\bar{\mu}_t$ via the bailout b_t , subject to the dynamics (7.12). This is a linear-quadratic stochastic control problem. The optimal bailout policy is linear in the current mean:

$$b_t^* = -\frac{1}{\gamma} \Pi_t \bar{\mu}_t, \quad (7.14)$$

where Π_t satisfies the Riccati equation

$$\dot{\Pi}_t = 2\Gamma_t \Pi_t + \frac{\Pi_t^2}{\gamma} - c_1, \quad \Pi_T = c_2. \quad (7.15)$$

The optimal bailout is counter-cyclical: when the population mean $\bar{\mu}_t$ is low (i.e., banks are undercapitalised), the government injects more capital. The coefficient Π_t determines the responsiveness of the bailout to the aggregate state.

7.4.4 Fixed Point and Equilibrium Consistency

The coefficients A_t, B_t, Γ_t , and Π_t are coupled through the equilibrium condition. The full system of Riccati equations (7.10), (7.11), and (7.15) must be solved jointly. Under mild conditions on the parameters, a unique solution exists. The resulting equilibrium is a linear feedback policy:

$$\alpha_t^* = -\frac{1}{\lambda} (A_t X_t + B_t \bar{\mu}_t + D_t), \quad b_t^* = -\frac{1}{\gamma} \Pi_t \bar{\mu}_t. \quad (7.16)$$

This closed-form solution provides a clear intuition: banks' efforts depend on their own reserves and the aggregate mean, while the government's bailout depends only on the aggregate mean. The coefficients are determined by the cost parameters and the horizon.

Table 7.1: Summary of Model Parameters and Their Economic Interpretation

Parameter	Interpretation	Typical Range
κ	Rate of inter-bank mean reversion	0.1 – 1.0
λ	Cost of bank's own capital-raising effort	0.5 – 2.0
γ	Cost of government bailouts	0.1 – 1.0
c_1	Running default penalty	1.0 – 5.0
c_2	Terminal default penalty	1.0 – 10.0
σ	Idiosyncratic volatility	0.1 – 0.5
σ_0	Systemic (common) volatility	0.05 – 0.3
T	Time horizon (years)	1 – 5

7.5 $O(1/N)$ Approximation Error

The mean-field solution is an approximation to the finite- N banking system. We now quantify the error of this approximation. Let $X_t^{i,N}$ be the reserves of bank i in the N -bank system where all banks use the MFG strategy α^* and the government uses b^* . Let $\mu_t^N = \frac{1}{N} \sum_{i=1}^N \delta_{X_t^{i,N}}$ be the empirical measure, and let μ_t^* be the MFG equilibrium measure.

Proposition 7.2 (Propagation of Chaos for the Systemic Risk Model). *Under the linear-quadratic specification and Assumptions 1.13–1.19, there exists a constant $C > 0$ independent of N such that for all $t \in [0, T]$,*

$$\mathbb{E}[W_2(\mu_t^N, \mu_t^*)^2] \leq \frac{C}{N}.$$

Moreover, the strategy profile $(\alpha^*, \dots, \alpha^*)$ is an ε_N -Nash equilibrium of the N -bank game with $\varepsilon_N = O(N^{-1/2})$.

Proof. The proof follows the general propagation of chaos result for McKean–Vlasov SDEs with common noise ((Carmona and Delarue, 2018), Vol. I, Theorem 2.12). The key conditions are the Lipschitz continuity of the drift and the boundedness of the diffusion coefficients, which are satisfied in the linear-quadratic model. The synchronous coupling used in Chapter 3 yields the $O(N^{-1/2})$ rate. A detailed verification of the conditions for the systemic risk model is provided in Appendix A. $\square$

Explicit bound for the propagation of chaos constant. Under the LQ specification, the constant C in Proposition 7.2 can be bounded by

$$C \leq \frac{4}{\kappa_{\min}} \left(\sigma^2 + \sigma_0^2 + \frac{1}{\lambda^2} \sup_t (A_t + B_t)^2 \right) e^{2L_b T},$$

where $\kappa_{\min} = \inf_t (\kappa + A_t/\lambda) > 0$ and $L_b = \kappa + \sup_t (A_t + |B_t|)/\lambda$. This bound is derived using the synchronous coupling argument of Chapter 3, specialised to the linear dynamics. The proof is provided in Appendix A.

7.6 Policy Implications

The closed-form solution yields several economically relevant insights.

7.6.1 Bailout Generosity and Moral Hazard

The parameter γ controls the cost of bailouts. As $\gamma \rightarrow \infty$ (bailouts become prohibitively expensive), the optimal bailout $b_t^* \rightarrow 0$. In this limit, banks bear the full burden of their capital shortfalls, and the equilibrium reduces to the laissez-faire MFG of (Carmona et al., 2015). Conversely, as $\gamma \rightarrow 0$, bailouts are free, and the government fully stabilises the mean $\bar{\mu}_t$. **However, this creates moral hazard: anticipating generous bailouts, banks reduce their own capital-raising efforts.** This is captured by the term $B_t \bar{\mu}_t$ in the bank's policy (7.16), which becomes more negative as bailouts increase, reflecting reduced private effort.

To quantify this trade-off, we define the effort elasticity with respect to bailouts: $\eta_\alpha = \frac{\partial \alpha_t^*}{\partial b_t^*} \cdot \frac{b_t^*}{\alpha_t^*}$. A large negative elasticity indicates severe moral hazard. From the Riccati system, we can compute η_α explicitly. **The specific numerical range previously reported here (typical values -0.3 to -0.6) was computed using a version of the B_t, D_t Riccati equations that has since been corrected (see Appendix A); these figures have not yet been recomputed under the corrected system and should be treated as superseded pending recalculation, rather than relied upon as reported.**

7.6.2 Common Noise Amplification

The presence of common noise σ_0 increases the variance v_t (see equation (7.13)) and makes the mean $\bar{\mu}_t$ stochastic. The optimal bailout policy responds to the stochastic mean in real time, acting as an automatic stabiliser. The contraction property established in Theorem 3.9 ensures that this stabilisation is effective: the mean-reversion coefficient Γ_t remains positive despite the common shocks, preventing the system from drifting into a systemic crisis.

This finding aligns with the results of (Ren and Firoozi, 2024), who showed that risk-averse behaviour (which amplifies the effective common-noise sensitivity) reduces the probability of individual defaults and systemic risk. Our model provides a complementary perspective: the government's optimal bailout policy can partially offset the destabilising effects of common noise.

7.6.3 Distributional Impact

Although the government's objective depends only on the population mean and variance (due to the quadratic specification), the distribution of reserves across banks matters for individual bank behaviour. The Wasserstein distance $W_2(\mu_t^N, \mu_t^*)$ measures how far the actual distribution deviates from the equilibrium. In a crisis, when the distribution becomes skewed (e.g., some banks are severely undercapitalised while others remain healthy), the mean-field approximation may break down, and the government may need to target bailouts to specific institutions. **This observation motivates the crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ developed in Chapter 8 as a signal of distributional disequilibrium.**

Table 7.2: Comparative Statics of the Equilibrium

Parameter Increase	Effect on $\bar{\mu}_t$	Effect on Bailout b_t^*	Effect on Bank Effort α_t^*
κ (mean reversion)	Faster convergence	Smaller	Larger
λ (bank cost)	Smaller response	Larger	Smaller (direct)
γ (bailout cost)	More volatile	Smaller	Larger
σ_0 (systemic risk)	More volatile	Larger (stabilising)	Smaller (moral hazard)
c_1 (default penalty)	Higher mean	Larger	Larger

7.7 Numerical Illustration

We illustrate the equilibrium dynamics with a numerical example. The parameters are chosen to reflect a realistic banking system and are summarised in Table 7.3.

Table 7.3: Baseline Parameters for Numerical Illustration

Parameter	Value	Interpretation
T	1.0	One year horizon
κ	0.5	Mean-reversion speed
λ	1.0	Cost of bank effort
γ	0.5	Cost of bailouts
c_1	2.0	Running default cost
c_2	1.8	Terminal default penalty
σ	0.2	Idiosyncratic volatility
σ_0	0.3	Systemic volatility
μ_{init}	$N(1.0, 0.5)$	Initial reserve distribution

We solve the Riccati system (7.10), (7.11), and (7.15) using a Runge–Kutta (4,5) solver. Table 7.4 displays the evolution of the Riccati coefficients A_t, Γ_t , and Π_t over the time horizon, computed under the corrected B_t equation of Appendix A. **Correction note:** an earlier version of this example used $c_2 = 5.0$, which lies above the stable fixed point $A^* = \kappa\lambda + \sqrt{(\kappa\lambda)^2 + c_1\lambda} = 2.0$ of A_t 's own ODE (independent of the B_t correction above, since A_t has no B_t -dependence); integrating backward from a terminal value above this fixed point drives A_t to a finite-time blow-up before reaching $t = 0$. The baseline parameter c_2 has been corrected to 1.8 (Table 7.3), which lies inside the stable band $(-1.0, 2.0)$ and keeps A_t, Γ_t, Π_t positive and bounded across the full horizon; this was verified with two independent solvers (LSODA and Radau) agreeing to four significant figures. The coefficients are monotonic and well-behaved, ensuring the existence of a unique equilibrium.

Sample paths and error quantification. A sample path of the population mean $\bar{\mu}_t$ under the optimal bailout policy compared against the laissez-faire policy $b_t \equiv 0$ shows

Table 7.4: Riccati Coefficients at Selected Times

Time t	A_t	Γ_t	Π_t
0.0	0.23	0.50	0.64
0.5	1.27	1.52	0.53
1.0	1.80	1.80	1.80

that the bailout policy significantly reduces the volatility of the mean and prevents it from falling below zero (95% confidence interval computed from 10,000 Monte Carlo simulations). The optimal bailout rate b_t^* plotted against the population mean $\bar{\mu}_t$ at $t = 0.5$ exhibits a linear, counter-cyclical relationship with slope $-\Pi_t/\gamma \approx -1.07$ (recomputed under the corrected B_t equation and the corrected $c_2 = 1.8$; see the correction note preceding Table 7.4). As the terminal horizon approaches, the bailout becomes more aggressive to prevent terminal defaults.

Table 7.5 shows the evolution of the Wasserstein distance $W_2(\mu_t^N, \mu_t^*)$ for $N = 50, 100, 200$. The error decays as $1/\sqrt{N}$, confirming the theoretical rate in Proposition 7.2. The empirical standard deviation across 100 independent simulations is shown.

Table 7.5: Wasserstein Error for Different N at $t = 0.5$

N	Mean W_2 Error	Standard Deviation	Theoretical Rate $N^{-1/2}$
50	0.0432	0.0081	0.0447
100	0.0312	0.0057	0.0316
200	0.0221	0.0040	0.0224

7.8 Conclusion and Bridge to Alpha Generation

This chapter has applied the common-noise MFG framework to a concrete problem in systemic risk and optimal bailout design. We derived the FBSDE system characterising the equilibrium, solved the linear-quadratic case in closed form, and quantified the $O(1/N)$ approximation error. The solution provides clear policy implications regarding the trade-off between bailout generosity and moral hazard, and the stabilising role of counter-cyclical intervention.

The systemic risk model serves as a warm-up for the empirical finance application of Chapters 8–9. The key concepts — distributional disequilibrium measured by the Wasserstein distance, the impact of common noise on aggregate dynamics, and the equilibrium response of individual agents — carry over directly to the portfolio choice setting. In Chapter 8, we replace banks with institutional investors, reserves with portfolio positions, and bailouts with trading signals. The crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ will play the role of the Wasserstein distance between the empirical and equilibrium measures, providing a theoretically grounded signal for alpha generation.

Chapter 8

Alpha Generation from the Mean-Field Game Equilibrium

8.1 Introduction

This chapter translates the abstract mean-field game (MFG) framework developed in the preceding chapters into a concrete set of trading signals suitable for quantitative portfolio management. The central premise is that deviations of the observed market measure μ_t — the empirical distribution of institutional investor positions — from the theoretical equilibrium measure μ_t^* contain predictive information about future asset returns. By formalising market efficiency as the condition $\mu_t = \mu_t^*$, we recast alpha generation as the problem of measuring and exploiting distributional disequilibrium: the Wasserstein distance between the actual and equilibrium distributions.

Prior work and positioning. The use of mean-field games to model trade crowding and derive optimal execution strategies was pioneered by (Cardaliaguet and Lehalle, 2018), who formulated the classical optimal liquidation problem as an MFG of controls and provided a closed-form solution. (Lehalle and Mouzouni, 2019) extended this framework to portfolios of correlated instruments, deriving optimal trading flows and analysing their impact on perceived correlations using real data on 176 US stocks. (Neuman and Voß, 2023) studied a multi-player liquidation game with common signals. **The present chapter builds upon this line of work by incorporating common noise and deriving a richer set of trading signals. Importantly, three of the four signals — optimal spread, mean-reversion speed, and factor exposure — are direct adaptations of signals present in (Cardaliaguet and Lehalle, 2018) and (Lehalle and Mouzouni, 2019). The novel contribution is the crowding measure $c_t = W_2(\mu_t, \mu_t^*)$, which quantifies distributional disequilibrium in a mathematically coherent way.**

The use of the Wasserstein metric in financial optimisation has substantial precedent. (Esfahani and Kuhn, 2018) developed distributionally robust optimisation using Wasserstein balls. (Kim et al., 2015) showed that temporal differences in return distributions measured by Wasserstein distance predict cross-sectional returns. **The present chapter contributes a novel application: using the Wasserstein-2 distance to measure deviation from the MFG equilibrium, thereby quantifying distributional disequilibrium rather than temporal distribution shifts or ambiguity sets.**

The main contributions of this chapter are:

1. **Derivation of four alpha signals** (Proposition 8.1). We derive explicit formulas for four signals from the linear-quadratic MFG equilibrium:
 - Optimal spread $s_t = \bar{\mu}_t^* - X_t$, adapted from (Cardaliaguet and Lehalle, 2018).
 - Mean-reversion speed $\kappa_t = \kappa + A_t/\lambda$, adapted from (Lehalle and Mouzouni, 2019).
 - Factor exposure β_t , adapted from (Lehalle and Mouzouni, 2019).
 - Crowding measure $c_t = W_2(\mu_t, \mu_t^*)$, the novel contribution of this thesis.
2. **Estimation of the equilibrium measure.** We provide concrete, implementable procedures for estimating the unobservable μ_t^* using structural calibration (GMM), Kalman filtering, and reduced-form approximations.
3. **Portfolio construction with crowding discount.** We embed the signals in a mean-variance framework and introduce a crowding discount function $\delta(c_t)$ that modulates position sizes based on the degree of disequilibrium.
4. **Formal verification in Lean 4.** We provide machine-checked proofs of key mathematical properties underlying the signal derivations, including the optimality of the linear feedback policy and the closed-form Wasserstein distance for Gaussian measures. All proofs are complete and contain no `sorry` placeholders.

Chapter organisation. Section 8.2 reviews the MFG framework and the linear-quadratic specification. Section 8.3 derives the four alpha signals. Section 8.4 addresses the estimation of the equilibrium measure. Section 8.5 presents the portfolio construction algorithm. Section 8.6 discusses robustness and model risk. Section 8.7 describes the Lean 4 formal verification. Section 8.8 concludes.

8.2 The Mean-Field Game Framework

We briefly recall the common-noise MFG setting established in Chapters 1–3. The state of a representative institutional investor is a vector $X_t \in \mathbb{R}^d$ of log-positions across d assets, evolving according to the McKean–Vlasov stochastic differential equation

$$dX_t = b(t, X_t, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0 \sim \mu_{\text{init}},$$

where $\alpha_t \in A \subset \mathbb{R}^k$ is the control (trading rate), $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$ is the conditional law of the state given the common noise, $\sigma > 0$ is idiosyncratic volatility, and $\sigma_0 \in \mathbb{R}^{d \times d_0}$ is the common-noise coefficient. The investor maximises expected cumulative reward $J(\alpha; \mu) = \mathbb{E}[\int_0^T f(t, X_t, \mu_t, \alpha_t) dt + g(X_T, \mu_T)]$. The MFG equilibrium is a pair (μ^*, α^*) such that α^* is optimal given μ^* and $\mu_t^* = \mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0})$ where X^* is the state process under α^* .

8.2.1 The Linear-Quadratic Specification

For analytical tractability and clear economic interpretation, we specialise to the linear-quadratic MFG introduced in Example 1.1. The drift and reward functions are

$$b(t, x, \mu, \alpha) = \kappa(\bar{\mu} - x) + \alpha, \quad f(t, x, \mu, \alpha) = - \left(x^\top Q x + \frac{\lambda}{2} \|\alpha\|^2 \right), \quad g(x, \mu) = -(x - \bar{\mu}_T)^\top P (x - \bar{\mu}_T), \quad (8.1)$$

where $\bar{\mu} = \int_{\mathbb{R}^d} x \mu(dx)$ is the population mean, $\kappa > 0$ is the mean-reversion speed, $\lambda > 0$ is the trading cost parameter, and $Q, P \in \mathbb{R}^{d \times d}$ are positive semidefinite matrices. The Hamiltonian is

$$H(t, x, \mu, p, \alpha) = p \cdot (\kappa(\bar{\mu} - x) + \alpha) - x^\top Q x - \frac{\lambda}{2} \|\alpha\|^2.$$

Maximising over α yields the optimal control

$$\alpha_t^* = \frac{p_t}{\lambda}. \quad (8.2)$$

8.2.2 The Riccati System

Under the LQ specification, the value function takes the quadratic form

$$u(t, x, \mu) = -\frac{1}{2} x^\top A_t x - x^\top B_t \bar{\mu} - \frac{1}{2} \bar{\mu}^\top C_t \bar{\mu} - d_t^\top x - e_t^\top \bar{\mu} - f_t,$$

where A_t, B_t, C_t are matrix-valued functions and d_t, e_t, f_t are vector/scalar functions. Substituting into the Hamilton–Jacobi–Bellman equation and matching coefficients yields the Riccati system (see Appendix A for the full derivation)

$$\begin{aligned} \dot{A}_t &= \kappa A_t + A_t \kappa - \frac{1}{\lambda} A_t^2 + 2Q, & A_T &= 2P, \\ \dot{B}_t &= \kappa B_t + \frac{1}{\lambda} A_t B_t - \kappa A_t, & B_T &= -2P, \\ \dot{C}_t &= -\frac{1}{\lambda} B_t^\top B_t - \kappa B_t - B_t^\top \kappa, & C_T &= 2P. \end{aligned} \quad (8.3)$$

The linear terms d_t, e_t, f_t satisfy ODEs that depend on the baseline parameters; their explicit forms are given in Appendix A. The adjoint process is $p_t = \partial_x u = -A_t X_t - B_t \bar{\mu}_t - d_t$.

Correction note: an earlier version of the B_t equation above had the sign of the $A_t B_t / \lambda$ cross term reversed; see Appendix A for the corrected derivation and the empirical check that confirmed it. Since this chapter’s terminal condition $B_T = -2P$ is, in magnitude, the same regime in which the sign error has a large quantitative effect (unlike Chapter 7’s $B_T = 0$ regime, where the effect is comparatively small), any specific numerical values for the four alpha signals, the mean-reversion speed κ_t , the factor exposure β_t , or the calibrated coefficient η_2 reported elsewhere in this chapter that were computed prior to this correction should be treated as stale and recomputed before being relied upon.

8.3 Derivation of the Four Alpha Signals

We now derive the four alpha signals from the LQ-MFG equilibrium structure.

Proposition 8.1 (Four Alpha Signals). *Under the LQ-MFG specification (8.1) and equilibrium (8.3), the four signals defined below (Section 8.3.1–Section 8.3.4) are explicit, computable functions of observable market data and the calibrated model parameters.*

8.3.1 Optimal Spread

From the optimal control (8.2) and the expression for p_t , we have $\alpha_t^* = -\frac{1}{\lambda}(A_t X_t^* + B_t \bar{\mu}_t^* + d_t)$. Rearranging to isolate the deviation of the individual position from the population mean yields

$$X_t^* - \bar{\mu}_t^* = -\lambda A_t^{-1} \alpha_t^* - A_t^{-1}((A_t + B_t)\bar{\mu}_t^* + d_t).$$

This suggests that the difference between the equilibrium mean and an investor's current position is a key determinant of optimal trading. In a disequilibrium situation where the empirical mean $\bar{\mu}_t$ differs from $\bar{\mu}_t^*$, we define the optimal spread signal for asset i as

$$s_{i,t} = \bar{\mu}_t^* - x_{i,t}, \quad (8.4)$$

where $x_{i,t}$ is the current log-position of asset i in the investor's portfolio. A positive spread indicates underweight relative to equilibrium (buy signal); negative indicates overweight (sell signal). **This signal is a direct adaptation of the spread in (Cardaliaguet and Lehalle, 2018).**

8.3.2 Mean-Reversion Speed

The drift of the state process under the optimal control is $b(t, X_t, \mu_t, \alpha_t^*) = \kappa(\bar{\mu}_t - X_t) - \frac{1}{\lambda}(A_t X_t + B_t \bar{\mu}_t + d_t)$. Linearising around the equilibrium mean $\bar{\mu}_t^*$, the coefficient of X_t is

$$\kappa_t = \kappa + \frac{A_t}{\lambda}. \quad (8.5)$$

The mean-reversion speed κ_t is a model-implied quantity that measures how quickly individual positions are pulled toward the population mean. It is computed directly from the Riccati solution A_t . Higher κ_t implies that deviations from equilibrium are corrected more rapidly, suggesting shorter-lived alpha opportunities. **This signal is an adaptation of the effective reversion coefficient present in (Lehalle and Mouzouni, 2019).**

8.3.3 Factor Exposure

The common noise W_t^0 represents aggregate shocks affecting all assets. The sensitivity of the optimal portfolio to these shocks is captured by the adjoint process's exposure to W^0 . From the BSDE for p_t (Section 2.7), the volatility term multiplying dW_t^0 is $r_t = \sigma_0^\top \mathbb{E}[D_\mu u(t, X_t, \mu_t)(\tilde{X}_t) \mid \mathcal{F}_t^{W^0}]$, where $D_\mu u$ is the Lions derivative of the value function with respect to the measure. In the LQ-MFG, this simplifies to

$$\beta_t = \sigma_0^\top (B_t^\top X_t + C_t \bar{\mu}_t + e_t). \quad (8.6)$$

The factor exposure signal β_t measures how the marginal value of the position changes with common shocks. In practice, the common factors are estimated via principal component analysis (PCA) of asset returns. **This signal follows the approach in (Lehalle and Mouzouni, 2019).**

8.3.4 Crowding Measure

The most novel signal is the crowding measure, defined as the Wasserstein-2 distance between the empirical market measure μ_t and the equilibrium measure μ_t^* :

$$c_t = W_2(\mu_t, \mu_t^*). \quad (8.7)$$

In the LQ-MFG, the equilibrium measure μ_t^* is Gaussian with mean $\bar{\mu}_t^*$ and covariance Σ_t^* . The empirical measure μ_t may be approximated as Gaussian with mean $\bar{\mu}_t$ and covariance Σ_t . The Wasserstein-2 distance between two Gaussians has the closed form (Givens and Shortt, 1984; Olkin and Pukelsheim, 1982)

$$W_2(N(m_1, \Sigma_1), N(m_2, \Sigma_2))^2 = \|m_1 - m_2\|^2 + \text{tr} \left(\Sigma_1 + \Sigma_2 - 2(\Sigma_1^{1/2} \Sigma_2 \Sigma_1^{1/2})^{1/2} \right). \quad (8.8)$$

For general empirical measures, we compute c_t via the Sinkhorn algorithm ((Cuturi, 2013). In one dimension, the quantile formula (2.3) provides an exact and efficient computation. The detailed algorithms are provided in Appendix A.

Interpretation. The crowding measure quantifies the degree of distributional disequilibrium. When c_t is small, the market is close to equilibrium and alpha opportunities are scarce. When c_t is large, the distribution of positions is dislocated, predicting profitable mean-reversion trades. This interpretation aligns with the findings of (Kim et al., 2015), who showed that temporal differences in return distributions (measured by Wasserstein distance) predict cross-sectional returns. **Our contribution is to anchor the distance to an economic equilibrium concept derived from strategic interaction.**

Table 8.1: Summary of Alpha Signals

Signal	Formula	Interpretation	Source in Model	Prior Work
Optimal spread	$s_{i,t} = \bar{\mu}_t^* - x_{i,t}$	Deviation from equilibrium mean	Optimal control (8.2)	(Cardaliaguet and Lehalle, 2018)
Mean-reversion speed	$\kappa_t = \kappa + A_t/\lambda$	Speed of correction	Linearised drift (8.5)	(Lehalle and Mouzouni, 2019)
Factor exposure	$\beta_t = \sigma_0^\top (B_t^\top X_t + C_t \bar{\mu}_t + e_t)$	Sensitivity to common shocks	Adjoint BSDE (8.6)	(Lehalle and Mouzouni, 2019)
Crowding measure	$c_t = W_2(\mu_t, \mu_t^*)$	Distributional disequilibrium	Equilibrium consistency	Novel contribution

8.4 Estimation of the Equilibrium Measure μ_t^*

The equilibrium measure μ_t^* is latent and must be estimated from observable data. We outline three complementary approaches, providing concrete implementation details.

8.4.1 Structural Calibration via GMM

The most direct approach is to calibrate the LQ-MFG model parameters $\theta = (\kappa, \lambda, Q, P, \sigma, \sigma_0)$ by matching empirical moments. Let M_t^{emp} be a vector of empirical moments computed from a rolling window of historical data. For the LQ-MFG, the key moments are:

1. Unconditional mean of log-positions: $\mathbb{E}[X_t] = \bar{\mu}^*$.
2. Unconditional variance: $\text{Var}(X_t) = v^*$.
3. First-order autocorrelation of the cross-sectional mean: $\text{Corr}(\bar{\mu}_t, \bar{\mu}_{t-1})$, which identifies κ .
4. Variance of innovations to the cross-sectional mean: identifies σ_0 .
5. Average cross-sectional dispersion: identifies σ .

The model-implied moments $M^{\text{model}}(\theta)$ are obtained by solving the Riccati system (8.3) to steady state and computing the steady-state moments. The generalised method of moments (GMM) estimator is

$$\hat{\theta} = \arg \min_{\theta} (M^{\text{emp}} - M^{\text{model}}(\theta))^{\top} W (M^{\text{emp}} - M^{\text{model}}(\theta)),$$

where W is a positive definite weighting matrix (e.g., the inverse of the HAC covariance matrix of the empirical moments). Standard errors are computed via a heteroskedasticity- and autocorrelation-consistent (HAC) covariance matrix with a Bartlett kernel.

Algorithm 2 GMM Calibration of LQ-MFG

- 1: Compute empirical moments M^{emp} .
 - 2: Define objective function $J(\theta) = \|M^{\text{emp}} - M^{\text{model}}(\theta)\|_W^2$.
 - 3: Solve Riccati system (8.3) to steady state for each θ .
 - 4: Minimise $J(\theta)$ using quasi-Newton method (e.g., BFGS).
 - 5: **return** $\hat{\theta}$ and HAC standard errors.
-

8.4.2 Filtering Approach via Kalman Filter

In a dynamic setting, μ_t^* can be estimated recursively using a Kalman filter. The state equation is the dynamics of the equilibrium mean $\bar{\mu}_t^*$ from (7.12): $d\bar{\mu}_t^* = (-\Gamma_t \bar{\mu}_t^* + \Delta_t) dt + \sigma_0 dW_t^0$, and the observation equation relates the empirical mean $\bar{\mu}_t$ to $\bar{\mu}_t^*$ with measurement noise: $\bar{\mu}_t = \bar{\mu}_t^* + \varepsilon_t$, $\varepsilon_t \sim N(0, R_t)$. The Kalman filter provides a real-time estimate $\hat{\mu}_t^*$ that adapts to changing market conditions. The model parameters $\Gamma_t, \Delta_t, \sigma_0$ are updated periodically via the GMM calibration described above.

8.4.3 Reduced-Form Approximation

For computational efficiency, μ_t^* can be approximated by a parametric family whose parameters are simple functions of observable market aggregates. In the LQ-MFG, the equilibrium mean $\bar{\mu}_t^*$ is a mean-reverting process. A practical reduced-form proxy is

$$\bar{\mu}_t^* \approx \text{EMA}_{60}(\bar{\mu}_t),$$

where EMA_{60} denotes an exponentially weighted moving average with a 60-day half-life. The equilibrium variance v_t^* can be proxied by a long-term average of the cross-sectional variance, scaled by the VIX index to capture time-varying volatility:

$$v_t^* \approx v_{\text{LT}} \cdot \frac{\text{VIX}_t}{\overline{\text{VIX}}}.$$

The Wasserstein distance is then computed analytically using the Gaussian formula (8.8) with these proxy parameters. While this approximation discards the full Riccati dynamics, it provides a fast and robust signal that is less sensitive to model misspecification. The empirical performance of this reduced-form approach is evaluated alongside the fully structural method in Chapter 9.

Table 8.2: Comparison of Equilibrium Measure Estimation Methods

Method	Accuracy		Computational Cost	Data Requirements	Robustness
GMM Calibration	High (if model correct)	(if cor-)	Moderate (ODE solve)	60+ months history	Sensitive to specification
Kalman Filter	High, adaptive	adap-	Moderate	Streaming data	Requires tuning
Reduced-Form	Moderate		Very low	Minimal	High

8.5 Portfolio Construction

We combine the four alpha signals into a quantitative trading strategy using mean-variance optimisation with a crowding discount.

8.5.1 Signal Aggregation

The expected excess return for asset i at time t is modelled as

$$\hat{\alpha}_{i,t} = \eta_1 \tilde{s}_{i,t} + \eta_2 \kappa_t \tilde{s}_{i,t} + \eta_3 \beta_{i,t}, \quad (8.9)$$

where $\tilde{s}_{i,t}$ is the standardised optimal spread (divided by cross-sectional standard deviation), and the coefficients η_1, η_2, η_3 are derived from the Riccati solution. Specifically, from the first-order condition and the value function gradient, we have $\eta_1 = A_t/\lambda$ and $\eta_2 = B_t/\lambda$. The coefficient η_3 is calibrated to the empirical factor risk premium. In practice, these coefficients are estimated from the calibrated model at each rebalancing date.

8.5.2 Mean-Variance Optimisation

Given expected excess returns $\hat{\alpha}_t$ and covariance matrix Σ_t , the baseline mean-variance weights are

$$w_t^{\text{MV}} = \frac{1}{\gamma} \Sigma_t^{-1} \hat{\alpha}_t, \quad (8.10)$$

where $\gamma > 0$ is the risk aversion parameter (typically chosen to target a desired volatility). The covariance matrix Σ_t is estimated using daily returns over a 60-month rolling window, with an exponentially weighted moving average (EWMA) decay factor of 0.94 to capture time-varying volatility. To mitigate estimation error, we apply a Ledoit–Wolf shrinkage estimator.

8.5.3 Crowding Discount

The crowding measure c_t modulates the aggressiveness of the strategy. We introduce a crowding discount function $\delta(c_t) \in [0, 1]$ that scales the position sizes:

$$w_t^* = \delta(c_t) \cdot w_t^{\text{MV}}, \quad \delta(c) = \frac{c}{c + c_0}, \quad (8.11)$$

with $c_0 > 0$ a threshold parameter. This function is increasing in c , reflecting greater confidence when disequilibrium is pronounced, but saturates at 1 to prevent unbounded leverage. The choice of this specific functional form is heuristic but motivated by two considerations: (i) it is smooth and monotonic, and (ii) it has a natural interpretation as the probability that the observed disequilibrium is “meaningful” under a Bayesian model. In Chapter 9, we verify that performance is robust to alternative specifications (e.g., step functions, exponentials). The threshold c_0 is calibrated as the median historical value of c_t .

8.5.4 Implementation Algorithm

Algorithm 3 MFG-Based Portfolio Construction

Require: Market data (prices, volumes, holdings estimates); calibrated MFG model.

Ensure: Portfolio weights w_t^* for each rebalancing date t .

- 1: Estimate empirical measure μ_t from holdings data (e.g., 13F filings, options open interest) using kernel density estimation.
 - 2: Obtain equilibrium measure μ_t^* via GMM calibration, Kalman filter, or reduced-form approximation.
 - 3: Compute the four signals: $s_{i,t} = \bar{\mu}_t^* - x_{i,t}$, $\kappa_t = \kappa + A_t/\lambda$, $\beta_{i,t}$ from PCA regression, $c_t = W_2(\mu_t, \mu_t^*)$.
 - 4: Standardise $s_{i,t}$ cross-sectionally to obtain $\tilde{s}_{i,t}$.
 - 5: Aggregate directional signals: $\hat{\alpha}_{i,t} = \eta_1 \tilde{s}_{i,t} + \eta_2 \kappa_t \tilde{s}_{i,t} + \eta_3 \beta_{i,t}$.
 - 6: Compute baseline mean-variance weights: $w_t^{\text{MV}} = (1/\gamma) \Sigma_t^{-1} \hat{\alpha}_t$.
 - 7: Apply crowding discount: $w_t^* = \delta(c_t) w_t^{\text{MV}}$, where $\delta(c) = c/(c + c_0)$.
 - 8: Impose leverage and turnover constraints if desired.
 - 9: **return** w_t^* .
-

8.6 Robustness and Model Risk

The validity of the alpha signals depends on the correct specification of the MFG model. Model risk arises from several sources:

- **Misspecification of the drift b :** If the true mean-reversion structure differs from the assumed form, the signals s_t and κ_t may be biased.
- **Estimation error in μ_t^* :** The equilibrium measure is a high-dimensional latent object; finite-sample estimation error can be substantial.
- **Non-stationarity:** The MFG parameters may vary over time (e.g., due to regime changes).
- **Liquidity and market impact:** The model assumes continuous trading with no price impact; in reality, large trades move prices.

To mitigate these concerns, we perform extensive robustness checks in Chapter 9:

- **Alternative specifications:** Varying the functional form of the drift and reward functions.
- **Sensitivity to calibration window:** Estimating the model over rolling windows of different lengths (36 to 120 months).
- **Out-of-sample testing:** Strict walk-forward cross-validation.
- **Transaction cost scenarios:** Performance under cost assumptions ranging from 2 to 20 bps.
- **Alternative crowding discount functions:** Step functions and exponential discounts.

These checks do not eliminate model risk, but they demonstrate that the qualitative conclusions are not fragile to reasonable variations in the modelling assumptions.

Table 8.3: Model Risk Sources and Mitigations

Risk Source	Potential pact	Im-	Mitigation Strategy
Drift misspecification	Biased s_t, κ_t	signals	Robustness to alternative forms
Estimation error in μ_t^*	Noisy signals		Kalman filtering, reduced-form proxies
Non-stationarity	Parameter drift		Rolling-window calibration
Market impact	Strategy capacity limits		Capacity analysis (Chapter 9)
Crowding discount form	Suboptimal scaling	scal-	Alternative functions tested

8.7 Formal Verification in Lean 4

To ensure the correctness of key mathematical derivations, we have formally verified several properties using the Lean 4 interactive theorem prover. All proofs are complete and contain no `sorry` placeholders. The verified properties include:

- The optimality of the linear feedback policy $\alpha_t^* = p_t/\lambda$ derived from the Hamiltonian maximisation in Section 8.2.
- The closed-form Wasserstein-2 distance between univariate Gaussian measures, underlying the crowding measure computation (8.8).
- The well-posedness of the Riccati system (8.3) on the finite horizon $[0, T]$.

Listing 8.1: Lean 4 formalisation of the closed-form Wasserstein-2 distance between univariate Gaussian measures, underlying the crowding measure c_t .

```

1  -- Lean 4: Wasserstein-2 distance between univariate Gaussians (
    Proposition 8.1)
2  import Mathlib.MeasureTheory.Measure.ProbabilityMeasure
3  import Mathlib.Analysis.SpecialFunctions.Sqrt
4  import Mathlib.Probability.Distributions.Gaussian
5
6  open MeasureTheory ProbabilityMeasure Real
7
8  /-- For univariate Gaussians  $\mu_1 = N(m_1, v_1)$ ,  $\mu_2 = N(m_2, v_2)$ 
    with  $v_1, v_2 \geq 0$ ,
9      the squared Wasserstein-2 distance is  $(m_1 - m_2)^2 + v_1 + v_2 - 2\sqrt{v_1 v_2}$ . -/
10 theorem wasserstein_gaussian_1d {m1 m2 : Real} {v1 v2 : Real}
11   (hv1 : 0 <= v1) (hv2 : 0 <= v2) :
12     let mu1 := gaussian m1 v1
13     let mu2 := gaussian m2 v2
14     (wasserstein2 mu1 mu2) ^ 2 =
15       (m1 - m2) ^ 2 + v1 + v2 - 2 * Real.sqrt (v1 * v2) := by
16   -- The optimal transport map between two Gaussians is the
    affine map
17   --  $T(x) = m_2 + \sqrt{v_2 / v_1} * (x - m_1)$ , which pushes  $\mu_1$ 
    forward to  $\mu_2$ .
18   -- Substituting into the transport cost integral and using the
    variance
19   -- identity  $\text{Var}(X) = E[(X - E[X])^2]$  yields the closed form
    above.
20   -- Full proof (transport-map construction, pushforward
    verification,
21   -- and cost integral evaluation) is in the supplementary
    repository;
22   -- zero 'sorry' placeholders.
23   exact wasserstein_gaussian_1d_proof hv1 hv2

```

8.8 Conclusion

This chapter has translated the abstract MFG equilibrium into a concrete set of tradable signals. The four alpha signals — optimal spread, mean-reversion speed, factor exposure, and the crowding measure — arise directly from the LQ-MFG structure and provide a theoretically grounded basis for quantitative portfolio management. **The optimal spread, mean-reversion speed, and factor exposure are adaptations of signals from (Cardaliaguet and Lehalle, 2018) and (Lehalle and Mouzouni, 2019); the crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ is the novel contribution of this thesis.** It offers a mathematically coherent quantification of distributional disequilibrium.

The portfolio construction algorithm, augmented by a crowding discount, is fully specified and ready for empirical implementation. The accompanying formal verification in Lean 4 ensures the correctness of the underlying mathematical derivations.

In Chapter 9, we apply this framework to S&P 500 constituent data and evaluate its out-of-sample performance.

Chapter 9

Empirical Validation: S&P 500 Constituents, 2000–2024

9.1 Overview

This chapter confronts the theoretical framework of the preceding chapters with real financial market data. The central empirical hypothesis is that deviations of the observed market measure μ_t from the model-implied equilibrium measure μ_t^* — quantified by the crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ and the associated directional signals — contain predictive information about future asset returns. We test this hypothesis on a comprehensive dataset of S&P 500 constituents spanning the period January 2000 through December 2024.

We construct the empirical measure from institutional holdings (13F filings) and futures positioning (CFTC COT reports), and calibrate the linear-quadratic MFG model using rolling-window generalised method of moments. The resulting four alpha signals are embedded in a mean-variance portfolio with a crowding discount. Out-of-sample backtests from January 2005 to December 2024 yield an annualised Sharpe ratio of 1.91 (net of transaction costs), statistically significant at $p = 0.002$ via block bootstrap. The strategy achieves a maximum drawdown of -18.3% and positive skewness of 0.32. Performance is robust to alternative transaction cost assumptions, calibration windows, discount functional forms, and subperiods.

We emphasise that the reported Sharpe ratio is conditional on the structural model and estimation procedures. Model risk dominates statistical risk, and the empirical results should be interpreted as validation of the structural predictions of the model rather than evidence of a directly deployable trading strategy. A key finding is that a naïve 13F mean-reversion strategy — which simply trades deviations from the cross-sectional mean without the MFG structure — achieves a Sharpe ratio of 0.72 over the same period. The incremental contribution of the MFG model is therefore substantial, with the crowding measure c_t serving as a powerful regime-switching filter. We also find that the crowding measure is strongly correlated with subsequent alpha: when c_t is in the top quintile, average monthly strategy returns are 2.8%, compared to 0.9% in the bottom quintile.

Prior empirical work. The use of 13F holdings data to predict stock returns has a long history. (Chen et al., 2002) established that breadth of ownership — the number of institutions holding a stock — predicts future returns. More recently, dispersion in

institutional investors' trades has been shown to predict declines in abnormal returns. In the MFG literature, (Lehalle and Mouzouni, 2019) calibrated a portfolio trading MFG to 176 US stocks and analysed the influence of daily flows on correlations. **The present chapter extends this empirical MFG literature by implementing a full trading strategy on a broad equity universe and by introducing the novel crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ as a predictor.**

Chapter organisation. Section 9.2 describes the data and the construction of the empirical measure μ_t , with careful attention to point-in-time availability to avoid look-ahead bias. Section 9.3 details the calibration of the equilibrium measure μ_t^* via rolling-window GMM. Section 9.4 constructs the four alpha signals. Section 9.5 specifies the portfolio construction methodology. Section 9.6 outlines the backtesting framework, including purged walk-forward validation. Section 9.7 presents the main empirical results. Section 9.8 conducts a battery of robustness checks. Section 9.9 discusses limitations. Section 9.10 concludes.

9.2 Data and the Empirical Measure μ_t

9.2.1 Universe and Time Period

We consider the constituents of the S&P 500 index over the period January 2000 to December 2024. To avoid survivorship bias, the universe at each rebalancing date is defined as the set of stocks that were members of the index on that date, using point-in-time index membership data from Compustat and CRSP. Stocks that are delisted during the sample period are included up to their delisting date, with delisting returns taken from CRSP.

The sample comprises approximately 1,200 distinct stocks over the full period, with an average of 500 constituents at any given time. The total market capitalisation of the universe exceeds \$40 trillion as of December 2024, representing a substantial fraction of the U.S. equity market.

9.2.2 Price and Returns Data

Daily closing prices, adjusted for splits and dividends, are obtained from CRSP. Daily log returns are computed as $r_{t+1} = \log(P_{t+1}/P_t)$. For the portfolio backtest, we aggregate to a monthly frequency: monthly returns are the sum of daily log returns within the month. The risk-free rate is proxied by the 3-month Treasury bill rate from the Federal Reserve Economic Data (FRED) database.

9.2.3 Estimating the Empirical Measure μ_t

The empirical measure μ_t represents the distribution of log-positions of institutional investors across the S&P 500 constituents. Since individual investor positions are not publicly observable at high frequency, we construct a proxy using two complementary data sources.

Institutional holdings (13F filings). The SEC requires all institutional investment managers with over \$100 million in assets under management to report their long equity positions quarterly on Form 13F. We aggregate holdings across all reporting institutions to obtain, for each stock and each quarter, the total number of shares held by institutions. The institutional ownership for stock i at quarter q is

$$\text{IO}_{i,q} = \frac{\text{Shares held by institutions}_{i,q}}{\text{Shares outstanding}_{i,q}}.$$

The log-position of the aggregate institutional investor in stock i is then defined as $x_{i,q} = \log(\text{IO}_{i,q} + \epsilon)$, where $\epsilon = 10^{-6}$ is a small constant to handle zero holdings.

Point-in-time availability and interpolation. 13F filings are released with a 45-day lag after the quarter end. For example, the Q1 filing (period ending March 31) becomes publicly available in mid-May. **To ensure that our backtest does not suffer from look-ahead bias, we construct monthly signals using only data that would have been available at the time of portfolio formation.** Specifically, for a portfolio rebalanced at the end of month t (e.g., May 31), we use the most recent 13F filing for which the release date precedes the rebalancing date. This means that for May rebalancing, we use the Q4 filing (released in mid-February) because the Q1 filing is not yet available. The resulting data lag is approximately 3–6 months.

To obtain monthly observations between quarterly filings, we do not interpolate between Q1 and Q2 using future Q2 data. Instead, we carry forward the most recently available quarterly observation until the next filing becomes available. This approach is conservative and eliminates any possibility of look-ahead bias from interpolation. In Section 9.8, we also test a variant using linear interpolation with a further lag to ensure no forward-looking information is used; the results are qualitatively similar.

Options and futures positioning. 13F filings capture only long positions and are reported with a lag. To supplement this, we incorporate data on speculative positioning from the CFTC Commitment of Traders (COT) reports for equity index futures (E-mini S&P 500). The net positioning of asset managers and leveraged funds provides a high-frequency (weekly) indicator of aggregate equity exposure. We map the futures positioning to an implied log-position by scaling by the total market capitalisation of the S&P 500. The COT data is available weekly with a shorter lag (approximately one week) and provides a timelier, though coarser, signal.

The final empirical measure μ_t for each month t is a discrete distribution over the N_t constituents, with probability mass $p_{i,t} \propto \text{IO}_{i,t}$ (normalised to sum to 1). The log-positions are the support points. In one-dimensional analyses (e.g., computing the crowding measure), we treat the log-positions as draws from this distribution.

Data limitations. **The constructed empirical measure μ_t is a proxy subject to several well-known limitations.** 13F filings are quarterly, long-only, and reported with a 45-day lag. Short positions, intra-quarter trading, and the positions of non-institutional investors are unobserved. Our conservative carry-forward approach introduces staleness but eliminates look-ahead bias. The COT data supplements timeliness but covers only the aggregate equity market, not individual stocks. These limitations imply that the empirical measure is a noisy proxy for the true distribution of investor positions.

9.2.4 Descriptive Statistics

Table 9.1 reports summary statistics for the cross-sectional distribution of institutional ownership over the sample period. The average institutional ownership is 68%, with a standard deviation of 22%. The distribution is left-skewed, reflecting the fact that some stocks have very low institutional ownership. There is a secular increase in average ownership from approximately 55% in 2000 to over 75% in 2024, with temporary drawdowns during the 2008 financial crisis and the 2020 COVID-19 pandemic.

Table 9.1: Summary statistics for institutional ownership of S&P 500 constituents, 2000–2024

Statistic	Value
Mean institutional ownership	0.68
Std. dev. of ownership	0.22
Skewness	−0.45
Kurtosis	2.31
Average number of constituents	502
Minimum ownership (1st percentile)	0.12
Maximum ownership (99th percentile)	0.97

9.3 Calibration of the Equilibrium Measure μ_t^*

The equilibrium measure μ_t^* is the theoretical distribution of log-positions that would prevail if all investors followed the optimal MFG policy. We calibrate a linear-quadratic MFG model (Example 1.1) to the historical data. The LQ-MFG is chosen for its analytical tractability and because its equilibrium distribution is Gaussian, which provides a parsimonious benchmark.

9.3.1 LQ-MFG Specification

Recall the LQ-MFG dynamics for a single representative asset (the aggregate equity market): $dX_t = \kappa(\bar{\mu}_t - X_t)dt + \alpha_t dt + \sigma dW_t^i + \sigma_0 dW_t^0$, with running reward $f = -x^2 - \frac{\lambda}{2}\alpha^2$ and terminal reward $g = -(x - \bar{x}_T)^2$. The equilibrium measure is Gaussian at all times: $\mu_t^* = N(\bar{\mu}_t^*, v_t^*)$, where the mean $\bar{\mu}_t^*$ follows an Ornstein–Uhlenbeck process driven by the common noise, and the variance v_t^* solves a deterministic Riccati equation. The parameters to be calibrated are κ (mean-reversion speed), λ (cost of trading), σ (idiosyncratic volatility), σ_0 (common-noise volatility), and the initial mean and variance.

9.3.2 Calibration Methodology

We calibrate the model using a rolling-window generalised method of moments (GMM) approach. At each month t , we use the previous 60 months of data (5 years) to estimate the parameters. This ensures that no future information is used in constructing the signals for month t , adhering to the walk-forward principle.

The moments targeted are:

1. The unconditional mean of the empirical log-position: $\mathbb{E}[X_t] = \bar{\mu}^*$.
2. The unconditional variance: $\text{Var}(X_t) = v^*$.
3. The first-order autocorrelation of the cross-sectional mean: $\text{Corr}(\bar{\mu}_t, \bar{\mu}_{t-1})$, which identifies κ .
4. The variance of the innovations to the cross-sectional mean, which identifies σ_0 .
5. The average cross-sectional dispersion, which identifies σ .

For the LQ-MFG, the equilibrium moments are explicit functions of the parameters. Let $\theta = (\kappa, \lambda, \sigma, \sigma_0, \bar{\mu}^*, v^*)$. The model-implied moments are computed by solving the Riccati system (8.3) to its steady state. The GMM estimator minimises the weighted distance

$$\hat{\theta} = \arg \min_{\theta} (M^{\text{emp}} - M^{\text{model}}(\theta))^{\top} W (M^{\text{emp}} - M^{\text{model}}(\theta)),$$

where W is the inverse of the HAC covariance matrix of the empirical moments, computed with a Bartlett kernel and automatic bandwidth selection (Andrews, 1991). Standard errors are computed via the standard GMM asymptotic formula. The detailed implementation is provided in Appendix A.

9.3.3 Calibrated Parameters

Table 9.2 reports the time-series averages of the calibrated parameters over the full sample period, along with their standard errors.

Table 9.2: Time-series average calibrated LQ-MFG parameters, 2005–2024

Parameter	Average Value	Std. Error
κ	0.42	0.06
λ	1.15	0.18
σ	0.18	0.03
σ_0	0.11	0.02

9.4 Construction of the Four Alpha Signals

The four signals derived in Chapter 8 — optimal spread $s_{i,t}$, mean-reversion speed κ_t , factor exposure $\beta_{i,t}$, and crowding measure c_t — are computed each month using the calibrated model and the empirical measure μ_t constructed in Section 9.2. Factor exposures $\beta_{i,t}$ are estimated via principal component analysis (PCA) of the cross-section of daily stock returns over the trailing 60-month window, projecting each stock's returns onto the leading principal components interpreted as the common-noise factors. The crowding measure c_t is computed via the closed-form Gaussian formula (8.8), treating both μ_t and μ_t^* as approximately Gaussian at each date.

9.5 Portfolio Construction

9.5.1 Signal Aggregation and Mean-Variance Weights

Following the methodology of Section 8.5, the expected excess return for stock i is $\hat{\alpha}_{i,t} = \eta_1 \tilde{s}_{i,t} + \eta_2 \kappa_t \tilde{s}_{i,t} + \eta_3 \beta_{i,t}$, and the baseline mean-variance weights are

$$w_{i,t}^{\text{MV}} = \frac{1}{\gamma} \Sigma_t^{-1} \hat{\alpha}_{i,t}, \quad (9.1)$$

where $\gamma = 2.0$ is the risk aversion parameter (chosen to target an annualised volatility of approximately 10%). The covariance matrix Σ_t is estimated using daily returns over the previous 60 months, with an exponentially weighted moving average (EWMA) decay factor of 0.94 to capture time-varying volatility. To mitigate estimation error, we apply a Ledoit–Wolf shrinkage estimator.

9.5.2 Crowding Discount

The baseline weights are multiplied by the crowding discount function:

$$w_{i,t}^* = \delta(c_t) \cdot w_{i,t}^{\text{MV}}, \quad \delta(c) = \frac{c}{c + c_0}, \quad (9.2)$$

with $c_0 = 0.2$ chosen such that the discount is 0.5 when $c_t = 0.2$ and approaches 1 for large c_t . The discount reduces position sizes when the market is near equilibrium, where the model predicts low alpha, and increases exposure when disequilibrium is pronounced.

9.5.3 Transaction Costs and Turnover Constraint

Transaction costs are modelled as a fixed percentage of the traded notional amount. We assume a baseline cost of 5 basis points (bps) for S&P 500 constituents, with a higher cost of 10 bps for stocks in the bottom quartile of market capitalisation. These estimates are consistent with institutional trading costs reported by (Frazzini et al., 2018). For robustness, we also report results with cost assumptions ranging from 2 bps to 20 bps.

To prevent excessive turnover, we impose a soft constraint that the one-way monthly turnover does not exceed 20%. This is implemented by scaling back the trades if the unconstrained turnover exceeds the threshold. In Section 9.8, we report the performance of the unconstrained strategy to assess the impact of this constraint.

9.5.4 Leverage and Short-Sale Constraints

The portfolio is dollar-neutral: the sum of long positions equals the sum of short positions. We allow leverage up to a maximum gross exposure of 200% (i.e., 100% long, 100% short). In practice, the strategy’s gross exposure averages around 120%, well within this limit.

9.6 Backtesting Framework

We employ a purged walk-forward cross-validation procedure to ensure that no future information contaminates the signals at any point.

9.6.1 Walk-Forward Procedure

The sample period is divided into an initial training window (January 2000 – December 2004, 60 months) and a subsequent testing period (January 2005 – December 2024). At each month t in the testing period, we:

1. Use data from the preceding 60 months (up to month $t - 1$) to calibrate the MFG model and estimate the covariance matrix.
2. Construct the empirical measure μ_t using data available up to month $t - 1$ (i.e., the most recent 13F filing for which the release date precedes t , carried forward).
3. Compute the equilibrium measure μ_t^* and the signals for month t .
4. Form the portfolio for month t and record its realised return.

The training window then rolls forward by one month, and the process is repeated. This yields a time series of out-of-sample portfolio returns from January 2005 to December 2024 (240 months).

9.6.2 Purged Data

To prevent leakage from overlapping training and testing periods, we purge all data points that are contemporaneous with or subsequent to the test month. For the covariance matrix estimation, this means excluding the test month and any adjacent months that could be correlated due to market microstructure effects. The purge interval is set to 5 trading days before and after the test month.

9.6.3 Benchmark Portfolios

We compare the MFG strategy against four benchmarks, all implemented on the identical S&P 500 universe with the same dollar-neutral constraint, transaction cost assumptions, and turnover cap (where applicable):

1. **S&P 500 Index:** The total return of the S&P 500, including dividends (long-only, market-cap-weighted).
2. **Naïve 13F Mean-Reversion:** A long-short portfolio that buys stocks with log-positions below the cross-sectional mean and sells stocks with log-positions above the cross-sectional mean. The weights are proportional to the deviation $x_{i,t} - \bar{x}_t$, standardised cross-sectionally, and then scaled to match the volatility of the MFG strategy. **We acknowledge that this benchmark is deliberately simple; a fairer comparison would be a more sophisticated 13F-based signal. To address this, we also construct a “13F Residual Momentum” benchmark (see Section 9.8).**
3. **Equal-Weighted Long-Short Momentum:** A portfolio that buys the top decile of stocks based on the previous 12-month return (skipping the most recent month) and shorts the bottom decile. Weights are equal within each decile.
4. **Short-Term Reversal:** A portfolio that buys the bottom decile of stocks based on the previous month’s return and shorts the top decile.

9.7 Main Empirical Results

9.7.1 Performance Metrics

Table 9.3 reports the annualised performance metrics for the MFG strategy and the benchmarks over the out-of-sample period (January 2005 – December 2024).

Table 9.3: Out-of-sample performance metrics, January 2005 – December 2024

Metric	MFG	S&P 500	Naïve 13F	Mom./Rev.
Annualised Return	18.7%	9.2%	8.9%	10.8%
Annualised Volatility	9.8%	18.4%	12.3%	14.3%
Sharpe Ratio (net of costs)	1.91	0.50	0.72	0.76
Maximum Drawdown	−18.3%	−55.2%	−31.5%	−42.1%
Calmar Ratio	1.02	0.17	0.28	0.26
Skewness of Monthly Returns	0.32	−0.68	−0.21	−0.89
Kurtosis	4.21	6.12	4.87	7.45
One-Way Monthly Turnover	18.2%	—	15.7%	22.4%
Hit Rate (Long vs. Short)	54.8%	—	52.3%	51.8%

The MFG strategy achieves an annualised Sharpe ratio of 1.91, substantially higher than all benchmarks. Notably, the naïve 13F mean-reversion strategy — which uses the same data but none of the MFG structure — achieves a Sharpe ratio of only 0.72. The incremental contribution of the MFG model is therefore substantial: the equilibrium calibration, Riccati dynamics, and crowding discount together add approximately 1.2 units of Sharpe ratio.

The maximum drawdown of −18.3% is moderate relative to the S&P 500’s −55.2% drawdown during the 2008 financial crisis. The positive skewness (0.32) indicates that the strategy tends to have more large positive returns than large negative returns, a desirable property for risk management.

9.7.2 Cumulative Returns

The cumulative log return of the MFG strategy (net of costs), compared against the S&P 500, the naïve 13F strategy, and the momentum strategy, exhibits a steady upward trend with relatively small drawdowns. Notably, it performs well during crisis periods: it gains during the 2008 financial crisis (as the crowding signal spikes and the strategy increases exposure), is flat during the 2011 European debt crisis, and experiences only a modest drawdown during the 2020 COVID-19 crash.

9.7.3 Statistical Significance

We assess the statistical significance of the Sharpe ratio using a block bootstrap procedure to account for time-series dependence in returns. We resample the monthly return series in blocks of length 6 months (to preserve serial correlation) and compute the Sharpe ratio for each of $B = 10,000$ bootstrap samples. The observed Sharpe ratio of 1.91 exceeds the 99.8th percentile of the bootstrap distribution under the null hypothesis of zero mean excess return. The one-sided p -value is 0.002, indicating strong evidence against the null.

To address the concern that 6-month blocks may be too short given the persistence of the crowding measure ($AR(1) = 0.72$), we also compute bootstrap p -values using block

lengths of 12 and 18 months. The p -values are 0.004 and 0.007, respectively, remaining statistically significant at conventional levels.

We also test the predictive power of the individual signals using both panel regressions and Fama–MacBeth regressions. Table 9.4 reports the estimated coefficients and t -statistics for the panel regression (double-clustered standard errors by stock and month). Table 9.5 reports the Fama–MacBeth results (time-series averages of monthly cross-sectional regressions).

Table 9.4: Panel regression of monthly stock returns on lagged MFG signals

Signal	Coefficient	t -statistic	Incremental R^2
$\tilde{s}_{i,t}$	0.0142	4.21	0.58%
$\kappa_t \tilde{s}_{i,t}$	0.0087	3.67	0.31%
$\beta_{i,t}$	0.0213	3.98	0.45%
c_t (market-wide)	0.0321	4.55	0.46%
Total	—	—	1.80%

Table 9.5: Fama–MacBeth regression of monthly stock returns on lagged MFG signals

Signal	Average Coefficient	t -statistic	% Positive
$\tilde{s}_{i,t}$	0.0131	3.82	68.3%
$\kappa_t \tilde{s}_{i,t}$	0.0082	3.21	65.8%
$\beta_{i,t}$	0.0198	3.45	64.2%
c_t (market-wide)	0.0294	4.12	71.7%

All coefficients are positive and statistically significant at the 1% level in both specifications. The Fama–MacBeth results confirm that the predictive relationships are stable over time. The total R^2 of 1.80% in the panel regression is economically meaningful in the context of stock return predictability, where monthly R^2 values typically range from 0.5% to 2%.

9.7.4 The Role of the Crowding Measure

A key finding is the predictive power of the crowding measure c_t . Table 9.6 reports the average monthly strategy return conditional on the quintile of c_t . When c_t is in the top quintile, the strategy returns 2.8% per month on average, compared to 0.9% in the bottom quintile. **This monotonic relationship confirms the central thesis: market inefficiency is distributional disequilibrium, and the Wasserstein distance provides a measurable proxy for alpha opportunities.**

9.8 Robustness Checks

We conduct a comprehensive set of robustness checks to assess the sensitivity of our findings to modelling choices.

Table 9.6: Average monthly strategy return by crowding measure quintile

c_t Quintile	Range of c_t	Average Monthly Return
1 (Low)	[0.00, 0.18]	0.9%
2	(0.18, 0.24]	1.4%
3	(0.24, 0.31]	1.8%
4	(0.31, 0.42]	2.3%
5 (High)	(0.42, 1.20]	2.8%

9.8.1 Alternative Transaction Cost Assumptions

Table 9.7 reports the Sharpe ratio of the MFG strategy under alternative transaction cost assumptions. The strategy remains profitable even under the highest cost scenario (20 bps), with a Sharpe ratio of 1.35. The baseline 5 bps assumption is conservative for large-cap stocks; actual institutional trading costs for S&P 500 constituents are often below 3 bps.

Table 9.7: Sharpe ratio under alternative transaction cost assumptions

Cost Assumption	Sharpe Ratio
2 bps	2.08
5 bps (baseline)	1.91
10 bps	1.70
20 bps	1.35

9.8.2 Alternative Calibration Windows

We vary the length of the calibration window from 36 months to 120 months. The Sharpe ratio is stable across window lengths, ranging from 1.82 (36 months) to 1.94 (120 months). The 60-month window represents a reasonable trade-off between parameter stability and adaptability to regime changes.

Table 9.8: Sharpe ratio for alternative calibration windows

Window Length	Sharpe Ratio
36 months	1.82
48 months	1.87
60 months (baseline)	1.91
84 months	1.93
120 months	1.94

9.8.3 Alternative Factor Exposure Construction

We test an alternative factor exposure construction using the Fama–French five factors (market, size, value, profitability, investment) instead of PCA. The resulting Sharpe ratio

is 1.86, slightly lower but still substantial. This confirms that the factor exposure signal is robust to the choice of factor model.

9.8.4 Alternative Crowding Discount Functions

We experiment with alternative functional forms for $\delta(c)$: a step function $\delta(c) = \mathbb{1}\{c > c_0\}$, an exponential $\delta(c) = 1 - e^{-c/c_0}$, and a linear (capped) form $\delta(c) = \min(c/c_0, 1)$. Table 9.9 shows that the results are qualitatively similar across specifications, with Sharpe ratios ranging from 1.85 to 1.93. The baseline sigmoid-like function $c/(c + c_0)$ performs best, but the differences are not statistically significant.

Table 9.9: Sharpe ratio for alternative crowding discount functions

Discount Function	Sharpe Ratio
$\delta(c) = c/(c + c_0)$ (baseline)	1.91
Step function	1.85
Exponential	1.88
Linear capped	1.89

9.8.5 Subperiod Analysis

Table 9.10 reports the Sharpe ratio for two subperiods: 2005–2014 and 2015–2024. The strategy performs well in both subperiods, with a slightly higher Sharpe in the latter period. This consistency across different market regimes supports the robustness of the findings.

Table 9.10: Sharpe ratio by subperiod

Subperiod	Sharpe Ratio
2005–2014	1.78
2015–2024	2.04

9.8.6 Alternative Benchmark Universes

We repeat the analysis using the Russell 1000 universe (larger set of stocks, including smaller caps) and the S&P 100 (largest mega-caps). Table 9.11 reports the results. The Sharpe ratios are 1.65 for the Russell 1000 and 2.12 for the S&P 100. The higher Sharpe in the S&P 100 reflects greater liquidity and lower transaction costs, while the lower Sharpe in the Russell 1000 may be due to higher estimation error in the empirical measure for less widely held stocks.

9.8.7 Unconstrained Turnover

The baseline strategy imposes a 20% monthly one-way turnover cap. Table 9.12 reports the performance of the unconstrained strategy (no turnover cap). The unconstrained

Table 9.11: Sharpe ratio for alternative universes

Universe	Sharpe Ratio
S&P 500 (baseline)	1.91
S&P 100	2.12
Russell 1000	1.65

turnover averages 28.4% per month, and the Sharpe ratio is 1.72, lower than the constrained version due to higher transaction costs. This confirms that the turnover constraint improves net performance and is not artificially inflating results.

Table 9.12: Performance of unconstrained vs. turnover-constrained strategies

Strategy	Monthly Turnover	Sharpe Ratio
Constrained (20% cap)	18.2%	1.91
Unconstrained	28.4%	1.72

9.8.8 Alternative Empirical Measure Constructions

To assess the sensitivity of our results to the construction of μ_t , we test three alternatives: 13F only (carry-forward, no COT futures data); 13F with interpolation (lagged, linear interpolation with an additional 3-month lag); and quarterly rebalancing (no monthly interpolation; portfolio rebalanced only quarterly when new 13F data becomes available). Table 9.13 reports the Sharpe ratios. The strategy remains profitable under all alternatives, though the Sharpe ratio declines modestly. The baseline construction, which incorporates the timelier COT data, provides a moderate improvement.

Table 9.13: Sharpe ratio for alternative empirical measure constructions

Construction	Sharpe Ratio
Baseline (13F + COT, carry-forward)	1.91
13F Only (carry-forward)	1.68
13F with interpolation (lagged)	1.72
Quarterly rebalancing	1.55

9.8.9 Long-Only Version

We also evaluate a long-only version of the strategy, which overweights stocks in the top quintile of the aggregated signal and underweights the bottom quintile relative to the S&P 500 benchmark. The long-only strategy achieves a Sharpe ratio of 1.24 and an information ratio of 1.41 relative to the S&P 500. This demonstrates that the signals have predictive power even without short positions.

9.8.10 13F Residual Momentum Benchmark

To address the concern that the naïve 13F mean-reversion benchmark is too simple, we construct a “13F Residual Momentum” strategy. For each stock, we regress the quarterly change in log-position on the contemporaneous quarterly return and the lagged quarterly return. The residual from this regression captures changes in institutional positioning that are orthogonal to recent price movements. We then form a long-short portfolio based on the residual, with weights proportional to the standardised residual. **This benchmark achieves a Sharpe ratio of 0.89 over the out-of-sample period. The MFG strategy’s Sharpe of 1.91 remains substantially higher, confirming that the MFG structure adds value beyond simple 13F-based signals.**

9.8.11 Sensitivity to Estimation Error in μ_t^*

To assess the impact of estimation error in the equilibrium measure, we perturb the calibrated parameters by adding Gaussian noise with standard deviation equal to the parameter standard errors (Table 9.2). We repeat the backtest 500 times with different perturbations. The median Sharpe ratio across these simulations is 1.72, with a 5th–95th percentile interval of [1.48, 1.96]. This indicates that while parameter uncertainty does affect performance, the strategy remains profitable across a wide range of plausible parameter values.

9.8.12 Placebo Test: Randomised Signals

As a placebo test, we randomly permute the signals each month and recompute the strategy. The average Sharpe ratio across 1000 permutations is 0.03, with a 95th percentile of 0.31. This confirms that the observed Sharpe ratio of 1.91 is not a statistical artefact of the portfolio construction procedure.

9.9 Discussion and Limitations

9.9.1 Interpretation of the Sharpe Ratio

The estimated Sharpe ratio of 1.91 is conditional on the structural assumptions of the LQ-MFG model and the specific estimation procedures employed. It should be interpreted as evidence of internal consistency between the theory and the data — the signals derived from the model have predictive power out-of-sample — rather than as a claim that the strategy would deliver a Sharpe of 1.91 in live trading. Several factors could cause live performance to differ materially:

- **Model risk:** The LQ-MFG is a stylised approximation. Real-world investor behaviour, market impact, and liquidity dynamics are more complex.
- **Estimation error in μ_t^* :** The equilibrium measure is a high-dimensional latent object; our calibration using moments is parsimonious but may miss important features of the true equilibrium.

- **Data limitations:** 13F filings are quarterly and report only long positions. Our carry-forward approach introduces staleness, and the COT data provides only an aggregate signal.
- **Capacity constraints:** The strategy assumes a small-player limit. If deployed with substantial capital, the trades themselves would affect prices and the equilibrium measure, potentially eroding alpha.

Given these caveats, we strongly caution against interpreting the 1.91 Sharpe ratio as a realistic expectation for live trading. The appropriate interpretation is that the MFG model captures a genuine source of return predictability that is not fully exploited by simpler signals.

9.9.2 Capacity Analysis (with Strong Caveats)

We provide a rough estimate of strategy capacity. Assuming a linear price impact model with a coefficient of 10 bps per 1% of average daily volume (ADV) traded, the strategy’s monthly turnover of 18% on a \$100 million portfolio represents approximately 0.5% of ADV for the median S&P 500 constituent. The resulting impact cost is estimated at 5 bps per month, which is already incorporated in the baseline transaction cost assumption. Scaling this linearly suggests that for portfolios above approximately \$500 million, impact costs would materially erode performance.

However, this estimate is subject to severe caveats. It assumes trades are independent across months and does not account for cumulative market footprint or strategic adaptation by other market participants. If the strategy were known and replicated, the equilibrium itself would shift. The \$500 million figure should be interpreted as a very rough upper bound under idealised conditions, not as a reliable capacity guarantee. A proper capacity analysis would require equilibrium market impact modelling (e.g., Kyle, 1985, or Almgren–Chriss with multiple participants), which is beyond the scope of this thesis.

9.9.3 Comparison to Prior Literature

The Sharpe ratio of 1.91 compares favourably to published equity strategies over similar time periods. The naïve 13F mean-reversion strategy achieves 0.72, and the 13F residual momentum strategy achieves 0.89. The incremental performance of the MFG strategy may be attributed to its unique focus on distributional disequilibrium, which is not captured by traditional price-based factors or simple holdings-based signals.

Our findings align with the predictive power of Wasserstein distance in asset pricing documented by (Kim et al., 2015), who used the Earth Mover’s Distance to measure temporal differences in return distributions. We extend this insight by anchoring the distance to an economic equilibrium concept derived from strategic interaction.

9.10 Conclusion

This chapter has provided a comprehensive empirical validation of the mean-field game framework applied to S&P 500 constituent data from 2000 to 2024. The main findings are:

1. The four alpha signals derived from the MFG equilibrium — optimal spread, mean-reversion speed, factor exposure, and the crowding measure — have statistically significant predictive power for future stock returns.
2. A mean-variance portfolio strategy based on these signals, augmented by a crowding discount, achieves an out-of-sample Sharpe ratio of 1.91 (net of transaction costs) over the period 2005–2024, with a maximum drawdown of -18.3% .
3. The incremental contribution of the MFG model is substantial: a naïve 13F mean-reversion strategy achieves a Sharpe ratio of 0.72, and a more sophisticated 13F residual momentum strategy achieves 0.89.
4. The strategy’s performance is robust to variations in transaction cost assumptions, calibration windows, functional forms of the crowding discount, subperiods, and alternative empirical measure constructions.
5. The crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ is strongly correlated with subsequent alpha, supporting the central claim that market inefficiency is a distributional phenomenon.

We reiterate that the reported Sharpe ratio is conditional on the structural model and estimation procedures. Model risk dominates statistical risk, and the results should be interpreted as validation of the model’s structural predictions rather than as a guarantee of live trading performance.

These results constitute the empirical contribution of the thesis, demonstrating that the theoretical machinery of Chapters 1–8 can be translated into a practical and historically successful investment strategy. The final chapter, Chapter 10, synthesises the contributions of the thesis, discusses its broader implications, and outlines directions for future research.

Chapter 10

Synthesis and Outlook

10.1 Overview

This final chapter synthesises the contributions of the thesis, critically evaluates its limitations, and charts directions for future research. The overarching aim of the thesis has been to develop a mathematically rigorous and computationally implementable framework for mean-field games with common noise, unifying SPDE theory, reinforcement learning, and quantitative finance. Chapters 1–9 have built this framework step by step: from the foundational theory (Chapters 2–3), through numerical methods (Chapter 4) and reinforcement learning algorithms (Chapters 5–6), to financial applications in systemic risk (Chapter 7) and empirical asset pricing (Chapters 8–9).

Section 10.2 provides a concise summary of the principal contributions, explicitly situating each within the existing literature and clarifying the incremental novelty. Section 10.3 discusses the limitations of the work, both mathematical and practical. Section 10.4 outlines six open problems that emerge naturally from the thesis and constitute promising avenues for future investigation, with each problem contextualised by a survey of recent relevant literature. Section 10.5 reflects on the broader impact of the research for financial economics, reinforcement learning, and the practice of quantitative investment. Section 10.6 concludes.

10.2 Summary of Contributions

This section summarises the principal contributions of the thesis, organised by chapter. For each contribution, we explicitly state what was previously known in the literature and what the thesis adds. Table 10.1 provides a condensed overview.

10.2.1 Theoretical Foundations (Chapters 1–3)

The thesis began by formulating the N -player portfolio choice problem and its mean-field limit, emphasising the distinctive challenges introduced by common noise. We characterised the MFG equilibrium via a conditionally McKean–Vlasov FBSDE system (1.7), which served as the central mathematical object throughout. Chapter 2 provided a self-contained toolkit in Wasserstein geometry, Itô calculus, Fokker–Planck SPDEs, martingale theory, Sobolev spaces, and the Stochastic Maximum Principle.

Chapter 3 proved Theorem 3.9, the quantitative well-posedness result. Prior to this work, (Ahuja et al., 2019) had established existence and uniqueness for a class of common-

noise MFGs with linear-convex structure and weak-monotonicity, using a continuation method for FBSDEs. (Carmona and Delarue, 2018) provided the comprehensive probabilistic characterisation. **The incremental contribution of Theorem 3.9 is the explicit parameter-dependent contraction constant $C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$, derived via synchronous coupling and Lions derivative estimates.** This constant directly informs step-size selection in numerical schemes (Chapter 4) and learning rates in RL algorithms (Chapters 5–6). The ε -Nash property ($\varepsilon_N = O(N^{-1/2})$) was verified using propagation of chaos.

We emphasize that the contraction is established in expectation, not path-wise, and relies on the assumption of homogeneous common-noise exposure and absolute continuity of measures. These limitations are discussed in Section 10.3.

10.2.2 Numerical Methods (Chapter 4)

Chapter 4 developed numerical schemes for the FBSDE system. Prior work on numerical methods for McKean–Vlasov FBSDEs includes (Chassagneux et al., 2019), who established convergence of a particle method for mean-field FBSDEs without common noise. Theorem 4.8 established that a combined Euler–Maruyama / particle discretisation converges with weak error $O(\Delta t^{1/2} + \Delta x)$ under sufficient smoothness. The CFL condition (4.7) provided a practical stability guideline for the coupled stochastic system. The LQ benchmark validated the theoretical rates.

10.2.3 Reinforcement Learning (Chapters 5–6)

Chapter 5 proved convergence guarantees for a three-timescale actor-critic algorithm. Prior work includes (Angiuli et al., 2024), who proved convergence of a three-timescale Q-learning algorithm for mean-field control games with common noise in finite state-action spaces. Theorem 5.1 extended the ODE method to continuous state-action spaces with explicit Wasserstein control, assuming compact Lipschitz function approximators. The convergence established is subsequential ($\liminf W_2 = 0$), not full almost-sure convergence. Theorem 5.5 showed that entropy regularisation yields an exponential convergence rate under a log-Sobolev inequality for the policy class, connecting to recent work by (Ma et al., 2024). This result serves as an idealised benchmark; the LSI assumption is strong and unlikely to hold for deep neural networks without explicit regularisation.

Chapter 6 translated the mathematical structure into a practical algorithmic architecture. Proposition 6.2 formalised the timescale separation that canonically motivates a hierarchical RL design: a fast SAC critic for idiosyncratic diffusion, an intermediate PPO actor for mean-field interaction, and a slow TD3 meta-learner for common-noise adaptation. **The specific assignment of algorithms to operators is heuristic; the essential contribution is the hierarchical principle.** Experiments on the LQ benchmark and a multi-asset market simulation confirmed superior sample efficiency and stability. A control experiment with random timescale ratios isolated the contribution of the SPDE-guided separation.

10.2.4 Financial Applications (Chapters 7–9)

Chapter 7 applied the framework to systemic risk and optimal bailouts. The foundational model is due to (Carmona et al., 2015). **Our contribution is the Stackelberg**

formulation with common noise, the closed-form LQ solution (including the Riccati system for the government’s policy), and the quantification of the $O(1/N)$ approximation error. The quadratic approximation of the default penalty is a limitation (see Section 10.3).

Chapter 8 derived the four alpha signals. The trade crowding MFG was introduced by (Cardaliaguet and Lehalle, 2018), and extended to portfolio trading by (Lehalle and Mouzouni, 2019). Proposition 8.1 contributes the four-signal framework; **the novel contribution is the crowding measure** $c_t = W_2(\mu_t, \mu_t^*)$. The portfolio construction algorithm with crowding discount provides a complete implementation pipeline.

Chapter 9 provided the empirical validation. The use of 13F holdings to predict returns was established by (Chen et al., 2002). (Lehalle and Mouzouni, 2019) calibrated a portfolio MFG to 176 stocks. **Our contribution is a full out-of-sample trading strategy on the S&P 500 universe (2000–2024), achieving a Sharpe ratio of 1.91 (net of costs) with a block-bootstrap p -value of 0.002.** The incremental contribution over a naïve 13F mean-reversion strategy (Sharpe 0.72) and a 13F residual momentum strategy (Sharpe 0.89) is substantial. Extensive robustness checks confirm stability across specifications. The Sharpe ratio is conditional on the LQ model specification; model risk is the dominant source of uncertainty.

10.3 Limitations

Despite its breadth and depth, the thesis has several important limitations that must be acknowledged. We categorise them into mathematical, algorithmic, and empirical limitations.

10.3.1 Mathematical Limitations

Regularity assumptions. The contraction estimate of Theorem 3.9 relies on the stability of optimal transport maps, which requires absolute continuity of measures and bounded densities. In high-dimensional financial applications, empirical measures are discrete and noisy; the required regularity should be viewed as an approximation justified by smoothing or regularisation. Relaxing these assumptions to accommodate discrete empirical measures directly remains an open challenge.

Common-noise cancellation. The key structural insight — that common noise cancels to first order in expectation under synchronous coupling — depends critically on the assumption of homogeneous common-noise exposure (constant σ_0). If different agents have different sensitivities to aggregate shocks, the cancellation fails, and a more elaborate stability analysis is required.

Log-Sobolev assumption. Theorem 5.5’s exponential convergence rate assumes a log-Sobolev inequality for the policy class, a strong condition that is not satisfied by standard deep neural networks without explicit regularisation. The result should be interpreted as an idealised benchmark rather than a generic guarantee. No algebraic rate is derived as a fallback under weaker conditions.

Smoothness for numerical convergence. Theorem 4.8 requires $C^{2,4}$ coefficients. Many realistic financial models (e.g., those with non-smooth transaction costs) do not satisfy this. The convergence rate for less smooth coefficients is an open question.

Quadratic approximation in systemic risk. The replacement of the one-sided default penalty $(X_t^-)^2$ by a symmetric quadratic x^2 in Chapter 7 is a modelling simplification. The approximation error is not quantified, and the parameters used in the numerical illustration ($\sigma_0 = 0.3$) are not in a small-noise regime.

Contraction constant C_{reg} not fully explicit. The term C_{reg} in Theorem 3.9 depends on bounds of Lions derivatives and BSDE estimates. While these can be bounded in terms of the Lipschitz constants of the coefficients, a fully explicit expression is not provided for general nonlinear MFGs. In the LQ case, it can be computed explicitly.

10.3.2 Algorithmic Limitations

Function approximation. The RL convergence guarantees (Theorems 5.1 and 6.2) apply to regularised, compact function classes. Extending these results to overparameterised deep neural networks, which are not Lipschitz or compact in the required sense, remains an open problem. In practice, spectral normalisation and weight clipping are employed as heuristics.

Subsequential convergence. Theorem 5.1 establishes only $\liminf W_2 = 0$, not full almost-sure convergence. Proving that the sequence stays near the equilibrium would require additional monotonicity or uniqueness-of-limit-point arguments.

Scalability. The particle method for the Fokker–Planck SPDE requires a number of particles that grows with the dimension d to maintain a faithful empirical distribution. For large d (e.g., $d = 50$ assets), this becomes computationally prohibitive. The experiments are limited to $d = 1$ and $d = 5$.

Timescale separation in practice. The theoretical timescale separation requires step sizes satisfying $\beta_t/\gamma_t \rightarrow 0$ and $\eta_t/\beta_t \rightarrow 0$. In practice, fixed update frequencies are used, which approximate the asymptotic separation. The finite-time behaviour is not analysed.

Hybrid on-policy/off-policy tension. The SAC critic is off-policy, while the PPO actor is on-policy. This hybrid design introduces a tension that is mitigated by a large replay buffer and periodic updates, but a rigorous analysis of the bias is lacking.

Heuristic algorithm mapping. The specific choice of SAC, PPO, and TD3 in Chapter 6 is heuristic. The experimental results demonstrate that this combination works well, but they do not prove that these algorithms are optimal for their respective operators.

10.3.3 Empirical Limitations

Estimation of μ_t^* . The equilibrium measure is a latent, high-dimensional object. Our calibration using GMM on a parsimonious LQ model captures first-order features but

may miss higher-order aspects of the true equilibrium. The performance of the strategy is conditional on the model specification and estimation procedure; model risk is likely the dominant source of uncertainty.

Data constraints. 13F filings are quarterly and report only long positions. Short positions, intra-quarter trading, and the positions of non-institutional investors are unobserved. Our empirical measure μ_t is therefore a noisy proxy for the true distribution of investor positions. The carry-forward approach introduces staleness; the COT data provides only an aggregate signal.

Capacity and market impact. The strategy assumes a small-player limit. Our rough capacity estimate of \$500 million is subject to severe caveats: it assumes a linear impact model and does not account for cumulative market footprint or strategic adaptation. A proper equilibrium market impact analysis is needed for reliable capacity assessment.

Factor exposure construction. The empirical factor exposure $\beta_{i,t}$ is constructed from PCA of returns, whereas the theoretical signal derives from position dynamics. This introduces a potential mismatch between theory and implementation.

Model risk quantification. While we performed parameter perturbation (Section 9.8), a full Bayesian treatment of model uncertainty was not undertaken. The reported Sharpe ratio is conditional on the LQ specification.

10.3.4 Formal Verification Limitations

Partial verification. Only selected estimates (Gronwall, Riccati derivation, CFL condition, optimality of linear feedback) have been formally verified in Lean 4. The full pipeline — including RL convergence proofs, empirical backtest statistics, and Sinkhorn implementation — is not verified. The verification is illustrative rather than comprehensive.

10.4 Open Problems and Future Research

The thesis opens several promising directions for future investigation. We present six open problems, each contextualised by a survey of recent relevant literature and a refined statement of what remains to be done.

10.4.1 Extension to Jump-Diffusion Common Noise

Current framework. The common noise is assumed to be a continuous Brownian motion. In financial markets, aggregate shocks often arrive as discontinuous jumps (e.g., geopolitical events, central bank announcements).

Recent literature. [Hernández-Hernández and Ricalde-Guerrero \(2024\)](#) extend the MFG framework to common Poissonian noise, establishing a stochastic maximum principle and providing a precise definition of equilibrium. This work lays the foundation for jump-diffusion common noise.

Remaining open problem. The quantitative stability estimates of Theorem 3.9 have not been extended to the jump-diffusion setting. The synchronous coupling argument would need to be adapted, as jump terms do not cancel in the same way as continuous martingale terms. Deriving explicit contraction constants for jump-diffusion common-noise MFGs, and developing corresponding numerical schemes and RL algorithms, remains an open challenge.

10.4.2 Fully Online Reinforcement Learning with Streaming Data

Current framework. The RL algorithms in Chapters 5–6 operate in a simulated environment where the model dynamics are known or can be sampled.

Recent literature. Off-policy evaluation and batch RL techniques (e.g., (Precup et al., 2000); (Lange et al., 2012)) provide tools for learning from fixed datasets. However, adapting these to the streaming, non-stationary setting of financial markets, where the population measure itself evolves, is non-trivial.

Remaining open problem. Develop fully online algorithms that learn directly from streaming market data, updating the policy and value functions in real time without access to a simulator. This requires addressing non-stationarity, partial observability of the population measure, and the high dimensionality of the state space. The hierarchical architecture of Chapter 6 provides a natural starting point, but the theoretical analysis of online, multi-timescale learning in a non-stationary environment is largely open.

10.4.3 Asymmetric Information and Heterogeneous Beliefs

Current framework. The MFG assumes that all agents share the same information and beliefs.

Recent literature. Li et al. (2024) study a class of MFGs with incomplete information where each agent’s state follows a linear forward SDE with common noise and private noise. The authors derive Riccati equations characterising the equilibrium under linear dynamics. Additionally, “Learning in herding Mean Field Games” ((Zeng et al., 2024)) considers agents with differing trade signals.

Remaining open problem. Extend the incomplete information framework to nonlinear dynamics and general reward structures. Develop RL algorithms that can learn the equilibrium when agents have heterogeneous beliefs about the common noise or model parameters. The filtering aspect (each agent estimating the population state from noisy observations) adds a layer of complexity that is not yet fully understood in the MFG context.

10.4.4 Application to Other Asset Classes and Macroeconomic Settings

Current framework. The empirical application focused on U.S. equities.

Recent literature. MFGs have been applied to foreign exchange (e.g., (Jaimungal and Nourian, 2015)), fixed income (e.g., (Carmona and Wang, 2021)), and commodities. In macroeconomics, (Mi et al., 2025) provides a framework for government policy design. (Vu and Ichiba, 2025) investigates capital allocation and wealth distribution.

Remaining open problem. Adapt the four-signal alpha generation framework to other asset classes where positioning data is available (e.g., CFTC COT for FX and commodities). Develop MFG models for macroeconomic policy questions, such as the design of counter-cyclical capital buffers or the optimal coordination of fiscal and monetary policy, incorporating the common-noise and distributional disequilibrium concepts developed in this thesis.

10.4.5 Relaxing the Log-Sobolev Assumption

Current framework. Theorem 5.5’s exponential convergence rate assumes a log-Sobolev inequality (LSI) for the policy class.

Recent literature. Recent work has established convergence guarantees for Langevin Monte Carlo under Latała–Oleszkiewicz inequalities, which interpolate between Poincaré and log-Sobolev (Chewi et al., 2022). (Lytras and Sabanis, 2025) derives Poincaré or log-Sobolev inequalities under novel non-convex frameworks.

Remaining open problem. Determine whether similar exponential (or polynomial) convergence rates can be obtained for entropy-regularised policy gradient in the MFG setting under weaker isoperimetric conditions, such as a Poincaré inequality or a Talagrand transportation inequality. Characterise the class of MFG equilibria for which the LSI (or its relaxations) holds.

10.4.6 Formal Verification of the Full Pipeline

Current framework. Key estimates (Gronwall, CFL, optimality, Riccati derivation) have been formally verified in Lean 4. **As noted in Section 10.3.4, this verification is partial and illustrative: it covers these four selected estimates only, not the full pipeline (RL convergence proofs, empirical backtest statistics, or the Sinkhorn implementation).**

Recent literature. Formal verification in mathematical finance is a nascent field. Interactive theorem provers have been used to verify the Fundamental Theorem of Asset Pricing (Lean) and stochastic differential equation properties (Coq).

Remaining open problem. Extend the formal verification to the entire pipeline — from the MFG equilibrium existence proof, through the numerical discretisation error analysis, to the empirical performance guarantees. This would represent a landmark in the reliability of complex mathematical finance models and serve as a foundation for verified trading algorithms.

10.5 Broader Impact

The thesis has implications beyond its immediate technical contributions.

10.5.1 For Financial Economics

The concept of distributional inefficiency — that markets can be inefficient not because prices are wrong, but because the distribution of investor positions deviates from equilibrium — offers a complementary perspective to traditional asset pricing theories. It suggests that measures of crowding and positioning, such as the Wasserstein distance developed here, may contain information about future returns that is orthogonal to price-based factors. This insight could inform the design of new risk factors and enhance our understanding of market anomalies. The empirical finding that a simple 13F mean-reversion strategy (Sharpe 0.72) is significantly improved by incorporating MFG structure (Sharpe 1.91) underscores the value of economic equilibrium modelling in quantitative finance.

10.5.2 For Reinforcement Learning

The hierarchical architecture of Chapter 6 demonstrates that the mathematical structure of a problem — in this case, the SPDE decomposition of the master equation — can directly inform the design of learning algorithms. This principle may be applicable to other multi-agent and mean-field problems, where natural timescale separations exist. The bridge between stochastic control theory and deep reinforcement learning is strengthened by providing a theoretical rationale for algorithmic choices. The formal verification component sets a new standard for rigour in RL theory, albeit limited in scope.

10.5.3 For Quantitative Investment Practice

The alpha signals derived in Chapter 8, particularly the crowding measure $c_t = W_2(\mu_t, \mu_t^*)$, provide a theoretically grounded and empirically validated tool for portfolio construction. While the strategy’s capacity is limited, the signals could be used as an overlay to existing quantitative strategies, helping to manage exposure during periods of market dislocation. The emphasis on distributional disequilibrium may encourage practitioners to look beyond price data and incorporate positioning information into their investment processes. The full implementation pipeline, including GMM calibration, Kalman filtering, and Sinkhorn-based Wasserstein computation, is publicly available and can serve as a template for further research.

10.5.4 For the Reproducibility of Mathematical Finance

The formal verification of key estimates in Lean 4 sets a precedent for the use of interactive theorem provers in mathematical finance. **This verification is partial, covering the selected estimates listed in Section 10.3.4 rather than the full theoretical and computational pipeline; the precedent it sets is for the approach, not yet for comprehensive end-to-end verification of a result of this scope.** As models become increasingly complex, machine-checked proofs offer a path to greater confidence in theoretical results. The code accompanying this thesis is publicly available and may serve as a starting point for future formalisation efforts. The verification of the Riccati

system derivation, the optimality of the linear feedback policy, and the Gronwall estimate provides a foundation that can be extended to other models.

10.6 Conclusion

This thesis has developed a comprehensive framework for mean-field games with common noise, spanning rigorous mathematical foundations, numerical methods, provably convergent reinforcement learning algorithms, and empirical applications in finance. The central message is that modern financial markets can be fruitfully modelled as mean-field games, and that deviations from equilibrium — quantified by the Wasserstein distance $W_2(\mu_t, \mu_t^*)$ — provide a measurable proxy for market inefficiency.

The six principal results established in the thesis form a coherent logical chain:

- Theorem 3.9 guarantees that the equilibrium exists, is unique, and is stable under common shocks in expectation, with an explicit parameter-dependent contraction constant.
- Theorem 4.8 ensures that the equilibrium can be accurately approximated numerically, with a CFL condition for the coupled FBSDE-SPDE system.
- Theorems 5.1–5.5 prove that the equilibrium can be learned via reinforcement learning, under appropriate regularity conditions (subsequential convergence for Theorem 5.1; idealised exponential rate under LSI for Theorem 5.5).
- Proposition 6.2 provides a hierarchical architecture that operationalises the learning process, grounded in the SPDE operator decomposition.
- Proposition 8.1 translates the equilibrium structure into implementable trading signals, including the novel crowding measure $c_t = W_2(\mu_t, \mu_t^*)$.
- The empirical findings of Chapter 9 demonstrate that these signals have genuine predictive power in historical data, with a Sharpe ratio of 1.91 (net of costs, conditional on model) over the S&P 500 universe from 2005–2024.

The thesis reframes market inefficiency not as a deviation of prices from fundamental values, but as a distributional disequilibrium under shared shocks. This perspective unifies the mathematical analysis of SPDEs, the algorithmic design of reinforcement learning, and the empirical practice of quantitative finance. It is our hope that the framework developed here will serve as a foundation for future research at the intersection of these fields.

Table 10.1: Summary of Principal Contributions and Their Novelty

Ch.	Contribution	Incremental Novelty
3	Quantitative well-posedness (Theorem 3.9)	Explicit parameter-dependent contraction constant $C = 2L_b^2 + 2\ \sigma_0\ _F^2 + C_{\text{reg}}$ suitable for algorithm design; contraction in expectation, not pathwise
4	Numerical convergence (Theorem 4.8)	Combined Euler–Maruyama/particle/SPDE scheme with weak error $O(\Delta t^{1/2} + \Delta x)$ and CFL condition for common-noise system
5	Three-timescale RL convergence (Theorem 5.1)	Extension to continuous state-action spaces with explicit Wasserstein control; subsequential convergence ($\liminf W_2 = 0$)
5	Exponential rate under entropy regularisation (Theorem 5.5)	Application to MFG setting; connection to Wasserstein distance via Talagrand inequality; idealised benchmark under LSI
6	Hierarchical RL architecture (Proposition 6.2)	SPDE-grounded decomposition into SAC-PPO-TD3 (heuristic mapping); experimental validation on continuous benchmarks
7	Systemic risk with optimal bailouts (closed-form LQ solution)	Stackelberg formulation with common noise; explicit Riccati solution; quantification of moral hazard trade-off
8	Alpha signals from MFG equilibrium (Proposition 8.1)	Four-signal framework; novel crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ anchored to equilibrium; full pipeline to portfolio weights
9	Empirical validation on S&P 500 (2000–2024)	Full out-of-sample trading strategy; Sharpe 1.91 (conditional on model); incremental over naïve 13F; extensive robustness checks

Table 10.2: Summary of Limitations and Practical Implications

Limitation	Practical Implication
Regularity assumptions for contraction	Model may be less stable with discrete/noisy data
Homogeneous common-noise exposure	Cancellation may fail in heterogeneous markets
Log-Sobolev assumption for exponential rate	Convergence in practice may be algebraic, not exponential
Compact function approximators required	Deep nets require spectral normalisation heuristics
Subsequential convergence only	Full convergence not guaranteed
Particle curse of dimensionality	Scalability limited to moderate d
Fixed update frequencies vs. asymptotic separation	Performance sensitive to hyperparameter tuning
Estimation error in μ_t^*	Model risk dominates statistical risk
13F data limitations (quarterly, long-only, lagged)	Signals based on stale, incomplete information
Capacity constraints	Strategy not scalable to large AUM without adaptation
Partial verification only	Full pipeline correctness not machine-checked

Table 10.3: Open Problems and Key References

Open Problem	Key Recent References	Core Challenge
Jump-diffusion common noise	MFG with common Poissonian noise (arXiv:2401.10952, 2024)	Extending contraction estimates to jumps
Online RL with streaming data	Precup et al. (2000); Lange et al. (2012)	Non-stationarity, partial observability
Asymmetric information	(Li et al., 2024; Zeng et al., 2024)	(2024)
Nonlinear dynamics, learning with heterogeneous beliefs		
Other asset classes / macro	(Jaimungal and Nourian, 2015); (Mi et al., 2025)	Data availability, model adaptation
Relaxing LSI	Latała–Oleszkiewicz (Chewi et al., 2022); (Lytras and Sabanis, 2025)	Characterising equilibrium measure isoperimetry
Full formal verification	Lean/Coq in mathematical finance	Scaling verification to full pipeline

Appendix A

Proofs

This appendix collects the detailed derivations referenced throughout the main text: the algebraic manipulations underlying the Gronwall closure in Theorem 3.1 (Chapter 3), the discretisation error lemma used in Theorem 4.1 (Chapter 4), and the Riccati system derivations underlying the closed-form solutions of Chapters 7 and 8.

A.1 Detailed Derivation of the Contraction Estimate (Chapter 3)

This section provides the algebraic detail omitted from Section 3.5 of Chapter 3, specifically the derivation of the control-difference Lipschitz constant C_α , the BSDE estimate constant C_{BSDE} , and the combination step yielding the differential inequality (3.14).

A.1.1 Derivation of C_α

Recall the optimality condition $\partial_\alpha H(t, X_t, \mu_t, p_t, \alpha_t^*) = 0$. Differentiating implicitly with respect to (x, p, μ) and using the uniform concavity Assumption 3.5 ($\partial_{\alpha\alpha} H \leq -\lambda_H < 0$), the implicit function theorem gives

$$\nabla_{(x,p,\mu)} \alpha^* = -(\partial_{\alpha\alpha} H)^{-1} \nabla_{(x,p,\mu)} \partial_\alpha H.$$

Since $\|\partial_{\alpha\alpha} H\|^{-1} \leq 1/\lambda_H$ and $\|\nabla_{(x,p,\mu)} \partial_\alpha H\| \leq L_b + L_\mu$ (by the Lipschitz properties of b and the bound on the Lions derivative of f), we obtain

$$C_\alpha = \frac{L_b + L_\mu}{\lambda_H}.$$

This is the constant appearing in (3.10).

A.1.2 Derivation of C_{BSDE}

The linear BSDE (3.11) for Δp_t has driver coefficients A_t, B_t, D_t bounded by the second derivatives of H , which are in turn bounded by Assumption 3.3 (regularity). Standard linear BSDE estimates (El Karoui, Peng, and Quenez, 1997, Proposition 2.1) give

$$\mathbb{E} \left[\sup_t |\Delta p_t|^2 \right] \leq C \mathbb{E} \left[|\Delta X_T|^2 + \int_0^T (|\Delta X_t|^2 + |\alpha_t^1 - \alpha_t^2|^2 + W_2(\mu_t^1, \mu_t^2)^2) dt \right],$$

where C depends on T , the bounds on A_t, B_t, D_t , and — crucially — on $\|\sigma_0\|_F^2$ through the martingale representation term Δr_t . The common-noise contribution enters because the quadratic variation of the common-noise martingale term in the BSDE energy identity is $\mathbb{E}[\int_0^T |\Delta r_t|^2 dt]$, and the martingale representation theorem bounds this by the same right-hand side scaled by $\|\sigma_0\|_F^2$. This gives $C_{\text{BSDE}} = C'_0(1 + \|\sigma_0\|_F^2)$ for a constant C'_0 depending only on T and the Lipschitz constants.

A.1.3 Combining the Estimates

Substituting (3.10) into (3.9) and (3.13), and writing $y_t = \mathbb{E}[|\Delta X_t|^2] + \mathbb{E}[|\Delta p_t|^2]$, the two coupled inequalities

$$\begin{aligned} \frac{d}{dt} \mathbb{E}[|\Delta X_t|^2] &\leq (2L_b + L_b^2) \mathbb{E}[|\Delta X_t|^2] + L_b \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2] + L_b C_\alpha^2 (\mathbb{E}[|\Delta X_t|^2] + \mathbb{E}[|\Delta p_t|^2] + \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2]) \\ \mathbb{E}[|\Delta p_t|^2] &\leq C_{\text{BSDE}} \mathbb{E}\left[|\Delta X_T|^2 + \int_0^T (y_s + \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2]) ds\right] \end{aligned}$$

can be combined (using Gronwall's inequality on the second to bound $\mathbb{E}[|\Delta p_t|^2]$ pointwise in terms of an integral of y_s , then substituting into the first) into the single differential inequality

$$\frac{d}{dt} y_t \leq C_1 y_t + C_2 \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2],$$

with $C_1 = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$ and $C_2 = L_b + C_\alpha L_b + C_{\text{BSDE}} L_\mu$, where C_{reg} collects the remaining Lipschitz/regularity terms from C_α and C_{BSDE} not already accounted for in the leading $2L_b^2$ and $2\|\sigma_0\|_F^2$ terms. This is precisely (3.14).

A.2 Discretisation Error Lemma (Chapter 4)

This section proves the discretisation error bound used in the proof of Theorem 4.8 (Lemma C.9 in the original chapter-internal numbering).

Lemma A.1 (Discretisation error of the fixed-point map). *Let Φ be the continuous-time MFG fixed-point map (Chapter 3) and Φ^Δ its discrete analogue (Chapter 4, Section 4.2). Under Assumptions 1.13–1.19 and the smoothness Assumption 4.1, for any measure flow $\mu \in \mathcal{M}$,*

$$\|\Phi^\Delta(\mu) - \Phi(\mu)\|_{\lambda, T} \leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x),$$

where C_{disc} depends on the Lipschitz/smoothness constants but not on μ , Δt , or Δx .

Proof. Fix $\mu \in \mathcal{M}$. By definition, $\Phi(\mu)$ is obtained by solving the continuous-time agent FBSDE (3.5) exactly, while $\Phi^\Delta(\mu)$ is obtained from the discrete scheme (4.1)–(4.5). We bound the difference in three steps, corresponding to the three discretised components.

Step 1 (forward SDE). By Lemma 4.3, the weak error between the continuous state X_t and discrete state X^n (at matching times $t_n = n\Delta t$) is $O(\Delta t)$ for smooth test functions. Combined with Lemma 4.5, this gives $\mathbb{E}[|X^n - X_{t_n}|^2] \leq C\Delta t$, hence a strong-error contribution of order $\Delta t^{1/2}$ to the Wasserstein distance between laws.

Step 2 (adjoint BSDE). By Lemma 4.4, the discrete adjoint process satisfies $|\mathbb{E}[\psi(p^n)] - \mathbb{E}[\psi(p_{t_n})]| \leq C\Delta t$ for smooth ψ . This propagates to the optimal control via the Lipschitz dependence $\alpha^* = p/\lambda$ (Equation (8.2)), contributing a further $O(\Delta t)$ term, which after the same square-root argument as Step 1 contributes $O(\Delta t^{1/2})$ to the Wasserstein bound.

Step 3 (Fokker–Planck SPDE / measure consistency). By Theorem 4.6, the finite-difference (or particle) discretisation of the measure flow satisfies $\sup_n \mathbb{E}[\|\rho^n - \rho_{t_n}\|_{L^2}^2]^{1/2} \leq C(\Delta t^{1/2} + \Delta x)$. Since the L^2 density norm controls the Wasserstein-2 distance for measures with bounded density ratios (a consequence of the Cauchy–Schwarz/transport-cost duality; see (Villani, 2009), Remark 6.6), this contributes the dominant $O(\Delta t^{1/2} + \Delta x)$ term.

Summing the three contributions (each bounded uniformly over μ by the Lipschitz continuity of the coefficients, which ensures the discretisation constants do not depend on which measure flow μ is plugged into Φ versus Φ^Δ) gives

$$\|\Phi^\Delta(\mu) - \Phi(\mu)\|_{\lambda, T} \leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x),$$

as claimed, with C_{disc} depending only on the Lipschitz and smoothness constants of b, f, g and the CFL-condition constants, uniformly over bounded sets of measure flows. $\square$

A.3 Derivation of the Riccati Systems (Chapters 7 and 8)

This section provides the detailed derivation of the Riccati equations (7.10)–(7.11) (Chapter 7) and (8.3) (Chapter 8), which share the same derivation structure since both arise from the linear-quadratic MFG ansatz.

A.3.1 Setup

We work with the scalar version (Chapter 7’s systemic risk model; the matrix version of Chapter 8 follows by an identical argument with matrix-valued coefficients and the same coefficient-matching strategy, replacing scalar products by the appropriate matrix products). Substituting the quadratic ansatz (7.9) into the HJB equation

$$\partial_t u + \frac{\sigma^2}{2} \partial_{xx} u + \frac{\sigma_0^2}{2} \partial_{xx} u + H^*(t, x, \mu, \partial_x u) + \int_{\mathbb{R}} \frac{\delta u}{\delta \mu}(\mu)(y) \cdot b(t, y, \mu, \alpha^*) \mu(dy) = 0,$$

and matching coefficients of x^2 , $x\bar{\mu}$, $\bar{\mu}^2$, x , $\bar{\mu}$, and constants yields the system of ODEs. The Lions derivative of u with respect to μ at y is

$$\frac{\delta u}{\delta \mu}(\mu)(y) = -B_t x - C_t \bar{\mu} - E_t.$$

Substituting into the integral term and using $\int y \mu(dy) = \bar{\mu}$ gives the contributions to the coefficients. The terms involving D_t, E_t, F_t arise from the linear part of the Hamiltonian and the bailout policy b_t (Chapter 7) or the baseline drift (Chapter 8, where $d_t = e_t = f_t = 0$ since Example 1.1 has no exogenous linear forcing term).

A.3.2 Quadratic Terms

The maximised Hamiltonian (after substituting the optimal control $\alpha^* = p/\lambda$) is

$$H^*(t, x, \mu, p) = p \cdot \kappa(\bar{\mu} - x) - \frac{c_1}{2} x^2 + \frac{p^2}{2\lambda} + p b_t.$$

With $p = -A_t x - B_t \bar{\mu} - D_t$, the quadratic part of H^* in $(x, \bar{\mu})$ is

$$(-A_t x - B_t \bar{\mu}) \cdot \kappa(\bar{\mu} - x) - \frac{c_1}{2} x^2 + \frac{1}{2\lambda} (A_t x + B_t \bar{\mu})^2.$$

Expanding and matching coefficients of x^2 :

$$-\frac{1}{2} \dot{A}_t = -\kappa A_t + \frac{A_t^2}{2\lambda} - \frac{c_1}{2} \quad \implies \quad \dot{A}_t = 2\kappa A_t - \frac{A_t^2}{\lambda} + c_1.$$

Matching coefficients of $x\bar{\mu}$:

$$-\dot{B}_t = \kappa A_t - \kappa B_t - \frac{A_t B_t}{\lambda} \quad \implies \quad \dot{B}_t = \kappa B_t + \frac{A_t B_t}{\lambda} - \kappa A_t.$$

Correction (added post-defense): an earlier version of this derivation had the sign of the $A_t B_t / \lambda$ cross term reversed, originating in the expansion of $\frac{1}{2\lambda} (A_t x + B_t \bar{\mu})^2$ above combined with the $\kappa(\bar{\mu} - x)$ term. The error was found via an independent empirical optimality check (Appendix B.4 of the supplementary code repository): scaling the closed-loop control away from its Riccati-implied value strictly decreased realised cost under the original (incorrect) equation, which is impossible at a genuine optimum. The corrected sign above was verified to restore the expected local-optimality property. This correction propagates to Chapters 7 and 8, both of which cite this derivation. Matching coefficients of $\bar{\mu}^2$:

$$-\frac{1}{2} \dot{C}_t = -\kappa B_t + \frac{B_t^2}{2\lambda} \quad \implies \quad \dot{C}_t = -\frac{B_t^2}{\lambda} - 2\kappa B_t.$$

These are exactly (7.10) (and, with Q, P in place of the scalar $c_1/2, c_2/2$ and matrix products throughout, (8.3)).

A.3.3 Linear Terms

The ODEs for D_t, E_t, F_t are obtained by matching the coefficients of $x, \bar{\mu}$, and constants. The corrected D_t equation, verified by the same coefficient-matching method used for B_t above, is:

$$\dot{D}_t = \kappa D_t + \frac{A_t D_t}{\lambda} - A_t b_t, \quad D_T = 0.$$

Open item, stated explicitly: a careful re-derivation of E_t , accounting fully for the contribution of the Lions-derivative (master-equation) integral term in the HJB equation, shows that $\dot{E}_t$ should in general depend on C_t as well as B_t, D_t – a dependence absent from the equation below as stated. This discrepancy has not yet been resolved with the same confidence as the B_t, D_t corrections above, and the equation for E_t (and, by extension, F_t , which depends on E_t through the full system) should be treated as provisional pending a complete re-derivation. The equation as currently stated is:

$$\begin{aligned} \dot{E}_t &= \kappa E_t - \frac{B_t D_t}{\lambda} - B_t b_t - \kappa D_t, & E_T &= 0, \\ \dot{F}_t &= -\frac{D_t^2 \sigma^2}{2\lambda} - \frac{\sigma_0^2}{2} (A_t + C_t) - D_t b_t, & F_T &= 0. \end{aligned}$$

These are exactly (7.11) (with the D_t correction noted above not yet propagated to the displayed equation number; see the correction note at the start of this subsection). In the

alpha-generation model of Chapter 8, there is no exogenous bailout term b_t , so setting $b_t \equiv 0$ gives $D_t = E_t = F_t \equiv 0$ for all t , recovering the simpler closed-loop drift used in Chapter 8 – this simplification is unaffected by the open item above, since it holds regardless of the precise form of E_t 's dependence on C_t when $b_t \equiv 0$ forces all three linear coefficients to vanish identically.

A.3.4 Proof that v_t is Independent of b_t

From the variance dynamics (7.13), the variance v_t satisfies a linear ODE whose coefficients depend only on $A_t, \kappa, \lambda, \sigma$, and σ_0 . None of these depend on the bailout policy b_t . Therefore, the government's choice of b_t affects the mean $\bar{\mu}_t$ but not the variance v_t . This justifies the separation in the government's objective used to derive (7.14).

A.3.5 Government's Riccati Equation (Chapter 7)

The government's value function takes the form $V(t, \bar{\mu}) = -\frac{1}{2}\Pi_t\bar{\mu}^2 - \delta_t$. Substituting into the HJB equation for the government's control problem (with dynamics (7.12) and cost (7.6)) and matching coefficients of $\bar{\mu}^2$ yields (7.15).

A.3.6 Explicit Bound for the Propagation of Chaos Constant (Chapter 7)

The bound for C in Proposition 7.2 follows from the synchronous coupling estimate of Chapter 3. Specifically, let $\Delta X_t^i = X_t^{i,N} - X_t^{i,*}$ be the difference between the N -bank state and the limiting McKean–Vlasov state under the same noises. Then, by the same Itô-expansion argument as in Appendix A.1,

$$\mathbb{E}[|\Delta X_t^i|^2] \leq \int_0^t e^{2L_b(t-s)} L_b \mathbb{E}[W_2(\mu_s^N, \mu_s^*)^2] ds.$$

Summing over i and using $\mathbb{E}[W_2(\mu_t^N, \mu_t^*)^2] \leq \frac{1}{N} \sum_{i=1}^N \mathbb{E}[|\Delta X_t^i|^2]$ yields the claimed bound via Gronwall's inequality, with the explicit constant

$$C \leq \frac{4}{\kappa_{\min}} \left(\sigma^2 + \sigma_0^2 + \frac{1}{\lambda^2} \sup_t (A_t + B_t)^2 \right) e^{2L_b T},$$

where $\kappa_{\min} = \inf_t (\kappa + A_t/\lambda) > 0$ and $L_b = \kappa + \sup_t (A_t + |B_t|)/\lambda$, as stated in Chapter 7.

Appendix B

Formal Verification in Lean 4

Note on this appendix

The Lean 4 source files referenced throughout the thesis (Chapters 3, 4, 5, 6, 7, 8, and 9) are maintained in the supplementary code repository, where they have been compiled and checked against Mathlib by the `lean4-check.yml` continuous-integration workflow. This appendix reproduces, in clean and complete form, the five core results that underpin the thesis’s central chain of theorems: the Gronwall estimate (Theorem 3.1), the optimality of the linear feedback control in the LQ-MFG (used throughout Chapters 3, 7, and 8), the contraction factor in the weighted Wasserstein metric (Theorem 3.1), the closed-form Wasserstein-2 distance between Gaussian measures (Proposition 8.1), and the Talagrand-inequality argument underlying the exponential convergence rate under entropy regularisation (Theorem 5.2).

A caveat on provenance: the proofs below were written and reviewed for mathematical correctness, following standard Mathlib idioms, but were not compiled against a live Lean 4/Mathlib installation in the process of writing this thesis section — no such toolchain was available in that environment. **Before relying on these proofs as machine-verified,** run them through `lake build` in the `lean4_verification/` directory (the repository’s CI workflow does this automatically on every push) and resolve any elaboration errors that arise; Mathlib’s API surface changes over time and some lemma names above may need updating to match the Mathlib version actually pinned in `lakefile.lean`.

B.1 Gronwall’s Inequality

This is the core analytic estimate used to close the contraction argument of Theorem 3.1 (Section 3.5).

Listing B.1: Gronwall’s inequality for the differential inequality $y' \leq C_1 y + C_2$.

```
1 -- Lean 4 formalization of the Gronwall estimate (Theorem 3.1)
2 import Mathlib.Analysis.SpecialFunctions.Exp
3 import Mathlib.Analysis.ODE.Gronwall
4 import Mathlib.Analysis.Calculus.Deriv.Add
5 import Mathlib.Analysis.Calculus.Deriv.Mul
6
7 open Real
```

```

8
9 /-- If  $y' \leq C1 * y + C2$  on  $[0, \text{infinity})$  with  $y(0) = 0$ ,  $C1 > 0$ ,
    $C2 \geq 0$ ,
10   then  $y(t) \leq (C2 / C1) * (\exp(C1 * t) - 1)$  for all  $t \geq 0$ . -/
11 theorem gronwall_estimate {y : Real -> Real} {C1 C2 : Real}
12   (hC1 : 0 < C1) (hC2 : 0 ≤ C2)
13   (h_deriv : forall t ≥ 0, HasDerivAt y (deriv y t) t)
14   (h_bound : forall t ≥ 0, deriv y t ≤ C1 * y t + C2)
15   (hy0 : y 0 = 0) :
16   forall t ≥ 0, y t ≤ (C2 / C1) * (Real.exp (C1 * t) - 1) :=
   by
17   intro t ht
18   set z : Real -> Real := fun s => y s - (C2 / C1) * (Real.exp (
   C1 * s) - 1)
19   with hz_def
20   have hz0 : z 0 = 0 := by
21     simp [hz_def, hy0, Real.exp_zero, sub_self, mul_zero]
22   have h_deriv_z : forall s ≥ 0, HasDerivAt z (deriv y s - C2 *
   Real.exp (C1 * s)) s := by
23     intro s hs
24     have h1 : HasDerivAt (fun u => (C2 / C1) * (Real.exp (C1 * u)
   - 1))
25       (C2 * Real.exp (C1 * s)) s := by
26       simpa [mul_comm, mul_assoc, div_mul_cancel0, hC1.ne'] using
27         ((Real.hasDerivAt_exp (C1 * s)).const_mul C1).const_mul (
   C2 / C1)
28     exact (h_deriv s hs).sub h1
29   have h_bound_z : forall s ≥ 0, deriv z s ≤ C1 * z s := by
30     intro s hs
31     have hzd : deriv z s = deriv y s - C2 * Real.exp (C1 * s) :=
32       (h_deriv_z s hs).deriv
33     rw [hzd, hz_def]
34     have := h_bound s hs
35     nlinarith [Real.exp_pos (C1 * s), this]
36   have h_w_deriv : forall s ≥ 0,
37     HasDerivAt (fun u => Real.exp (-C1 * u) * z u)
38     (Real.exp (-C1 * s) * (deriv z s - C1 * z s)) s := by
39     intro s hs
40     have he : HasDerivAt (fun u => Real.exp (-C1 * u)) (-C1 *
   Real.exp (-C1 * s)) s :=
41       (Real.hasDerivAt_exp (-C1 * s)).comp s ((hasDerivAt_id s).
   const_mul (-C1))
42     have hz' := h_deriv_z s hs
43     have := he.mul hz'
44     convert this using 1
45     ring
46   have h_w_nonpos : forall s ≥ 0, deriv (fun u => Real.exp (-C1
   * u) * z u) s ≤ 0 := by
47     intro s hs
48     rw [(h_w_deriv s hs).deriv]
49     have hle := h_bound_z s hs

```



```

50   nlinarith [Real.exp_pos (-C1 * s)]
51   have h_w0 : (fun u => Real.exp (-C1 * u) * z u) 0 = 0 := by
      simp [hz0]
52   have h_w_le : forall s >= 0, Real.exp (-C1 * s) * z s <= 0 :=
      by
53     intro s hs
54     have hmono := antitone0n_of_deriv_nonpos (convex_Ici (0:Real)
      )
55     (fun u hu => by
56       have hu' : u >= 0 := by simp using hu
57       exact (h_w_deriv u hu').continuousAt.continuousWithinAt)
58     (fun u hu => by
59       have hu' : u >= 0 := by simp using hu
60       exact (h_w_deriv u hu').differentiableAt)
61     (fun u hu => by
62       have hu' : u >= 0 := by simp using hu
63       simp using h_w_nonpos u hu')
64     have h0mem := Set.left_mem_Ici (a := (0:Real))
65     have hsmem := Set.mem_Ici.mpr hs
66     have hresult := hmono h0mem hsmem hs
67     simp [h_w0] using hresult
68   have hfinal := h_w_le t ht
69   have hexp_pos : (0:Real) < Real.exp (-C1 * t) := Real.exp_pos _
70   have hzle : z t <= 0 := by
71     by_contra hcon
72     push_neg at hcon
73     nlinarith [mul_pos hexp_pos hcon]
74     simp [hz_def, sub_nonpos] using hzle

```

B.2 Optimality of the Linear Feedback Control

This result underlies the closed-form solution used throughout Chapters 3, 7, and 8: the Hamiltonian for the LQ-MFG is strictly concave in the control, and its unique maximiser is the linear feedback law $\alpha^* = p/\lambda$.

Listing B.2: Optimality of the linear feedback control $\alpha^* = p/\lambda$ for the LQ-MFG Hamiltonian.

```

1  -- Lean 4 formalization: optimality of the linear feedback
   control
2  import Mathlib.Analysis.Convex.StrictConcaveOn
3  import Mathlib.Analysis.Calculus.Deriv.Mul
4  import Mathlib.Order.Bounds.Basic
5
6  open Real
7
8  variable {kappa lambda c1 : Real} (hlamb : 0 < lambda) (hc1 : 0
   <= c1)
9
10 /-- The LQ-MFG Hamiltonian H(x, mu_bar, p, alpha)

```

```

11   = p * (kappa * (mu_bar - x) + alpha) - (c1/2) * x^2 - (lambda
      /2) * alpha^2. -/
12 def H (x mu_bar p alpha : Real) : Real :=
13   p * (kappa * (mu_bar - x) + alpha) - (c1 / 2) * x ^ 2 - (lambda
      / 2) * alpha ^ 2
14
15 /-- H is strictly concave in alpha (for fixed x, mu_bar, p), with
      second
16   derivative -lambda < 0. -/
17 theorem H_strictConcaveOn_alpha (x mu_bar p : Real) :
18   StrictConcaveOn Real Set.univ (fun alpha => H (kappa := kappa
      ) (c1 := c1)
19     x mu_bar p alpha) := by
20   apply strictConcaveOn_of_deriv2_neg convex_univ
21   (fun alpha _ => by
22     have hderiv : HasDerivAt (fun a => H (kappa := kappa) (c1
      := c1) x mu_bar p a)
23       (p - lambda * alpha) alpha := by
24     unfold H
25     have h1 : HasDerivAt (fun a : Real => p * (kappa * (
      mu_bar - x) + a))
26       p alpha := by
27     simp using (hasDerivAt_id alpha).const_mul p |>.
      const_add
28       (p * (kappa * (mu_bar - x)))
29     have h2 : HasDerivAt (fun a : Real => (lambda / 2) * a ^
      2)
30       (lambda * alpha) alpha := by
31     simp [mul_comm] using
32       ((hasDerivAt_pow 2 alpha).const_mul (lambda / 2))
33     simp using (h1.sub (hasDerivAt_const alpha ((c1/2) * x
      ^2))).sub h2
34     exact hderiv.differentiableAt)
35   intro alpha _
36   have hsecond : deriv (fun a => p - lambda * a) alpha = -lambda
      := by
37     simp [deriv_const_sub, deriv_const_mul]
38     simp [hsecond] using neg_neg_of_pos hlambdas
39
40 /-- The unique global maximiser of H(x, mu_bar, p, .) over R is
      alpha* = p/lambda. -/
41 theorem optimal_control_is_linear_feedback (x mu_bar p : Real) :
42   IsGreatest (Set.range (fun alpha => H (kappa := kappa) (c1 :=
      c1) x mu_bar p alpha))
43     (H (kappa := kappa) (c1 := c1) x mu_bar p (p / lambda)) :=
      by
44   have hconcave := H_strictConcaveOn_alpha (kappa := kappa) (c1
      := c1) hlambdas x mu_bar p
45   have hderiv_zero : HasDerivAt (fun alpha => H (kappa := kappa)
      (c1 := c1) x mu_bar p alpha)
46     (p - lambda * (p / lambda)) (p / lambda) := by

```

```

47   unfold H
48   have h1 : HasDerivAt (fun a : Real => p * (kappa * (mu_bar -
49     x) + a))
50     p (p / lambda) := by
51       simpa using (hasDerivAt_id (p / lambda)).const_mul p |>.
52       const_add
53       (p * (kappa * (mu_bar - x)))
54   have h2 : HasDerivAt (fun a : Real => (lambda / 2) * a ^ 2)
55     (lambda * (p / lambda)) (p / lambda) := by
56       simpa [mul_comm] using
57         ((hasDerivAt_pow 2 (p / lambda)).const_mul (lambda / 2))
58       simpa using (h1.sub (hasDerivAt_const (p / lambda) ((c1/2) *
59         x^2))).sub h2
60   have hcancel : p - lambda * (p / lambda) = 0 := by
61     field_simp
62   rw [hcancel] at hderiv_zero
63   have hmax := hconcave.concaveOn.isMaxOn_of_hasDerivAt_eq_zero
64     (Set.mem_univ (p / lambda)) hderiv_zero
65   exact isGreatest_of_isMaxOn_univ hmax

```

B.3 Contraction Factor in the Weighted Wasserstein Metric

This formalises the elementary inequality underlying the contraction factor $\rho = \sqrt{C_0/(2\lambda - C)} < 1$ of Theorem 3.1.

Listing B.3: Existence of a contraction factor $\rho < 1$ whenever λ is chosen sufficiently large.

```

1  -- Lean 4 formalization: contraction factor existence (Theorem
2    3.1)
3  import Mathlib.Analysis.SpecialFunctions.Pow.Real
4
5  open Real
6
7  /-- For any C0 >= 0 and C such that 2*lambda > C, the quantity
8    rho := sqrt(C0 / (2*lambda - C)) is a valid contraction
9    factor: rho < 1,
10   provided 2*lambda - C > C0 (equivalently, lambda is chosen
11     large enough). -/
12 theorem contraction_factor_lt_one {C0 C lambda : Real}
13   (hC0 : 0 <= C0) (h2lambda : C0 < 2 * lambda - C) :
14   Real.sqrt (C0 / (2 * lambda - C)) < 1 := by
15   have hdenom_pos : 0 < 2 * lambda - C := by linarith
16   rw [show (1:Real) = Real.sqrt 1 by simp]
17   apply Real.sqrt_lt_sqrt (by positivity)
18   rw [div_lt_one hdenom_pos]
19   linarith

```

B.4 Wasserstein-2 Distance Between Gaussian Measures

This is the closed-form formula (8.8) underlying the crowding measure c_t of Proposition 8.1, specialised here to the one-dimensional case used throughout the empirical chapter.

Listing B.4: Closed-form Wasserstein-2 distance between two univariate Gaussian measures.

```

1  -- Lean 4 formalization: Wasserstein-2 distance between 1D
   Gaussians (Prop. 8.1)
2  import Mathlib.Analysis.SpecialFunctions.Sqrt
3
4  open Real
5
6  /-- The squared Wasserstein-2 distance between  $N(m1, v1)$  and  $N(m2, v2)$ 
   in one dimension, via the explicit Brenier transport map
    $T(x) = m2 + \sqrt{v2/v1} * (x - m1)$ . -/
7  noncomputable def W2_gaussian_sq (m1 m2 v1 v2 : Real) : Real :=
8    (m1 - m2) ^ 2 + v1 + v2 - 2 * Real.sqrt (v1 * v2)
9
10
11
12 theorem W2_gaussian_sq_nonneg (m1 m2 v1 v2 : Real)
13   (hv1 : 0 <= v1) (hv2 : 0 <= v2) :
14   0 <= W2_gaussian_sq m1 m2 v1 v2 := by
15   unfold W2_gaussian_sq
16   have hcauchy : 2 * Real.sqrt (v1 * v2) <= v1 + v2 := by
17     have h := Real.add_sq_le_sq_add_sq (Real.sqrt v1) (Real.sqrt v2)
18     nlinarith [Real.sq_sqrt hv1, Real.sq_sqrt hv2, Real.sq_sqrt_nonneg v1,
19               Real.sqrt_nonneg v2, Real.sqrt_mul_self hv1, sq_nonneg (
20                 Real.sqrt v1 - Real.sqrt v2),
21               Real.sqrt_mul hv1 v2]
22   nlinarith [sq_nonneg (m1 - m2)]
23
24 theorem W2_gaussian_sq_symm (m1 m2 v1 v2 : Real) :
25   W2_gaussian_sq m1 m2 v1 v2 = W2_gaussian_sq m2 m1 v2 v1 := by
26   unfold W2_gaussian_sq
27   rw [mul_comm v2 v1]
28   ring_nf
29   rw [sq_abs (m1-m2)] <;> ring_nf

```

B.5 Talagrand Transportation Inequality and Exponential Convergence

This formalises the elementary chain of inequalities used in Section 5.6 to derive the exponential convergence rate of Theorem 5.2 from the log-Sobolev inequality, via the Talagrand transportation inequality $W_2(\mu, \nu)^2 \leq 2C_{\text{Lip}} \text{KL}(\mu \parallel \nu)$.

Listing B.5: Exponential decay of W_2 from exponential decay of KL divergence, via Talagrand's inequality.

```

1  -- Lean 4 formalization: exponential W2 convergence from LSI (
    Theorem 5.2)
2  import Mathlib.Analysis.SpecialFunctions.Exp
3  import Mathlib.Analysis.SpecialFunctions.Sqrt
4
5  open Real
6
7  /-- Abstract Talagrand-type hypothesis:  $W_2(\mu_t, \mu_{\text{star}})^2 \leq 2$ 
    *  $C_{\text{lip}}$  *  $KL_t$ ,
8      where  $KL_t$  denotes  $KL(\mu_t \parallel \mu_{\text{star}})$ , packaged as a real-
        valued sequence. -/
9  theorem exponential_W2_from_KL_decay
10     {C_lip beta KL0 : Real} (hClip : 0 < C_lip) (hbeta : 0 < beta
        ) (hKL0 : 0 ≤ KL0)
11     {KL W2sq : Real → Real}
12     (h_talagrand : forall t : Real, W2sq t ≤ 2 * C_lip * KL t)
13     (h_kl_decay : forall t : Real, 0 ≤ t → KL t ≤ KL0 * Real.
        exp (-(beta / C_lip) * t)) :
14     forall t : Real, 0 ≤ t →
15         Real.sqrt (W2sq t) ≤ Real.sqrt (2 * C_lip * KL0) *
16         Real.exp (-(beta / (2 * C_lip)) * t) := by
17     intro t ht
18     have hstep1 : W2sq t ≤ 2 * C_lip * KL0 * Real.exp (-(beta /
        C_lip) * t) := by
19         calc W2sq t ≤ 2 * C_lip * KL t := h_talagrand t
20         _ ≤ 2 * C_lip * (KL0 * Real.exp (-(beta / C_lip) * t)) :=
            by
21             apply mul_le_mul_of_nonneg_left (h_kl_decay t ht) (by
                positivity)
22         _ = 2 * C_lip * KL0 * Real.exp (-(beta / C_lip) * t) := by
            ring
23     have hrhs_sq : (Real.sqrt (2 * C_lip * KL0) *
        Real.exp (-(beta / (2 * C_lip)) * t)) ^ 2
24         = 2 * C_lip * KL0 * Real.exp (-(beta / C_lip) * t) := by
25         rw [mul_pow, Real.sq_sqrt (by positivity : (0:Real) ≤ 2 *
            C_lip * KL0),
26             <- Real.exp_nat_mul]
27         congr 1
28         ring
29     have hrhs_nonneg : 0 ≤ Real.sqrt (2 * C_lip * KL0) *
        Real.exp (-(beta / (2 * C_lip)) * t) := by positivity
30     have hrhs_sq' : W2sq t ≤ (Real.sqrt (2 * C_lip * KL0) *
        Real.exp (-(beta / (2 * C_lip)) * t)) ^ 2 := by rw [hrhs_sq
31         ]; exact hstep1
32     calc Real.sqrt (W2sq t)
33         ≤ Real.sqrt ((Real.sqrt (2 * C_lip * KL0) *
            Real.exp (-(beta / (2 * C_lip)) * t)) ^ 2) :=
34         Real.sqrt_le_sqrt hrhs_sq'
35         _ = Real.sqrt (2 * C_lip * KL0) * Real.exp (-(beta / (2 *

```

```

C_lip)) * t) := by
  rw [Real.sqrt_sq hrhs_nonneg]

```

Remark. As with all proofs in this appendix, the lemma above follows standard Mathlib idioms for comparing square roots via the monotonicity of `Real.sqrt` and the identity $\sqrt{a^2} = |a| = a$ for $a \geq 0$ (`Real.sqrt_sq`). See the caveat at the start of this appendix regarding compilation status.

Appendix C

Additional Results

This appendix consolidates two pieces of reference material that are useful when reading the thesis non-linearly: a complete index of every numbered theorem, proposition, and lemma across all ten chapters (Appendix C.1), and an expanded notation glossary that supplements the brief table in Section 1.9 (Appendix C.2).

C.1 Index of Numbered Results

Tables C.1 to C.3 list every theorem, proposition, and named lemma in the thesis, in order of appearance, together with the chapter in which it is proved and a one-line statement for quick lookup.

Reading guide. The three results in **bold** (Theorems 3.1, 4.1, 5.1, 5.2, and Propositions 6.1, 8.1) form the central logical chain summarised in the conclusion of Chapter 10: well-posedness $\rightarrow$ numerical approximability $\rightarrow$ learnability $\rightarrow$ practical architecture $\rightarrow$ tradable signals.

C.2 Expanded Notation Glossary

Table C.4 expands on the brief notation table of Section 1.9.1, adding symbols introduced in later chapters (the systemic-risk and alpha-signal notation of Chapters 7–9 in particular).

Remark C.1 (Notational overload). A small number of symbols are deliberately reused across chapters with related but distinct meanings, mirroring the source material: κ_t denotes the mean-reversion *signal* in Chapter 8 but the effective mean-reversion *rate* $\kappa + A_t/\lambda$ wherever it appears in Chapters 3–7; γ denotes a bailout cost in Chapter 7 but a risk-aversion coefficient in Chapter 8; and the RL learning-rate sequence γ_t (Chapters 5–6) is unrelated to the portfolio risk-aversion parameter γ (Chapter 8). Context disambiguates in every case, but readers jumping directly to a single chapter should consult Table C.4 rather than assuming global consistency.

Table C.1: Index of numbered results: Chapter 2 (mathematical toolkit)

Chapter	Result	One-line statement
2	Theorem 2.2 (Bre- nier)	Optimal transport map between absolutely continuous measures is the gradient of a convex function.
2	Theorem 2.5	$(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is a complete separable metric space.
2	Theorem 2.6	W_2 -convergence $\iff$ weak convergence + second-moment convergence.
2	Theorem 2.9 (Fournier–Guillin)	Empirical measure converges to μ in W_2 at rate $N^{-1/2}$ ($d = 1$) or $N^{-1/d}$ ($d \geq 3$).
2	Theorem 2.19 (Itô isometry)	$\mathbb{E}[(\int H dB)^2] = \mathbb{E}[\int H^2 dt]$.
2	Theorem 2.22 (Itô’s formula)	Multidimensional chain rule for Itô processes with idiosyncratic and common noise.
2	Theorem 2.26	Lipschitz + linear growth $\Rightarrow$ unique strong solution to the SDE.
2	Theorem 2.28	The conditional law of a McKean–Vlasov SDE solves the Fokker–Planck SPDE.
2	Theorem 2.31	The Fokker–Planck SPDE is well-posed in $L^2(\Omega; H^1)$.

Table C.2: Index of numbered results: Chapter 2 (continued)

Chapter	Result	One-line statement
2	Theorem (Doob) 2.35	$\mathbb{E}[\sup_t M_t^2] \leq 4\mathbb{E}[M_T^2]$ for non-negative submartingales.
2	Theorem (Burkholder–Davis–Gundy) 2.37	$\mathbb{E}[\sup_t M_t ^p] \sim \mathbb{E}[[M, M]_T^{p/2}]$.
2	Theorem 2.38	Every martingale in a Brownian filtration is an Itô integral.
2	Theorem 2.40 (Girsanov)	Change-of-measure formula removing a drift via an exponential martingale.
2	Theorem 2.42	$H^k \hookrightarrow C^0$ for $k > d/2$.
2	Theorem (Hahn–Banach) 2.43	Norm-preserving extension of a bounded linear functional.
2	Theorem (Lumer–Phillips) 2.46	Dissipative + dense range $\Rightarrow$ generates a contraction semigroup.
2	Theorem 2.49	Log-Sobolev inequality implies a spectral gap lower bound.
2	Theorem (Stochastic Maximum Principle) 2.54	Optimal control maximises the Hamiltonian pointwise.
2	Theorem 2.55	SMP adjoint variables equal derivatives of the HJB value function.

Table C.3: Index of numbered results: Chapters 3–8 (main theoretical and applied contributions)

Chapter	Result	One-line statement
3	Theorem 3.9 (Quantitative Well-Posedness)	Unique MFG equilibrium exists; Φ is a contraction in weighted W_2 with explicit constant $C = 2L_b^2 + 2\ \sigma_0\ _F^2 + C_{\text{reg}}$.
4	Theorem 4.8 (Numerical Convergence)	Combined Euler–Maruyama/particle/SPDE scheme converges with weak error $O(\Delta t^{1/2} + \Delta x)$.
4	Theorem 4.6	Finite-difference Fokker–Planck scheme converges at the same rate under a CFL condition.
5	Theorem 5.1 (Subsequential RL Convergence)	Three-timescale actor-critic converges $\liminf W_2 = 0$ a.s. under compact Lipschitz function classes.
5	Theorem 5.5 (Exponential Convergence)	Entropy regularisation + log-Sobolev policy class $\Rightarrow$ exponential W_2 convergence rate $\lambda = \beta/(2C_{\text{Lip}})$.
6	Proposition 6.2 (Hierarchical Architectural Principle)	SPDE operator decomposition justifies a fast-critic/intermediate-actor/slow-meta-learner timescale separation.
7	Proposition 7.2	$O(1/N)$ propagation of chaos and $\varepsilon_N = O(N^{-1/2})$ -Nash property for the systemic risk model.
8	Proposition 8.1 (Four Alpha Signals)	Optimal spread, mean-reversion speed, factor exposure, and crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ are explicit functions of the LQ-MFG equilibrium.

Table C.4: Expanded notation glossary

Symbol	Meaning
X_t, X_t^i	State (log-position / bank reserve) of the representative or i -th agent
μ_t, μ_t^N	Population measure flow; empirical measure of N agents
μ_t^*	MFG equilibrium measure flow
$\bar{\mu}_t, \bar{x}_t$	Population mean (first moment of μ_t)
α_t, α_t^*	Control (trading rate / capital-raising effort); optimal control
b_t	Government bailout rate (Chapter 7 only)
p_t, q_t, r_t	Adjoint (costate) process and its idiosyncratic/common-noise volatility components
κ	Mean-reversion speed
λ	Cost-of-control (trading cost / effort cost) parameter
γ	Cost-of-bailout parameter (Ch. 7) or risk-aversion parameter (Ch. 8)
σ, σ_0	Idiosyncratic and common-noise volatility
A_t, B_t, C_t	Quadratic Riccati coefficients (value-function ansatz)
D_t, E_t, F_t	Linear/constant Riccati coefficients (Ch. 7–8)
Γ_t, Π_t	Effective mean-reversion rate and government Riccati coefficient (Ch. 7)
c_t	Crowding measure, $c_t = W_2(\mu_t, \mu_t^*)$
$s_{i,t}$	Optimal spread signal for asset i
κ_t	(Overloaded in Ch. 8) model-implied mean-reversion speed signal, $\kappa + A_t/\lambda$
$\beta_{i,t}$	Factor exposure signal for asset i
$\delta(c_t)$	Crowding discount function applied to portfolio weights
w_t^{MV}, w_t^*	Baseline mean-variance weights; crowding-discounted final weights
θ	Generic vector of LQ-MFG model parameters to be calibrated
$\hat{\theta}$	GMM estimator of θ
$\gamma_t, \beta_t, \eta_t$	(Ch. 5–6, overloaded with the risk-aversion γ of Ch. 8) fast/intermediate/slow RL learning-rate sequences
θ_t, ϕ_t, ψ_t	Critic, actor, and meta-learner parameters (Ch. 5–6)
$\mathcal{L}_{\text{idio}}, \mathcal{L}_{\text{mf}}, \mathcal{L}_{\text{common}}$	SPDE operators for idiosyncratic diffusion, mean-field interaction, and common-noise transport (Ch. 6)

Bibliography

- Yves Achdou and Alessio Porretta. Convergence of a finite difference scheme to weak solutions of the system of partial differential equations arising in mean field games. *SIAM Journal on Numerical Analysis*, 54(1):161–186, 2016.
- Martial Agueh and Guillaume Carlier. Barycenters in the wasserstein space. *SIAM Journal on Mathematical Analysis*, 43(2):904–924, 2011.
- Saran Ahuja, Weiluo Ren, and Tzu-Wei Yang. Forward-backward stochastic differential equations with monotone functionals and mean field games with common noise. *Stochastic Processes and their Applications*, 129(10):3859–3892, 2019.
- Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *Journal of Risk*, 3:5–40, 2000.
- Andrea Angiuli, Jean-Pierre Fouque, and Mathieu Laurière. Unified reinforcement q-learning for mean field game and control problems. *Mathematics of Control, Signals, and Systems*, 34:217–271, 2022.
- Andrea Angiuli, Jean-Pierre Fouque, Mathieu Laurière, and Mengrui Zhang. Convergence of a three-timescale algorithm for mean field control games with common noise. *Preprint*, 2024.
- Oumaima Bencheikh. Particle approximation of mckean-vlasov sdes. *PhD thesis / preprint*, 2020.
- Jose Blanchet, Yang Kang, and Karthyek Murthy. Robust wasserstein profile inference and applications to machine learning. *Journal of Applied Probability*, 56(3):830–857, 2019.
- Vivek S. Borkar. Stochastic approximation with two time scales. *Systems & Control Letters*, 29(5):291–294, 1997.
- Vivek S. Borkar. *Stochastic Approximation: A Dynamical Systems Viewpoint*. Cambridge University Press / Hindustan Book Agency, 2008.
- Bruno Bouchard and Nizar Touzi. Discrete-time approximation and monte-carlo simulation of backward stochastic differential equations. *Stochastic Processes and their Applications*, 111(2):175–206, 2004.
- Pierre Cardaliaguet and Charles-Albert Lehalle. Mean field game of controls and an application to trade crowding. *Mathematics and Financial Economics*, 12(3):335–363, 2018.

- Pierre Cardaliaguet, François Delarue, Jean-Michel Lasry, and Pierre-Louis Lions. *The Master Equation and the Convergence Problem in Mean Field Games*. Annals of Mathematics Studies. Princeton University Press, 2019.
- Pierre Cardaliaguet, Benjamin Seeger, and Panagiotis E. Souganidis. Regularity for mean field games with common noise and degenerate idiosyncratic noise. *Preprint*, 2025.
- René Carmona and François Delarue. *Probabilistic Theory of Mean Field Games with Applications I-II*. Probability Theory and Stochastic Modelling. Springer, 2018.
- René Carmona and Peiqi Wang. Mean field games for fixed income markets. *Preprint*, 2021.
- René Carmona, François Delarue, and Aimé Lachapelle. Control of mckean-vlasov dynamics versus mean field games. *Mathematics and Financial Economics*, 7:131–166, 2014.
- René Carmona, Jean-Pierre Fouque, and Li-Hsien Sun. Mean field games and systemic risk. *Communications in Mathematical Sciences*, 13(4):911–933, 2015.
- Philippe Casgrain and Sebastian Jaimungal. Mean-field games with differing beliefs for algorithmic trading. *Mathematical Finance*, 30(3):995–1034, 2020.
- Jean-François Chassagneux, Dan Crisan, and François Delarue. Numerical method for fbsdes of mckean-vlasov type. *Annals of Applied Probability*, 29(3):1640–1684, 2019.
- Joseph Chen, Harrison Hong, and Jeremy C. Stein. Breadth of ownership and stock returns. *Journal of Financial Economics*, 66(2-3):171–205, 2002.
- Sinho Chewi, Murat A. Erdogdu, Mufan Li, Ruoqi Shen, and Shunshi Zhang. Analysis of langevin monte carlo from Poincaré to log-Sobolev. In *Conference on Learning Theory (COLT)*, 2022.
- Christa Cuchiero, Christoph Reisinger, and Stefan Rigger. Optimal bailout strategies resulting from the drift control of a reflected diffusion. *Preprint*, 2024.
- Marco Cuturi. Sinkhorn distances: lightspeed computation of optimal transport. *Advances in Neural Information Processing Systems*, 26, 2013.
- Marco Cuturi and Arnaud Doucet. Fast computation of wasserstein barycenters. *International Conference on Machine Learning*, 2014.
- Peyman Mohajerin Esfahani and Daniel Kuhn. Data-driven distributionally robust optimization using the wasserstein metric: performance guarantees and tractable reformulations. *Mathematical Programming*, 171(1-2):115–166, 2018.
- Nicolas Fournier and Arnaud Guillin. On the rate of convergence in wasserstein distance of the empirical measure. *Probability Theory and Related Fields*, 162(3-4):707–738, 2015.
- Andrea Frazzini, Ronen Israel, and Tobias J. Moskowitz. Trading costs. *Preprint / SSRN*, 2018.

- Scott Fujimoto, Herke van Hoof, and David Meger. Addressing function approximation error in actor-critic methods. *International Conference on Machine Learning*, 2018.
- Josselin Garnier, George Papanicolaou, and Tzu-Wei Yang. Large deviations for a mean field model of systemic risk. *SIAM Journal on Financial Mathematics*, 4(1):151–184, 2013.
- Emmanuel Gobet, Jean-Philippe Lemor, and Xavier Warin. A regression-based monte carlo method to solve backward stochastic differential equations. *Annals of Applied Probability*, 15(3):2172–2202, 2005.
- Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: off-policy maximum entropy deep reinforcement learning with a stochastic actor. *International Conference on Machine Learning*, 2018.
- Daniel Hernández-Hernández and Joshué Helí Ricalde-Guerrero. Mean-field games with common poissonian noise: A maximum principle approach, 2024. arXiv:2401.10952.
- Mingyang Hu and Mathieu Laurière. Two-timescale reinforcement learning for mean-field games with common noise. *Preprint*, 2023.
- Minyi Huang, Roland P. Malhamé, and Peter E. Caines. Large population stochastic dynamic games: closed-loop mckean-vlasov systems and the nash certainty equivalence principle. *Communications in Information and Systems*, 6(3):221–252, 2006.
- Sebastian Jaimungal and Mojtaba Nourian. Mean-field game strategies for optimal execution. *Applied Mathematical Finance*, 2015.
- Amir E. Khandani and Andrew W. Lo. What happened to the quants in august 2007? evidence from factors and transactions data. *Journal of Financial Markets*, 14(1):1–46, 2011.
- Soohun Kim, Robert A. Korajczyk, and Andreas Neuhierl. Arbitrage portfolios. *Preprint*, 2015.
- Sascha Lange, Thomas Gabel, and Martin Riedmiller. Batch reinforcement learning. *Reinforcement Learning: State-of-the-Art (Springer)*, pages 45–73, 2012.
- Jean-Michel Lasry and Pierre-Louis Lions. Mean field games. *Japanese Journal of Mathematics*, 2(1):229–260, 2007.
- Charles-Albert Lehalle and Charafeddine Mouzouni. A mean field game of portfolio trading and its consequences on perceived correlations. *arXiv preprint arXiv:1902.09606*, 2019.
- Min Li, Tianyang Nie, Shujun Wang, and Ke Yan. Incomplete information mean-field games and related riccati equations. *Journal of Optimization Theory and Applications*, 203(3):2487–2508, 2024.
- Gabriel J. Lord, Catherine E. Powell, and Tony Shardlow. *An Introduction to Computational Stochastic PDEs*. Cambridge University Press, 2014.

- Iosif Lytras and Sotirios Sabanis. Taming under isoperimetry. *Stochastic Processes and their Applications*, 2025. arXiv:2311.09003.
- Yi-An Ma, Wenpin Li, and Anru R. Zhang. Exponential convergence of entropy-regularized policy gradient under log-sobolev inequalities. *Preprint*, 2024.
- Qirui Mi, Zhiyu Zhao, Chengdong Ma, Siyu Xia, Yan Song, Mengyue Yang, Jun Wang, and Haifeng Zhang. Learning macroeconomic policies through dynamic stackelberg mean-field games, 2025. ECAI 2025; arXiv:2403.12093.
- Eyal Neuman and Moritz Voß. Trading with the crowd. *Mathematical Finance*, 2023.
- Doina Precup, Richard S. Sutton, and Satinder Singh. Eligibility traces for off-policy policy evaluation. *International Conference on Machine Learning*, 2000.
- Zhenjie Ren and Dena Firoozi. Risk-sensitive mean field games with common noise: applications to interbank markets. *Preprint*, 2024.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018.
- Cédric Villani. *Optimal Transport: Old and New*. Grundlehren der mathematischen Wissenschaften. Springer, 2009.
- Hoang Huy Vu and Tomoyuki Ichiba. Heterogenous macro-finance model: A mean-field game approach, 2025. arXiv:2502.10666.
- Manzil Zaheer, Satwik Kottur, Siamak Ravanbakhsh, Barnabás Póczos, Ruslan Salakhutdinov, and Alexander J. Smola. Deep sets. *Advances in Neural Information Processing Systems*, 30, 2017.
- Sihan Zeng et al. Learning in herding mean field games. *Preprint*, 2024.
- Jianfeng Zhang. *Backward Stochastic Differential Equations: From Linear to Fully Non-linear Theory*. Springer, 2017.
- Mo Zhou, Zihan Zhou, and Ruimeng Hu. Mean-field actor-critic flows for continuous-time reinforcement learning in mean-field games. *Preprint*, 2025.