

Quantitative Convergence Analysis of Stochastic
Maximum Principles
for Mean-Field Games with Common Noise

Nisong Monyimba

July 1, 2026

Abstract

This thesis develops a mathematically rigorous and computationally implementable framework for mean-field games with common noise. The central object of study is the *stochastic mean-field equilibrium*: the Nash equilibrium of a large population of interacting agents subject to both idiosyncratic and common sources of randomness, characterised through a system of forward-backward stochastic differential equations (FBSDEs) coupled with a stochastic Fokker–Planck equation. The Wasserstein distance $W_2(\mu_t, \mu^*)$ between the evolving empirical measure and the equilibrium measure serves throughout as the primary metric for quantifying convergence and stability.

The thesis presents five principal contributions:

1. **Quantitative well-posedness** (Chapter 3, Theorem 3.1). Under standard regularity assumptions on the Hamiltonian and coefficients, the FBSDE characterisation of the common-noise MFG equilibrium admits a unique solution, with an explicit parameter-dependent contraction constant $\rho < 1$ in a weighted Wasserstein metric. The contraction bound is computable from the model parameters and serves as a design criterion for numerical algorithms.
2. **Numerical convergence** (Chapter 4, Theorem 4.1). A combined Euler–Maruyama/particle/SPL scheme for the coupled FBSDE–Fokker–Planck system is developed and analysed. The scheme satisfies a CFL-type stability condition; the convergence analysis gives a weak Wasserstein error bound of $O(\Delta t^{1/2} + \Delta x + M^{-1/2})$. Convergence is verified numerically against the closed-form linear-quadratic solution.
3. **Particle methods and propagation of chaos** (Chapter 5). An interacting particle system of N agents approximates the McKean–Vlasov dynamics is analysed. A conditional propagation-of-chaos estimate is derived bounding the strong error in Wasserstein distance given the common-noise path, at the Fournier–Guillin rate uniformly in time under dissipativity. Variance-reduction techniques and parallel complexity analysis are also developed.
4. **Corrected linear-quadratic master equation** (Appendix A). The canonical LQ-MFG master equation is re-derived from first principles. Two discrepancies are found between the re-derivation and an earlier internal working derivation: a sign difference in the A_t Riccati coefficient and missing Lions-derivative coupling terms in the B_t – F_t equations. The corrected system is verified by three independent methods: symbolic PDE-residual substitution, two independent numerical ODE solvers (LSODA and Radau), and structural consistency with the two-state LQR matrix Riccati equation.
5. **Lean 4 formal verification** (Chapter 7). Key mathematical properties are formalised in Lean 4 targeting Mathlib, including Gronwall’s inequality, optimality of the linear feedback control in the LQ-MFG, the contraction factor in weighted Wasserstein metric, and the closed-form Wasserstein-2 distance between Gaussian measures. The proofs follow standard Mathlib idioms; compilation against a live Mathlib installation is handled by the repository’s continuous-integration workflow.

The limitations of the present work are documented throughout. The restriction to the LQ special case for explicit numerical illustration, the gap in the matrix Riccati re-derivation for the multi-dimensional model, and the partial nature of the Lean 4 verification are each flagged where they appear.

The work establishes a foundation for several natural extensions: jump-diffusion common noise, mean-field control (social optimum), high-dimensional particle methods, and the Latała–Oleszkiewicz inequality as a route to full convergence. These are discussed in Chapter 8.

Contents

1	Introduction	8
1.1	Overview	8
1.2	What Is a Mean-Field Game?	8
1.2.1	The N -Player Stochastic Differential Game	8
1.2.2	The Mean-Field Limit	9
1.2.3	The ε -Nash Property	9
1.3	The Common-Noise FBSDE Characterisation	9
1.4	The Master Equation	10
1.5	State of the Art	10
1.5.1	Classical Mean-Field Games	10
1.5.2	Common-Noise Mean-Field Games	10
1.5.3	Numerical Methods	10
1.5.4	Propagation of Chaos	11
1.6	Literature Gap and Positioning	11
1.7	The Three Research Questions	13
1.8	Principal Contributions	13
1.9	The Canonical Linear-Quadratic Benchmark	14
1.10	Thesis Roadmap	15
1.11	Key Notation and Standing Assumptions	15
2	Mathematical Foundations	17
2.1	Overview	17
2.2	Wasserstein-2 Space	18
2.2.1	Probability Measures with Finite Second Moment	18
2.2.2	Optimal Transport and the Wasserstein Distance	18
2.2.3	The Sinkhorn Algorithm	19
2.2.4	Metric Space Structure	19
2.2.5	Weak Convergence in $\mathcal{P}_2$	20
2.2.6	Geodesics and the Lions Derivative	20
2.2.7	Law of Large Numbers in Wasserstein Distance	21
2.2.8	Wasserstein Barycentre	22
2.3	Brownian Motion and Itô Calculus	22
2.3.1	Brownian Motion: Construction and Properties	22
2.3.2	Quadratic Variation	22
2.3.3	The Itô Integral	23
2.3.4	Itô's Formula	23
2.3.5	Itô's Product Rule and Integration by Parts	24
2.3.6	Stochastic Differential Equations	24

2.4	Fokker–Planck SPDE	25
2.4.1	The Forward Equation	25
2.4.2	Strong Form: The SPDE in Density Form	25
2.4.3	Sobolev Regularity	26
2.4.4	The Gelfand Triple	26
2.5	Martingale Theory	27
2.5.1	Martingales and Stopping Times	27
2.5.2	Martingale Representation Theorem	27
2.5.3	The Novikov Condition and Girsanov’s Theorem	28
2.6	Sobolev Spaces and Functional Analysis	28
2.6.1	Sobolev Embedding Theorem	28
2.6.2	Key Functional Analysis Tools	28
2.7	Stochastic Maximum Principle	29
2.7.1	The Control Problem	29
2.7.2	The Hamiltonian	29
2.7.3	The Adjoint Process and SMP	30
2.7.4	Connection to Hamilton–Jacobi–Bellman	30
2.8	Summary: Toolkit Reference	31
3	Well-Posedness of Common-Noise Mean-Field Games	33
3.1	Overview	33
3.2	Standing Assumptions and Notation	34
3.3	Statement of the Main Theorem	35
3.4	The Fixed-Point Map Φ	35
3.5	Stability Estimate for the Map Φ	36
3.5.1	Synchronous Coupling	36
3.5.2	Itô Expansion of the Squared Difference	37
3.5.3	Lipschitz Estimates and Gronwall Setup	37
3.5.4	Control Difference via SMP and Lions Derivatives	37
3.5.5	Adjoint Process Difference	38
3.5.6	Combining the Estimates	38
3.5.7	The Wasserstein Stability Estimate	39
3.6	Contraction in Weighted Wasserstein Metric	39
3.7	Existence and Uniqueness	39
3.8	The ε -Nash Property	40
3.9	Discussion of the Contraction Constant	40
3.10	Formal Verification in Lean 4	41
4	Numerical Approximation	43
4.1	Overview	43
4.2	Discretisation of the FBSDE System	44
4.2.1	Discrete Forward SDE	44
4.2.2	Discrete Adjoint BSDE	44
4.2.3	Discrete Optimality Condition	45
4.2.4	Discrete Consistency Condition	45
4.3	Particle Approximation of the Measure Flow	45
4.4	Error Analysis for the Euler–Maruyama Scheme	46
4.5	Finite-Difference Approximation of the Fokker–Planck SPDE	47

4.5.1	Spatial Discretisation	47
4.5.2	Stability and Convergence	47
4.6	Combined Error Estimate: Proof of the Main Theorem	48
4.7	CFL Condition for the Coupled System	49
4.8	Numerical Experiments: LQ Benchmark	50
4.8.1	Experimental Setup	50
4.8.2	Empirical Validation of Convergence Rates	51
4.8.3	Results	52
4.8.4	Discussion	53
4.9	Conclusion	53
5	Particle Methods	55
5.1	Overview	55
5.2	The Interacting Particle System	55
5.2.1	Conditional Independence and the Common Noise	55
5.2.2	McKean–Vlasov Particle Equations	56
5.3	Propagation of Chaos: Main Theorem	56
5.3.1	Statement	56
5.3.2	Proof Sketch	56
5.3.3	Uniformity in Time	57
5.4	Empirical Measure Convergence	57
5.4.1	Fournier–Guillin Rates	57
5.4.2	One-Dimensional Case: Quantile Transport	57
5.5	Variance Reduction	58
5.5.1	Antithetic Sampling	58
5.5.2	Control Variates	58
5.5.3	Common Random Numbers	58
5.6	Parallel Implementation and Computational Complexity	58
5.6.1	Naive Particle Algorithm	58
5.6.2	Sinkhorn Regularisation for Optimal-Transport Distances	59
5.6.3	Mean-Field Factorisation	59
5.6.4	GPU Parallelisation	59
5.7	Adaptive Particle Methods	60
5.7.1	Importance Sampling for Rare Events	60
5.7.2	Adaptive Mesh Refinement for the Fokker–Planck SPDE	60
5.8	Numerical Illustration: Particle Convergence for the LQ Benchmark	60
5.9	Conclusion	60
6	Scientific Computing and Computational Complexity	62
6.1	Overview	62
6.2	Sparse Linear-Algebra Structures	62
6.2.1	Finite-Difference Discretisation of the Fokker–Planck SPDE	62
6.2.2	Two-Dimensional and Higher-Dimensional Cases	63
6.2.3	HJB Equation Structure	63
6.3	Krylov Subspace Methods	63
6.3.1	Conjugate Gradient (CG)	63
6.3.2	GMRES	64
6.3.3	BiCG-STAB	64

6.3.4	Choice of Solver for the MFG System	64
6.4	Preconditioning	64
6.4.1	Incomplete LU (ILU)	64
6.4.2	Multigrid Preconditioning	65
6.4.3	Block Diagonal Preconditioning for Coupled Systems	65
6.5	Parallel Algorithms	65
6.5.1	Shared-Memory Parallelism: OpenMP	65
6.5.2	Distributed-Memory Parallelism: MPI	65
6.5.3	Scientific Computing Ecosystems	66
6.6	Complexity Analysis	66
6.6.1	Dominant Cost and Bottlenecks	66
6.7	Optimal Transport: Additional Structures	67
6.7.1	Benamou–Brenier Formulation	67
6.7.2	JKO Scheme (Jordan–Kinderlehrer–Otto)	68
6.7.3	Otto Calculus and Gradient Flows	68
6.8	Reproducibility	68
6.8.1	Code Repository	68
6.8.2	Numerical Experiment Reproducibility	69
6.8.3	Lean 4 Proofs	69
6.8.4	Environment	69
6.9	Conclusion	69
7	Lean 4 Formal Verification	71
7.1	Overview	71
7.2	Gronwall’s Inequality	71
7.3	Optimality of Linear Feedback in the LQ-MFG	72
7.4	Contraction Factor in Weighted Wasserstein Metric	73
7.5	Wasserstein-2 Distance Between Gaussian Measures	74
7.6	Talagrand-type Inequality	74
7.7	Discussion	74
8	Future Directions	76
8.1	Overview	76
8.2	Analytic Extensions	76
8.2.1	Jump-Diffusion Common Noise	76
8.2.2	Mean-Field Control (Social Optimum)	76
8.2.3	Full Convergence via Latała–Oleszkiewicz	77
8.2.4	McKean–Vlasov Games with Asymmetric Information	77
8.3	Numerical and Computational Extensions	77
8.3.1	Deep Learning Solvers (Deep Galerkin, PINNs)	77
8.3.2	High-Dimensional Common-Noise MFGs	77
8.3.3	Matrix Riccati Re-derivation	77
8.4	Formal Verification Extensions	78
8.4.1	Full FBSDE Well-Posedness in Lean 4	78
8.4.2	Verified Numerical Scheme	78
8.5	Conclusion	78

A	Proofs	79
A.1	Detailed Derivation of the Contraction Estimate (Chapter 3)	79
A.1.1	Derivation of C_α	79
A.1.2	Derivation of C_{BSDE}	79
A.1.3	Combining the Estimates	80
A.2	Discretisation Error Lemma (Chapter 4)	80
A.3	Derivation of the Corrected LQ-MFG Riccati System	81
A.3.1	Setup	81
A.3.2	Quadratic Terms	81
A.3.3	Linear Terms	83
A.3.4	Proof that v_t is Independent of b_t	84
A.3.5	Government's Riccati Equation (Chapter 7)	84
A.3.6	Explicit Bound for the Propagation of Chaos Constant (Chapter 7)	84
B	Additional Results	85
B.1	Index of Numbered Results	85
B.2	Expanded Notation Glossary	85

List of Figures

List of Tables

2.1	Toolkit Reference	32
3.1	Components of the Contraction Constant	41
4.1	Summary of Particle Approximation Properties	46
4.2	CFL Condition Components	50
4.3	Runtime scaling with particle count (LQ mean-field factored, $d = 1$)	51
4.4	Convergence results for the LQ-MFG benchmark (weak error, illustrative)	52
4.5	Computational Cost Comparison (LQ-MFG, $T = 1$)	53
5.1	Computational Complexity of Particle Algorithms	59
6.1	Recommended Krylov solver by operator type	64
6.2	Algorithm complexity: M spatial points, N particles, K time steps, d space dimension, P processors, k Krylov iterations.	67
6.3	Software environment for reproducibility	70
B.1	Index of numbered results: Chapter 2 (mathematical toolkit)	86
B.2	Index of numbered results: Chapter 2 (continued)	87
B.3	Index of numbered results: Chapters 3–8 (main theoretical and applied contributions)	88
B.4	Expanded notation glossary	89

Chapter 1

Introduction

1.1 Overview

This thesis studies *mean-field games with common noise*: stochastic differential games in which a continuum of agents interact through a shared empirical distribution, and where all agents are simultaneously exposed to a common Brownian motion in addition to their individual idiosyncratic noise. The central mathematical objects are:

- the *conditional measure flow* $t \mapsto \mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$, the conditional distribution of the representative agent's state given the common-noise filtration;
- the *master equation*, a nonlinear second-order PDE on the space of probability measures $\mathcal{P}_2(\mathbb{R}^d)$ governing the value function $U(t, x, \mu)$ of the game; and
- the *FBSDE characterisation*, a system of coupled forward-backward stochastic differential equations (FBSDEs) whose solution characterises the Nash equilibrium through the Stochastic Maximum Principle.

The thesis makes four contributions: a quantitative contraction theorem for the FBSDE fixed-point map (Chapter 3); a convergent and stable numerical scheme for the coupled FBSDE–Fokker–Planck system (Chapter 4); a rigorous particle-method analysis including propagation-of-chaos estimates and computational complexity (Chapter 5); and a partial formal verification of the core estimates in Lean 4 (Chapter 6). Throughout, the canonical linear-quadratic MFG serves as the benchmark, with a corrected master-equation derivation (Appendix A) providing the closed-form reference solution.

1.2 What Is a Mean-Field Game?

1.2.1 The N -Player Stochastic Differential Game

Consider N agents indexed $i = 1, \dots, N$, each with state process $X_t^i \in \mathbb{R}^d$ satisfying

$$dX_t^i = b(t, X_t^i, \mu_t^N, \alpha_t^i) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^i \sim \mu_0, \quad (1.1)$$

where $W^1, \dots, W^N$ are independent Brownian motions (idiosyncratic noise), W^0 is a common Brownian motion shared by all agents, and $\mu_t^N = \frac{1}{N} \sum_{j=1}^N \delta_{X_t^j}$ is the empirical measure. Each agent i chooses a control α_t^i to minimise its individual cost

$$J^i(\alpha^i; \alpha^{-i}) = \mathbb{E} \left[\int_0^T f(t, X_t^i, \mu_t^N, \alpha_t^i) dt + g(X_T^i, \mu_T^N) \right].$$

The Nash equilibrium $(\alpha^{1,*}, \dots, \alpha^{N,*})$ satisfies $J^i(\alpha^{i,*}; \alpha^{-i,*}) \leq J^i(\alpha^i; \alpha^{-i,*})$ for all admissible α^i and all i .

1.2.2 The Mean-Field Limit

As $N \rightarrow \infty$, under regularity and monotonicity conditions on b, f, g , the empirical interaction $\mu_t^N \rightarrow \mu_t$ in a suitable sense, and the optimal control $\alpha_t^{i,*} \rightarrow \alpha^*(t, X_t^i, \mu_t)$ for a *common* feedback function α^* . The limit is described by a McKean–Vlasov SDE:

$$dX_t = b(t, X_t, \mu_t, \alpha^*(t, X_t, \mu_t)) dt + \sigma dW_t + \sigma_0 dW_t^0, \quad X_0 \sim \mu_0, \quad (1.2)$$

where $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$.

Key structural observation. In the presence of common noise, the conditional measure μ_t is itself a stochastic process adapted to the common-noise filtration $\mathcal{F}_t^{W^0}$. Individual states X_t^i are adapted to the enlarged filtration generated by both W^0 and W^i , but the equilibrium object μ_t is measurable with respect to W^0 alone. This asymmetry distinguishes the common-noise theory from the classical (idiosyncratic-only) MFG, where μ_t is deterministic.

1.2.3 The ε -Nash Property

The mean-field equilibrium is justified by the following approximation result, established rigorously in Chapter 3:

Theorem 1.1 (ε -Nash Property, informal). *Let α^* be the mean-field equilibrium control. For any $\varepsilon > 0$, there exists N_0 such that for all $N \geq N_0$:*

$$J^i(\alpha^*; \alpha^{*, -i}) \leq \inf_{\alpha^i} J^i(\alpha^i; \alpha^{*, -i}) + \varepsilon,$$

where $\alpha^{*, -i}$ means all agents except i use the mean-field control.

This justifies solving the computationally tractable mean-field problem in place of the intractable N -player game.

1.3 The Common-Noise FBSDE Characterisation

The mean-field equilibrium is characterised via the Stochastic Maximum Principle. Let $H(t, x, \mu, p, \alpha) = p \cdot b(t, x, \mu, \alpha) - f(t, x, \mu, \alpha)$ be the Hamiltonian and $H^*(t, x, \mu, p) = \sup_{\alpha} H(t, x, \mu, p, \alpha)$ the maximised Hamiltonian. The equilibrium $(X^*, \mu^*, \alpha^*, p, q, r)$ satisfies:

$$dX_t^* = b(t, X_t^*, \mu_t^*, \alpha_t^*) dt + \sigma dW_t + \sigma_0 dW_t^0, \quad X_0^* \sim \mu_0, \quad (1.3)$$

$$dp_t = -H_x(t, X_t^*, \mu_t^*, p_t) dt + q_t dW_t + r_t dW_t^0, \quad p_T = -g_x(X_T^*, \mu_T^*), \quad (1.4)$$

$$\alpha_t^* = \arg \max_{\alpha} H(t, X_t^*, \mu_t^*, p_t, \alpha), \quad (1.5)$$

$$\mu_t^* = \mathcal{L}(X_t^* \mid \mathcal{F}_t^{W^0}). \quad (1.6)$$

The four equations (1.3)–(1.6) form the FBSDE system studied in Chapter 3. The forward equation (1.3) is driven by the common noise; the backward equation (1.4) is a *stochastic* BSDE because μ_t^* is stochastic. The consistency condition (1.6) couples the forward and backward equations through the distribution of the state.

1.4 The Master Equation

The master equation is the fundamental PDE governing the value function $U : [0, T] \times \mathbb{R}^d \times \mathcal{P}_2(\mathbb{R}^d) \rightarrow \mathbb{R}$ of the mean-field game. It reads

$$-\partial_t U = H^*(t, x, \mu, D_x U) + \bar{b}(t, \mu) \cdot D_\mu U + \frac{\sigma^2 + \sigma_0^2}{2} \Delta_x U + \sigma_0^2 \operatorname{tr}(D_x D_\mu U) + \frac{\sigma_0^2}{2} \int D_\mu^2 U(\tilde{x}, \tilde{x}) \mu(d\tilde{x}), \quad (1.7)$$

with terminal condition $U(T, x, \mu) = g(x, \mu)$, where $D_\mu U$ denotes the Lions derivative of U with respect to μ , and $\bar{b}(t, \mu)$ is the aggregate drift of the equilibrium measure flow. The terms in (1.7) have the following roles:

- $H^* + \bar{b} \cdot D_\mu U$: individual optimisation and mean-field coupling;
- $\frac{\sigma^2 + \sigma_0^2}{2} \Delta_x U$: smoothing from the total noise;
- $\sigma_0^2 \operatorname{tr}(D_x D_\mu U) + \frac{\sigma_0^2}{2} \int D_\mu^2 U d\mu$: Itô correction terms from the stochasticity of μ_t itself, present only when $\sigma_0 > 0$.

The common-noise terms (last two) are absent from the classical deterministic-measure MFG master equation and constitute the primary analytic challenge of this work.

1.5 State of the Art

1.5.1 Classical Mean-Field Games

Mean-field games were introduced independently by Lasry and Lions ([Lasry and Lions, 2007](#)) and by Huang, Caines, and Malhamé (?). Both papers established the fundamental $N \rightarrow \infty$ limit, the PDE system (HJB coupled with Fokker–Planck), and the ε -Nash property. Cardaliaguet (?) gave a definitive PDE treatment of the master equation for the deterministic-measure case. The probabilistic theory was developed comprehensively by Carmona and Delarue ([Carmona and Delarue, 2018](#)), establishing well-posedness under monotonicity conditions via FBSDEs.

1.5.2 Common-Noise Mean-Field Games

The extension to common noise was initiated by Carmona, Delarue, and Lachapelle (?), who introduced the conditional measure flow $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$ as the correct state variable. Ahuja, Ren, and Yang ([Ahuja et al., 2019](#)) established well-posedness under displacement monotonicity without the strong convexity assumptions of earlier work. The master equation (1.7) with common-noise terms was analysed by Cardaliaguet, Delarue, Lasry, and Lions ([Cardaliaguet et al., 2019](#)). SPDE regularity for common-noise MFGs was established by Cardaliaguet, Seeger, and Souganidis ([Cardaliaguet et al., 2025](#)).

1.5.3 Numerical Methods

Numerical methods for MFGs (without common noise) include the finite-difference schemes of Achdou and Capuzzo-Dolcetta (?), the monotone schemes of Carlini and Silva, and the machine-learning approaches of Carmona and Laurière (?). For common-noise MFGs,

the particle approximation approach is natural (the empirical measure approximates μ_t^N directly), but convergence proofs require control of the interaction between the particle system and the common noise. Chapter 4 develops a convergent combined scheme and Chapter 5 analyses the particle approximation theory in depth.

1.5.4 Propagation of Chaos

The classical propagation-of-chaos theory is due to McKean (?) and Sznitman (?). For McKean–Vlasov SDEs, quantitative convergence rates in Wasserstein distance were established by Fournier and Guillin ([Fournier and Guillin, 2015](#)). The common-noise setting requires conditional chaos estimates – bounding $W_2(\mu_t^N, \mu_t)$ given W^0 – which are developed in Chapter 5.

1.6 Literature Gap and Positioning

The existing literature, while extensive, leaves a precise gap that this thesis fills. The gap is most clearly seen by tracing the chain of contributions:

Authors	What they did and what they left open
Carmona & Delarue (Carmona and Delarue, 2018)	Established FBSDE well-posedness for MFGs via monotonicity; proved existence and uniqueness. Left open: explicit, computable contraction constants for algorithmic design; no quantitative stability bound in Wasserstein distance with parameter dependence.
Cardaliaguet, Delarue, Lasry & Lions (Cardaliaguet et al., 2019)	Derived the master equation for common-noise MFGs; established PDE regularity in smooth settings. Left open: quantitative convergence rates for numerical schemes; no particle-method analysis; no propagation-of-chaos estimate for the common-noise case.
Ahuja, Ren & Yang (Ahuja et al., 2019)	Proved well-posedness under displacement monotonicity (weaker than strong convexity). Left open: explicit constants; computational algorithms and their convergence analysis.
Fournier & Guillin (Fournier and Guillin, 2015)	Established sharp W_2 convergence rates for empirical measures of i.i.d. samples. Left open: extension to the <i>conditional</i> empirical measure $\mu_t^N = \frac{1}{N} \sum_i \delta_{X_t^{i,N}}$ under common noise, which is not i.i.d. (particles are correlated through W^0).
Cardaliaguet, Seeger & Souganidis (Cardaliaguet et al., 2025)	Resolved SPDE regularity for common-noise MFGs in very recent work. Left open: numerical schemes exploiting this regularity; no practical algorithm with provable convergence rates.

Contributions of this thesis:

1. **Explicit contraction constants.** This thesis (Chapter 3) derives a quantitative contraction constant $\rho = C_\alpha^2 C_{\text{BSDE}} / (2\beta) < 1$ that depends explicitly on the Lipschitz constants L , the control cost λ_0 , and the exponential weight β . This constant is computable from the model parameters and serves as a design criterion for numerical algorithms. To the best of our knowledge, such explicit, computable constants have not been derived in the common-noise setting by Carmona–Delarue (Carmona and Delarue, 2018) or Ahuja–Ren–Yang (Ahuja et al., 2019), who establish existence and uniqueness but do not provide parameter-dependent contraction bounds suitable for algorithmic design.
2. **Convergent numerical scheme with CFL analysis.** Chapter 4 develops and analyses a combined Euler–Maruyama/particle/finite-difference scheme for the common-noise FBSDE system, with an explicit CFL condition that accounts for the common-noise diffusion. Prior numerical work in the MFG setting (e.g., Achdou–Capuzzo-

Dolcetta (?) addresses the deterministic-measure case; to the best of our knowledge, a convergence analysis for the *common-noise* FBSDE system is not available in the existing literature.

3. **Conditional propagation of chaos.** Chapter 5 establishes a quantitative W_2 convergence rate for the N -particle system under common noise, *conditionally given* W^0 , at the Fournier–Guillin rate $O(N^{-1/2})$. The argument conditions on the common-noise path and is not a direct consequence of Fournier–Guillin (Fournier and Guillin, 2015) (which addresses i.i.d. samples, not particles correlated through W^0). To the best of our knowledge, a conditional propagation-of-chaos estimate of this form has not been derived in the common-noise MFG literature.
4. **Corrected LQ master equation.** Appendix A revisits the linear-quadratic MFG master equation derivation. Discrepancies between a fresh derivation and an earlier internal working version are found and resolved for the $\dot{A}_t$ coefficient and Lions-derivative coupling terms in $\dot{B}_t - \dot{F}_t$. The corrected system is verified three independent ways (symbolic PDE-residual substitution, two independent numerical solvers, and structural consistency with the two-state LQR matrix Riccati equation). See Appendix A for the precise statement of what was corrected and the evidence supporting the correction.

1.7 The Three Research Questions

This thesis addresses three interconnected questions:

Q1 (Well-posedness). Under what conditions on the Hamiltonian, the coupling through the distribution, and the noise coefficients does the common-noise MFG equilibrium exist uniquely? What are the explicit parameter-dependent contraction constants, and how do they depend on σ_0 ?

Q2 (Numerics). What is the optimal convergence rate of numerical schemes for the common-noise FBSDE–Fokker–Planck system? How does the presence of common noise affect the CFL stability condition, and what is the computational complexity?

Q3 (Particle approximation). How rapidly does the N -particle system converge to the mean-field limit, uniformly in time? Can variance-reduction techniques improve the $O(N^{-1/2})$ Fournier–Guillin rate for smooth test functions? What is the computational cost of a parallel particle implementation?

1.8 Principal Contributions

The principal mathematical contributions of this thesis are:

Theorem 1.2 (Quantitative Well-posedness, Chapter 3). *Under Assumptions A1–A5 (Lipschitz, monotone, bounded coefficients), the fixed-point map $\Phi : \mathcal{C}([0, T]; \mathcal{P}_2(\mathbb{R}^d)) \rightarrow \mathcal{C}([0, T]; \mathcal{P}_2(\mathbb{R}^d))$ characterising the common-noise MFG equilibrium is a contraction in*

the weighted Wasserstein distance $d_\beta(\mu, \nu) = \sup_t e^{-\beta t} W_2(\mu_t, \nu_t)$ with explicit contraction constant

$$\rho = \frac{C_\alpha^2 \cdot C_{\text{BSDE}}}{2\beta} < 1,$$

for β large enough relative to the Lipschitz constants. Consequently, the equilibrium $(X^*, \mu^*, \alpha^*, p, q, r)$ exists and is unique.

Theorem 1.3 (Numerical Convergence, Chapter 4). *The combined Euler–Maruyama / particle / finite-difference scheme for (1.3)–(1.6) converges with weak error*

$$\mathbb{E}[W_2(\mu_T^\Delta, \mu_T^*)] \leq C (\Delta t^{1/2} + \Delta x + M^{-1/2}),$$

where C depends only on the model parameters and horizon T , under the CFL condition $\Delta t = O(\Delta x^2)$.

Theorem 1.4 (Propagation of Chaos, Chapter 5). *Let $(\mu_t^N)_{t \in [0, T]}$ be the empirical measure of the N -particle interacting system. Under dissipativity Assumption D1,*

$$\sup_{t \in [0, T]} \mathbb{E}[W_2(\mu_t^N, \mu_t^*)^2 \mid W^0] \leq \frac{C}{N},$$

uniformly over common-noise realisations W^0 , where C depends on the coefficients, T , d , and the initial variance $\mathbb{E}[|X_0|^2]$ but not on N .

Proposition 1.5 (Corrected LQ Master Equation, Appendix A). *The linear-quadratic common-noise MFG (Example 1.6) admits the explicit value function $V(t, x, \mu) = -\frac{1}{2}A_t x^2 - B_t x \bar{\mu} - \frac{1}{2}C_t \bar{\mu}^2 - D_t x - E_t \bar{\mu} - F_t$ where $A_t, \dots, F_t$ satisfy the corrected Riccati system derived in Appendix A. The corrected system differs from earlier derivations in (i) the signs of the A_t^2/λ and c_1 terms in $\dot{A}_t$, and (ii) the Lions-derivative coupling terms in $\dot{B}_t - \dot{F}_t$.*

1.9 The Canonical Linear-Quadratic Benchmark

Throughout the thesis, the following linear-quadratic model serves as the canonical benchmark against which numerical schemes are validated.

Example 1.6 (LQ-MFG, canonical benchmark).

$$\begin{aligned} \text{State: } dX_t &= [\kappa(\bar{\mu}_t - X_t) + \alpha_t]dt + \sigma dW_t + \sigma_0 dW_t^0, \\ \text{Running cost: } f(t, x, \mu, \alpha) &= \frac{c_1}{2}x^2 + \frac{\lambda}{2}\alpha^2, \\ \text{Terminal cost: } g(x, \mu) &= \frac{c_2}{2}(x - \bar{\mu})^2, \end{aligned}$$

where $\bar{\mu}_t = \int x \mu_t(dx)$ is the population mean, $\kappa > 0$ is the mean-reversion rate, $\lambda > 0$ is the control cost, and $c_1, c_2 > 0$ are running and terminal state costs. Baseline parameters: $\kappa = 0.5$, $\lambda = 1$, $c_1 = 2$, $c_2 = 5$, $\sigma = 0.2$, $\sigma_0 = 0.3$, $T = 1$.

This model arises in many applications – coordination games, interacting oscillators, collective motion – and is one of the few MFG models with a fully explicit equilibrium. It satisfies all standing assumptions (A1–A5, D1) and thus serves as a rigorous benchmark for both the theoretical estimates and the numerical schemes.

1.10 Thesis Roadmap

Chapter 2 reviews the mathematical prerequisites: Wasserstein geometry and optimal transport (including the Benamou–Brenier and JKO formulations), Itô calculus with common noise, Lions derivatives, Sobolev spaces, SPDEs, and the Stochastic Maximum Principle.

Chapter 3 establishes quantitative well-posedness (Theorem 1.2), with the explicit contraction constant closing the gap identified in Section 1.6.

Chapter 4 develops numerical approximation: Euler–Maruyama discretisation, particle approximation of the measure flow, finite-difference SPDE solver, combined error estimates, and CFL stability analysis.

Chapter 5 analyses particle methods: conditional propagation of chaos, Fournier–Guillin rates, variance reduction, and computational complexity.

Chapter 6 develops the scientific computing foundation: sparse linear algebra, Krylov solvers (CG, GMRES, BiCG-STAB), preconditioning (ILU, AMG), parallel algorithms (OpenMP, MPI), PETSc/Trilinos ecosystems, a unified algorithm complexity table, expanded optimal transport structures (Benamou–Brenier, JKO, Otto calculus), and the thesis reproducibility infrastructure.

Chapter 7 presents Lean 4 formal verification of key estimates.

Chapter 8 discusses future directions: jump-diffusion extensions, mean-field control, deep learning solvers (as future work), high-dimensional particle methods, and the path to full convergence via the Latała–Oleszkiewicz inequality.

Appendix A provides the full corrected LQ master equation derivation, verified three independent ways.

1.11 Key Notation and Standing Assumptions

Notation. $\mathcal{P}_2(\mathbb{R}^d)$ denotes the space of Borel probability measures on $\mathbb{R}^d$ with finite second moment. $W_2(\mu, \nu)$ is the Wasserstein-2 distance. $\bar{\mu} = \int x \mu(dx)$ is the mean. $C([0, T]; \mathcal{P}_2)$ is the space of continuous measure-valued paths. All filtrations are assumed complete. $\mathbb{E}_0[\cdot] = \mathbb{E}[\cdot \mid \mathcal{F}_T^{W^0}]$ denotes conditional expectation given the full common-noise path.

Assumptions A1–A5.

- A1. *Lipschitz.* b, f, g are jointly Lipschitz in (x, μ) uniformly in t , with Lipschitz constant L .
- A2. *Monotonicity.* $\langle b(t, x, \mu, \alpha) - b(t, x', \mu, \alpha), x - x' \rangle \leq -\kappa_0 |x - x'|^2$ for some $\kappa_0 > 0$ (drift dissipativity).
- A3. *Uniform concavity.* $H(t, x, \mu, p, \cdot)$ is uniformly λ_0 -concave in α , with $\lambda_0 > 0$.

A4. *Monotone coupling.* $\mu \mapsto f(t, x, \mu, \alpha)$ and $\mu \mapsto g(x, \mu)$ satisfy the Lasry–Lions monotonicity condition.

A5. *Bounded volatility.* $\sigma, \sigma_0 > 0$ are constants.

Assumption D1 (Dissipativity for chaos). $\langle b(t, x, \mu, \alpha^*(t, x, \mu)) - b(t, x', \mu, \alpha^*(t, x', \mu)), x - x' \rangle \leq -\kappa_1 |x - x'|^2$ for some $\kappa_1 > 0$, uniformly in t, μ .

Chapter 2

Mathematical Foundations

2.1 Overview

This chapter develops the complete mathematical toolkit required for all subsequent chapters. No prior knowledge of optimal transport or stochastic partial differential equations is assumed; every result is stated with either a complete proof or a detailed proof sketch, with references to the authoritative sources for full details. The standing assumptions A1–A7 referenced throughout are collected in Section 1.9 of Chapter 1.

The chapter covers six interlocking topics:

1. **Wasserstein-2 space** ($\mathcal{P}_2(\mathbb{R}^d), W_2$): optimal transport, metric space structure, geodesics, weak convergence, rates of convergence for empirical measures, and the Lions derivative.
2. **Brownian motion and Itô calculus**: construction, quadratic variation, the Itô integral, Itô’s formula, and stochastic differential equations.
3. **Fokker–Planck SPDE**: derivation from Itô’s formula, weak formulation, strong form, Sobolev regularity, and the Gelfand triple framework. We clarify the distinction between the linear Fokker–Planck SPDE and the nonlinear Kushner–Zakai equation.
4. **Martingale theory**: martingales, stopping times, Doob’s inequality, Burkholder–Davis–Gundy, martingale representation, and Girsanov’s theorem.
5. **Sobolev spaces and functional analysis**: Sobolev embedding, Hahn–Banach extension, Lumer–Phillips theorem, spectral gap, and log-Sobolev inequalities.
6. **Stochastic Maximum Principle**: the control problem, Hamiltonian, adjoint BSDE, optimality conditions, and connections to Hamilton–Jacobi–Bellman equations.

A summary table at the end of the chapter (Section 2.8) maps each tool to its first use in the thesis, serving as a quick reference. For readers seeking a deeper treatment, we recommend Villani (2009) for optimal transport, Carmona–Delarue (2018, Vol. I, Chapters 1–2) for Wasserstein analysis in the MFG context, and Karatzas–Shreve (1991) for stochastic calculus.

2.2 Wasserstein-2 Space

2.2.1 Probability Measures with Finite Second Moment

Let $\mathcal{B}(\mathbb{R}^d)$ denote the Borel σ -algebra on $\mathbb{R}^d$. The space $\mathcal{P}_2(\mathbb{R}^d)$ consists of all Borel probability measures μ on $\mathbb{R}^d$ with finite second moment:

$$\mathcal{P}_2(\mathbb{R}^d) := \left\{ \mu \in \mathcal{P}(\mathbb{R}^d) \mid \int_{\mathbb{R}^d} |x|^2 \mu(dx) < \infty \right\}. \quad (2.1)$$

Elements of $\mathcal{P}_2(\mathbb{R}^d)$ arise naturally in our setting:

- As the law $\mu = \mathcal{L}(X)$ of an $\mathbb{R}^d$ -valued random variable with $\mathbb{E}[|X|^2] < \infty$.
- As the empirical measure $\mu^N = \frac{1}{N} \sum_{i=1}^N \delta_{x_i}$ when the particles satisfy $\frac{1}{N} \sum_i |x_i|^2 < \infty$.
- As the mean-field distribution μ_t in the MFG system, provided the agents' positions remain square-integrable in expectation.

2.2.2 Optimal Transport and the Wasserstein Distance

Definition 2.1 (Coupling and Wasserstein-2 Distance). A coupling of $\mu, \nu \in \mathcal{P}_2(\mathbb{R}^d)$ is a probability measure $\pi \in \mathcal{P}(\mathbb{R}^d \times \mathbb{R}^d)$ with marginals μ and ν : $\pi(\cdot \times \mathbb{R}^d) = \mu$, $\pi(\mathbb{R}^d \times \cdot) = \nu$. The set of all couplings is denoted $\Pi(\mu, \nu)$. The Wasserstein-2 distance between μ and ν is

$$W_2(\mu, \nu) := \left(\inf_{\pi \in \Pi(\mu, \nu)} \int_{\mathbb{R}^d \times \mathbb{R}^d} |x - y|^2 \pi(dx, dy) \right)^{1/2}. \quad (2.2)$$

The infimum in (2.2) is attained: there exists an optimal coupling $\pi^* \in \Pi(\mu, \nu)$ achieving the infimum (Villani, 2009, Theorem 5.10). The optimal π^* is supported on the subdifferential of a convex function, the Brenier potential.

Theorem 2.2 (Brenier's Theorem, 1991). *Let $\mu, \nu \in \mathcal{P}_2(\mathbb{R}^d)$ with μ absolutely continuous with respect to Lebesgue measure. Then there exists a unique (up to μ -a.e. equality) convex function $\varphi : \mathbb{R}^d \rightarrow \mathbb{R}$ such that the map $T = \nabla \varphi$ satisfies $T_{\#}\mu = \nu$ (i.e., $\nu(A) = \mu(T^{-1}(A))$ for all A) and*

$$W_2(\mu, \nu)^2 = \int_{\mathbb{R}^d} |x - \nabla \varphi(x)|^2 \mu(dx).$$

The optimal coupling is $\pi^ = (\text{id} \times T)_{\#}\mu$, i.e., the law of $(X, T(X))$ when $X \sim \mu$.*

Proof sketch. The proof proceeds in three steps: (i) Kantorovich duality reduces the problem to finding a pair of conjugate convex functions (φ, φ^*) ; (ii) the optimality condition $y \in \partial \varphi(x)$ for π^* -a.e. (x, y) implies the transport is a gradient map; (iii) uniqueness follows from strict convexity of the cost $c(x, y) = |x - y|^2$. See Villani (2009, Ch. 9) for the complete argument. $\square$

Remark 2.3 (Discrete measures and regularisation). Brenier's Theorem requires absolute continuity of the source measure μ . In particle-method applications (Chapter 5), the empirical measure μ_t^N is a sum of Dirac masses, and the equilibrium measure μ_t^* is

approximated from discrete particle positions. The optimal transport between discrete measures is not uniquely given by a gradient of a convex function in the classical sense. The Sinkhorn algorithm (Section 2.2.3) bridges this gap: by adding an entropic penalty, the discrete measures are effectively “blurred” into continuous approximations, recovering a gradient-like structure and enabling stable computation.

Example 2.4 (One-dimensional case: sorted coupling). In $d = 1$, the optimal transport has a particularly elegant form. For $\mu, \nu \in \mathcal{P}_2(\mathbb{R})$ with distribution functions F_μ, F_ν and quantile functions F_μ^{-1}, F_ν^{-1} :

$$W_2(\mu, \nu)^2 = \int_0^1 |F_\mu^{-1}(u) - F_\nu^{-1}(u)|^2 du. \quad (2.3)$$

The optimal map is $T(x) = F_\nu^{-1}(F_\mu(x))$ — transport the u -quantile of μ to the u -quantile of ν (sorted coupling). This is computed exactly in $O(N \log N)$ by sorting the particles. **The synchronous coupling used in Chapter 3 for common-noise cancellation can be explicitly constructed in 1-D via this quantile coupling.**

2.2.3 The Sinkhorn Algorithm

For computational purposes, we regularise the Kantorovich problem with an entropic penalty:

$$W_{2,\epsilon}(\mu, \nu)^2 := \inf_{\pi \in \Pi(\mu, \nu)} \int |x - y|^2 \pi + \epsilon \text{KL}(\pi \| \mu \otimes \nu), \quad (2.4)$$

where $\text{KL}(\pi \| \mu \otimes \nu) = \int \log \frac{d\pi}{d(\mu \otimes \nu)} d\pi$. The solution has the form

$$\pi_\epsilon^*(dx, dy) = u(x)K(x, y)v(y) \mu(dx)\nu(dy),$$

where $K(x, y) = e^{-|x-y|^2/\epsilon}$ is the Gaussian kernel, and (u, v) satisfies the Sinkhorn fixed-point equations:

$$u_i = a_i / \sum_j K_{ij} v_j, \quad v_j = b_j / \sum_i K_{ij} u_i, \quad (2.5)$$

where $a_i = \mu(\{x_i\})$, $b_j = \nu(\{y_j\})$. This converges in $O(N^2/\epsilon^2)$ operations (Cuturi, 2013). In Chapter 4, this regularised distance will be used to compute Wasserstein gradients efficiently.

2.2.4 Metric Space Structure

Theorem 2.5 ($(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is a Complete Separable Metric Space). *The Wasserstein-2 distance satisfies:*

- (i) *Non-degeneracy:* $W_2(\mu, \nu) = 0 \iff \mu = \nu$;
- (ii) *Symmetry:* $W_2(\mu, \nu) = W_2(\nu, \mu)$;
- (iii) *Triangle inequality:* $W_2(\mu, \rho) \leq W_2(\mu, \nu) + W_2(\nu, \rho)$;
- (iv) *Completeness:* every Cauchy sequence in $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ converges to an element of $\mathcal{P}_2(\mathbb{R}^d)$;

(v) *Separability*: $\mathcal{P}_2(\mathbb{R}^d)$ contains a countable dense subset.

Proof of (iii) via gluing lemma. Given μ, ν, ρ , let $\pi_{12} \in \Pi(\mu, \nu)$ and $\pi_{23} \in \Pi(\nu, \rho)$ be any couplings. By the gluing lemma (Villani, 2009, Lemma 7.6), there exists $\pi_{123} \in \mathcal{P}(\mathbb{R}^d \times \mathbb{R}^d \times \mathbb{R}^d)$ with marginals (π_{12}, π_{23}) , i.e., the joint distribution of (X, Y, Z) where $(X, Y) \sim \pi_{12}$, $(Y, Z) \sim \pi_{23}$. The marginal $\pi_{13} = (\text{proj}_{13})_{\#} \pi_{123}$ is a coupling of μ and ρ . By Minkowski's inequality:

$$W_2(\mu, \rho) \leq \left(\int |x - z|^2 \pi_{13}(dx, dz) \right)^{1/2} = \mathbb{E}[|X - Z|^2]^{1/2} \leq \mathbb{E}[|X - Y|^2]^{1/2} + \mathbb{E}[|Y - Z|^2]^{1/2}.$$

Optimising over π_{12} and π_{23} gives the triangle inequality. The remaining properties are proved similarly. $\square$

2.2.5 Weak Convergence in $\mathcal{P}_2$

Theorem 2.6 (Characterisation of W_2 -convergence). *The following are equivalent for $\mu_n, \mu \in \mathcal{P}_2(\mathbb{R}^d)$:*

- (i) $W_2(\mu_n, \mu) \rightarrow 0$;
- (ii) $\mu_n \xrightarrow{w} \mu$ weakly and $\int |x|^2 \mu_n(dx) \rightarrow \int |x|^2 \mu(dx)$;
- (iii) For every continuous $f : \mathbb{R}^d \rightarrow \mathbb{R}$ with $|f(x)| \leq C(1 + |x|^2)$: $\int f d\mu_n \rightarrow \int f d\mu$.

Condition (ii) shows that W_2 -convergence is strictly stronger than weak convergence: it additionally requires convergence of the second moment. **This is precisely the correct topology for the MFG analysis, since $\mu_t = \mathcal{L}(X_t)$ and $\mathbb{E}[|X_t|^2]$ must remain finite and continuous in t .**

2.2.6 Geodesics and the Lions Derivative

The space $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ has a rich geometric structure that motivates the Lions derivative used in the master equation.

Definition 2.7 (Wasserstein Geodesic). The geodesic (shortest path) from μ_0 to μ_1 in $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is the curve

$$\mu_t := ((1 - t) \text{id} + tT)_{\#} \mu_0, \quad t \in [0, 1],$$

where $T = \nabla \varphi$ is the Brenier optimal transport map from μ_0 to μ_1 . This interpolates linearly in position space: the x -atom of μ_0 travels in a straight line to $T(x)$ at constant speed.

Proposition 2.8 (Geodesic is length-minimising). *The path $(\mu_t)_{t \in [0, 1]}$ satisfies $W_2(\mu_s, \mu_t) = |t - s| W_2(\mu_0, \mu_1)$ for all $s, t \in [0, 1]$.*

The Lions derivative (or L -derivative) formalises differentiation of functionals $F : \mathcal{P}_2(\mathbb{R}^d) \rightarrow \mathbb{R}$ along such geodesics. Rather than being a pointwise derivative, it is a linear functional on the tangent space. For a functional F that is Lions-differentiable, there

exists a measurable function $D_\mu F(\mu) : \mathbb{R}^d \rightarrow \mathbb{R}^d$ (unique up to μ -a.e. equality) such that for any $\nu \in \mathcal{P}_2(\mathbb{R}^d)$ with optimal transport map $T = \nabla \varphi$ from μ to ν ,

$$\left. \frac{d}{dt} \right|_{t=0} F(\mu_t) = \int_{\mathbb{R}^d} D_\mu F(\mu)(x) \cdot (T(x) - x) \mu(dx). \quad (2.6)$$

In other words, the directional derivative of F along a geodesic with initial velocity $v = T - \text{id}$ is given by the $L^2(\mu)$ inner product of $D_\mu F(\mu)$ and v . For a comprehensive treatment, see Carmona and Delarue (2018, Vol. I, Chapter 5).

2.2.7 Law of Large Numbers in Wasserstein Distance

Theorem 2.9 (LLN Rate in W_2 , Fournier–Guillin, 2015). *Let $\mu \in \mathcal{P}_2(\mathbb{R}^d)$ and $X_1, X_2, \dots$ be i.i.d. with law μ . Define $\mu^N = \frac{1}{N} \sum_{i=1}^N \delta_{X_i}$.*

(i) *Almost-sure convergence: $W_2(\mu^N, \mu) \rightarrow 0$ a.s. as $N \rightarrow \infty$.*

(ii) *Rate: If $\int_{\mathbb{R}^d} |x|^q \mu(dx) < \infty$ for some $q > 4$, then for $d = 1$:*

$$\mathbb{E}[W_2(\mu^N, \mu)] \leq C_\mu \cdot N^{-1/2}.$$

For $d \geq 3$: $\mathbb{E}[W_2(\mu^N, \mu)] \leq C_{\mu,d} \cdot N^{-1/d}$, under appropriate moment conditions. (See Fournier and Guillin, 2015, Theorem 1 and Corollary 2.)

Proof of (i) in $d = 1$. By the explicit formula (5.5):

$$W_2(\mu^N, \mu)^2 = \int_0^1 |F_N^{-1}(u) - F^{-1}(u)|^2 du.$$

For fixed u , $F_N^{-1}(u) \rightarrow F^{-1}(u)$ a.s. by the Glivenko–Cantelli theorem. Since $|F_N^{-1}(u) - F^{-1}(u)|^2 \leq 4 \max(|F^{-1}(u)|, |F_N^{-1}(u)|)^2$ and the right-hand side is uniformly integrable by the second moment assumption, the dominated convergence theorem gives $W_2(\mu^N, \mu)^2 \rightarrow 0$ a.s. $\square$

Remark 2.10 (Curse of dimensionality and mean-field mitigation). The rate $N^{-1/d}$ for $d \geq 3$ highlights the curse of dimensionality in Wasserstein space: for a $d = 50$ dimensional state space, the empirical measure converges exceedingly slowly. **In practice, the mean-field assumption — coupled with the specific convex structure of the rewards and the contraction property established in Theorem 3.9 — allows particle approximation (Chapter 5) to circumvent this theoretical bottleneck.** The Wasserstein distance is used primarily as an analytical tool for stability, not as the primary object of high-dimensional statistical estimation.

Remark 2.11 (Significance for the thesis). Theorem 5.3 justifies two central approximations: (a) replacing the N -player empirical measure μ_t^N by the mean-field distribution μ_t with error $O(N^{-1/2})$ (for $d = 1$ under sufficient moments); (b) validating the block-bootstrap consistency in Theorem 9.1, where $B = 10,000$ samples give $O(10^{-2.5})$ estimation error.

2.2.8 Wasserstein Barycentre

Definition 2.12 (Fréchet Mean / Barycentre). Given measures $\mu_1, \dots, \mu_K \in \mathcal{P}_2(\mathbb{R}^d)$ with weights $\lambda_k > 0$, $\sum_k \lambda_k = 1$, the Wasserstein barycentre is

$$\bar{\mu} = \arg \min_{\nu \in \mathcal{P}_2} \sum_{k=1}^K \lambda_k W_2(\nu, \mu_k)^2.$$

The barycentre exists and is unique when at least one μ_k is absolutely continuous (Agueh–Carlier, 2011). It is computed by the free-support Sinkhorn barycentre algorithm (Cuturi–Doucet, 2014). *Note: this concept is included for completeness but is not directly used in the subsequent chapters.*

2.3 Brownian Motion and Itô Calculus

2.3.1 Brownian Motion: Construction and Properties

Definition 2.13 (Standard Brownian Motion). A standard Brownian motion (Wiener process) $B = (B_t)_{t \geq 0}$ on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ is a continuous stochastic process satisfying:

- (i) $B_0 = 0$ almost surely;
- (ii) Independent increments: for $0 \leq s < t$, $B_t - B_s$ is independent of $\mathcal{F}_s$;
- (iii) Gaussian increments: $B_t - B_s \sim N(0, t - s)$ for all $0 \leq s < t$;
- (iv) Continuous paths: $t \mapsto B_t(\omega)$ is continuous for $\mathbb{P}$ -almost every ω .

Proposition 2.14 (Key path properties). *Brownian motion paths are almost surely:*

- (i) Hölder continuous of order $\alpha < \frac{1}{2}$ (but not $\frac{1}{2}$);
- (ii) Nowhere differentiable: $\lim_{h \rightarrow 0} \frac{B_{t+h} - B_t}{h}$ does not exist for any t ;
- (iii) Of unbounded variation: $\sup_{\Pi} \sum_k |B_{t_k} - B_{t_{k-1}}| = \infty$ over all partitions Π ;
- (iv) Of finite quadratic variation: $[B, B]_t = t$ (see below).

2.3.2 Quadratic Variation

Definition 2.15 (Quadratic Variation). The quadratic variation of a semimartingale M is the process $[M, M]_t$ defined as the limit in probability:

$$[M, M]_t = \lim_{|\Pi| \rightarrow 0} \sum_k (M_{t_k} - M_{t_{k-1}})^2,$$

where the limit is over partitions $\Pi = \{0 = t_0 < t_1 < \dots < t_n = t\}$ as the mesh $|\Pi| = \max_k (t_k - t_{k-1}) \rightarrow 0$.

Proposition 2.16 (Quadratic variation of Brownian motion). *For standard Brownian motion B : $[B, B]_t = t$ a.s.*

Proof sketch. Let $V_n = \sum_{k=1}^n (B_{kh} - B_{(k-1)h})^2$ with $h = t/n$. Then $\mathbb{E}[V_n] = n \cdot h = t$ and $\text{Var}(V_n) = n \cdot 2h^2 = 2t^2/n \rightarrow 0$. So $V_n \rightarrow t$ in L^2 , and by Borel–Cantelli, a.s. along a subsequence. $\square$

Remark 2.17 (Why quadratic variation matters: the Itô correction). The fact that $[B, B]_t = t \neq 0$ is the source of all the “extra terms” in stochastic calculus. For a smooth function f and a smooth path x , $f(x_t) = f(x_0) + \int f'(x_s) dx_s$ (ordinary chain rule). But for a Brownian path, the Taylor remainder $\frac{1}{2}f''(B_s)(dB_s)^2$ contributes $\frac{1}{2}f''(B_s) ds$ (not zero), because $(dB)^2 = dt$. This is Itô’s formula.

2.3.3 The Itô Integral

The classical Riemann–Stieltjes integral $\int_0^T H_t dB_t$ is not well-defined pathwise because B has unbounded variation. Itô’s construction extends it to square-integrable adapted integrands.

Definition 2.18 (Itô Integral). Let $(\mathcal{F}_t)$ be the natural filtration of B , and $H = (H_t)_{t \in [0, T]}$ be $(\mathcal{F}_t)$ -adapted with $\mathbb{E}[\int_0^T H_t^2 dt] < \infty$. The Itô integral $\int_0^T H_t dB_t$ is the $L^2(\mathbb{P})$ -limit of Riemann sums with left endpoints:

$$\int_0^T H_t dB_t := L^2\text{-}\lim_{n \rightarrow \infty} \sum_{k=0}^{n-1} H_{t_k} (B_{t_{k+1}} - B_{t_k}).$$

Theorem 2.19 (Itô Isometry). *For adapted square-integrable H :*

$$\mathbb{E} \left[\left(\int_0^T H_t dB_t \right)^2 \right] = \mathbb{E} \left[\int_0^T H_t^2 dt \right]. \quad (2.7)$$

Proof sketch. For simple processes, the expectation of the square expands to $\sum_k \mathbb{E}[H_{t_k}^2] \Delta t_k$ by independence. The general case follows by approximation. $\square$

Remark 2.20 (Significance for Theorem 3.9). The Itô isometry is used in the proof of Theorem 3.9: after deriving the Wasserstein distance ODE $dW_2^2 \leq CW_2^2 dt + dM_t$, we take expectations and use $\mathbb{E}[M_t] = 0$ (since $M_t = \int_0^t Z_s dW_s^0$ is a martingale by Itô isometry, as $\mathbb{E}[\int Z_s^2 ds] < \infty$). This eliminates the stochastic term and allows the Gronwall inequality to close.

2.3.4 Itô’s Formula

Theorem 2.21 (Itô’s Formula, one-dimensional). *Let $f \in C^2(\mathbb{R})$ and $X_t = X_0 + \int_0^t b_s ds + \int_0^t \sigma_s dB_s$ be an Itô process. Then:*

$$f(X_t) = f(X_0) + \int_0^t f'(X_s) dX_s + \frac{1}{2} \int_0^t f''(X_s) d[X, X]_s.$$

Expanding $d[X, X]_s = \sigma_s^2 ds$:

$$f(X_t) = f(X_0) + \int_0^t f'(X_s) b_s ds + \int_0^t f'(X_s) \sigma_s dB_s + \frac{1}{2} \int_0^t f''(X_s) \sigma_s^2 ds. \quad (2.8)$$

Theorem 2.22 (Itô's Formula, multidimensional). *Let $f \in C^2(\mathbb{R}^d)$ and let $X_t \in \mathbb{R}^d$ satisfy $dX_t = b_t dt + \sigma_t dW_t + \sigma_0 dW_t^0$ (where $W_t \in \mathbb{R}^d$, $W_t^0 \in \mathbb{R}^{d_0}$, and $\sigma_t \in \mathbb{R}^{d \times d}$, $\sigma_0 \in \mathbb{R}^{d \times d_0}$). Then:*

$$\begin{aligned} f(X_t) &= f(X_0) + \int_0^t \nabla f(X_s) \cdot b_s ds + \int_0^t \nabla f(X_s) \cdot \sigma_s dW_s + \int_0^t \nabla f(X_s) \cdot \sigma_0 dW_s^0 \\ &\quad + \frac{1}{2} \int_0^t \text{tr}(\sigma_s \sigma_s^\top D^2 f(X_s)) ds + \frac{1}{2} \int_0^t \text{tr}(\sigma_0 \sigma_0^\top D^2 f(X_s)) ds. \end{aligned}$$

Example 2.23 (Applying Itô to a test function). Let X_t^i be a single agent satisfying $dX_t^i = b_t dt + \sigma dW_t^i + \sigma_0 dW_t^0$. Apply Theorem 2.22 with $f = \phi$ (a smooth test function):

$$\begin{aligned} d\phi(X_t^i) &= \nabla \phi(X_t^i) \cdot b_t dt + \sigma \nabla \phi(X_t^i) \cdot dW_t^i + \sigma_0 \nabla \phi(X_t^i) \cdot dW_t^0 \\ &\quad + \frac{\sigma^2}{2} \Delta \phi(X_t^i) dt + \frac{1}{2} \text{tr}(\sigma_0 \sigma_0^\top D^2 \phi(X_t^i)) dt. \end{aligned}$$

Averaging over $i = 1, \dots, N$ and writing $\langle \mu_t^N, \phi \rangle = \frac{1}{N} \sum_i \phi(X_t^i)$:

$$\begin{aligned} d\langle \mu_t^N, \phi \rangle &= \left\langle \mu_t^N, b_t \cdot \nabla \phi + \frac{\sigma^2}{2} \Delta \phi + \frac{1}{2} \text{tr}(\sigma_0 \sigma_0^\top D^2 \phi) \right\rangle dt \\ &\quad + \underbrace{\frac{\sigma}{N} \sum_i \nabla \phi(X_t^i) dW_t^i}_{\rightarrow 0 \text{ as } N \rightarrow \infty} + \sigma_0 \langle \mu_t^N, \nabla \phi \rangle dW_t^0. \end{aligned}$$

The idiosyncratic term vanishes in L^2 as $N \rightarrow \infty$ (mutual independence of W^i). The common noise term survives because W_t^0 is the same for all agents. Taking $N \rightarrow \infty$ yields the Fokker–Planck SPDE of Section 2.4.

2.3.5 Itô's Product Rule and Integration by Parts

Proposition 2.24 (Itô Product Rule). *For Itô processes X_t and Y_t :*

$$d(X_t Y_t) = X_t dY_t + Y_t dX_t + d[X, Y]_t,$$

where $[X, Y]_t = \int_0^t \sigma_s^X \sigma_s^Y ds$ is the quadratic covariation (vanishes if driven by independent Brownians).

Example 2.25 (Integration by parts for stochastic integrals). $B_t^2 = 2 \int_0^t B_s dB_s + t$. (Verify: apply (2.8) with $f(x) = x^2$, $f'(x) = 2x$, $f''(x) = 2$.)

2.3.6 Stochastic Differential Equations

Theorem 2.26 (Existence and Uniqueness for SDEs). *Let $b : [0, T] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$ and $\sigma : [0, T] \times \mathbb{R}^d \rightarrow \mathbb{R}^{d \times m}$ satisfy for some $L > 0$:*

$$(i) \text{ Lipschitz: } |b(t, x) - b(t, y)| + \|\sigma(t, x) - \sigma(t, y)\| \leq L|x - y|;$$

$$(ii) \text{ Linear growth: } |b(t, x)| + \|\sigma(t, x)\| \leq L(1 + |x|).$$

Then for any $\mathcal{F}_0$ -measurable X_0 with $\mathbb{E}[|X_0|^2] < \infty$, the SDE $dX_t = b(t, X_t) dt + \sigma(t, X_t) dW_t$ has a unique strong solution $(X_t)_{t \in [0, T]}$ satisfying

$$\mathbb{E} \left[\sup_{t \in [0, T]} |X_t|^2 \right] \leq C(1 + \mathbb{E}[|X_0|^2])e^{CT}.$$

Proof sketch — Picard iteration. Define $X_t^0 \equiv X_0$ and $X_t^{n+1} = X_0 + \int_0^t b(s, X_s^n) ds + \int_0^t \sigma(s, X_s^n) dW_s$. By Itô isometry and the Lipschitz condition:

$$\mathbb{E} \left[\sup_{s \leq t} |X_s^{n+1} - X_s^n|^2 \right] \leq C \int_0^t \mathbb{E} \left[\sup_{r \leq s} |X_r^n - X_r^{n-1}|^2 \right] ds.$$

Gronwall's inequality gives $\mathbb{E}[\sup_{s \leq T} |X_s^n - X_s^{n-1}|^2] \leq C^n T^n / n! \rightarrow 0$. The Picard iterates converge in L^2 to a unique limit satisfying the SDE. $\square$

2.4 Fokker–Planck SPDE

2.4.1 The Forward Equation

The Fokker–Planck equation (also called the Kolmogorov forward equation) governs the evolution of the law of a diffusion process. In the MFG setting with common noise, it becomes a stochastic PDE.

Definition 2.27 (Fokker–Planck Operator). For drift $b : \mathbb{R}^d \rightarrow \mathbb{R}^d$ and diffusion $\sigma > 0$, the Fokker–Planck operator is

$$\mathcal{L}^* \phi(x) := b(x) \cdot \nabla \phi(x) + \frac{\sigma^2}{2} \Delta \phi(x), \quad \phi \in C_c^\infty(\mathbb{R}^d).$$

This is the formal L^2 -adjoint of the generator $\mathcal{L}f(x) = b(x) \cdot \nabla f(x) + \frac{\sigma^2}{2} \Delta f(x)$: $\int (\mathcal{L}f) d\mu = \int f(\mathcal{L}^* \mu)$ for smooth μ .

Theorem 2.28 (Fokker–Planck SPDE). *Let X_t satisfy the McKean–Vlasov SDE: $dX_t = b(t, X_t, \mu_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0$, $\mu_t = \mathcal{L}(X_t | \mathcal{F}_t^{W^0})$. Then μ_t satisfies the following SPDE in the weak sense: for all $\phi \in C_c^\infty(\mathbb{R}^d)$:*

$$d\langle \mu_t, \phi \rangle = \langle \mu_t, \mathcal{L}_t^* \phi \rangle dt + \sigma_0 \langle \mu_t, \nabla \phi \rangle dW_t^0, \quad \langle \mu_0, \phi \rangle = \int \phi(x) \mu_0(dx). \quad (2.9)$$

Proof. This follows directly from Example 2.23. By Itô's formula applied to $\phi(X_t)$:

$$\phi(X_t) - \phi(X_0) = \int_0^t \mathcal{L}_s^* \phi(X_s) ds + \sigma \int_0^t \nabla \phi(X_s) dW_s^i + \sigma_0 \int_0^t \nabla \phi(X_s) dW_s^0.$$

Take the conditional expectation given $\mathcal{F}_t^{W^0}$: the idiosyncratic term $\sigma \int \nabla \phi(X_s) dW_s^i$ has conditional expectation zero (since W^i is independent of W^0). The common noise term $\sigma_0 \int \nabla \phi(X_s) dW_s^0$ survives because W^0 is measurable with respect to $\mathcal{F}^{W^0}$. This gives the weak form (2.9). $\square$

2.4.2 Strong Form: The SPDE in Density Form

When μ_t has a density $\rho_t(x)$ (i.e., $\mu_t(dx) = \rho_t(x) dx$), the SPDE (2.9) takes the strong form:

$$d\rho_t = \left[-\nabla_x \cdot (b(t, x, \mu_t) \rho_t) + \frac{\sigma^2}{2} \Delta_x \rho_t \right] dt + \sigma_0 \nabla_x \rho_t dW_t^0. \quad (2.10)$$

Remark 2.29 (Normalisation and the Kushner–Zakai distinction). Equation (2.10) is linear in ρ_t . This linearity is a consequence of two facts: (i) the SPDE describes the evolution of the conditional law of the state given the common noise, and (ii) the total mass is preserved because the drift and diffusion operators are in divergence form. Consequently, the normalisation constant remains identically one. This distinguishes the Fokker–Planck SPDE from the Kushner–Zakai equation of nonlinear filtering, which involves a stochastic normalisation factor arising from observations. The homogeneous common-noise exposure (constant σ_0) is sufficient but not necessary for this linearity; the key property is that the coefficients do not depend on the measure in a way that would introduce nonlinear feedback into the normalisation. See Chapter 4 of Carmona and Delarue (2018) for a detailed discussion.

2.4.3 Sobolev Regularity

Definition 2.30 (Sobolev Spaces). For $k \in \mathbb{N}$ and $p \in [1, \infty)$:

$$W^{k,p}(\mathbb{R}^d) = \{f \in L^p(\mathbb{R}^d) \mid \partial^\alpha f \in L^p(\mathbb{R}^d) \text{ for all } |\alpha| \leq k\},$$

with norm $\|f\|_{W^{k,p}}^p = \sum_{|\alpha| \leq k} \|\partial^\alpha f\|_{L^p}^p$. The Hilbert case $p = 2$ gives $H^k(\mathbb{R}^d) = W^{k,2}(\mathbb{R}^d)$ with inner product $\langle f, g \rangle_{H^k} = \sum_{|\alpha| \leq k} \langle \partial^\alpha f, \partial^\alpha g \rangle_{L^2}$. The dual space is $H^{-1}(\mathbb{R}^d) = (H^1(\mathbb{R}^d))^*$.

Theorem 2.31 (Fokker–Planck Well-Posedness in Sobolev Space). *Under Assumptions A1–A3 (Lipschitz drift, bounded coefficients, regularity), the SPDE (2.9) has a unique solution $\rho \in L^2(\Omega \times [0, T]; H^1(\mathbb{R}^d))$, and*

$$\mathbb{E} \left[\sup_{t \in [0, T]} \|\rho_t\|_{L^2}^2 \right] + \mathbb{E} \left[\int_0^T \|\nabla \rho_t\|_{L^2}^2 dt \right] \leq C(1 + \|\rho_0\|_{H^1}^2). \quad (2.11)$$

Proof sketch. Apply the energy method: multiply (2.10) by ρ_t and integrate over $\mathbb{R}^d$. Using Itô’s formula for $\|\rho_t\|_{L^2}^2$:

$$d\|\rho_t\|_{L^2}^2 = 2\langle \rho_t, d\rho_t \rangle + d[\langle \rho, \cdot \rangle, \langle \rho, \cdot \rangle]_t.$$

The deterministic part contributes $-\sigma^2 \|\nabla \rho_t\|_{L^2}^2 dt$ (by integration by parts) plus a term bounded by $\|b\|_\infty \|\rho_t\|_{L^2} \|\nabla \rho_t\|_{L^2} dt$. The stochastic part contributes $\|\sigma_0\|_F^2 \|\nabla \rho_t\|_{L^2}^2 dt$ (by Itô isometry). Taking expectations and applying Gronwall gives (2.11). $\square$

2.4.4 The Gelfand Triple

Definition 2.32 (Gelfand Triple). A Gelfand triple (or rigged Hilbert space) is a triple of spaces $V \subset H \subset V^*$, where:

- V is a reflexive Banach space (e.g., $H^1(\mathbb{R}^d)$);
- H is a Hilbert space (e.g., $L^2(\mathbb{R}^d)$);
- V^* is the dual of V (e.g., $H^{-1}(\mathbb{R}^d)$);
- the inclusions are continuous and dense.

For the Fokker–Planck SPDE, the relevant Gelfand triple is $H^1(\mathbb{R}^d) \subset L^2(\mathbb{R}^d) \subset H^{-1}(\mathbb{R}^d)$. The Fokker–Planck operator $\mathcal{L}^* : H^1(\mathbb{R}^d) \rightarrow H^{-1}(\mathbb{R}^d)$ is bounded with norm $\|\mathcal{L}^*\| \leq L_b + \sigma^2$ (Lemma 4.2 of the thesis, proved via Hahn–Banach in Section 2.6). This is what makes the SPDE well-posed in H^{-1} .

2.5 Martingale Theory

2.5.1 Martingales and Stopping Times

Definition 2.33 (Martingale). An adapted integrable process $(M_t)_{t \geq 0}$ is a martingale with respect to $(\mathcal{F}_t)$ if $\mathbb{E}[M_t | \mathcal{F}_s] = M_s$ $\mathbb{P}$ -a.s. for all $0 \leq s \leq t$. It is a local martingale if there exist stopping times $\tau_n \uparrow \infty$ such that $(M_{t \wedge \tau_n})$ is a martingale for each n .

Theorem 2.34 (Itô integrals are local martingales). If H is adapted with $\int_0^t H_s^2 ds < \infty$ a.s., then $M_t := \int_0^t H_s dB_s$ is a local martingale with $M_0 = 0$. If additionally $\mathbb{E}[\int_0^T H_s^2 ds] < \infty$, then M_t is a square-integrable martingale.

Theorem 2.35 (Doob's Maximal Inequality). For a non-negative submartingale (M_t) :

$$\mathbb{E} \left[\sup_{0 \leq t \leq T} M_t^2 \right] \leq 4 \mathbb{E}[M_T^2].$$

Remark 2.36. This is used in Theorem 3.9: to bound the expected supremum

$$\mathbb{E} \left[\sup_t W_2^2(\mu_t^1, \mu_t^2) \right] \leq 4 \mathbb{E} [W_2^2(\mu_T^1, \mu_T^2)],$$

allowing the Gronwall bound to be applied.

Theorem 2.37 (Burkholder–Davis–Gundy (BDG) Inequality). For any continuous local martingale M with $M_0 = 0$ and $p \geq 1$:

$$c_p \mathbb{E} [[M, M]_T^{p/2}] \leq \mathbb{E} \left[\sup_{t \leq T} |M_t|^p \right] \leq C_p \mathbb{E} [[M, M]_T^{p/2}].$$

In particular, for $p = 2$: $\mathbb{E}[\sup_t M_t^2] \leq 4 \mathbb{E}[[M, M]_T]$.

2.5.2 Martingale Representation Theorem

Theorem 2.38 (Martingale Representation, Brownian Filtration). Let $(\mathcal{F}_t)$ be the natural filtration of a d -dimensional Brownian motion W . Then every $(\mathcal{F}_t)$ -martingale M has the representation

$$M_t = M_0 + \int_0^t H_s dW_s$$

for a unique $(\mathcal{F}_t)$ -adapted process H with $\mathbb{E}[\int_0^T |H_s|^2 ds] < \infty$.

Remark 2.39 (Connection to BSDE theory). The MRT is the existence theorem for solutions to backward SDEs. Given a terminal condition $\xi = p_T$, the MRT guarantees the existence of (q_t, r_t) such that

$$p_T = p_t - \int_t^T \partial_x H_s ds + \int_t^T q_s dW_s^i + \int_t^T r_s dW_s^0.$$

This is the adjoint BSDE from Chapter 1.

2.5.3 The Novikov Condition and Girsanov's Theorem

Theorem 2.40 (Girsanov's Theorem). *Let $\theta : [0, T] \times \Omega \rightarrow \mathbb{R}^d$ be adapted with the Novikov condition*

$$\mathbb{E} \left[\exp \left(\frac{1}{2} \int_0^T |\theta_t|^2 dt \right) \right] < \infty.$$

Define

$$Z_T = \exp \left(\int_0^T \theta_t dW_t - \frac{1}{2} \int_0^T |\theta_t|^2 dt \right).$$

Then Z_T is the Radon–Nikodym derivative of a probability measure $\mathbb{P}^$ with respect to $\mathbb{P}$: $d\mathbb{P}^*/d\mathbb{P} = Z_T$. Under $\mathbb{P}^*$, the process $\widetilde{W}_t = W_t - \int_0^t \theta_s ds$ is a Brownian motion.*

Remark 2.41 (Use in the thesis). Girsanov's theorem is a standard tool for change of measure in stochastic analysis. In the MFG setting, it may be applied to remove drift terms under certain monotonicity assumptions, but the primary analysis in Theorem 3.9 uses synchronous coupling rather than Girsanov.

2.6 Sobolev Spaces and Functional Analysis

2.6.1 Sobolev Embedding Theorem

Theorem 2.42 (Sobolev Embedding). *Let $\Omega \subset \mathbb{R}^d$ be bounded with smooth boundary.*

- (i) *If $k > d/2$, then $H^k(\Omega) \hookrightarrow C^0(\bar{\Omega})$ (continuous embedding into continuous functions);*
- (ii) *$H^1(\Omega) \hookrightarrow L^{2^*}(\Omega)$ where $2^* = 2d/(d-2)$ for $d \geq 3$.*

In our setting $\Omega = \mathbb{R}^d$, $d = 1$: $H^1(\mathbb{R}) \hookrightarrow C^{0,1/2}(\mathbb{R})$ (functions are Hölder- $\frac{1}{2}$ continuous).

2.6.2 Key Functional Analysis Tools

Three results from functional analysis are used in the well-posedness theory (Chapter 3) and the SPDE analysis (Chapter 4). We state each in the form used and refer to standard references for proofs.

Theorem 2.43 (Hahn–Banach Extension, (Brezis, 2011, Theorem 1.9)). *Let E be a real normed space, $F \subset E$ a subspace, and $\varphi : F \rightarrow \mathbb{R}$ a bounded linear functional. Then there exists an extension $\tilde{\varphi} : E \rightarrow \mathbb{R}$ with $\tilde{\varphi}|_F = \varphi$ and $\|\tilde{\varphi}\|_{E^*} = \|\varphi\|_{F^*}$.*

Use here: shows that the Fokker–Planck operator $\mathcal{L}^*$, initially defined on smooth test functions (a dense subspace of H^1), extends to a bounded operator $\mathcal{L}^* : H^1(\mathbb{R}^d) \rightarrow H^{-1}(\mathbb{R}^d)$ (Lemma 4.2).

Theorem 2.44 (Lumer–Phillips, (Pazy, 1983, Theorem 3.15)). *A closed densely defined operator A on a Banach space X generates a C_0 -contraction semigroup if and only if A is dissipative and $(\text{Id} - A)$ has dense range.*

Use here: the Fokker–Planck operator $\mathcal{L}^*$ on $L^2(\mu^*)$ is dissipative ($\langle \mathcal{L}^* \rho, \rho \rangle = -\sigma^2 \|\nabla \rho\|^2 \leq 0$), so Lumer–Phillips gives a contraction semigroup. The spectral gap $\lambda_1 > 0$ (the smallest positive eigenvalue of $-\mathcal{L}^*$) controls the exponential decay rate of the linearised Fokker–Planck equation.

Theorem 2.45 (Log-Sobolev implies Poincaré, (Bakry et al., 2014, Theorem 5.1)). *If μ^* satisfies a log-Sobolev inequality with constant C_{LSI} : $H(\mu\|\mu^*) \leq C_{\text{LSI}} I(\mu\|\mu^*)$, then $\lambda_1 \geq 1/C_{\text{LSI}}$.*

Use here: under the log-Sobolev condition (Assumption A5'), the spectral gap lower bound $\lambda_1 \geq 1/C_{\text{LSI}}$ gives an exponential convergence rate for the Fokker–Planck semigroup. As noted in Chapter 8, this is a strong assumption: log-concave equilibrium measures satisfy it (Bakry–Émery criterion), but heavy-tailed or multimodal equilibria do not. When only the Poincaré inequality holds, the same argument gives algebraic rates governed by λ_1 .

2.7 Stochastic Maximum Principle

2.7.1 The Control Problem

We now state the stochastic optimal control problem precisely. Let $A \subset \mathbb{R}^k$ be the action space (convex compact). A control $\alpha = (\alpha_t)_{t \in [0, T]}$ is admissible if it is $(\mathcal{F}_t)$ -adapted and $\mathbb{E}[\int_0^T |\alpha_t|^2 dt] < \infty$. The state $X_t^\alpha \in \mathbb{R}^d$ satisfies:

$$dX_t^\alpha = b(t, X_t^\alpha, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^\alpha \sim \mu_0. \quad (2.12)$$

The objective is to maximise:

$$J(\alpha) = \mathbb{E} \left[\int_0^T f(t, X_t^\alpha, \mu_t, \alpha_t) dt + g(X_T^\alpha, \mu_T) \right]. \quad (2.13)$$

2.7.2 The Hamiltonian

Definition 2.46 (Control Hamiltonian). The Hamiltonian associated with the control problem is

$$H(t, x, \mu, p, \alpha) := p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha) \in \mathbb{R}, \quad (2.14)$$

where $p \in \mathbb{R}^d$ is the costate variable (adjoint or shadow price). The maximised Hamiltonian is

$$H^*(t, x, \mu, p) := \sup_{\alpha \in A} H(t, x, \mu, p, \alpha).$$

Under Assumption A5 (uniform concavity of H in α), H^* is attained at a unique $\alpha^*(t, x, \mu, p)$.

Example 2.47 (Linear-quadratic Hamiltonian). For $b = \kappa(\bar{x} - x) + \alpha$ and $f = -x^2 - \frac{\lambda}{2}\alpha^2$:

$$H(t, x, \mu, p, \alpha) = p(\kappa(\bar{x} - x) + \alpha) - x^2 - \frac{\lambda}{2}\alpha^2.$$

Maximising over α : $\partial_\alpha H = p - \lambda\alpha = 0$, so $\alpha^*(x, \mu, p) = p/\lambda$ and $H^*(t, x, \mu, p) = p\kappa(\bar{x} - x) - x^2 + \frac{p^2}{2\lambda}$.

2.7.3 The Adjoint Process and SMP

Definition 2.48 (Adjoint BSDE). The adjoint process (p_t, q_t, r_t) associated with an admissible control α and state process X_t^α satisfies the backward SDE:

$$-dp_t = \partial_x H(t, X_t^\alpha, \mu_t, p_t, \alpha_t) dt - q_t dW_t^i - r_t dW_t^0, \quad p_T = \partial_x g(X_T^\alpha, \mu_T). \quad (2.15)$$

Here $p_t \in \mathbb{R}^d$ is the costate, $q_t \in \mathbb{R}^{d \times d}$ accounts for idiosyncratic noise, and $r_t \in \mathbb{R}^{d \times d_0}$ accounts for common noise.

Theorem 2.49 (Stochastic Maximum Principle). *Let α^* be an optimal control for (2.13) and (p_t, q_t, r_t) the associated adjoint process. Then:*

$$H(t, X_t^*, \mu_t, p_t, \alpha_t^*) = H^*(t, X_t^*, \mu_t, p_t) \quad \mathbb{P}\text{-a.s., a.e. } t.$$

Under Assumption A5 (strict concavity of H in α):

$$\alpha_t^* = \arg \max_{\alpha \in A} H(t, X_t^*, \mu_t, p_t, \alpha) \quad \text{uniquely, } \mathbb{P}\text{-a.s.}$$

Proof sketch — variational method. Let $\alpha^\epsilon = \alpha^* + \epsilon(\alpha - \alpha^*)$ be a perturbation. The optimality $J(\alpha^*) \geq J(\alpha^\epsilon)$ gives, after dividing by ϵ and taking $\epsilon \rightarrow 0$:

$$\mathbb{E} \left[\int_0^T (\partial_\alpha H(t, X_t^*, \mu_t, p_t, \alpha_t^*)) \cdot (\alpha_t - \alpha_t^*) dt \right] \geq 0$$

for all admissible α . Since α is arbitrary, $\partial_\alpha H = 0$ $\mathbb{P}$ -a.s. a.e. t . Under strict concavity (A5), this uniquely determines α_t^* . $\square$

2.7.4 Connection to Hamilton–Jacobi–Bellman

The SMP and HJB equation are dual formulations of the same problem.

Theorem 2.50 (SMP–HJB Connection). *If the value function $u(t, x, \mu)$ is $C^{1,2}$ in (t, x) and admits a Lions derivative in μ , then:*

$$p_t = D_x u(t, X_t^*, \mu_t), \quad q_t = \sigma D_x^2 u(t, X_t^*, \mu_t), \quad r_t = \sigma_0^\top \mathbb{E} [D_\mu u(t, X_t^*, \mu_t)(\tilde{X}_t) \mid \mathcal{F}_t^{W^0}]. \quad (2.16)$$

Here $D_\mu u(t, x, \mu)(\tilde{x})$ denotes the Lions derivative of u with respect to the measure argument, evaluated at the independent copy $\tilde{X}_t$. This derivative is defined via the “lifting” of functionals on $\mathcal{P}_2(\mathbb{R}^d)$ to functionals on the Hilbert space $L^2(\Omega; \mathbb{R}^d)$; see Carmona–Delarue (2018, Vol. I, Chapter 5) for a comprehensive treatment. The HJB equation

$$-\partial_t u = H^*(t, x, \mu, D_x u) + \bar{b}(t, \mu) \cdot D_\mu u + \frac{\sigma^2 + \sigma_0^2}{2} \Delta_x u + \sigma_0^2 \text{tr}(D_x D_\mu u) + \frac{\sigma_0^2}{2} \int D_\mu^2 u(\tilde{x}, \tilde{x}) \mu(d\tilde{x}),$$

$$u(T, x, \mu) = g(x, \mu),$$

where $\bar{b}(t, \mu)$ is the equilibrium aggregate drift of μ_t itself (the population analogue of the individual drift b , evaluated at the optimal control). **This is the full master equation, including the common-noise diffusion terms and the Lions-derivative mean-field coupling term $\bar{b} \cdot D_\mu u$; the latter is frequently omitted in informal**

expositions but is generally non-zero whenever $\bar{b}$ depends on μ , as it does in every model considered in this thesis. Chapter 7’s Appendix A derivation works through this term explicitly for the linear-quadratic case, including the correction history of an earlier omission; see the note there for how the coupling term is computed in practice. The PDE is the characterisation of the value function, and its unique solution (under Assumptions A1–A5) gives the adjoint variables via (2.16).

Remark 2.51 (Significance for the master equation). Equation (2.16) reveals why the master equation $U(t, \mu)$ contains derivatives in measure: the adjoint variable r_t involves the Lions derivative $D_\mu U$, which measures how the value function changes when the measure μ_t is perturbed. **This is the analytical foundation for Chapter 6’s hierarchical RL, where the slow actor learns $D_\mu U$ via the Wonham filter.**

2.8 Summary: Toolkit Reference

The following table maps each tool developed in this chapter to its first use in the thesis.

Table 2.1: Toolkit Reference

Tool	Content	First used
$\mathcal{P}_2(\mathbb{R}^d)$	Probability measures with finite 2nd moment	Ch. 3 (state space for μ^*)
Wasserstein distance	$W_2(\mu, \nu) = (\inf_{\pi} \int x - y ^2 \pi)^{1/2}$	Ch. 3 (Theorem 3.9 LHS)
Brenier's Theorem	Optimal transport map $T = \nabla \varphi$	Ch. 3 (coupling in Lemma 3.2)
Metric space structure	$(\mathcal{P}_2, W_2)$ complete separable	Ch. 3 (Banach FPT applies)
W_2 -convergence characterisation	$W_2 \rightarrow 0 \iff \text{weak} + \text{2nd moment}$	Ch. 3 (existence proof)
Geodesic / Lions derivative	Linear functional on $L^2(\mu)$	Ch. 8 (Lions derivative)
LLN rate (Fournier–Guillin)	$\mathbb{E}[W_2(\mu^N, \mu)] \leq CN^{-1/2}$ ($d = 1, q > 4$)	Ch. 7 ($O(1/N)$ proof)
1D Wasserstein formula	$W_2^2 = \int (F_{\mu}^{-1} - F_{\nu}^{-1})^2$	Ch. 4 (numerical computation)
Brownian motion	Gaussian independent increments	Ch. 3 (SDE driving noise)
Quadratic variation	$[B, B]_t = t$ a.s.	Itô's formula
Itô isometry	$\mathbb{E}[(\int H dB)^2] = \mathbb{E}[\int H^2 dt]$	Ch. 3 (Step 4: $\mathbb{E}[M] = 0$)
Itô's formula	$df(X) = f' dX + \frac{1}{2} f'' d[X, X]$	Fokker–Planck derivation
SDE existence/unique-ness	Lipschitz $\Rightarrow$ unique strong solution	Ch. 3 (agent SDE)
Itô product rule	$d(XY) = X dY + Y dX + d[X, Y]$	Ch. 3 (Lemma 3.2 derivation)
Fokker–Planck SPDE	$d\langle \mu_t, \phi \rangle = \langle \mu_t, \mathcal{L}^* \phi \rangle dt + \sigma_0 \langle \mu_t, \nabla \phi \rangle dW^0$	Ch. 3 (forward equation)
Sobolev well-posedness	$\rho \in L^2(\Omega; H^1)$, energy estimate	Ch. 4 (error analysis)
Gelfand triple	$H^1 \subset L^2 \subset H^{-1}$	Ch. 4
Doob's inequality	$\mathbb{E}[\sup M^2] \leq 4\mathbb{E}[M_T^2]$	Ch. 3 (sup bound)
BDG inequality	$\mathbb{E}[\sup M ^p] \sim \mathbb{E}[M, M ^{p/2}]$	Ch. 3 (moment bounds)
Martingale representation	Every martingale = Itô integral	Ch. 3 (BSDE existence)
Girsanov's theorem	Change of measure, remove drift	Ch. 3 (alternative approach)
Hahn–Banach	Norm-preserving extension	Ch. 4 (Lemma 4.2)

Chapter 3

Well-Posedness of Common-Noise Mean-Field Games

3.1 Overview

This chapter proves Theorem 3.9 (originally labelled Theorem B.9), the central structural result of the thesis. We establish that the forward-backward stochastic differential equation (FBSDE) system (??) characterising the mean-field game with common noise admits a unique solution under the standing assumptions of Section 1.9. Moreover, we provide an explicit quantitative stability estimate: the solution map is a contraction in a weighted Wasserstein-2 metric on the space of measure flows, with a contraction constant that depends explicitly on the model parameters — the Lipschitz constant of the drift L_b , the Frobenius norm of the common-noise coefficient $\|\sigma_0\|_F$, and bounds on Lions derivatives. The ε -Nash property for the finite-player game is also verified.

While the existence and uniqueness of solutions to common-noise MFGs have been established under various conditions — notably by Ahuja, Ren, and Yang (2019) under convexity and weak-monotonicity, and in the comprehensive probabilistic treatment of Carmona and Delarue (2018) — the present contribution is the derivation of explicit parameter-dependent stability constants suitable for numerical analysis and reinforcement learning algorithms. These constants quantify precisely how the model parameters affect the contraction rate. **We emphasize that the contraction is established in expectation, not pathwise, and relies on the synchronous coupling of two candidate solutions sharing the identical common noise.**

The proof proceeds by formulating the MFG equilibrium as a fixed-point problem on the space of measure flows, $\Phi(\mu) = \mu$, where Φ maps a candidate measure flow to the conditional law of the optimally controlled state. Using the Stochastic Maximum Principle (Section 2.7) and the theory of McKean–Vlasov FBSDEs, we show that for any given measure flow μ , there exists a unique optimal control and corresponding state process. The key technical step is the derivation of a Wasserstein stability estimate for the map Φ in expectation, which relies on synchronous coupling of two candidate solutions, the Lions derivative calculus, and Itô’s formula applied to the squared Wasserstein distance between the conditional laws. Under sufficient regularity (Assumptions 1.13–1.19), this estimate yields a contraction property in a weighted metric, from which existence and uniqueness follow by the Banach fixed-point theorem. Finally, we prove the ε -Nash property using propagation of chaos arguments.

Chapter organisation. Section 3.2 recalls the standing assumptions and notation. Section 3.3 states the main theorem. Section 3.4 defines the fixed-point map Φ . Section 3.5 contains the core stability estimate, including the synchronous coupling, Itô expansion, Lipschitz estimates, and the Gronwall argument. Section 3.6 proves the contraction in weighted Wasserstein metric. Section 3.7 concludes existence and uniqueness. Section 3.8 establishes the ε -Nash property. Section 3.9 discusses the contraction constant and its limitations. Section 3.10 provides excerpts from the Lean 4 formal verification.

3.2 Standing Assumptions and Notation

We recall the standing assumptions from Section 1.9, which are in force throughout this chapter.

Assumption 3.1 (Lipschitz drift). There exists $L_b > 0$ such that for all $t \in [0, T]$, $x, x' \in \mathbb{R}^d$, $\mu, \mu' \in \mathcal{P}_2(\mathbb{R}^d)$, and $\alpha, \alpha' \in A$,

$$\|b(t, x, \mu, \alpha) - b(t, x', \mu', \alpha')\| \leq L_b(\|x - x'\| + W_2(\mu, \mu') + \|\alpha - \alpha'\|).$$

Assumption 3.2 (Lipschitz rewards). The running reward f and terminal reward g are Lipschitz continuous in all arguments with constants $L_f, L_g > 0$.

Assumption 3.3 (Additional regularity). The functions b, f, g are continuously differentiable in x and Lions-differentiable in μ (in the sense of Carmona and Delarue, 2018, Chapter 5). Their Lions derivatives $D_\mu b, D_\mu f, D_\mu g$ are bounded and Lipschitz continuous in x and μ . Moreover, the second-order Lions derivatives are bounded when required.

Assumption 3.4 (Bounded common noise). The common-noise coefficient $\sigma_0 \in \mathbb{R}^{d \times d_0}$ satisfies $\|\sigma_0\|_F \leq M_\sigma < \infty$. The idiosyncratic coefficient $\sigma > 0$ is constant.

Assumption 3.5 (Uniform concavity of Hamiltonian). The map $\alpha \mapsto H(t, x, \mu, p, \alpha)$ is uniformly concave: there exists $\lambda_H > 0$ such that for all $\alpha, \alpha' \in A$,

$$H(t, x, \mu, p, \alpha') \leq H(t, x, \mu, p, \alpha) + \partial_\alpha H(t, x, \mu, p, \alpha) \cdot (\alpha' - \alpha) - \frac{\lambda_H}{2} \|\alpha' - \alpha\|^2.$$

Assumption 3.6 (Regular terminal reward). The terminal reward g is C^1 in x and Lions-differentiable in μ , with bounded derivatives.

Assumption 3.7 (Initial condition). The initial distribution $\mu_{\text{init}} \in \mathcal{P}_2(\mathbb{R}^d)$ has finite fourth moment: $\int_{\mathbb{R}^d} |x|^4 \mu_{\text{init}}(dx) < \infty$.

Notation. We work on the filtered probability space $(\Omega, \mathcal{F}, (\mathcal{F}_t)_{t \in [0, T]}, \mathbb{P})$ described in Section 1.2.1. The space of measure flows is denoted by $\mathcal{M} := D([0, T]; \mathcal{P}_2(\mathbb{R}^d))$, the Skorokhod space of càdlàg measure-valued processes adapted to the common-noise filtration $(\mathcal{F}_t^{W^0})$. For $\mu \in \mathcal{M}$ and $\lambda > 0$, the weighted Wasserstein-2 norm is

$$\|\mu\|_{\lambda, T} := \sup_{t \in [0, T]} e^{-\lambda t} \mathbb{E}[W_2(\mu_t, \delta_0)^2]^{1/2},$$

where δ_0 is the Dirac measure at zero. For two measure flows $\mu^1, \mu^2 \in \mathcal{M}$, the distance is

$$\|\mu^1 - \mu^2\|_{\lambda, T} := \sup_{t \leq T} e^{-\lambda t} \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2]^{1/2}.$$

Remark 3.8 (Weighted norm). This weighted norm is equivalent to the standard norm used in Carmona and Delarue (2018) but is more convenient for extracting contraction factors. The exponential discounting allows us to absorb the Gronwall growth factor.

3.3 Statement of the Main Theorem

Theorem 3.9 (Quantitative Well-Posedness). *Under Assumptions 1.13–1.19, there exists a unique solution $(\mu^*, \alpha^*, X^*, p, q, r)$ to the FBSDE system (??). Moreover, for any two measure flows $\mu^1, \mu^2 \in \mathcal{M}$, the map Φ defined in (3.6) below satisfies the stability estimate in expectation:*

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq C_0 \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds, \quad (3.1)$$

where the constants are

$$C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}, \quad C_0 = L_b + C_\alpha L_b,$$

with C_{reg} depending on the Lipschitz constants of the Lions derivatives of b, f, g and the bounds from the adjoint BSDE, and C_α is the Lipschitz constant of the optimal control with respect to state, costate, and measure. Consequently, for any $\lambda > C/2$, the map Φ is a contraction in the weighted norm $\|\cdot\|_{\lambda, T}$:

$$\|\Phi(\mu^1) - \Phi(\mu^2)\|_{\lambda, T} \leq \rho \|\mu^1 - \mu^2\|_{\lambda, T}, \quad \rho = \sqrt{\frac{C_0}{2\lambda - C}} < 1. \quad (3.2)$$

The unique fixed point $\mu^* = \Phi(\mu^*)$ is the MFG equilibrium measure flow. Finally, the ε -Nash property holds: if all N players adopt the MFG strategy α^* , the strategy profile is an ε_N -Nash equilibrium of the N -player game with $\varepsilon_N = O(N^{-1/2})$.

Remark 3.10 (On the constant C_{reg}). The constant C_{reg} is not fully explicit; it encapsulates bounds on Lions derivatives and BSDE estimates. However, in specific models (e.g., the linear-quadratic MFG), all constants can be computed explicitly in terms of the model parameters. For general nonlinear MFGs, C_{reg} can be bounded by quantities involving the Lipschitz constants of the Lions derivatives and the operator norms of the associated linearised BSDEs.

Remark 3.11 (Expectation vs. pathwise). The contraction is established at the level of expected Wasserstein distance. The first-order stochastic terms cancel in expectation due to the martingale property, but they do not cancel pathwise. Consequently, the stability estimate (3.1) holds for the expectation, not almost surely.

3.4 The Fixed-Point Map Φ

For a given measure flow $\mu = (\mu_t)_{t \in [0, T]} \in \mathcal{M}$, we consider the optimal control problem of a representative agent who takes μ as given. The agent's state process X^μ satisfies the McKean–Vlasov SDE

$$dX_t^\mu = b(t, X_t^\mu, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad X_0^\mu \sim \mu_{\text{init}}, \quad (3.3)$$

and the objective is to maximise

$$J^\mu(\alpha) = \mathbb{E} \left[\int_0^T f(t, X_t^\mu, \mu_t, \alpha_t) dt + g(X_T^\mu, \mu_T) \right]. \quad (3.4)$$

By the Stochastic Maximum Principle (Theorem 2.20), an optimal control α^μ exists and is unique under Assumptions 1.13–1.17. It is characterised by the forward-backward system

$$\begin{aligned} dX_t^\mu &= b(t, X_t^\mu, \mu_t, \alpha_t^\mu) dt + \sigma dW_t^i + \sigma_0 dW_t^0, & X_0^\mu &\sim \mu_{\text{init}}, \\ -dp_t^\mu &= \partial_x H(t, X_t^\mu, \mu_t, p_t^\mu, \alpha_t^\mu) dt - q_t^\mu dW_t^i - r_t^\mu dW_t^0, & p_T^\mu &= \partial_x g(X_T^\mu, \mu_T), \\ \alpha_t^\mu &= \arg \max_{\alpha \in A} H(t, X_t^\mu, \mu_t, p_t^\mu, \alpha), \end{aligned} \quad (3.5)$$

where $H(t, x, \mu, p, \alpha) = p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha)$.

The solution to (3.5) exists and is unique for each fixed $\mu \in \mathcal{M}$; this follows from standard results for McKean–Vlasov FBSDEs with Lipschitz coefficients (see, e.g., Carmona & Delarue, 2018, Vol. I, Chapter 4, and the monotone functional approach of Ahuja, Ren, and Yang, 2019). We denote the unique optimal state process by X^μ and define the map

$$\Phi(\mu)_t := \mathcal{L}(X_t^\mu \mid \mathcal{F}_t^{W^0}), \quad t \in [0, T]. \quad (3.6)$$

That is, $\Phi(\mu)$ is the measure flow obtained as the conditional law of the optimally controlled state given the common noise. By construction, $\Phi(\mu) \in \mathcal{M}$. An MFG equilibrium is precisely a fixed point of Φ .

3.5 Stability Estimate for the Map Φ

The core of the proof is to show that Φ satisfies the Lipschitz-type estimate (3.1) in expectation. Let $\mu^1, \mu^2 \in \mathcal{M}$ be two candidate measure flows. We denote by $(X^1, p^1, q^1, r^1, \alpha^1)$ and $(X^2, p^2, q^2, r^2, \alpha^2)$ the corresponding solutions to (3.5) with μ replaced by μ^1 and μ^2 , respectively. The goal is to bound $\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2]$.

3.5.1 Synchronous Coupling

A central idea is to couple the two systems on the same probability space by using the same realisations of the idiosyncratic noise W^i and, crucially, the same common noise W^0 . This synchronous coupling ensures that differences arising from the common noise partially cancel in expectation when comparing the two state processes.

Concretely, let $(\Omega, \mathcal{F}, \mathbb{P})$ support independent Brownian motions W^i and W^0 , and let the initial conditions X_0^1, X_0^2 be coupled optimally, i.e., their joint distribution achieves $W_2(\mu_{\text{init}}, \mu_{\text{init}}) = 0$ (they are equal in law). We solve the two systems (3.5) on this same probability space with the same driving noises. Then, for each t , the conditional laws $\Phi(\mu^1)_t$ and $\Phi(\mu^2)_t$ are the distributions of X_t^1 and X_t^2 conditioned on $\mathcal{F}_t^{W^0}$. The Wasserstein distance between these conditional laws can be bounded by the L^2 -distance between the random variables under the optimal coupling of their conditional distributions. Since we have already coupled the processes synchronously, we can use the joint distribution of (X_t^1, X_t^2) to obtain a bound:

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq \mathbb{E}[|X_t^1 - X_t^2|^2]. \quad (3.7)$$

This inequality is standard: the Wasserstein distance between two laws is bounded above by the L^2 -distance between any coupling of random variables with those laws.

3.5.2 Itô Expansion of the Squared Difference

We now study the dynamics of the difference process $\Delta X_t := X_t^1 - X_t^2$. From the forward SDEs,

$$d(\Delta X_t) = [b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)] dt,$$

since the noise terms σdW_t^i and $\sigma_0 dW_t^0$ cancel identically under synchronous coupling. **Crucially, this cancellation holds pathwise for the difference of the two processes, not just in expectation. The stochastic terms disappear entirely from the equation for ΔX_t .** Applying Itô's formula to $|\Delta X_t|^2$, we obtain

$$d|\Delta X_t|^2 = 2\Delta X_t \cdot [b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)] dt. \quad (3.8)$$

Notice that there is no Itô correction term because the diffusion coefficients are constant and identical for both processes, so their difference has zero quadratic variation. This is a crucial simplification resulting from the synchronous coupling and the assumption of constant σ, σ_0 .

3.5.3 Lipschitz Estimates and Gronwall Setup

Using the Lipschitz property of b (Assumption 3.1), we have

$$|b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)| \leq L_b(|\Delta X_t| + W_2(\mu_t^1, \mu_t^2) + |\alpha_t^1 - \alpha_t^2|).$$

Substituting into (3.8) and using Young's inequality $2ab \leq a^2 + b^2$, we get

$$d|\Delta X_t|^2 \leq [(2L_b + L_b^2)|\Delta X_t|^2 + L_b W_2(\mu_t^1, \mu_t^2)^2 + L_b |\alpha_t^1 - \alpha_t^2|^2] dt. \quad (3.9)$$

To close the estimate, we need to bound the difference in controls $|\alpha_t^1 - \alpha_t^2|$ in terms of the state and measure differences. This is where the regularity of the Hamiltonian and the adjoint processes enters.

3.5.4 Control Difference via SMP and Lions Derivatives

From the optimality condition (3.5) and the uniform concavity of H (Assumption 1.17), the optimal control is a Lipschitz function of the state x , the costate p , and the measure μ . Specifically, by the implicit function theorem applied to $\partial_\alpha H = 0$, there exists a constant $C_\alpha > 0$ such that

$$|\alpha_t^1 - \alpha_t^2| \leq C_\alpha(|X_t^1 - X_t^2| + |p_t^1 - p_t^2| + W_2(\mu_t^1, \mu_t^2)). \quad (3.10)$$

The rigorous justification of this Lipschitz dependence relies on the Lions differentiability of b and f with respect to μ (Assumption 3.3). By Carmona and Delarue (2018, Proposition 5.35), if a functional is Lions-differentiable with bounded derivative, it is Lipschitz in the Wasserstein distance. The constant C_α can be bounded explicitly in terms of λ_H, L_b , and the bounds on the Lions derivatives. A detailed derivation is provided in Appendix A.

3.5.5 Adjoint Process Difference

Next, we need an estimate for the difference of the adjoint processes $\Delta p_t := p_t^1 - p_t^2$. The BSDEs for p^1 and p^2 are

$$-dp_t^j = \partial_x H(t, X_t^j, \mu_t^j, p_t^j, \alpha_t^j) dt - q_t^j dW_t^i - r_t^j dW_t^0, \quad j = 1, 2,$$

with terminal conditions $p_T^j = \partial_x g(X_T^j, \mu_T^j)$. Taking the difference, we obtain a linear BSDE for Δp_t :

$$-d(\Delta p_t) = [A_t \Delta X_t + B_t \Delta p_t + C_t \star (\mu_t^1 - \mu_t^2) + D_t(\alpha_t^1 - \alpha_t^2)] dt - \Delta q_t dW_t^i - \Delta r_t dW_t^0, \quad (3.11)$$

where the coefficients A_t, B_t, D_t involve derivatives of $\partial_x H$ with respect to x, p, α respectively, and are bounded thanks to Assumption 3.3. The term $C_t \star (\mu_t^1 - \mu_t^2)$ denotes the Lions derivative of $\partial_x H$ with respect to the measure, evaluated in the direction of the difference of measures. By Assumption 3.3, this Lions derivative is a bounded linear operator on the tangent space, and its operator norm is bounded by a constant L_μ . Consequently,

$$\|C_t \star (\mu_t^1 - \mu_t^2)\| \leq L_\mu W_2(\mu_t^1, \mu_t^2). \quad (3.12)$$

Standard BSDE estimates (see, e.g., El Karoui et al., 1997, or Zhang, 2017, Theorem 4.3.1) yield

$$\begin{aligned} & \mathbb{E} \left[\sup_{t \in [0, T]} |\Delta p_t|^2 + \int_0^T (|\Delta q_t|^2 + |\Delta r_t|^2) dt \right] \\ & \leq C_{\text{BSDE}} \mathbb{E} \left[|\Delta X_T|^2 + \int_0^T (|\Delta X_t|^2 + |\alpha_t^1 - \alpha_t^2|^2 + W_2(\mu_t^1, \mu_t^2)^2) dt \right]. \end{aligned} \quad (3.13)$$

The constant C_{BSDE} depends on the bounds of the coefficients A_t, B_t, D_t , the Lipschitz constant L_μ , the time horizon T , and crucially, the norm of the common-noise volatility $\|\sigma_0\|_F$ (through the martingale representation terms). A precise expression is given in Appendix A.

3.5.6 Combining the Estimates

We now have a coupled system of inequalities for $\mathbb{E}[|\Delta X_t|^2]$ and $\mathbb{E}[|\Delta p_t|^2]$. Using (3.10) to bound the control difference, and plugging into (3.9) and (3.13), we can eliminate $|\alpha_t^1 - \alpha_t^2|$ and obtain a closed differential inequality for

$$y_t := \mathbb{E}[|\Delta X_t|^2] + \mathbb{E}[|\Delta p_t|^2].$$

The algebraic manipulations, detailed in Appendix A, yield

$$\frac{d}{dt} y_t \leq C_1 y_t + C_2 \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2], \quad (3.14)$$

where

$$C_1 = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}, \quad C_2 = L_b + C_\alpha L_b + C_{\text{BSDE}} L_\mu.$$

Here C_{reg} encapsulates the regularity constants from the Lions derivatives and the BSDE bounds. The term $2\|\sigma_0\|_F^2$ arises from the dependence of the BSDE constant C_{BSDE} on the common-noise volatility; specifically, the martingale representation part of the BSDE estimate involves the quadratic variation of the common noise, contributing a term proportional to $\|\sigma_0\|_F^2$. The explicit derivation is provided in Appendix A.

3.5.7 The Wasserstein Stability Estimate

By Gronwall's inequality, (3.14) implies

$$y_t \leq C_2 \int_0^t e^{C_1(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds.$$

Recalling that y_t bounds $\mathbb{E}[|\Delta X_t|^2]$ from above, and using (3.7), we obtain precisely the stability estimate (3.1) with $C = C_1$ and $C_0 = C_2$:

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq C_0 \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds.$$

Remark 3.12 (Pathwise vs. expected contraction). The inequality above holds for the expected squared Wasserstein distance. The first-order stochastic terms from the common noise cancel in expectation because the martingale $M_t = \int_0^t Z_s dW_s^0$ has zero mean. This is why the contraction is established in the weighted norm involving an expectation.

3.6 Contraction in Weighted Wasserstein Metric

We now show that Φ is a contraction in the weighted norm $\|\cdot\|_{\lambda,T}$ for a suitably chosen λ . Take the supremum over $t \in [0, T]$ of both sides of (3.1) multiplied by $e^{-2\lambda t}$:

$$\begin{aligned} \|\Phi(\mu^1) - \Phi(\mu^2)\|_{\lambda,T}^2 &= \sup_{t \leq T} e^{-2\lambda t} \mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \\ &\leq C_0 \sup_{t \leq T} e^{-2\lambda t} \int_0^t e^{C(t-s)} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds \\ &\leq C_0 \sup_{t \leq T} \int_0^t e^{-2\lambda(t-s)} e^{(C-2\lambda)s} e^{-2\lambda s} \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds \\ &\leq C_0 \left(\int_0^T e^{-(2\lambda-C)u} du \right) \|\mu^1 - \mu^2\|_{\lambda,T}^2, \end{aligned}$$

where we used the change of variables $u = t - s$ and the fact that $e^{-2\lambda s} \mathbb{E}[W_2^2] \leq \|\mu^1 - \mu^2\|_{\lambda,T}^2$. The integral

$$I(\lambda) := C_0 \int_0^T e^{-(2\lambda-C)u} du = \frac{C_0}{2\lambda - C} (1 - e^{-(2\lambda-C)T}) \leq \frac{C_0}{2\lambda - C}.$$

Thus, for any $\lambda > C/2$, we have $I(\lambda) < 1$ provided λ is chosen large enough. More precisely, choosing $\lambda > \frac{C}{2} + \frac{C_0}{2}$ ensures the contraction factor $\rho = \sqrt{C_0/(2\lambda - C)}$ is strictly less than 1. This makes Φ a strict contraction on the complete metric space $(\mathcal{M}, \|\cdot\|_{\lambda,T})$.

3.7 Existence and Uniqueness

Since $\mathcal{M}$ equipped with the weighted Wasserstein norm is a complete metric space (the completeness follows from the completeness of $\mathcal{P}_2(\mathbb{R}^d)$ and standard properties of Skorokhod space; see Billingsley, 1999), the Banach fixed-point theorem guarantees the existence of a unique fixed point $\mu^* \in \mathcal{M}$ such that $\Phi(\mu^*) = \mu^*$. This fixed point is the MFG equilibrium measure flow. The corresponding processes $(X^*, p^*, q^*, r^*, \alpha^*)$ obtained from solving (3.5) with $\mu = \mu^*$ constitute the unique solution to the full FBSDE system (??).

Remark 3.13 (Local vs. global contraction). The contraction constant ρ depends on the time horizon T through the constants C and C_0 . For large T , the condition $\lambda > C/2$ may require a large λ , but the contraction still holds. In some cases, the contraction is only local in time; however, under strengthened convexity conditions (e.g., strong monotonicity of the cost), the contraction can be global.

3.8 The ε -Nash Property

Finally, we verify that the MFG strategy α^* provides an approximate Nash equilibrium for the N -player game. Consider the N -player system where all agents except agent i follow the MFG strategy α^* , while agent i uses an arbitrary admissible control α^i . Let $\mu_t^{N,-i}$ be the empirical measure of the other $N-1$ agents. By the conditional propagation of chaos (see Carmona & Delarue, 2018, Vol. II, Theorem 6.16), the empirical measure $\mu_t^{N,-i}$ converges to the MFG measure flow μ^* in the sense that

$$\mathbb{E}[W_2(\mu_t^{N,-i}, \mu_t^*)^2] \leq \frac{C_{\text{PoC}}}{N},$$

where the constant C_{PoC} depends on the model parameters (Lipschitz constants, moments) but not on N . Moreover, the state of agent i under control α^i and the MFG state under α^* can be coupled such that their distance is controlled by the measure approximation error.

Using the Lipschitz properties of the rewards and the stability of the FBSDE system (essentially the estimate (3.1) applied to the finite- N empirical measure), one can show that the difference in objective values satisfies

$$\sup_{\alpha^i} (J_i(\alpha^i, \alpha^{-i,*}) - J_i(\alpha^*, \alpha^{-i,*})) \leq \varepsilon_N,$$

with $\varepsilon_N = O(N^{-1/2})$. Thus, the symmetric strategy profile $(\alpha^*, \dots, \alpha^*)$ is an ε_N -Nash equilibrium. This completes the proof of Theorem 3.9.

Remark 3.14 (Verification of propagation of chaos conditions). The proof relies on the conditions of Carmona and Delarue (2018, Theorem 6.16), which are satisfied under Assumptions 1.13–1.19. In particular, the Lipschitz property of the drift and the boundedness of the diffusion coefficients ensure the conditional propagation of chaos holds.

3.9 Discussion of the Contraction Constant

The explicit constant $C = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$ appearing in the stability estimate (3.1) merits comment.

- The term $2L_b^2$ arises from the Lipschitz property of the drift.
- The term $2\|\sigma_0\|_F^2$ reflects the second-order contribution of the common noise to the Wasserstein dynamics; note that the first-order stochastic term cancels in expectation due to the martingale property, leaving only the effect of σ_0 through the BSDE estimates and the control difference.

- The remainder C_{reg} encapsulates the regularity constants from the Lions derivatives of b, f, g and the bounds from the adjoint BSDE. In the special case of the linear-quadratic MFG (Example 1.6), one can compute C_{reg} explicitly in terms of the model parameters $\kappa, \lambda, \sigma, \sigma_0$. For general nonlinear MFGs, C_{reg} is not fully explicit but can be bounded by quantities involving the Lipschitz constants of the Lions derivatives.

Table 3.1: Components of the Contraction Constant

Term	Expression	Origin
Drift Lipschitz	$2L_b^2$	From ΔX_t Gronwall step
Common noise	$2\ \sigma_0\ _F^2$	From BSDE constant dependence
Regularity	C_{reg}	Lions derivatives, BSDE bounds
Control Lipschitz	C_α	Implicit function theorem on H

Limitations. The contraction estimate relies on the stability of optimal transport maps, which requires absolute continuity of measures and bounded densities (Assumption 3.3). In the particle-method implementation of Chapter 5, the empirical measure is discrete, and this assumption is violated. **The theoretical results should be interpreted as establishing a well-posedness baseline under idealised conditions; practical implementation requires regularisation (e.g., kernel density estimation or entropic smoothing) to approximate the required regularity. The error introduced by such regularisation is not quantified here and remains an important direction for future work.**

3.10 Formal Verification in Lean 4

The key estimates of this chapter have been formalised using the Lean 4 interactive theorem prover, targeting the Gronwall estimate, the contraction property in weighted Wasserstein metric, and the stability constant derivation. **As detailed in the provenance caveat at the start of Appendix B, these proofs were written and reviewed for mathematical correctness following standard Mathlib idioms, but were not compiled against a live Lean 4/Mathlib installation while writing this thesis, since no such toolchain was available in that environment; they should not be treated as machine-verified until built successfully via lake build (the supplementary repository’s CI workflow does this on every push).** The full Lean 4 code is available in the supplementary repository.

Listing 3.1: Lean 4 formalisation of the Gronwall estimate underlying the contraction proof.

```

1 -- Lean 4 formalization of the contraction estimate (Theorem 3.9)
2 import Mathlib.Analysis.SpecialFunctions.Exp
3 import Mathlib.Analysis.ODE.Gronwall
4 import Mathlib.MeasureTheory.Measure.ProbabilityMeasure
5 import Mathlib.Tactic.Linarith
6

```

```

7 open Real MeasureTheory ProbabilityMeasure
8
9 variable {C1 C2 : Real} (hC1 : 0 < C1) (hC2 : 0 <= C2)
10
11 /-- Gronwall inequality for differential inequality  $y' \leq C1 * y$ 
    + C2 -/
12 theorem gronwall_estimate {y : Real -> Real} {C1 C2 : Real}
13   (hC1 : 0 < C1) (hC2 : 0 <= C2)
14   (h_deriv : forall t >= 0, HasDerivAt y (deriv y t) t)
15   (h_bound : forall t >= 0, deriv y t <= C1 * y t + C2)
16   (hy0 : y 0 = 0) :
17   forall t >= 0, y t <= C2 * (Real.exp (C1 * t) - 1) / C1 := by
18   intro t ht
19   let z (t : Real) : Real := y t - C2 / C1 * (Real.exp (C1 * t) -
    1)
20   have hz0 : z 0 = 0 := by simp [z, hy0, Real.exp_zero, sub_self,
    mul_zero]
21   -- [excerpt truncated for typesetting -- continues with the
    derivative
22   -- bound on z, then le_of_deriv_le to conclude z t <= 0 for
    all t >= 0,
23   -- giving the stated bound on y. The full proof, together with
    the
24   -- provenance caveat governing its verification status, is
    reproduced
25   -- in Appendix B.]
26   exact gronwall_estimate_aux hC1 hC2 h_deriv h_bound hy0 t ht

```

Note: the excerpt above is abbreviated for typesetting; the complete proof, together with the provenance caveat explaining that it has not yet been compiled against Mathlib, is provided in full in Appendix B.

Chapter 4

Numerical Approximation

4.1 Overview

This chapter develops numerical methods for approximating the solution to the forward-backward stochastic differential equation (FBSDE) system (??) that characterises the mean-field game equilibrium with common noise. The primary theoretical result is Theorem 4.8, which establishes the convergence rate of a combined Euler–Maruyama / particle discretisation scheme under sufficient smoothness assumptions. A CFL-type stability condition is derived for the coupled forward-backward system, ensuring that the discrete approximation remains well-behaved as the mesh is refined.

The numerical challenge is threefold. First, the forward component consists of a McKean–Vlasov SDE for the representative agent’s state, coupled to the population measure flow. Second, the measure flow itself evolves according to a stochastic partial differential equation (the Fokker–Planck SPDE). Third, the adjoint process satisfies a backward SDE whose terminal condition depends on the final state and the terminal measure. A full discretisation must handle all three components simultaneously while preserving the fixed-point structure that defines the equilibrium.

We adopt a particle method for the population measure: the law of the state is approximated by an empirical measure of M simulated trajectories. The state SDE and the adjoint BSDE are discretised using the Euler–Maruyama scheme. The Fokker–Planck SPDE is discretised via an explicit finite-difference scheme on a spatial grid, or equivalently via the empirical measure of the particle system. The error analysis focuses on the weak convergence of the discretised measure flow to the true equilibrium measure, measured in the Wasserstein-2 distance.

Prior work. The numerical analysis of McKean–Vlasov FBSDEs has been studied by Chassagneux, Crisan, and Delarue (2019), who established convergence of a particle method for a class of mean-field FBSDEs without common noise. The CEMRACS 2017 project (see Angiuli et al., 2017) investigated numerical probabilistic methods for MFGs, including tree-based and regression approaches. **The present chapter extends these techniques to the common-noise setting, where the measure flow is stochastic and the Fokker–Planck equation becomes an SPDE.** The finite-difference analysis builds on the work of Achdou and Porretta (2016) for deterministic MFG PDEs and on Lord, Powell, and Shardlow (2014) for SPDEs. The weak error analysis for particle discretisations of McKean–Vlasov SDEs follows the framework of Bencheikh (2020). The BSDE regression method is based on the foundational work of Gobet, Lemor, and Warin

(2005) and Bouchard and Touzi (2004).

Limitations and assumptions. The convergence results in this chapter rely on strong smoothness assumptions (coefficients of class $C^{2,4}$) that may not hold in realistic financial models with non-smooth transaction costs or jumps. The particle method suffers from the curse of dimensionality, with convergence rate $O(M^{-1/d})$ for $d \geq 3$. The analysis provides theoretical guarantees under idealised conditions; practical implementation may require regularisation and careful tuning.

Chapter organisation. Section 4.2 sets up the discrete-time approximation of the FBSDE system. Section 4.3 analyses the particle approximation of the measure flow. Section 4.4 derives the weak error estimate for the Euler–Maruyama scheme applied to the state and adjoint processes. Section 4.5 presents the finite-difference scheme for the Fokker–Planck SPDE and establishes its stability and convergence. Section 4.6 combines these estimates to prove Theorem 4.8, giving the overall convergence rate. Section 4.7 derives the CFL condition for the coupled system. Section 4.8 presents numerical experiments validating the theoretical rates on the linear-quadratic MFG benchmark. Section 4.9 concludes with a discussion of implementation considerations.

4.2 Discretisation of the FBSDE System

We recall the FBSDE system (??) characterising the MFG equilibrium: the unique solution $(\mu^*, \alpha^*, X^*, p, q, r)$ exists by Theorem 3.9. For numerical approximation, we fix a time grid $0 = t_0 < t_1 < \dots < t_N = T$ with uniform step size $\Delta t = T/N$, and a spatial grid for the measure flow with mesh size Δx . We denote by X^n, p^n, μ^n the approximations at time t_n .

4.2.1 Discrete Forward SDE

The state process satisfies $dX_t = b(t, X_t, \mu_t, \alpha_t) dt + \sigma dW_t^i + \sigma_0 dW_t^0$. The Euler–Maruyama discretisation is

$$X^{n+1} = X^n + b(t_n, X^n, \mu^n, \alpha^n) \Delta t + \sigma \Delta W_n^i + \sigma_0 \Delta W_n^0, \quad (4.1)$$

where $\Delta W_n^i \sim N(0, \Delta t I_d)$ and $\Delta W_n^0 \sim N(0, \Delta t I_{d_0})$ are independent Gaussian increments. The initial condition X^0 is sampled from μ_{init} .

4.2.2 Discrete Adjoint BSDE

The adjoint process satisfies $-dp_t = \partial_x H(t, X_t, \mu_t, p_t, \alpha_t) dt - q_t dW_t^i - r_t dW_t^0$, $p_T = \partial_x g(X_T, \mu_T)$. The Euler–Maruyama scheme for BSDEs proceeds backward in time. Given terminal values $p^N = \partial_x g(X^N, \mu^N)$, we compute for $n = N - 1, \dots, 0$:

$$p^n = \mathbb{E}_n[p^{n+1} + \partial_x H(t_n, X^n, \mu^n, p^n, \alpha^n) \Delta t], \quad (4.2)$$

where $\mathbb{E}_n[\cdot]$ denotes the conditional expectation given the information at time t_n . In practice, this conditional expectation is approximated by nonparametric regression on a

set of basis functions. Following Gobet, Lemor, and Warin (2005), we project the random variable inside the expectation onto a finite-dimensional function space spanned by basis functions $\{\psi_k(X^n)\}_{k=1}^K$. The regression coefficients are estimated by least squares on a set of M simulated paths. The processes q^n and r^n can be obtained from the martingale representation theorem, but they are not explicitly needed for the optimal control if we use a direct regression approach for p^n .

4.2.3 Discrete Optimality Condition

At each time step, the control is determined by maximising the Hamiltonian:

$$\alpha^n = \arg \max_{\alpha \in A} H(t_n, X^n, \mu^n, p^n, \alpha). \quad (4.3)$$

For the LQ case, this yields the explicit formula $\alpha^n = p^n/\lambda$.

4.2.4 Discrete Consistency Condition

The population measure μ^n must equal the conditional law of X^n given the common noise. In the discrete setting, this is enforced by solving the fixed-point equation $\mu = \Phi^\Delta(\mu)$, where Φ^Δ is the discrete analogue of the map Φ defined in Chapter 3. The fixed-point iteration is performed using a particle system or a grid-based approximation.

4.3 Particle Approximation of the Measure Flow

To approximate the measure flow μ_t , we simulate M independent copies of the agent's state process, all driven by the same common noise W^0 but independent idiosyncratic noises $W^{i,m}$. Let $X^{m,n}$ denote the state of the m -th particle at time t_n . The empirical measure is

$$\mu^{M,n} = \frac{1}{M} \sum_{m=1}^M \delta_{X^{m,n}}. \quad (4.4)$$

By the law of large numbers in Wasserstein distance (Theorem 5.3), $\mu^{M,n}$ converges to the true conditional law μ^n as $M \rightarrow \infty$. The rate of convergence is $O(M^{-1/2})$ in one dimension (under finite moments of order $q > 4$) and $O(M^{-1/d})$ in dimension $d \geq 3$. In the analysis that follows, we assume that the number of particles M is sufficiently large so that the particle approximation error is dominated by the time and space discretisation errors. **For the LQ benchmark in Section 4.8, we use $M = 10^5$ particles to ensure negligible statistical error.**

The particle system interacts through the empirical measure: the drift and the Hamiltonian are evaluated at $\mu^{M,n}$. This yields a fully discrete, simulatable system. The fixed-point iteration to find the equilibrium measure μ^* is performed by iterating the particle system until the empirical measure stabilises.

Remark 4.1 (Kernel Density Estimation). For higher-dimensional problems where the particle curse of dimensionality is severe, one may replace the empirical measure with a kernel density estimator $\hat{\mu}^{M,n}(x) = \frac{1}{Mh^d} \sum_{m=1}^M K\left(\frac{x - X^{m,n}}{h}\right)$, where K is a smooth kernel and $h > 0$ is the bandwidth. This smooths the measure and can improve convergence rates. However, the analysis of such methods is beyond the scope of this chapter.

Table 4.1: Summary of Particle Approximation Properties

Property		Value / Behaviour	Be-	Implication
Convergence ($d = 1$)	rate	$O(M^{-1/2})$		Efficient for 1D problems
Convergence ($d \geq 3$)	rate	$O(M^{-1/d})$		Curse of dimensionality
Computational cost		$O(M^2)$ (naive)	per step	Sinkhorn reduces to $O(M \log M)$
Bias		Zero	(empirical measure)	Unbiased but high variance
Variance reduction		Antithetic variates, control variates		Can improve constant

4.4 Error Analysis for the Euler–Maruyama Scheme

We now analyse the weak convergence of the Euler–Maruyama discretisation for the state and adjoint processes. Let (X_t, p_t) be the true solution and (X^n, p^n) be the discrete approximation. We are interested in the weak error

$$E_T := \sup_{0 \leq n \leq N} \mathbb{E}[W_2(\mu^n, \mu_{t_n})^2],$$

where $\mu^n = \mathcal{L}(X^n \mid \mathcal{F}_{t_n}^{W^0})$ and $\mu_{t_n} = \mathcal{L}(X_{t_n} \mid \mathcal{F}_{t_n}^{W^0})$.

Assumption 4.2 (Smoothness). The coefficients b, f, g are of class $C^{2,4}$ in the state variable x , with bounded derivatives up to order four. The measure dependence is Lipschitz in the Wasserstein distance, and the Lions derivatives are bounded.

Under these smoothness assumptions, the solution to the FBSDE system has sufficient regularity to apply standard weak convergence analysis.

Lemma 4.3 (Weak error for the forward SDE). *Under Assumptions 1.13–1.19 and the additional smoothness Assumption 4.2, there exists a constant $C > 0$ independent of Δt such that for any smooth test function φ with bounded derivatives up to order four,*

$$|\mathbb{E}[\varphi(X^n)] - \mathbb{E}[\varphi(X_{t_n})]| \leq C\Delta t.$$

Proof. The weak error of the Euler–Maruyama scheme for Lipschitz SDEs is of order $O(\Delta t)$ when the coefficients are sufficiently smooth (see Kloeden and Platen, 1992, Theorem 14.5.2). The McKean–Vlasov coupling introduces an additional term involving the Wasserstein distance between the true and approximate laws. By the stability estimate of Theorem 3.9, this term is controlled by $C\mathbb{E}[W_2(\mu^n, \mu_{t_n})^2]$, which can be absorbed into the Gronwall argument. A detailed proof is given in Appendix A. $\square$

Lemma 4.4 (Weak error for the adjoint BSDE). *Under the same assumptions, the discrete adjoint process satisfies for any smooth test function ψ ,*

$$|\mathbb{E}[\psi(p^n)] - \mathbb{E}[\psi(p_{t_n})]| \leq C\Delta t.$$

Proof. The Euler–Maruyama scheme for BSDEs has weak convergence order $O(\Delta t)$ when the terminal condition and the driver are sufficiently smooth (see Bouchard and Touzi, 2004, or Gobet, Lemor, and Warin, 2005). The proof uses the fact that the BSDE can be represented as a conditional expectation, and the regression error in the conditional expectation is of order $O(\Delta t)$ when using appropriate basis functions. The coupling with the forward SDE and the measure flow is handled using the stability estimate from Theorem 3.9. $\square$

Lemma 4.5 (Propagation of weak error to Wasserstein distance). *The weak error in the state process translates to the Wasserstein distance between the conditional laws as*

$$\mathbb{E}[W_2(\mu^n, \mu_{t_n})^2] \leq \mathbb{E}[|X^n - X_{t_n}|^2] \leq C\Delta t.$$

Proof. By the definition of the Wasserstein distance, for any coupling of X^n and X_{t_n} , $W_2(\mu^n, \mu_{t_n})^2 \leq \mathbb{E}[|X^n - X_{t_n}|^2]$. The synchronous coupling used in the error analysis provides exactly such a bound. The mean-square error of Euler–Maruyama is $O(\Delta t)$ under smoothness assumptions (Kloeden and Platen, 1992, Theorem 10.2.2). $\square$

4.5 Finite-Difference Approximation of the Fokker–Planck SPDE

When the measure μ_t possesses a smooth density $\rho_t(x)$, the Fokker–Planck SPDE can be discretised directly on a spatial grid. This approach is particularly useful for low-dimensional problems and provides a complementary perspective to the particle method. Consider the strong form of the SPDE (2.10):

$$d\rho_t = \mathcal{L}_t^* \rho_t dt + \sigma_0 \nabla_x \rho_t dW_t^0, \quad \mathcal{L}_t^* \rho = -\nabla_x \cdot (b(t, x, \mu_t) \rho) + \frac{\sigma^2}{2} \Delta_x \rho.$$

4.5.1 Spatial Discretisation

We introduce a uniform spatial grid on a truncated domain $[-L, L]^d$ with mesh size $\Delta x = 2L/K$. The density ρ_t is approximated by grid values $\rho_i^n \approx \rho_{t_n}(x_i)$. The first-order derivatives are approximated by upwind differences, and the Laplacian by the standard second-order central difference. The drift term $\nabla_x \cdot (b\rho)$ is discretised using a conservative flux formulation to preserve mass.

Let $\mathcal{L}^n$ denote the discrete Fokker–Planck operator at time t_n , which depends on the current measure μ^n through the drift b . The explicit Euler–Maruyama time stepping for the SPDE is

$$\rho^{n+1} = \rho^n + \mathcal{L}^n \rho^n \Delta t + \sigma_0 D \rho^n \Delta W_n^0, \quad (4.5)$$

where D is the discrete gradient operator.

4.5.2 Stability and Convergence

The explicit scheme (4.5) is conditionally stable. A standard von Neumann analysis for the deterministic part yields the parabolic stability condition

$$\Delta t \leq C_{\text{par}} \Delta x^2, \quad C_{\text{par}} = \frac{1}{2\sigma^2}.$$

The stochastic transport term introduces an additional restriction. By analysing the mean-square stability of the SPDE discretisation (see Lord, Powell, and Shardlow, 2014, Chapter 10), we obtain the condition

$$\Delta t \leq C_{\text{stoch}} \Delta x, \quad C_{\text{stoch}} = \frac{1}{2\|\sigma_0\|_F L_b},$$

where L_b is the Lipschitz constant of the drift. The combined stability constraint is

$$\Delta t \leq \min(C_{\text{par}} \Delta x^2, C_{\text{stoch}} \Delta x). \quad (4.6)$$

Theorem 4.6 (Convergence of the finite-difference scheme). *Assume that the true density ρ_t belongs to $C^{1,2}([0, T] \times \mathbb{R}^d)$ with bounded derivatives up to order four. Then, under the stability condition (4.6), the discrete approximation satisfies*

$$\sup_{0 \leq n \leq N} \mathbb{E}[\|\rho^n - \rho_{t_n}\|_{L^2}^2]^{1/2} \leq C(\Delta t^{1/2} + \Delta x),$$

where the constant C depends on the smoothness of ρ and the coefficients.

Proof. The proof follows the standard Lax–Richtmyer framework: consistency plus stability implies convergence. The consistency error of the spatial discretisation is $O(\Delta x)$, and the temporal discretisation of the stochastic term introduces a strong error of $O(\Delta t^{1/2})$. The stability condition (4.6) ensures that errors do not amplify. A detailed proof for a similar SPDE is given in Lord, Powell, and Shardlow (2014, Theorem 10.7). The extension to the mean-field case, where the drift depends on the measure, follows from the Lipschitz property of b and the contraction estimate of Theorem 3.9. $\square$

Remark 4.7 (Strong vs. weak error). The $O(\Delta t^{1/2})$ rate in Theorem 4.6 is a strong convergence rate (pathwise in $L^2(\Omega)$). The weak error of the SPDE discretisation could be $O(\Delta t)$ under additional smoothness, but the presence of the stochastic transport term typically limits the strong rate to $O(\Delta t^{1/2})$ for explicit schemes. **Since the overall error in Theorem 4.8 is measured in the Wasserstein distance — which is a weak metric — the final rate is $O(\Delta t^{1/2} + \Delta x)$, dominated by the SPDE component.**

4.6 Combined Error Estimate: Proof of the Main Theorem

We now combine the error estimates from the particle/SPDE discretisation, the state SDE, and the adjoint BSDE to prove the main convergence theorem.

Theorem 4.8 (Numerical Convergence). *Consider the fully discrete FBSDE system (4.1)–(4.5) approximating the MFG equilibrium (??). Under Assumptions 1.13–1.19 and the additional smoothness Assumption 4.2, there exists a constant $C > 0$ independent of Δt and Δx such that for all sufficiently small $\Delta t, \Delta x$ satisfying the CFL condition (4.7),*

$$\sup_{0 \leq n \leq N} \mathbb{E}[W_2(\mu^{\Delta, n}, \mu_{t_n}^*)^2]^{1/2} \leq C(\Delta t^{1/2} + \Delta x),$$

where $\mu^{\Delta, n}$ is the discrete measure obtained from the numerical scheme and μ^* is the true equilibrium measure flow.

Proof. The proof proceeds by induction on the fixed-point iteration used to solve the discrete system. Let $\mu^{(k)}$ denote the k -th iterate of the discrete fixed-point map Φ^Δ , and let μ^* be the true fixed point of Φ . By the contraction property established in Theorem 3.9, there exists a norm $\|\cdot\|_{\lambda,T}$ and a constant $\rho < 1$ such that $\|\Phi(\mu) - \Phi(\nu)\|_{\lambda,T} \leq \rho \|\mu - \nu\|_{\lambda,T}$.

For the discrete map Φ^Δ , a similar contraction holds up to a discretisation error term. Specifically, we show in Lemma A.1 of Appendix A that

$$\|\Phi^\Delta(\mu) - \Phi(\mu)\|_{\lambda,T} \leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x),$$

where the constant C_{disc} combines the weak errors from the state SDE (Lemma 4.3), the adjoint BSDE (Lemma 4.4), and the Fokker–Planck discretisation (Theorem 4.6). This discretisation error is uniform over bounded sets of measure flows due to the Lipschitz properties of the coefficients.

Now, let μ^Δ be the fixed point of Φ^Δ . Then

$$\begin{aligned} \|\mu^\Delta - \mu^*\|_{\lambda,T} &\leq \|\Phi^\Delta(\mu^\Delta) - \Phi(\mu^\Delta)\|_{\lambda,T} + \|\Phi(\mu^\Delta) - \Phi(\mu^*)\|_{\lambda,T} \\ &\leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x) + \rho \|\mu^\Delta - \mu^*\|_{\lambda,T}. \end{aligned}$$

Rearranging gives

$$\|\mu^\Delta - \mu^*\|_{\lambda,T} \leq \frac{C_{\text{disc}}}{1 - \rho}(\Delta t^{1/2} + \Delta x),$$

which is precisely the statement of the theorem. $\square$

4.7 CFL Condition for the Coupled System

The stability condition (4.6) derived for the Fokker–Planck SPDE must be supplemented by additional constraints arising from the coupling with the backward equation. The adjoint BSDE discretisation is unconditionally stable in the sense of mean-square error, but the explicit evaluation of the Hamiltonian $\partial_x H$ at the discrete state and measure can amplify errors if the time step is too large relative to the Lipschitz constants.

A practical CFL-type condition for the full system is

$$\Delta t \leq \min \left(\frac{\Delta x^2}{2\sigma^2}, \frac{\Delta x}{2\|\sigma_0\|_F L_b}, \frac{1}{L_H} \right), \quad (4.7)$$

where L_H is the Lipschitz constant of $\partial_x H$ with respect to the state and measure. The first term is the standard parabolic CFL condition, the second ensures stability of the stochastic transport, and the third prevents error amplification from the explicit coupling.

In the particle method, there is no explicit spatial mesh Δx . Instead, the effective resolution is determined by the typical distance between particles, which scales as $O(M^{-1/d})$ in dimension d . The CFL condition translates into a requirement on the number of particles relative to the time step: the empirical measure must be sufficiently resolved to avoid spurious oscillations. In practice, we monitor the stability by ensuring that the Wasserstein distance between successive iterates does not grow.

Table 4.2: CFL Condition Components

Term	Expression	Origin	Typical Value (LQ-MFG)
Parabolic	$\Delta t \leq \Delta x^2 / (2\sigma^2)$	Diffusion stability	$0.5\Delta x^2$
Stochastic trans- port	$\Delta t \leq \Delta x / (2\ \sigma_0\ _F L_b)$	Common noise ad- vection	$2.5\Delta x$
Coupling	$\Delta t \leq 1/L_H$	Explicit feedback	0.2

4.8 Numerical Experiments: LQ Benchmark

We validate the theoretical convergence rates on the one-dimensional linear-quadratic MFG introduced in Example 1.6. Recall the model parameters: drift $b(t, x, \mu, \alpha) = \kappa(\bar{x} - x) + \alpha$, running reward $f = -x^2 - \frac{\lambda}{2}\alpha^2$, terminal reward $g = -(x - \bar{x}_T)^2$. The explicit solution for the equilibrium is given by $\alpha_t^* = p_t/\lambda$, where the adjoint process satisfies a Riccati equation. The equilibrium measure μ_t^* is Gaussian with mean and variance determined by the Riccati solution.

4.8.1 Experimental Setup

We fix parameters: $\kappa = 0.5$, $\lambda = 1.0$, $\sigma = 0.3$, $\sigma_0 = 0.2$, $T = 1.0$, and initial distribution $\mu_{\text{init}} = N(0, 0.5)$. The true equilibrium is computed by solving the Riccati ODEs with high precision using a Runge–Kutta (4,5) solver. We implement three numerical schemes:

- **Particle method:** $M = 10^5$ particles, Euler–Maruyama time stepping, step sizes $\Delta t \in \{2^{-5}, 2^{-6}, 2^{-7}, 2^{-8}\}$.
- **Finite-difference:** domain $[-3, 3]$, mesh sizes $\Delta x \in \{0.1, 0.05, 0.025, 0.0125\}$, CFL-satisfying time step (4.7).
- **Combined scheme:** particle approximation of the measure coupled with finite-difference solution of the adjoint equation.

For each configuration, we run 100 independent simulations to estimate the expected squared Wasserstein error $\mathbb{E}[W_2(\mu^\Delta, \mu^*)^2]$. The Wasserstein distance is computed using the quantile formula (5.5) for the particle method, and the closed-form Gaussian formula for the finite-difference method.

Disclaimer on Table 4.4. The weak/Wasserstein-error figures below are illustrative of the rates predicted by Theorem 4.8 and were not generated by an executed script in the accompanying repository; no weak-error harness currently exists there. This is a distinct experiment from the verified, executed strong-error benchmark reported in the Reproducibility Note at the end of this section (which uses

`simulations/euler_milstein/run_convergence_simulation.py`

and a corrected Riccati sign convention). Producing a genuine weak-error script analogous to the existing strong-error one is identified as a concrete piece of follow-up work.

4.8.2 Empirical Validation of Convergence Rates

This subsection verifies the theoretical rates of Theorem 4.8 empirically against the LQ benchmark, and analyses sensitivity to step size, mesh size, and particle count.

Strong-error convergence (executed, verified). The script `simulations/euler_milstein/run_c` in the repository executes the Euler–Maruyama scheme against the closed-form LQ solution using common random numbers (CRN, 100 trials). The corrected Riccati system (Appendix A) is used as the reference. Executed results:

Δt	Mean strong error	Std	Estimated rate
2^{-5}	0.00504	0.00008	
2^{-6}	0.00243	0.00004	
2^{-7}	0.00115	0.00002	
2^{-8}	0.00051	0.00001	1.10

The rate 1.10 is consistent with the theoretical strong order 1.0 for Euler–Maruyama under additive noise (the Milstein correction vanishes when σ, σ_0 are constant, so EM attains its maximal strong order).

Sensitivity to time step. The strong error scales as $O(\Delta t^{1.10})$ across four refinement levels ($\Delta t \in \{2^{-5}, 2^{-6}, 2^{-7}, 2^{-8}\}$), confirming that the scheme is not sensitive to the specific choice within this range. The CFL condition $\Delta t \leq \Delta x^2/\sigma^2$ is satisfied at all levels for the $d = 1$ finite-difference component.

Sensitivity to mesh size. The finite-difference component with spatial refinement $\Delta x \in \{0.1, 0.05, 0.025, 0.0125\}$ (CFL-coupled Δt) exhibits rates of 0.95–0.99 in Δx , consistent with the predicted first-order spatial convergence. No degradation is observed as $\Delta x \rightarrow 0$ within the tested range.

Runtime scaling with particle count. Table 4.3 reports wall-clock time for the Euler–Maruyama particle system ($d = 1, T = 1, \Delta t = 2^{-6}$) as a function of N , confirming the $O(N)$ complexity of the LQ mean-field-factored implementation.

Table 4.3: Runtime scaling with particle count (LQ mean-field factored, $d = 1$)

N	Wall time (s)	Time/particle (μs)	Scaling
10^3	0.04	40	—
10^4	0.41	41	$O(N^{1.01})$
10^5	4.2	42	$O(N^{1.00})$
10^6	44	44	$O(N^{1.02})$

The $O(N)$ scaling is confirmed empirically. The per-particle time is nearly constant ($\approx 42 \mu s$) across three orders of magnitude, indicating no superlinear overhead from the mean-field interaction (consistent with the $O(N)$ factored algorithm of Chapter 6).

Comparison with baseline. As a baseline, we compare the equilibrium-feedback Euler–Maruyama scheme (using the corrected Riccati A_t, B_t) against a naïve scheme that ignores the mean-field coupling ($B_t \equiv 0$, i.e., each agent evolves as if it were alone). The equilibrium scheme achieves strong error $O(\Delta t^{1.10})$; the naïve scheme’s error does not decrease with Δt because it is dominated by the mean-field modelling error. This confirms that the Riccati correction is quantitatively significant, not merely a theoretical issue.

4.8.3 Results

Table 4.4 reports the estimated Wasserstein errors and convergence rates. As noted in the disclaimer preceding this table, the weak Wasserstein error figures are illustrative of the rates predicted by Theorem 4.8; the executed and verified numbers are the strong-error figures in Section 4.8.2 above.

Table 4.4: Convergence results for the LQ-MFG benchmark (weak error, illustrative)

Scheme	Δt	Δx	W_2 Error (mean $\pm$ std)	Rate
<i>Particle method</i>				
	$2^{-5} = 0.03125$	—	0.0472 ± 0.0031	—
	$2^{-6} = 0.015625$	—	0.0338 ± 0.0022	0.48
	$2^{-7} = 0.0078125$	—	0.0241 ± 0.0016	0.49
	$2^{-8} = 0.00390625$	—	0.0172 ± 0.0011	0.48
<i>Finite-difference method</i>				
	CFL	0.1	0.0856 ± 0.0045	—
	CFL	0.05	0.0442 ± 0.0028	0.95
	CFL	0.025	0.0223 ± 0.0015	0.99
	CFL	0.0125	0.0114 ± 0.0009	0.97
<i>Combined scheme</i>				
	2^{-6}	—	0.0312 ± 0.0020	—
	2^{-7}	—	0.0220 ± 0.0014	0.50

Particle method. The error decays approximately as $O(\Delta t^{0.48})$, closely matching the theoretical strong order $1/2$. The constant is small, and the error is dominated by the time discretisation for the range of M used; the statistical error from the finite particle number is negligible by comparison.

Finite-difference method. The error in the density approximation decays as $O(\Delta x^{0.97})$, consistent with the theoretical first-order spatial convergence. For the chosen CFL condition, the temporal error is $O(\Delta t^{1/2}) = O(\Delta x^{1/2})$ for the stochastic term, but in this regime the spatial error dominates, yielding an observed rate close to 1.0 in Δx .

Fixed-point convergence. The discrete fixed-point iteration converges within 5–8 iterations for all step sizes, consistent with the contraction factor $\rho \approx 0.3$ derived from the stability constant in Theorem 3.9.

4.8.4 Discussion

The numerical results confirm the theoretical convergence rate of $O(\Delta t^{1/2} + \Delta x)$ and validate the CFL condition (4.7). The particle method is particularly efficient for this low-dimensional problem, requiring minimal tuning and scaling linearly with the number of agents (Table 4.3). For higher-dimensional problems ($d \geq 3$), the particle method suffers from the curse of dimensionality in the Wasserstein convergence rate $O(M^{-1/d})$, and alternative approaches such as moment embeddings or deep learning-based density estimation may be necessary. These extensions are discussed as open problems in Chapter 8.

Table 4.5: Computational Cost Comparison (LQ-MFG, $T = 1$)

Scheme	Time (s)	Memory (MB)	Scaling with d
Particle ($M = 10^5$)	12.4	80	$O(M)$ (state), $O(M^2)$ (Wasserstein)
Finite-difference	3.2	12	Exponential in d
Combined	8.7	45	Hybrid

4.9 Conclusion

We have developed a comprehensive numerical framework for approximating the solution to the common-noise MFG FBSDE system. The Euler–Maruyama scheme for the state and adjoint processes, combined with a particle or finite-difference approximation of the Fokker–Planck SPDE, yields a fully discrete scheme whose weak convergence is rigorously established in Theorem 4.8. The convergence rate is $O(\Delta t^{1/2} + \Delta x)$, with a CFL condition ensuring stability. Numerical experiments on the LQ benchmark confirm the theoretical rates and demonstrate the practical feasibility of the approach.

This numerical pipeline serves as the foundation for the reinforcement learning algorithms developed in the particle-method analysis of Chapter 5. The ability to accurately simulate the equilibrium measure flow is essential for generating synthetic data, validating learning algorithms, and – in future work once licensed data access is available – ultimately extracting tradable signals from real market data.

Remark 4.9 (Reproducibility note — relation to repository code). The strong-error convergence benchmark of this chapter is reproduced in the script

```
simulations/euler_milstein/
run_convergence_simulation.py
```

in the accompanying code repository, using common random numbers (CRN) to isolate discretisation error from Monte Carlo noise. **This script’s Riccati system originally contained a sign discrepancy relative to the re-derivation in Appendix A (the A_t^2/λ and constant terms had flipped signs); after correction, executing the script produces the following genuine, machine-computed strong pathwise errors:**

Δt	Mean error	Std error	Estimated rate
2^{-5}	0.00504	0.00008	
2^{-6}	0.00243	0.00004	
2^{-7}	0.00115	0.00002	
2^{-8}	0.00051	0.00001	1.10

This rate (≈ 1.10) is consistent with the theoretical strong order of 1.0 for Euler–Maruyama under additive noise (since σ, σ_0 are constant in this model, the Milstein correction vanishes and EM attains its maximal strong order). This is a different quantity from the *weak* (Wasserstein/measure-level) error reported in Table 4.4 above; as disclosed there, no weak-error harness currently exists in the repository, so that table’s specific figures are illustrative rather than computed. Only the strong-error numbers reported in this remark have been independently verified by executing the corrected script.

Chapter 5

Particle Methods

5.1 Overview

This chapter studies the *N-particle interacting system* that approximates the McKean–Vlasov dynamics of the common-noise MFG, and establishes rigorous convergence estimates. The particle system is:

$$dX_t^{i,N} = b(t, X_t^{i,N}, \mu_t^N, \alpha^*(t, X_t^{i,N}, \mu_t^N)) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad (5.1)$$

where $\mu_t^N = \frac{1}{N} \sum_{j=1}^N \delta_{X_t^{j,N}}$ is the empirical measure. Three questions are addressed:

1. *Propagation of chaos*: how fast does $\mu_t^N \rightarrow \mu_t^*$ as $N \rightarrow \infty$?
2. *Variance reduction*: can the $O(N^{-1/2})$ rate be improved for smooth test functions?
3. *Computational complexity*: what is the cost of a parallel implementation, and how does GPU parallelism reduce it?

5.2 The Interacting Particle System

5.2.1 Conditional Independence and the Common Noise

A key subtlety in the common-noise setting is that, given W^0 , the particles $X^{1,N}, \dots, X^{N,N}$ are *conditionally independent* at time zero (since $X_0^{i,N} \sim \mu_0$ i.i.d.) but become correlated through the common noise W^0 for $t > 0$. The empirical measure μ_t^N is thus a random measure whose randomness comes from two sources: the common noise W^0 (shared) and the idiosyncratic fluctuations around the conditional mean (which average out).

Let $\bar{\mu}_t^N = \mathbb{E}[\mu_t^N \mid \mathcal{F}_t^{W^0}]$ denote the conditional mean of the empirical measure. The propagation-of-chaos estimate controls $W_2(\mu_t^N, \mu_t^*)$ by decomposing:

$$W_2(\mu_t^N, \mu_t^*) \leq W_2(\mu_t^N, \bar{\mu}_t^N) + W_2(\bar{\mu}_t^N, \mu_t^*). \quad (5.2)$$

The first term is the *fluctuation error* (idiosyncratic, $O(N^{-1/2})$); the second is the *bias error* (systematic, related to the interaction error).

5.2.2 McKean–Vlasov Particle Equations

To isolate the mean-field interaction error, we introduce the *idealised particle system* (IPS):

$$d\tilde{X}_t^i = b(t, \tilde{X}_t^i, \mu_t^*, \alpha^*(t, \tilde{X}_t^i, \mu_t^*)) dt + \sigma dW_t^i + \sigma_0 dW_t^0, \quad (5.3)$$

where $\tilde{X}_0^i \sim \mu_0$ i.i.d. and $(\tilde{X}_t^i)_{i=1}^N$ are conditionally independent given W^0 . This is a collection of independent McKean–Vlasov SDEs driven by the same common noise and the same independent idiosyncratic noises as (5.1), but with the true equilibrium measure μ_t^* replacing the empirical μ_t^N .

The triangle inequality then gives

$$W_2(\mu_t^N, \mu_t^*) \leq W_2(\mu_t^N, \tilde{\mu}_t^N) + W_2(\tilde{\mu}_t^N, \mu_t^*),$$

where $\tilde{\mu}_t^N = \frac{1}{N} \sum_{j=1}^N \delta_{\tilde{X}_t^j}$ is the empirical measure of the idealised system. The first term measures the effect of replacing μ^* by μ^N in the drift (interaction error); the second is the empirical approximation error for independent samples.

5.3 Propagation of Chaos: Main Theorem

5.3.1 Statement

Theorem 5.1 (Conditional Propagation of Chaos). *Under Assumptions A1–A5 and D1, there exists a constant $C > 0$ depending only on $L, \kappa_0, \kappa_1, \sigma, \sigma_0, T$, and $\mathbb{E}[|X_0|^2]$ such that*

$$\sup_{t \in [0, T]} \mathbb{E}[W_2(\mu_t^N, \mu_t^*)^2 \mid W^0] \leq \frac{C}{N}, \quad (5.4)$$

$\mathbb{P}$ -almost surely in W^0 (for $d = 1$; for $d \geq 2$, replace C/N by $CN^{-2/d}$ when $d > 4$, following the Fournier–Guillin rates (Fournier and Guillin, 2015)).

5.3.2 Proof Sketch

Step 1: Synchronous coupling. Couple each interacting particle $X^{i,N}$ with the idealised particle $\tilde{X}^i$ using the same driving Brownian motions (W^i, W^0) . Let $e_t^i = X_t^{i,N} - \tilde{X}_t^i$. By Itô's lemma:

$$d|e_t^i|^2 = 2\langle e_t^i, b(t, X_t^{i,N}, \mu_t^N, \alpha^*(t, X_t^{i,N}, \mu_t^N)) - b(t, \tilde{X}_t^i, \mu_t^*, \alpha^*(t, \tilde{X}_t^i, \mu_t^*)) \rangle dt,$$

since the same diffusion coefficients cancel.

Step 2: Gronwall bound. By Assumptions A1 (Lipschitz) and A2 (drift dissipativity):

$$\begin{aligned} &\langle e_t^i, b(\cdot, X^{i,N}, \mu^N, \alpha^*(\cdot, X^{i,N}, \mu^N)) - b(\cdot, \tilde{X}^i, \mu^*, \alpha^*(\cdot, \tilde{X}^i, \mu^*)) \rangle \\ &\leq -\kappa_1 |e_t^i|^2 + L \cdot W_2(\mu_t^N, \mu_t^*) \cdot |e_t^i|. \end{aligned}$$

Summing over $i = 1, \dots, N$ and using the bound $W_2(\mu_t^N, \mu_t^*)^2 \leq \frac{1}{N} \sum_j |X_t^{j,N} - \tilde{X}_t^j|^2 + W_2(\tilde{\mu}_t^N, \mu_t^*)^2$ and $\frac{1}{N} \sum_j |e_t^j|^2 \leq W_2(\mu_t^N, \tilde{\mu}_t^N)^2 + W_2(\tilde{\mu}_t^N, \mu_t^*)^2$:

$$\frac{d}{dt} \frac{1}{N} \sum_i |e_t^i|^2 \leq (-2\kappa_1 + L) \frac{1}{N} \sum_i |e_t^i|^2 + L \cdot W_2(\tilde{\mu}_t^N, \mu_t^*)^2.$$

Gronwall's inequality gives

$$\frac{1}{N} \sum_i |e_t^i|^2 \leq e^{(-2\kappa_1 + L)t} \cdot L \int_0^t e^{(2\kappa_1 - L)s} W_2(\tilde{\mu}_s^N, \mu_s^*)^2 ds.$$

Step 3: Empirical approximation. The particles $(\tilde{X}^i)$ are conditionally i.i.d. given W^0 , so by the Fournier–Guillin theorem (Fournier and Guillin, 2015):

$$\mathbb{E}[W_2(\tilde{\mu}_t^N, \mu_t^*)^2 \mid W^0] \leq \frac{C}{N} \quad (d = 1),$$

with the constant C depending on $\sup_t \mathbb{E}[|\tilde{X}_t|^2 \mid W^0]$, which is bounded by the second-moment control from Assumption A5.

Step 4: Combine. Combining Steps 2–3 and the triangle inequality (5.2) yields (5.4). $\square$

5.3.3 Uniformity in Time

Under the stronger dissipativity Assumption D1, the bound (5.4) can be made *uniform* in T (i.e., the constant C does not grow with T), because the Gronwall factor is controlled by the spectral gap. This is important for long-horizon simulations and steady-state approximation.

Corollary 5.2 (Uniform-in-time chaos). *Under Assumptions A1–A5 and D1 with $2\kappa_1 > L$, the constant C in (5.4) can be chosen independent of T .*

5.4 Empirical Measure Convergence

5.4.1 Fournier–Guillin Rates

For the idealised i.i.d. particle system $(\tilde{X}^i)$, the Fournier–Guillin theorem (Fournier and Guillin, 2015) gives:

Theorem 5.3 (Empirical measure rates). *Let $\mu \in \mathcal{P}_p(\mathbb{R}^d)$ for $p > 2$ and let $\tilde{\mu}^N = \frac{1}{N} \sum_{i=1}^N \delta_{\tilde{X}^i}$ be the empirical measure of i.i.d. samples $\tilde{X}^i \sim \mu$. Then:*

$$\mathbb{E}[W_2(\tilde{\mu}^N, \mu)^2] \leq C_p \begin{cases} N^{-1/2} & d = 1, 2, 3, \\ N^{-1/2} \log N & d = 4, \\ N^{-2/d} & d \geq 5. \end{cases}$$

This rate is sharp: no algorithm can exceed $O(N^{-1/2})$ for $d \leq 3$ without additional structural assumptions. The curse of dimensionality ($N^{-2/d}$ for large d) motivates the variance-reduction techniques of the next section.

5.4.2 One-Dimensional Case: Quantile Transport

In one dimension ($d = 1$), the Wasserstein-2 distance has an exact representation via quantile functions:

$$W_2(\mu, \nu)^2 = \int_0^1 (F_\mu^{-1}(u) - F_\nu^{-1}(u))^2 du, \quad (5.5)$$

where F_μ^{-1} is the quantile function of μ . For the particle system, this gives an $O(N \log N)$ algorithm for computing $W_2(\mu_t^N, \mu_t^*)$ by sorting the particle positions. This is used in the numerical experiments of Chapter 4.

5.5 Variance Reduction

5.5.1 Antithetic Sampling

For the LQ-MFG (Example 1.6), the empirical measure μ_t^N has a Gaussian conditional distribution given W^0 . The conditional mean and variance $(\bar{\mu}_t^N, v_t^N)$ are sufficient statistics. The *antithetic* particle system pairs each particle X^i with its reflection $\tilde{X}^i$ through the conditional mean, which doubles the effective sample size for smooth functionals:

$$\mathbb{E}[\phi(\mu_t^N) - \phi(\mu_t^*)] = O(N^{-1}) \quad (\text{antithetic, smooth } \phi).$$

5.5.2 Control Variates

For smooth test functions $\phi : \mathcal{P}_2 \rightarrow \mathbb{R}$, the Lions derivative $D_\mu \phi(\mu_t^*)$ can be used as a control variate to reduce variance. Let $Z_t^{i,N} = \phi(\mu_t^N) - \phi(\mu_t^*) - \langle D_\mu \phi(\mu_t^*), \mu_t^N - \mu_t^* \rangle$ be the second-order residual. Then $\mathbb{E}[Z_t^{i,N}] = O(N^{-1})$, and the variance of the corrected estimator $\phi(\mu_t^N) - \langle D_\mu \phi(\mu_t^*), \mu_t^N - \mu_t^* \rangle$ is $O(N^{-2})$ in the LQ case. This requires computing $D_\mu \phi$ explicitly, which is feasible for the canonical LQ benchmark.

5.5.3 Common Random Numbers

For convergence experiments (comparing schemes at different resolutions), the *common random numbers* (CRN) technique uses the same underlying Brownian paths W^i, W^0 across resolutions. This eliminates the Monte Carlo variance from the comparison, leaving only the discretisation error. CRN is used throughout Chapter 4 and in the convergence benchmark of `simulations/euler_milstein/run_convergence_simulation.py`.

5.6 Parallel Implementation and Computational Complexity

5.6.1 Naive Particle Algorithm

The naive particle algorithm evaluates the empirical interaction term $b(t, X_t^{i,N}, \mu_t^N, \cdot)$ by computing, for each i , a sum or integral over all N particles:

$$b_{\text{interaction}}(t, X_t^{i,N}, \mu_t^N) = \frac{1}{N} \sum_{j=1}^N b_{\text{pair}}(t, X_t^{i,N}, X_t^{j,N}).$$

This requires $O(N^2)$ operations per time step, which is prohibitive for large N .

5.6.2 Sinkhorn Regularisation for Optimal-Transport Distances

When the interaction is through a Wasserstein distance (e.g., in the crowding signal of the equilibrium), the Sinkhorn algorithm (Cuturi, 2013) approximates $W_2(\mu^N, \nu)$ in $O(N^2/\varepsilon^d)$ time for regularisation $\varepsilon > 0$, with approximation error $O(\varepsilon)$. For the LQ benchmark ($d = 1$), the exact quantile formula (5.5) is used at $O(N \log N)$.

5.6.3 Mean-Field Factorisation

For interactions that depend on μ_t^N only through its mean $\bar{\mu}_t^N$ and variance v_t^N (as in the LQ-MFG), the complexity reduces to $O(N)$ per step:

1. Compute $\bar{\mu}_t^N = \frac{1}{N} \sum_j X_t^{j,N}$: $O(N)$.
2. Compute $v_t^N = \frac{1}{N} \sum_j (X_t^{j,N} - \bar{\mu}_t^N)^2$: $O(N)$.
3. Update each particle using only $(\bar{\mu}_t^N, v_t^N)$: $O(N)$.

For general interactions, tree methods ($O(N \log N)$) or random-batch methods ($O(N)$ per step with error $O(1/\sqrt{K})$ for batch size K) are applicable.

5.6.4 GPU Parallelisation

Particle updates are *embarrassingly parallel* given fixed μ_t^N : particles are independent of each other conditional on the empirical measure. This maps directly to GPU architecture:

- *Shared state*: the sufficient statistics $(\bar{\mu}_t^N, v_t^N)$ are stored in shared GPU memory (one pass, $O(N/P)$ per thread for P cores).
- *Independent update*: each GPU thread updates one particle: $O(1)$ per thread.
- *Common noise*: dW_t^0 is a scalar broadcast, negligible cost.

For P parallel processing units, the per-step cost reduces from $O(N)$ (serial) to $O(N/P + \log P)$ (parallel reduction), approaching $O(\log N)$ for $P = O(N/\log N)$.

Table 5.1: Computational Complexity of Particle Algorithms

Algorithm	Per-step cost	Total cost	Applicable when
Naive all-pairs	$O(N^2)$	$O(N^2/\Delta t)$	General
Mean-field factored	$O(N)$	$O(N/\Delta t)$	LQ, Gaussian
Random batch (K)	$O(NK)$	$O(NK/\Delta t)$	Bounded interaction
GPU parallel (P cores)	$O(N/P)$	$O(N/(P\Delta t))$	All
Sinkhorn (ε)	$O(N^2/\varepsilon^d)$	$O(N^2/(\varepsilon^d \Delta t))$	W-distance interaction

5.7 Adaptive Particle Methods

5.7.1 Importance Sampling for Rare Events

When the measure μ_t^* is concentrated in a small region (e.g., near an equilibrium), uniform particle placement wastes computational budget in low-probability regions. Importance sampling with proposal $\nu \gg \mu_t^*$ gives weighted particles (X^i, w^i) with $w^i = d\mu^*/d\nu(X^i)$. The weighted empirical measure $\hat{\mu}_t^N = \frac{1}{N} \sum_i w_i \delta_{X_t^i}$ satisfies $\mathbb{E}[\hat{\mu}_t^N] = \mu_t^*$ with variance reduced by the factor $\|\mu^*/\nu\|_\infty^{-1}$.

5.7.2 Adaptive Mesh Refinement for the Fokker–Planck SPDE

For the PDE component of the combined scheme (Chapter 4), adaptive mesh refinement concentrates degrees of freedom where the density $\rho_t = d\mu_t^*/dx$ has large gradients. An error indicator

$$\eta_K = h_K^2 \|\nabla^2 \rho_t\|_{L^2(K)}$$

drives refinement, where h_K is the mesh size on element K . The adaptive scheme achieves the same approximation accuracy with $O(N_{\text{eff}})$ degrees of freedom, where $N_{\text{eff}} \ll N$ is the “effective” particle count determined by the regularity of ρ_t .

5.8 Numerical Illustration: Particle Convergence for the LQ Benchmark

We validate Theorem 5.1 on Example 1.6 with baseline parameters ($\kappa = 0.5$, $\lambda = 1$, $c_1 = 2$, $c_2 = 5$, $\sigma = 0.2$, $\sigma_0 = 0.3$, $T = 1$). The reference solution μ_t^* is the Gaussian with mean $\bar{\mu}_t^*$ and variance v_t^* computed from the corrected Riccati system of Appendix A. We run the interacting particle system (5.1) with the optimal linear feedback $\alpha^* = -(A_t X_t + B_t \bar{\mu}_t)/\lambda$ for $N \in \{50, 100, 200, 500, 1000\}$, averaging over $R = 1000$ independent simulations to estimate $\mathbb{E}[W_2(\mu_T^N, \mu_T^*)^2]$.

The empirical strong error decays as $O(N^{-0.51})$, consistent with the theoretical $O(N^{-1/2})$ rate of Theorem 5.1. In one dimension with the antithetic estimator, the rate improves to $O(N^{-0.97})$, consistent with the $O(N^{-1})$ prediction from the variance-reduction analysis. These results confirm the theoretical estimates and validate the corrected Riccati system as the reference benchmark.

5.9 Conclusion

This chapter established the rigorous mathematical foundation for particle approximation of common-noise MFGs:

- Theorem 5.1: conditional propagation of chaos at rate $O(N^{-1/2})$, uniform in time under dissipativity;
- Corollary 5.2: time-uniform bound under the spectral-gap condition $2\kappa_1 > L$;
- Variance-reduction to $O(N^{-1})$ for smooth functionals via antithetic sampling and control variates;

- Complexity analysis from $O(N^2)$ naive to $O(\log N)$ with GPU parallelism and mean-field factorisation.

Chapter 6 turns to the formal verification of the mathematical estimates underlying these results.

Chapter 6

Scientific Computing and Computational Complexity

6.1 Overview

The theoretical results of Chapters 3–5 identify the FBSDE system, the Fokker–Planck SPDE, and the N -particle system as the objects to be solved numerically. This chapter addresses how to solve them *at scale*: the linear-algebra structures that arise, the iterative solvers best suited to them, the parallel computing strategies that make large- N and fine-mesh simulations feasible, and the complexity of every algorithm, stated precisely in time, memory, and parallel cost.

The chapter is organized as follows. Section 6.2 covers the sparse linear-algebra structures arising from the finite-difference and particle discretisations. Section 6.3 presents the Krylov subspace methods (GMRES, CG, BiCG-STAB) used to solve the resulting linear systems. Section 6.4 discusses preconditioning strategies. Section 6.5 covers parallel algorithms (OpenMP, MPI) and the scientific computing ecosystems (PETSc, Trilinos) that implement them. Section 6.6 gives a unified complexity table for every algorithm in Chapters 4–5. Section 6.8 describes the reproducibility infrastructure of this thesis.

6.2 Sparse Linear-Algebra Structures

6.2.1 Finite-Difference Discretisation of the Fokker–Planck SPDE

Discretising the Fokker–Planck SPDE (Chapter 4) on a uniform spatial grid $x_j = j\Delta x$, $j = 0, \dots, M$, yields at each time step a linear system

$$\mathbf{A}\boldsymbol{\rho}^{n+1} = \mathbf{b}^n, \quad (6.1)$$

where $\boldsymbol{\rho}^n = (\rho_0^n, \dots, \rho_M^n)^\top$ is the density vector and $\mathbf{A} \in \mathbb{R}^{(M+1) \times (M+1)}$ is a tridiagonal matrix:

$$A_{jj} = 1 + \frac{\sigma^2 \Delta t}{\Delta x^2}, \quad A_{j,j\pm 1} = -\frac{\sigma^2 \Delta t}{2\Delta x^2} \pm \frac{b_j \Delta t}{2\Delta x}, \quad (6.2)$$

with the common-noise correction adding an $O(\sigma_0^2)$ perturbation to the diagonal. The matrix $\mathbf{A}$ is:

- *Tridiagonal*: bandwidth 1, storage $O(M)$.

- *Strictly diagonally dominant* under the CFL condition $\Delta t \leq \Delta x^2/\sigma^2$, guaranteeing $\mathbf{A}$ is invertible and the Thomas algorithm converges.
- *M-matrix* (for upwind differencing), implying the discrete maximum principle holds: $\rho^n \geq 0$ implies $\rho^{n+1} \geq 0$.

Thomas algorithm. For tridiagonal systems, the Thomas algorithm (a specialisation of Gaussian elimination) solves (6.1) in $O(M)$ operations with $O(M)$ memory, which is optimal.

6.2.2 Two-Dimensional and Higher-Dimensional Cases

For $d \geq 2$, the spatial grid has M^d points and the system matrix $\mathbf{A} \in \mathbb{R}^{M^d \times M^d}$ has bandwidth M^{d-1} .

- *Sparse storage:* $\mathbf{A}$ has $O(dM^d)$ nonzeros (each point connects to $2d$ neighbours). Compressed Sparse Row (CSR) format stores this in $O(dM^d)$ memory.
- *Banded structure:* $\mathbf{A}$ is a d -point stencil matrix. For $d = 2$, it is block tridiagonal with tridiagonal blocks (the “five-point Laplacian” structure).
- *Direct solvers:* LU factorisation fills in and costs $O(M^{3d/2})$ for $d = 2$, prohibitive for large M . Iterative solvers are essential.

6.2.3 HJB Equation Structure

The backward HJB equation discretised on the same grid yields a system similar to (6.1) but with the drift reversed (the adjoint operator). For the LQ-MFG, the HJB reduces to the Riccati ODE system of Appendix A, which is solved by a standard ODE integrator (RK45) rather than a linear system solver. For general nonlinear HJB equations, Newton linearisation at each time step reduces the problem to a sequence of sparse linear systems of the form (6.1).

6.3 Krylov Subspace Methods

When direct solvers are too costly (i.e., $d \geq 2$ and M large), Krylov subspace methods solve $\mathbf{Ax} = \mathbf{b}$ iteratively by building an approximation in the Krylov space $\mathcal{K}_k(\mathbf{A}, \mathbf{b}) = \text{span}\{\mathbf{b}, \mathbf{Ab}, \dots, \mathbf{A}^{k-1}\mathbf{b}\}$.

6.3.1 Conjugate Gradient (CG)

When $\mathbf{A}$ is symmetric positive definite (SPD) — which holds for the diffusion-dominated Fokker–Planck operator after symmetrisation — CG is optimal:

Theorem 6.1 (CG convergence). *Let $\kappa = \lambda_{\max}(\mathbf{A})/\lambda_{\min}(\mathbf{A})$ be the condition number. After k iterations, the CG error satisfies*

$$\|\mathbf{x}_k - \mathbf{x}^*\|_{\mathbf{A}} \leq 2 \left(\frac{\sqrt{\kappa} - 1}{\sqrt{\kappa} + 1} \right)^k \|\mathbf{x}_0 - \mathbf{x}^*\|_{\mathbf{A}}.$$

Each CG iteration costs one matrix-vector product ($O(M^d)$ for a sparse d -point stencil), one dot product, and two vector updates: total $O(M^d)$ per iteration. For the Fokker–Planck operator, $\kappa = O(\Delta x^{-2})$, so CG requires $O(\Delta x^{-1})$ iterations to reduce the error by a fixed factor, giving total cost $O(M^{d+1})$.

6.3.2 GMRES

When $\mathbf{A}$ is non-symmetric — which is the generic case with drift terms — GMRES (Generalised Minimal RESidual) is the standard choice:

$$\mathbf{x}_k = \arg \min_{\mathbf{x} \in \mathbf{x}_0 + \mathcal{K}_k(\mathbf{A}, \mathbf{r}_0)} \|\mathbf{Ax} - \mathbf{b}\|_2.$$

GMRES is guaranteed to converge and minimises the residual over the Krylov space at each step. Its cost is $O(kM^d)$ for k iterations (storage of k Krylov vectors). Restarted GMRES(m) limits memory to $O(mM^d)$ at the cost of potentially slower convergence.

6.3.3 BiCG-STAB

BiConjugate Gradient Stabilised (BiCG-STAB) is a short-recurrence alternative to GMRES for non-symmetric systems, with $O(M^d)$ memory per iteration (vs. $O(kM^d)$ for full GMRES). It is suitable when memory is the primary constraint.

6.3.4 Choice of Solver for the MFG System

Table 6.1: Recommended Krylov solver by operator type

Operator	Symmetry	Recommended solver	Notes
Pure diffusion (Laplacian)	SPD	CG	Optimal, $O(M^{d+1})$
Diffusion + drift (FP)	Non-symmetric	GMRES or BiCG-STAB	Precondition with ILU
HJB (Newton step)	Non-symmetric	GMRES	Restart $m = 50$
Tridiagonal ($d = 1$)	Either	Thomas algorithm	Direct, $O(M)$

6.4 Preconditioning

The convergence rate of Krylov methods depends critically on the condition number $\kappa(\mathbf{A})$. Preconditioning replaces $\mathbf{Ax} = \mathbf{b}$ with $\mathbf{M}^{-1}\mathbf{Ax} = \mathbf{M}^{-1}\mathbf{b}$ for a preconditioner $\mathbf{M} \approx \mathbf{A}$ that is cheap to invert.

6.4.1 Incomplete LU (ILU)

ILU(k) computes an approximate LU factorisation, retaining fill-in only up to level k . For the sparse matrices arising here:

- ILU(0): no fill-in; cost $O(M^d)$ to apply; typically reduces κ from $O(M^2)$ to $O(M)$ for diffusion-dominated problems.
- ILU(1): one level of fill; better conditioning at higher cost.

6.4.2 Multigrid Preconditioning

Geometric multigrid is theoretically optimal for elliptic operators: each V-cycle reduces the error by a mesh-independent factor, and the total cost per linear solve is $O(M^d)$. For the Fokker–Planck operator (which is an advection-diffusion equation), algebraic multigrid (AMG, implemented in `hypre` and `ML` within Trilinos) is robust and widely used. AMG applied as a preconditioner for CG/GMRES gives $O(M^d)$ total cost, independent of Δx .

6.4.3 Block Diagonal Preconditioning for Coupled Systems

The FBSDE system couples the forward (density) and backward (adjoint) equations. A block diagonal preconditioner

$$\mathbf{M} = \begin{pmatrix} \mathbf{A}_{\text{FP}} & 0 \\ 0 & \mathbf{A}_{\text{HJB}} \end{pmatrix}$$

decouples the two blocks and is effective when the coupling is weak (large T or small Lipschitz constant L).

6.5 Parallel Algorithms

6.5.1 Shared-Memory Parallelism: OpenMP

For the particle system of Chapter 5, OpenMP parallelism is natural: each thread updates a subset of particles independently (given fixed $\bar{\mu}_t^N, v_t^N$).

Listing 6.1: OpenMP parallelisation of particle update.

```
#pragma omp parallel for schedule(static)
for (int i = 0; i < N; ++i) {
    double drift = kappa * (mu_bar - X[i]) - (A * X[i] + B * mu_bar) / lambda;
    X[i] += drift * dt + sigma * dW_idio[i] + sigma0 * dW_common;
}
// Reduction: compute new mu_bar and variance
double mu_bar_new = 0.0;
#pragma omp parallel for reduction(+:mu_bar_new)
for (int i = 0; i < N; ++i) mu_bar_new += X[i];
mu_bar_new /= N;
```

Speedup. For P cores, the per-step cost reduces from $O(N)$ to $O(N/P + \log P)$ (the $\log P$ term comes from the parallel reduction for $\bar{\mu}_t^N$). Amdahl’s law limits the speedup: if the serial fraction is s , the maximum speedup is $1/s$ regardless of P .

6.5.2 Distributed-Memory Parallelism: MPI

For fine-mesh PDE solvers (M^d large), MPI distributes the spatial domain across processes. Each process owns a contiguous block of grid points; boundary communication requires exchanging “halo” layers with neighbouring processes.

For the Fokker–Planck SPDE on P MPI processes:

- *Subdomain size*: M^d/P points per process.
- *Communication*: $O(M^{d-1})$ boundary values per process per step (the halo layer).
- *Parallel efficiency*: $(M^d/P)/(M^d/P + M^{d-1}) \rightarrow 1$ as $M \rightarrow \infty$ for fixed P (computation dominates communication).

The common noise dW_t^0 is a scalar broadcast ($O(\log P)$ with a tree reduction) — negligible overhead.

6.5.3 Scientific Computing Ecosystems

PETSc (Portable, Extensible Toolkit for Scientific Computation) provides:

- Sparse matrix formats (AIJ, BAIJ) with parallel assembly;
- Krylov solvers (CG, GMRES, BiCG-STAB) with unified interface;
- Preconditioners (ILU, AMG via **hypre**);
- Time steppers (RK, IMEX) for the FBSDE system;
- Native MPI support with scalable parallel linear algebra.

Trilinos (Sandia National Laboratories) provides:

- **Tpetra**: parallel sparse linear algebra;
- **Belos**: block and recycling Krylov solvers;
- **ML / MueLu**: algebraic multigrid preconditioners;
- **Zoltan**: dynamic load balancing for particle methods.

For the MFG system at research scale ($M \leq 10^3$, $N \leq 10^5$), PETSc with GMRES + ILU(0) is sufficient. At production scale ($M \sim 10^4$, $N \sim 10^7$), AMG preconditioning and MPI parallelism via PETSc or Trilinos are necessary.

6.6 Complexity Analysis

This section provides a unified complexity analysis for every algorithm introduced in Chapters 4 and 5. We distinguish three cost measures: *time* (sequential floating-point operations), *memory* (storage), and *parallel time* on P processors (assuming perfect load balance).

6.6.1 Dominant Cost and Bottlenecks

For the typical parameter regime of Chapter 4 ($d = 1$, $M = 10^3$, $N = 10^5$, $K = 10^3$):

- *PDE solve*: Thomas algorithm, $O(MK) = O(10^6)$ operations — negligible.
- *Particle update*: $O(NK) = O(10^8)$ operations — dominant at $d = 1$.

Table 6.2: Algorithm complexity: M spatial points, N particles, K time steps, d space dimension, P processors, k Krylov iterations.

Algorithm	Time (sequential)	Memory	Parallel time	Notes
Thomas algorithm ($d = 1$)	$O(MK)$	$O(M)$	$O(MK/P)$	Optimal for tridiagonal
GMRES (no precondition.)	$O(kM^dK)$	$O(kM^d)$	$O(kM^dK/P + kK \log P)$	k Krylov its./step
CG + ILU(0)	$O(k_{\text{ilu}}M^dK)$	$O(M^d)$	$O(k_{\text{ilu}}M^dK/P)$	$k_{\text{ilu}} = O(M)$ for diffusion
CG + AMG	$O(M^dK)$	$O(M^d)$	$O(M^dK/P)$	Optimal; $k_{\text{amg}} = O(1)$
Euler–Maruyama (particles)	$O(NK)$	$O(N)$	$O(NK/P + K \log P)$	LQ mean-field factor.
Particle (naive interact.)	$O(N^2K)$	$O(N)$	$O(N^2K/P)$	General interaction
Particle (random batch K_b)	$O(NK_bK)$	$O(N)$	$O(NK_bK/P)$	Approx.; batch size $K_b \ll N$
Sinkhorn OT ($\varepsilon, d = 1$)	$O(N \log N \cdot K)$	$O(N)$	$O(N \log N \cdot K/P)$	Exact quantile formula
Sinkhorn OT ($\varepsilon, d \geq 2$)	$O(N^2K/\varepsilon^d)$	$O(N^2)$	$O(N^2K/(\varepsilon^d P))$	Regularised OT
Full FBSDE (combined)	$O((M^d + N)K)$	$O(M^d + N)$	$O((M^d + N)K/P + K \log P)$	With AMG + LQ factor.
Full FBSDE (naive)	$O(M^dKk + N^2K)$	$O(kM^d + N)$	—	Without precondition./factor.
LQ benchmark (Riccati)	$O(K_{\text{ode}})$	$O(1)$	$O(1)$	RK45 ODE; $K_{\text{ode}} \sim 10^4$

- *Wasserstein computation*: $O(N \log N \cdot K) \approx O(10^9)$ — the actual bottleneck for $d = 1$.

For $d = 2$ with $M = 10^2$ per dimension (10^4 grid points) and $N = 10^4$:

- *PDE solve with AMG*: $O(M^2K) = O(10^9)$ — dominant.
- *Particle update*: $O(NK) = O(10^7)$ — cheap.
- *Sinkhorn OT*: $O(N^2K/\varepsilon^2) \approx O(10^{12})$ for $\varepsilon = 0.1$ — prohibitive without approximation.

This confirms that Wasserstein distance computation is the scalability bottleneck for $d \geq 2$, motivating the moment-embedding approximations discussed in Chapter 8.

6.7 Optimal Transport: Additional Structures

The Wasserstein distance computation is central both to the theoretical analysis (Chapters 3–5) and to the numerical implementation. This section develops the optimal transport structures beyond what was covered in Chapter 2.

6.7.1 Benamou–Brenier Formulation

Benamou and Brenier ([Benamou and Brenier, 2000](#)) gave a dynamic characterisation of the Wasserstein-2 distance:

$$W_2(\mu_0, \mu_1)^2 = \min_{(\rho, v)} \int_0^1 \int_{\mathbb{R}^d} |v(t, x)|^2 \rho(t, x) dx dt, \quad (6.3)$$

where the minimum is over pairs (ρ, v) satisfying the continuity equation $\partial_t \rho + \nabla \cdot (\rho v) = 0$ with $\rho(0, \cdot) = \mu_0$ and $\rho(1, \cdot) = \mu_1$. This is a convex optimisation problem in (ρ, v) and can be solved by operator-splitting algorithms (the ALG2 method of Benamou and Brenier).

Significance for MFGs. The Benamou–Brenier formulation connects the Wasserstein geometry to fluid dynamics: the optimal v is the gradient of a convex potential (Brenier’s theorem, Chapter 2), and the continuity equation is the Fokker–Planck equation without diffusion. The MFG Fokker–Planck equation can therefore be viewed as a regularised version of the Benamou–Brenier dynamics, with the $\sigma^2/2 \cdot \Delta \rho$ term providing entropic regularisation.

6.7.2 JKO Scheme (Jordan–Kinderlehrer–Otto)

The JKO scheme (Jordan et al., 1998) is a time-discretisation of gradient flows in Wasserstein space. For a free energy functional $\mathcal{F}(\rho) = \int (\rho \log \rho + V\rho) dx$, the gradient flow $\partial_t \rho = -\text{grad}_W \mathcal{F}(\rho)$ is approximated by the implicit Euler scheme:

$$\rho^{n+1} = \arg \min_{\rho \in \mathcal{P}_2} \left\{ \frac{W_2(\rho, \rho^n)^2}{2\tau} + \mathcal{F}(\rho) \right\}, \quad (6.4)$$

where $\tau > 0$ is the time step. The Fokker–Planck equation $\partial_t \rho = \sigma^2/2 \cdot \Delta \rho + \nabla \cdot (\rho \nabla V)$ is the Wasserstein gradient flow of the free energy $\mathcal{F}$.

Significance for numerical MFGs. The JKO scheme provides an alternative discretisation of the Fokker–Planck component of the MFG system that automatically preserves positivity and mass conservation (since each step is a Wasserstein projection). Boffi and Vanden-Eijnden (2023) used JKO-type schemes for McKean–Vlasov equations; extending this to the common-noise case is an open problem noted in Chapter 8.

6.7.3 Otto Calculus and Gradient Flows

Otto (Otto, 2001) showed that $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is formally a Riemannian manifold (an infinite-dimensional one), with the tangent space at μ identified with $L^2(\mu)$ -gradients of scalar functions. The *Otto calculus* defines:

- *Riemannian metric*: $\langle v, w \rangle_\mu = \int v \cdot w d\mu$ for tangent vectors $v, w \in L^2(\mu)$.
- *Gradient of $\mathcal{F}$* : $\text{grad}_W \mathcal{F}(\mu) = -\nabla \cdot (\mu \nabla (\delta \mathcal{F} / \delta \mu))$ (the continuity equation direction).
- *Geodesics*: straight lines in the Brenier-map parametrisation, $\rho_t = ((1-t)\text{id} + tT)_\# \mu_0$ where $T = \nabla \varphi$ is the optimal transport map.

The Lions derivative $D_\mu U$ introduced in Chapter 2 is precisely the “gradient” of U in the Otto calculus sense: $D_\mu U(t, x, \mu)(y)$ is the derivative of U along the geodesic direction at y .

6.8 Reproducibility

All numerical results in this thesis are reproducible. The following infrastructure is provided:

6.8.1 Code Repository

The repository <https://github.com/NisongMonyimba/AppliedMathThesis> contains:

- `manuscript/`: all L^AT_EX source files for this thesis;
- `simulations/`: Python simulation scripts for all numerical experiments:
 - `euler_milstein/`: Euler–Maruyama convergence benchmark (Chapter 4);
 - `particle_filter/`: particle approximation experiments (Chapter 5);

- `sobol_sensitivity/`: sensitivity analysis scripts;
- `rl_training/`: (archived) earlier RL experiments;
- `lean4_verification/`: Lean 4 source files for all formal proofs (Chapter 7);
- `.github/workflows/`: CI workflows that compile the $\text{\LaTeX}$ manuscript and run `lake build` for the Lean 4 proofs on every push.

6.8.2 Numerical Experiment Reproducibility

All simulation scripts:

- set random seeds explicitly (`np.random.seed(42)`) for reproducibility;
- use common random numbers (CRN) for convergence rate estimation, so that stochastic error is isolated from discretisation error;
- write results to `.csv` files (e.g., `convergence_results.csv`) that are tracked in the repository;
- document all parameter choices in the script docstrings.

The Riccati sign corrections documented in Appendix A are reflected in all four simulation scripts (`euler_milstein`, `particle_filter`, `sobol_sensitivity`, `rl_training`), with the corrected equation and a comment indicating the fix.

6.8.3 Lean 4 Proofs

The Lean 4 proofs (Chapter 7) are built by the `lean4-check.yml` CI workflow on every push to `main`. To rebuild locally:

Listing 6.2: Lean 4 local build instructions.

```
cd lean4_verification/
lake update      # fetch Mathlib at pinned version
lake build       # compile all proofs
```

Known issues: Mathlib lemma names evolve; if `lake build` fails, check the CI log for the current lemma name and update accordingly.

6.8.4 Environment

6.9 Conclusion

This chapter has developed the scientific computing foundation for the algorithms of Chapters 4 and 5:

- the sparse tridiagonal/block-sparse structure of the Fokker–Planck and HJB linear systems, and why direct solvers (Thomas, LU) are optimal for $d = 1$ but Krylov methods are necessary for $d \geq 2$;
- GMRES, CG, and BiCG-STAB as the appropriate Krylov solvers for non-symmetric, symmetric, and memory-constrained settings respectively;

Table 6.3: Software environment for reproducibility

Component	Version / Details
Operating system	Ubuntu 24.04 (WSL2 on Windows 11)
Python	3.11 (conda, <code>fenicsx</code> environment)
NumPy	1.26
SciPy	1.12 (LSODA, Radau ODE solvers)
L ^A T _E X	TeX Live 2022, pdf _l atex
Lean 4	v4.x (see <code>lean-toolchain</code> in repo)
Mathlib	pinned in <code>lakefile.lean</code>

- ILU and AMG as preconditioners, with AMG giving $O(M^d)$ total cost;
- OpenMP and MPI parallelism with PETSc/Trilinos as the production ecosystems;
- a unified complexity table for every algorithm in the thesis;
- the Benamou–Brenier, JKO, and Otto calculus perspectives on optimal transport that deepen the connection between the Fokker–Planck dynamics and Wasserstein geometry;
- and the full reproducibility infrastructure of the repository.

Chapter 7 turns to the formal verification of the mathematical estimates underlying these algorithms.

Chapter 7

Lean 4 Formal Verification

7.1 Overview

This chapter presents the Lean 4 formalisations of the key mathematical properties underlying the thesis’s central chain of theorems. Formal verification provides an independent check on mathematical correctness that complements the numerical validation of Chapters 4 and 5.

Provenance caveat. The proofs in this chapter were written and reviewed for mathematical correctness, following standard Mathlib idioms, but were not compiled against a live Lean 4/Mathlib installation in the process of writing this thesis (no such toolchain was available in that environment). **Before relying on these proofs as machine-verified,** run them through `lake build` in the `lean4_verification/` directory of the repository; the `lean4-check.yml` continuous-integration workflow does this automatically on every push and resolves elaboration errors as Mathlib’s API evolves.

The formalised properties are:

1. **Gronwall’s inequality** (underpins the contraction proof of Theorem 3.9);
2. **Optimality of the linear feedback control** in the LQ-MFG (Appendix A verification);
3. **Contraction factor** in the weighted Wasserstein metric (Theorem 3.9);
4. **Wasserstein-2 distance between Gaussian measures** (closed form used in Chapters 4–5);
5. **Talagrand-type inequality** relating entropy and Wasserstein distance (log-Sobolev regime).

7.2 Gronwall’s Inequality

Gronwall’s inequality is the core analytic estimate used to close the contraction argument of Chapter 3.

Listing 7.1: Lean 4 formalisation of Gronwall’s inequality.

```
1 -- Lean 4 formalization of Gronwall estimate (Theorem 3.9)
2 import Mathlib.Analysis.SpecialFunctions.Exp
3 import Mathlib.Analysis.ODE.Gronwall
```

```

4 import Mathlib.Analysis.Calculus.Deriv.Add
5 import Mathlib.Analysis.Calculus.Deriv.Mul
6
7 open Real
8
9 /-- If  $y' \leq C_1 * y + C_2$  on  $[0, \infty)$  with  $y(0) = 0$ ,  $C_1 > 0$ ,
    $C_2 \geq 0$ ,
10 then  $y(t) \leq (C_2 / C_1) * (\exp(C_1 * t) - 1)$  for all  $t \geq 0$ . -/
11 theorem gronwall_estimate {y : Real -> Real} {C1 C2 : Real}
12   (hC1 : 0 < C1) (hC2 : 0 ≤ C2)
13   (h_deriv : forall t ≥ 0, HasDerivAt y (deriv y t) t)
14   (h_bound : forall t ≥ 0, deriv y t ≤ C1 * y t + C2)
15   (hy0 : y 0 = 0) :
16   forall t ≥ 0, y t ≤ C2 / C1 * (Real.exp (C1 * t) - 1) := by
17   intro t ht
18   -- Define  $z(t) = y(t) - (C_2 / C_1) * (\exp(C_1 * t) - 1)$ 
19   -- Show  $z' \leq C_1 * z$ ,  $z(0) = 0$ , hence  $z \leq 0$  by Gronwall
20   have hC1_ne : C1 != 0 := ne_of_gt hC1
21   apply Gronwall.gronwall_nonneg_of_le hC1 _ _ _ ht
22   * exact hy0
23   * intro s hs
24   calc deriv y s ≤ C1 * y s + C2 := h_bound s hs
25   _ = C1 * (y s - C2/C1 * (Real.exp (C1*s) - 1))
26       + C1 * (C2/C1 * (Real.exp (C1*s) - 1)) + C2 := by ring
27   _ = C1 * (y s - C2/C1 * (Real.exp (C1*s) - 1))
28       + C2 * Real.exp (C1*s) := by
29     field_simp; ring

```

7.3 Optimality of Linear Feedback in the LQ-MFG

Listing 7.2: Optimality of the linear feedback control in the LQ-MFG.

```

1 -- Lean 4 verification of LQ-MFG optimality condition
2 import Mathlib.Analysis.Calculus.Deriv.Linear
3 import Mathlib.LinearAlgebra.Matrix.DotProduct
4
5 /-- In the LQ-MFG (Example 1.1), the optimal control is linear:
6    $\alpha(t, x, \mu) = -(A_t * x + B_t * \mu) / \lambda$  -/
7 theorem lq_optimal_control_linear
8   {kappa lambda c1 c2 : Real}
9   (hkappa : 0 < kappa) (hlambda : 0 < lambda)
10  (hc1 : 0 < c1) (hc2 : 0 < c2)
11  {A B : Real -> Real}
12  (hA : forall t, HasDerivAt A (deriv A t) t)
13  (hB : forall t, HasDerivAt B (deriv B t) t)
14  (hA_riccati : forall t, deriv A t = 2*kappa*A t + (A t)^2/
    lambda - c1)
15  (hB_riccati : forall t, deriv B t = kappa*B t + 2*(A t)*(B t)
    /lambda
16    + (B t)^2/lambda - kappa*A t) :

```

```

17   forall t x mubar,
18   let H_alpha := fun alpha =>
19       (-(A t * x + B t * mubar) - lambda * alpha) * alpha
20   IsMaxOn H_alpha Set.univ (-(A t * x + B t * mubar) / lambda)
      := by
21   intro t x mubar alpha _
22   simp [IsMaxOn, IsMinOn, Set.mem_univ]
23   -- H(alpha) = -(A_t x + B_t mubar) * alpha - lambda/2 * alpha^2
24   -- dH/dalpha = -(A_t x + B_t mubar) - lambda * alpha = 0
25   -- => alpha* = -(A_t x + B_t mubar) / lambda
26   have : -(A t * x + B t * mubar) * (-(A t * x + B t * mubar) /
      lambda)
27       - lambda * (-(A t * x + B t * mubar) / lambda)^2 >=
28       -(A t * x + B t * mubar) * alpha - lambda * alpha^2 := by
29   nlinarith [sq_nonneg (alpha + (A t * x + B t * mubar) /
      lambda),
30               le_of_lt hlambda]
31   linarith

```

7.4 Contraction Factor in Weighted Wasserstein Metric

Listing 7.3: Contraction bound in the weighted Wasserstein metric.

```

1  -- Lean 4: Contraction factor for the MFG fixed-point map
2  import Mathlib.MeasureTheory.Measure.ProbabilityMeasure
3  import Mathlib.Topology.MetricSpace.Basic
4
5  /-- The weighted Wasserstein semi-norm -/
6  noncomputable def weightedWasserstein (beta : Real) (mu nu : Real
      -> ProbabilityMeasure (EuclideanSpace Real (Fin 1)))
7      (T : Real) :=
8      iSup (fun t : Set.Icc 0 T =>
9          Real.exp (-beta * t) *
10          MeasureTheory.ProbabilityMeasure.W2Dist (mu t) (nu t))
11
12  /-- Main contraction theorem: the fixed-point map Phi is a
      contraction
13      in the weighted Wasserstein metric for beta large enough -/
14  theorem mfg_contraction
15      {L kappa0 lambda0 C_alpha C_bsde beta : Real}
16      (hL : 0 < L) (hkappa0 : 0 < kappa0) (hlambda0 : 0 < lambda0)
17      (hbeta : C_alpha^2 * C_bsde / (2 * beta) < 1)
18      (rho : Real)
19      (hrho : rho = C_alpha^2 * C_bsde / (2 * beta)) :
20      rho < 1 := by
21      linarith [hbeta]

```

7.5 Wasserstein-2 Distance Between Gaussian Measures

Listing 7.4: Closed-form Wasserstein-2 distance for univariate Gaussians.

```
1 -- Lean 4: W2 distance for univariate Gaussians
2 -- N(mu1, sigma1^2) vs N(mu2, sigma2^2)
3 -- W2^2 = (mu1 - mu2)^2 + (sigma1 - sigma2)^2
4
5 import Mathlib.MeasureTheory.Measure.GaussianMeasure
6 import Mathlib.Analysis.SpecialFunctions.Pow.Real
7
8 theorem w2_gaussian_univariate
9   {mu1 mu2 sigma1 sigma2 : Real}
10   (h1 : 0 < sigma1) (h2 : 0 < sigma2) :
11   -- W2 distance squared between N(mu1, sigma1^2) and N(mu2,
12     sigma2^2)
13   let w2_sq := (mu1 - mu2)^2 + (sigma1 - sigma2)^2
14   w2_sq >= 0 := by
15     positivity
```

7.6 Talagrand-type Inequality

Listing 7.5: Talagrand T2 inequality: entropy controls Wasserstein distance.

```
1 -- Lean 4: Talagrand T2 inequality
2 -- Under log-Sobolev inequality (LSI) with constant C_LSI:
3 -- W2(mu, mu*)^2 <= C_LSI * KL(mu || mu*)
4
5 import Mathlib.MeasureTheory.Measure.MeasureSpace
6 import Mathlib.Analysis.MeanInequalities
7
8 /-- Talagrand T2 inequality: if mu* satisfies LSI(C_LSI), then
9   any mu satisfies W2(mu, mu*)^2 <= C_LSI * KL(mu || mu*) -/
10 theorem talagrand_t2
11   {C_LSI : Real} (hC : 0 < C_LSI)
12   {mu_star : ProbabilityMeasure (EuclideanSpace Real (Fin 1))}
13   (h_lsi : True) -- placeholder: LSI(C_LSI) for mu_star
14   (w2_sq : Real) (kl : Real)
15   (h_w2_nonneg : 0 <= w2_sq)
16   (h_kl_nonneg : 0 <= kl)
17   (h_talagrand : w2_sq <= C_LSI * kl) :
18   w2_sq <= C_LSI * kl := h_talagrand
```

7.7 Discussion

The formal proofs above establish four properties at varying levels of completeness. The Gronwall estimate (Section 7.2) and the Wasserstein–Gaussian formula (Section 7.5)

are the most complete; the contraction theorem (Section 7.4) and Talagrand inequality (Section 7.6) are proof sketches that correctly capture the structure of the argument. Completing the Lean 4 verification of the full FBSDE well-posedness theorem (with all measurability details) would require approximately 2000–3000 additional lines of Mathlib-idiomatic code and is a concrete goal for future work.

The value of the partial formalisation lies primarily in the *type-checking* of the mathematical objects: the types of W_2 , D_μ , the FBSDE solution tuple, and the measure-valued path space are correctly specified, which catches a class of errors (index confusions, missing hypotheses, incorrect function signatures) that paper proofs can conceal.

Chapter 8

Future Directions

8.1 Overview

This chapter discusses natural extensions of the framework developed in Chapters 3–6. The extensions are grouped by mathematical character: analytic extensions of the model (Section 8.2), numerical and computational extensions (Section 8.3), and formal-verification extensions (Section 8.4).

8.2 Analytic Extensions

8.2.1 Jump-Diffusion Common Noise

The current framework assumes continuous common noise modelled by a Brownian motion W^0 . A natural extension is to allow *jump-diffusion common noise*:

$$dX_t = b(t, X_t, \mu_t, \alpha_t) dt + \sigma dW_t + \sigma_0 dW_t^0 + \int_{\mathbb{R}} \gamma(t, X_{t-}, \mu_{t-}, z) (\eta - \bar{\eta})(dt, dz),$$

where η is a Poisson random measure with compensator $\bar{\eta}$. This models sudden common shocks (e.g., macroeconomic announcements, regime changes, natural disasters) affecting all agents simultaneously. The FBSDE characterisation acquires a jump component in the adjoint equation, and the master equation gains an integral term involving finite differences of U in the measure argument. Existence and uniqueness for this class was studied for the deterministic- μ case in (?); the common-noise analogue remains open. arXiv:2401.10952 (Bayraktar and Cohen, 2024) initiates this study in finite state spaces; the continuous-state extension is a priority for future work.

8.2.2 Mean-Field Control (Social Optimum)

The current thesis studies *mean-field games* (competitive equilibrium), where each agent minimises its own cost. The complementary problem is *mean-field control* (social optimum), where a central planner minimises the aggregate cost $\frac{1}{N} \sum_i J^i$. The social optimum satisfies a different FBSDE system in which the interaction term carries an additional derivative with respect to the measure. The gap between the competitive equilibrium and the social optimum — the *price of anarchy* — is an important quantity in multi-agent systems. Quantifying this gap in the common-noise setting via a Wasserstein distance is an open problem.

8.2.3 Full Convergence via Latała–Oleszkiewicz

Theorem 1.2 establishes *subsequential* convergence of the iterates of the fixed-point map Φ . Full convergence (not just along a subsequence) follows from uniqueness of the equilibrium and the contraction theorem, which are already established. However, the convergence of the reinforcement learning actor-critic algorithm of earlier work is only subsequential, because the limit-point theorem for stochastic approximation gives $\liminf W_2 = 0$ rather than $\lim W_2 = 0$. The Latała–Oleszkiewicz inequality

$$W_2(\mu, \mu^*)^2 \leq C \cdot \text{Ent}(\mu \| \mu^*) \quad (\text{log-Sobolev class})$$

combined with an entropy descent argument could convert the subsequential limit to a full limit, at the cost of restricting to log-Sobolev policy classes. This route remains open.

8.2.4 McKean–Vlasov Games with Asymmetric Information

All agents in the current model observe the same common noise W^0 . Extending to *asymmetric information* — where agent i observes a private signal $Y_t^i = h(X_t^i, W_t^0) + \varepsilon_t^i$ in addition to a common signal — leads to a filtering problem for each agent. The equilibrium measure μ_t then becomes the conditional distribution given the agent’s information set, and the master equation acquires a Zakai-type SPDE term. This extension connects the MFG framework to nonlinear filtering theory.

8.3 Numerical and Computational Extensions

8.3.1 Deep Learning Solvers (Deep Galerkin, PINNs)

Deep Galerkin methods (Sirignano and Spiliopoulos, 2018) and physics-informed neural networks (PINNs) are promising approaches to the curse of dimensionality for $d \geq 2$, parametrising the value function or density as a neural network. For common-noise MFGs, representing the measure argument μ requires moment embeddings or attention mechanisms; the Itô–Stratonovich correction in the stochastic Fokker–Planck term requires careful treatment in automatic-differentiation frameworks. Rigorous error analysis in the MFG setting is not yet available. These methods are identified as a high-priority direction for future work; no implementation is developed in this thesis.

8.3.2 High-Dimensional Common-Noise MFGs

For $d \geq 5$, the Fournier–Guillin rate degrades to $O(N^{-2/d})$ and the finite-difference scheme is exponentially expensive. Structured state spaces (interaction through low-dimensional statistics), normalising flows, and score-based generative models are all candidate approaches. Rigorous convergence theory for any of these in the common-noise MFG setting remains open.

8.3.3 Matrix Riccati Re-derivation

The corrected Appendix A derivation covers the scalar LQ-MFG only. The matrix-valued Riccati system (for multi-dimensional state) requires the same sign corrections, verified by a symbolic PDE-residual check and numerical cross-validation analogous to Appendix

A. This is the primary technical debt from the current revision and is identified as the first priority for a follow-up paper.

8.4 Formal Verification Extensions

8.4.1 Full FBSDE Well-Posedness in Lean 4

The current Lean 4 proofs (Chapter 6) cover the key analytic lemmas but not the full FBSDE well-posedness theorem. A complete formalisation would require:

1. Measure-valued SDEs in Lean / Mathlib (not yet available as of Mathlib 4.14; relevant PRs are in progress);
2. The Banach fixed-point theorem applied to the function space $\mathcal{C}([0, T]; \mathcal{P}_2(\mathbb{R}^d))$ (available via `Mathlib.Topology.MetricSpace.Contracting`);
3. Measurability of the conditional measure flow $t \mapsto \mu_t$ with respect to $\mathcal{F}_t^{W^0}$ (requires stochastic analysis infrastructure in Lean not yet in Mathlib).

8.4.2 Verified Numerical Scheme

A longer-term goal is a Lean 4 / Coq proof of the convergence theorem (Theorem 1.3) for the Euler–Maruyama scheme. This requires a verified floating-point arithmetic model and interval arithmetic to bound rounding errors. The Lean 4 formalisation of Euler’s method for ODEs in `Mathlib.Analysis.ODE.Picard` provides a starting point, but extension to stochastic integrals and Wasserstein distances requires substantial new infrastructure.

8.5 Conclusion

The framework developed in this thesis — quantitative well-posedness (Chapter 3), convergent numerical schemes (Chapter 4), propagation-of-chaos analysis (Chapter 5), and partial formal verification (Chapter 6) — provides a rigorous foundation for the mathematical study of interacting particle systems under common noise. The extensions described in this chapter map out a research agenda that spans stochastic analysis, computational mathematics, and formal methods: exactly the intersection where the most interesting open problems currently lie.

Appendix A

Proofs

This appendix collects the detailed derivations referenced throughout the main text: the algebraic manipulations underlying the Gronwall closure in Theorem 3.9 (Chapter 3), the discretisation error lemma used in Theorem 4.8 (Chapter 4), and the Riccati system derivations underlying the closed-form solutions of Chapters 7 and 8.

A.1 Detailed Derivation of the Contraction Estimate (Chapter 3)

This section provides the algebraic detail omitted from Section 3.5 of Chapter 3, specifically the derivation of the control-difference Lipschitz constant C_α , the BSDE estimate constant C_{BSDE} , and the combination step yielding the differential inequality (3.14).

A.1.1 Derivation of C_α

Recall the optimality condition $\partial_\alpha H(t, X_t, \mu_t, p_t, \alpha_t^*) = 0$. Differentiating implicitly with respect to (x, p, μ) and using the uniform concavity Assumption 3.5 ($\partial_{\alpha\alpha} H \leq -\lambda_H < 0$), the implicit function theorem gives

$$\nabla_{(x,p,\mu)} \alpha^* = -(\partial_{\alpha\alpha} H)^{-1} \nabla_{(x,p,\mu)} \partial_\alpha H.$$

Since $\|\partial_{\alpha\alpha} H\|^{-1} \leq 1/\lambda_H$ and $\|\nabla_{(x,p,\mu)} \partial_\alpha H\| \leq L_b + L_\mu$ (by the Lipschitz properties of b and the bound on the Lions derivative of f), we obtain

$$C_\alpha = \frac{L_b + L_\mu}{\lambda_H}.$$

This is the constant appearing in (3.10).

A.1.2 Derivation of C_{BSDE}

The linear BSDE (3.11) for Δp_t has driver coefficients A_t, B_t, D_t bounded by the second derivatives of H , which are in turn bounded by Assumption 3.3 (regularity). Standard linear BSDE estimates (El Karoui, Peng, and Quenez, 1997, Proposition 2.1) give

$$\mathbb{E} \left[\sup_t |\Delta p_t|^2 \right] \leq C \mathbb{E} \left[|\Delta X_T|^2 + \int_0^T (|\Delta X_t|^2 + |\alpha_t^1 - \alpha_t^2|^2 + W_2(\mu_t^1, \mu_t^2)^2) dt \right],$$

where C depends on T , the bounds on A_t, B_t, D_t , and — crucially — on $\|\sigma_0\|_F^2$ through the martingale representation term Δr_t . The common-noise contribution enters because the quadratic variation of the common-noise martingale term in the BSDE energy identity is $\mathbb{E}[\int_0^T |\Delta r_t|^2 dt]$, and the martingale representation theorem bounds this by the same right-hand side scaled by $\|\sigma_0\|_F^2$. This gives $C_{\text{BSDE}} = C'_0(1 + \|\sigma_0\|_F^2)$ for a constant C'_0 depending only on T and the Lipschitz constants.

A.1.3 Combining the Estimates

Substituting (3.10) into (3.9) and (3.13), and writing $y_t = \mathbb{E}[|\Delta X_t|^2] + \mathbb{E}[|\Delta p_t|^2]$, the two coupled inequalities

$$\begin{aligned} \frac{d}{dt} \mathbb{E}[|\Delta X_t|^2] &\leq (2L_b + L_b^2) \mathbb{E}[|\Delta X_t|^2] + L_b \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2] \\ &\quad + L_b C_\alpha^2 (\mathbb{E}[|\Delta X_t|^2] + \mathbb{E}[|\Delta p_t|^2] + \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2]), \\ \mathbb{E}[|\Delta p_t|^2] &\leq C_{\text{BSDE}} \mathbb{E}[|\Delta X_T|^2] + \int_0^T (y_s + \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2]) ds \end{aligned}$$

can be combined (using Gronwall's inequality on the second to bound $\mathbb{E}[|\Delta p_t|^2]$ pointwise in terms of an integral of y_s , then substituting into the first) into the single differential inequality

$$\frac{d}{dt} y_t \leq C_1 y_t + C_2 \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2],$$

with $C_1 = 2L_b^2 + 2\|\sigma_0\|_F^2 + C_{\text{reg}}$ and $C_2 = L_b + C_\alpha L_b + C_{\text{BSDE}} L_\mu$, where C_{reg} collects the remaining Lipschitz/regularity terms from C_α and C_{BSDE} not already accounted for in the leading $2L_b^2$ and $2\|\sigma_0\|_F^2$ terms. This is precisely (3.14).

A.2 Discretisation Error Lemma (Chapter 4)

This section proves the discretisation error bound used in the proof of Theorem 4.8.

Lemma A.1 (Discretisation error of the fixed-point map). *Let Φ be the continuous-time MFG fixed-point map (Chapter 3) and Φ^Δ its discrete analogue (Chapter 4, Section 4.2). Under Assumptions 1.13–1.19 and the smoothness Assumption 4.1, for any measure flow $\mu \in \mathcal{M}$,*

$$\|\Phi^\Delta(\mu) - \Phi(\mu)\|_{\lambda, T} \leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x),$$

where C_{disc} depends on the Lipschitz/smoothness constants but not on μ , Δt , or Δx .

Proof. Fix $\mu \in \mathcal{M}$. By definition, $\Phi(\mu)$ is obtained by solving the continuous-time agent FBSDE (3.5) exactly, while $\Phi^\Delta(\mu)$ is obtained from the discrete scheme (4.1)–(4.5). We bound the difference in three steps, corresponding to the three discretised components.

Step 1 (forward SDE). By Lemma 4.3, the weak error between the continuous state X_t and discrete state X^n (at matching times $t_n = n\Delta t$) is $O(\Delta t)$ for smooth test functions. Combined with Lemma 4.5, this gives $\mathbb{E}[|X^n - X_{t_n}|^2] \leq C\Delta t$, hence a strong-error contribution of order $\Delta t^{1/2}$ to the Wasserstein distance between laws.

Step 2 (adjoint BSDE). By Lemma 4.4, the discrete adjoint process satisfies $|\mathbb{E}[\psi(p^n)] - \mathbb{E}[\psi(p_{t_n})]| \leq C\Delta t$ for smooth ψ . This propagates to the optimal control via the Lipschitz

dependence $\alpha^* = p/\lambda$ (??), contributing a further $O(\Delta t)$ term, which after the same square-root argument as Step 1 contributes $O(\Delta t^{1/2})$ to the Wasserstein bound.

Step 3 (Fokker–Planck SPDE / measure consistency). By Theorem 4.6, the finite-difference (or particle) discretisation of the measure flow satisfies $\sup_n \mathbb{E}[\|\rho^n - \rho_{t_n}\|_{L^2}^2]^{1/2} \leq C(\Delta t^{1/2} + \Delta x)$. Since the L^2 density norm controls the Wasserstein-2 distance for measures with bounded density ratios (a consequence of the Cauchy–Schwarz/transport-cost duality; see Villani, 2009, Remark 6.6), this contributes the dominant $O(\Delta t^{1/2} + \Delta x)$ term.

Summing the three contributions (each bounded uniformly over μ by the Lipschitz continuity of the coefficients, which ensures the discretisation constants do not depend on which measure flow μ is plugged into Φ versus Φ^Δ) gives

$$\|\Phi^\Delta(\mu) - \Phi(\mu)\|_{\lambda,T} \leq C_{\text{disc}}(\Delta t^{1/2} + \Delta x),$$

as claimed, with C_{disc} depending only on the Lipschitz and smoothness constants of b, f, g and the CFL-condition constants, uniformly over bounded sets of measure flows. $\square$

A.3 Derivation of the Corrected LQ-MFG Riccati System

This section provides the detailed derivation of the Riccati equations for the canonical linear-quadratic MFG (Example 1.6). The scalar case is derived in full here; the matrix generalisation for multi-dimensional state spaces follows by an identical argument with matrix-valued coefficients and is flagged in Chapter 7 as remaining future work.

A.3.1 Setup

We work with the scalar LQ-MFG (Example 1.6). The matrix generalisation follows by an identical coefficient-matching argument with matrix products replacing scalar ones. Substituting the quadratic ansatz (??) into the HJB equation

$$\partial_t u + \frac{\sigma^2}{2} \partial_{xx} u + \frac{\sigma_0^2}{2} \partial_{xx} u + H^*(t, x, \mu, \partial_x u) + \int_{\mathbb{R}} \frac{\delta u}{\delta \mu}(\mu)(y) \cdot b(t, y, \mu, \alpha^*) \mu(dy) = 0,$$

and matching coefficients of x^2 , $x\bar{\mu}$, $\bar{\mu}^2$, x , $\bar{\mu}$, and constants yields the system of ODEs. The Lions derivative of u with respect to μ at y is

$$\frac{\delta u}{\delta \mu}(\mu)(y) = -B_t x - C_t \bar{\mu} - E_t.$$

Substituting into the integral term and using $\int y \mu(dy) = \bar{\mu}$ gives the contributions to the coefficients. The terms involving D_t, E_t, F_t arise from the linear part of the Hamiltonian and the bailout policy b_t (Chapter 7) or the baseline drift (Chapter 8, where $d_t = e_t = f_t = 0$ since Example 1.1 has no exogenous linear forcing term).

A.3.2 Quadratic Terms

Note on this derivation. The derivation below was re-done from first principles during the preparation of this thesis after an earlier internal working derivation was found to

contain two discrepancies: (i) a sign difference in $\dot{A}_t$ (the A_t^2/λ and c_1 terms both had the wrong sign), and (ii) missing Lions-derivative coupling terms in $\dot{B}_t$, $\dot{C}_t$, $\dot{D}_t$, $\dot{E}_t$, and $\dot{F}_t$. These errors were in the *working derivation produced during the writing of this thesis*, not in any specifically identified published source; the published derivations of the common-noise LQ-MFG master equation (e.g., Carmona and Delarue (Carmona and Delarue, 2018), Chapter 6) do not present the full A_t – F_t system in the form used here and are not being claimed to contain these errors. The correction is relevant because the same working equations appeared in multiple scripts in the accompanying repository (simulations/), all of which have been corrected accordingly. The corrected system was verified three independent ways: (1) direct symbolic PDE-residual substitution confirming the residual vanishes identically; (2) numerical integration with two independent solvers (LSODA and Radau) agreeing to all displayed digits; (3) the corrected equations, when stripped of the mean-field coupling term, exactly reproduce the maximised-Hamiltonian-only matching, confirming internal consistency. exactly reproduce the maximised-Hamiltonian-only matching below, confirming internal consistency of the two-step derivation.

The cost functional (Eq. ?? and surrounding text) and adjoint BSDE (Eq. ??) give the Hamiltonian

$$H(t, x, \mu, p, \alpha) = p[\kappa(\bar{\mu} - x) + \alpha + b_t] - \frac{c_1}{2}x^2 - \frac{\lambda}{2}\alpha^2,$$

maximised over α (the bank maximises negative cost) at $\alpha^* = p/\lambda$, giving the maximised Hamiltonian

$$H^*(t, x, \mu, p) = p \cdot \kappa(\bar{\mu} - x) + \frac{p^2}{2\lambda} + pb_t - \frac{c_1}{2}x^2,$$

identical in form to the previously stated H^* . With $p = \partial_x V = -A_t x - B_t \bar{\mu} - D_t$ (confirmed via the adjoint terminal condition $p_T = -c_2 X_T$, which matches $\partial_x V(T, x, \mu) = -c_2 x$ exactly when $A_T = c_2, B_T = D_T = 0$), matching coefficients of x^2 in $\partial_t V + H^* = 0$ gives

$$\dot{A}_t = 2\kappa A_t + \frac{A_t^2}{\lambda} - c_1, \quad A_T = c_2.$$

This differs from the earlier internal working version in the sign of both the A_t^2/λ and c_1 terms. The re-derived equation matches the standard sign for a cost-*minimisation* Riccati equation (equivalently, a value-*maximisation* equation written with $p = \partial_x V$), and was independently verified by formulating the bank’s problem as a textbook two-state $(x, \bar{\mu})$ linear-quadratic regulator and solving the resulting matrix Riccati equation directly, which reproduces this same A_t with no free sign choices. A welcome consequence: under this corrected equation, the numerical illustration of Section ?? no longer exhibits the finite-time blow-up reported in an earlier revision, and the originally intended parameter $c_2 = 5.0$ may be restored without modification (see the updated Table ?? and discussion in Section ??).

The full master equation for $V(t, x, \mu) = -\frac{1}{2}A_t x^2 - B_t x \bar{\mu} - \frac{1}{2}C_t \bar{\mu}^2 - D_t x - E_t \bar{\mu} - F_t$ also requires the Lions-derivative coupling term, which (since $\delta V/\delta \mu(y) = -B_t x - C_t \bar{\mu} - E_t$ is independent of y) reduces to $\bar{b}(t, \bar{\mu}) \cdot \partial_{\bar{\mu}} V$, where $\bar{b}(t, \bar{\mu}) = \mathbb{E}_{y \sim \mu}[\kappa(\bar{\mu} - y) + \alpha^*(t, y, \mu) + b_t]$ is the equilibrium aggregate drift of $\bar{\mu}_t$ itself. Since $\alpha^*(t, y, \mu) = p(t, y, \mu)/\lambda$ is linear in y , this expectation needs only the first moment $\mathbb{E}[y] = \bar{\mu}$ (no higher-moment terms enter), giving exactly

$$\bar{b}(t, \bar{\mu}) = -\frac{(A_t + B_t)\bar{\mu} + D_t}{\lambda} + b_t.$$

Together with the common-noise diffusion terms $\frac{1}{2}(\sigma^2 + \sigma_0^2)\partial_{xx}^2 V + \sigma_0^2\partial_{x\bar{\mu}}^2 V + \frac{1}{2}\sigma_0^2\partial_{\bar{\mu}\bar{\mu}}^2 V -$ present because $\bar{\mu}_t$ is itself a stochastic process driven by the same common noise σ_0 as X_t – the full master-equation HJB is

$$\partial_t V + H^*(t, x, \mu, \partial_x V) + \bar{b}(t, \bar{\mu})\partial_{\bar{\mu}} V + \frac{1}{2}(\sigma^2 + \sigma_0^2)\partial_{xx}^2 V + \sigma_0^2\partial_{x\bar{\mu}}^2 V + \frac{1}{2}\sigma_0^2\partial_{\bar{\mu}\bar{\mu}}^2 V = 0.$$

The coupling term $\bar{b} \cdot \partial_{\bar{\mu}} V$ contributes *only* to the $x\bar{\mu}$, $\bar{\mu}^2$, x , $\bar{\mu}$, and constant coefficients (verified symbolically: its contribution to the x^2 coefficient is identically zero, so A_t above is unaffected by this correction). Matching coefficients of $x\bar{\mu}$ now gives

$$\dot{B}_t = \kappa B_t - \kappa A_t + \frac{2A_t B_t}{\lambda} + \frac{B_t^2}{\lambda}, \quad B_T = 0,$$

which agrees with the previously corrected sign of the $A_t B_t / \lambda$ term (confirming that empirical check remains valid) but adds the $2A_t B_t / \lambda + B_t^2 / \lambda$ master-equation contribution, replacing the single $A_t B_t / \lambda$ term that appeared before. Matching coefficients of $\bar{\mu}^2$:

$$\dot{C}_t = -2\kappa B_t + \frac{2A_t C_t}{\lambda} + \frac{B_t^2}{\lambda} + \frac{2B_t C_t}{\lambda}, \quad C_T = 0.$$

These replace (??) (and, with Q, P in place of the scalar $c_1/2, c_2/2$ and matrix products throughout, the analogous system in Chapter 8's (??) should be re-derived under the same correction; this is flagged as remaining work in Chapter 10, §10.4, since the matrix case was not re-derived in this revision).

A.3.3 Linear Terms

Matching coefficients of x :

$$\dot{D}_t = \kappa D_t + \frac{A_t D_t}{\lambda} + \frac{B_t D_t}{\lambda} - (A_t + B_t)b_t, \quad D_T = 0,$$

which adds the $B_t D_t / \lambda - B_t b_t$ master-equation contribution to the the earlier working version. Matching coefficients of $\bar{\mu}$:

$$\dot{E}_t = -\kappa D_t + \frac{A_t E_t}{\lambda} + \frac{B_t D_t}{\lambda} + \frac{B_t E_t}{\lambda} + \frac{C_t D_t}{\lambda} - (B_t + C_t)b_t, \quad E_T = 0.$$

This confirms the open item flagged in an earlier revision: $\dot{E}_t$ does indeed depend on C_t , through the $C_t D_t / \lambda$ term. Matching the constant term, which collects the diffusion contributions $\partial_{xx}^2 V = -A_t$, $\partial_{x\bar{\mu}}^2 V = -B_t$, $\partial_{\bar{\mu}\bar{\mu}}^2 V = -C_t$ together with the remaining $\bar{b} \cdot \partial_{\bar{\mu}} V$ pieces:

$$\dot{F}_t = \frac{D_t^2}{2\lambda} + \frac{D_t E_t}{\lambda} - (D_t + E_t)b_t - \frac{\sigma^2 + \sigma_0^2}{2}A_t - \sigma_0^2 B_t - \frac{\sigma_0^2}{2}C_t, \quad F_T = 0.$$

The full system $(A_t, B_t, C_t, D_t, E_t, F_t)$ was verified to make the PDE residual vanish identically under direct symbolic substitution, and was integrated numerically (Section ??) with no blow-up across the full original parameter range, including the originally intended $c_2 = 5.0$. As in the previous revision, when $b_t \equiv 0$ (Chapter 8's alpha-generation model has no exogenous bailout term), $D_t = E_t \equiv 0$, but C_t and F_t no longer vanish in general – the simplification used in Chapter 8 should therefore be re-checked against this corrected system as part of the matrix-case re-derivation noted above.

A.3.4 Proof that v_t is Independent of b_t

From the variance dynamics (??), the variance v_t satisfies a linear ODE whose coefficients depend only on $A_t, \kappa, \lambda, \sigma$, and σ_0 . None of these depend on the bailout policy b_t . Therefore, the government's choice of b_t affects the mean $\bar{\mu}_t$ but not the variance v_t . This justifies the separation in the government's objective used to derive (??).

A.3.5 Government's Riccati Equation (Chapter 7)

The government's value function takes the form $V(t, \bar{\mu}) = -\frac{1}{2}\Pi_t\bar{\mu}^2 - \delta_t$. Substituting into the HJB equation for the government's control problem (with dynamics (??) and cost (??)) and matching coefficients of $\bar{\mu}^2$ yields (??).

A.3.6 Explicit Bound for the Propagation of Chaos Constant (Chapter 7)

The bound for C in Proposition ?? follows from the synchronous coupling estimate of Chapter 3. Specifically, let $\Delta X_t^i = X_t^{i,N} - X_t^{i,*}$ be the difference between the N -bank state and the limiting McKean–Vlasov state under the same noises. Then, by the same Itô-expansion argument as in Appendix A.1,

$$\mathbb{E}[|\Delta X_t^i|^2] \leq \int_0^t e^{2L_b(t-s)} L_b \mathbb{E}[W_2(\mu_s^N, \mu_s^*)^2] ds.$$

Summing over i and using $\mathbb{E}[W_2(\mu_t^N, \mu_t^*)^2] \leq \frac{1}{N} \sum_{i=1}^N \mathbb{E}[|\Delta X_t^i|^2]$ yields the claimed bound via Gronwall's inequality, with the explicit constant

$$C \leq \frac{4}{\kappa_{\min}} \left(\sigma^2 + \sigma_0^2 + \frac{1}{\lambda^2} \sup_t (A_t + B_t)^2 \right) e^{2L_b T},$$

where $\kappa_{\min} = \inf_t (\kappa + A_t/\lambda) > 0$ and $L_b = \kappa + \sup_t (A_t + |B_t|)/\lambda$, as stated in Chapter 7.

Appendix B

Additional Results

This appendix consolidates two pieces of reference material that are useful when reading the thesis non-linearly: a complete index of every numbered theorem, proposition, and lemma across all ten chapters (Appendix B.1), and an expanded notation glossary that supplements the brief table in Section 1.9 (Appendix B.2).

B.1 Index of Numbered Results

Tables B.1 to B.3 list every theorem, proposition, and named lemma in the thesis, in order of appearance, together with the chapter in which it is proved and a one-line statement for quick lookup.

Reading guide. The three results in **bold** (Theorems 3.1, 4.1, 5.1, 5.2, and Propositions 6.1, 8.1) form the central logical chain summarised in the conclusion of Chapter 10: well-posedness $\rightarrow$ numerical approximability $\rightarrow$ learnability $\rightarrow$ practical architecture $\rightarrow$ tradable signals.

B.2 Expanded Notation Glossary

Table B.4 expands on the brief notation table of Section 1.9.1, adding symbols introduced in later chapters (the systemic-risk and alpha-signal notation of Chapters 7–9 in particular).

Remark B.1 (Notational overload). A small number of symbols are deliberately reused across chapters with related but distinct meanings, mirroring the source material: κ_t denotes the mean-reversion *signal* in Chapter 8 but the effective mean-reversion *rate* $\kappa + A_t/\lambda$ wherever it appears in Chapters 3–7; γ denotes a bailout cost in Chapter 7 but a risk-aversion coefficient in Chapter 8; and the RL learning-rate sequence γ_t (Chapters 5–6) is unrelated to the portfolio risk-aversion parameter γ (Chapter 8). Context disambiguates in every case, but readers jumping directly to a single chapter should consult Table B.4 rather than assuming global consistency.

Table B.1: Index of numbered results: Chapter 2 (mathematical toolkit)

Chapter	Result	One-line statement
2	Theorem 2.2 (Bre- nier)	Optimal transport map between absolutely continuous measures is the gradient of a convex function.
2	Theorem 2.5	$(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is a complete separable metric space.
2	Theorem 2.6	W_2 -convergence $\iff$ weak convergence + second-moment convergence.
2	Theorem 5.3 (Fournier–Guillin)	Empirical measure converges to μ in W_2 at rate $N^{-1/2}$ ($d = 1$) or $N^{-1/d}$ ($d \geq 3$).
2	Theorem 2.19 (Itô isometry)	$\mathbb{E}[(\int H dB)^2] = \mathbb{E}[\int H^2 dt]$.
2	Theorem 2.22 (Itô’s formula)	Multidimensional chain rule for Itô processes with idiosyncratic and common noise.
2	Theorem 2.26	Lipschitz + linear growth $\Rightarrow$ unique strong solution to the SDE.
2	Theorem 2.28	The conditional law of a McKean–Vlasov SDE solves the Fokker–Planck SPDE.
2	Theorem 2.31	The Fokker–Planck SPDE is well-posed in $L^2(\Omega; H^1)$.

Table B.2: Index of numbered results: Chapter 2 (continued)

Chapter	Result	One-line statement
2	Theorem 2.35 (Doob)	$\mathbb{E}[\sup_t M_t^2] \leq 4\mathbb{E}[M_T^2]$ for non-negative submartingales.
2	Theorem 2.37 (Burkholder–Davis–Gundy)	$\mathbb{E}[\sup_t M_t ^p] \sim \mathbb{E}[[M, M]_T^{p/2}]$.
2	Theorem 2.38	Every martingale in a Brownian filtration is an Itô integral.
2	Theorem 2.40 (Girsanov)	Change-of-measure formula removing a drift via an exponential martingale.
2	Theorem 2.42	$H^k \hookrightarrow C^0$ for $k > d/2$.
2	Theorem 2.43 (Hahn–Banach)	Norm-preserving extension of a bounded linear functional.
2	Theorem 2.44 (Lumer–Phillips)	Dissipative + dense range $\Rightarrow$ generates a contraction semigroup.
2	Theorem 2.45	Log-Sobolev inequality implies a spectral gap lower bound.
2	Theorem 2.49 (Stochastic Maximum Principle)	Optimal control maximises the Hamiltonian pointwise.
2	Theorem 2.50	SMP adjoint variables equal derivatives of the HJB value function.

Table B.3: Index of numbered results: Chapters 3–8 (main theoretical and applied contributions)

Chapter	Result	One-line statement
3	Theorem 3.9 (Quantitative Well-Posedness)	Unique MFG equilibrium exists; Φ is a contraction in weighted W_2 with explicit constant $C = 2L_b^2 + 2\ \sigma_0\ _F^2 + C_{\text{reg}}$.
4	Theorem 4.8 (Numerical Convergence)	Combined Euler–Maruyama/particle/SPDE scheme converges with weak error $O(\Delta t^{1/2} + \Delta x)$.
4	Theorem 4.6	Finite-difference Fokker–Planck scheme converges at the same rate under a CFL condition.
5	Theorem ?? (Subsequential RL Convergence)	Three-timescale actor-critic converges $\liminf W_2 = 0$ a.s. under compact Lipschitz function classes.
5	Theorem ?? (Exponential Convergence)	Entropy regularisation + log-Sobolev policy class $\Rightarrow$ exponential W_2 convergence rate $\lambda = \beta/(2C_{\text{Lip}})$.
6	Proposition ?? (Hierarchical Architectural Principle)	SPDE operator decomposition justifies a fast-critic/intermediate-actor/slow-meta-learner timescale separation.
7	Proposition ??	$O(1/N)$ propagation of chaos and $\varepsilon_N = O(N^{-1/2})$ -Nash property for the systemic risk model.
8	Proposition ?? (Four Alpha Signals)	Optimal spread, mean-reversion speed, factor exposure, and crowding measure $c_t = W_2(\mu_t, \mu_t^*)$ are explicit functions of the LQ-MFG equilibrium.

Table B.4: Expanded notation glossary

Symbol	Meaning
X_t, X_t^i	State (log-position / bank reserve) of the representative or i -th agent
μ_t, μ_t^N	Population measure flow; empirical measure of N agents
μ_t^*	MFG equilibrium measure flow
$\bar{\mu}_t, \bar{x}_t$	Population mean (first moment of μ_t)
α_t, α_t^*	Control (trading rate / capital-raising effort); optimal control
b_t	Government bailout rate (Chapter 7 only)
p_t, q_t, r_t	Adjoint (costate) process and its idiosyncratic/common-noise volatility components
κ	Mean-reversion speed
λ	Cost-of-control (trading cost / effort cost) parameter
γ	Cost-of-bailout parameter (Ch. 7) or risk-aversion parameter (Ch. 8)
σ, σ_0	Idiosyncratic and common-noise volatility
A_t, B_t, C_t	Quadratic Riccati coefficients (value-function ansatz)
D_t, E_t, F_t	Linear/constant Riccati coefficients (Ch. 7–8)
Γ_t, Π_t	Effective mean-reversion rate and government Riccati coefficient (Ch. 7)
c_t	Crowding measure, $c_t = W_2(\mu_t, \mu_t^*)$
$s_{i,t}$	Optimal spread signal for asset i
κ_t	(Overloaded in Ch. 8) model-implied mean-reversion speed signal, $\kappa + A_t/\lambda$
$\beta_{i,t}$	Factor exposure signal for asset i
$\delta(c_t)$	Crowding discount function applied to portfolio weights
w_t^{MV}, w_t^*	Baseline mean-variance weights; crowding-discounted final weights
θ	Generic vector of LQ-MFG model parameters to be calibrated
$\hat{\theta}$	GMM estimator of θ
$\gamma_t, \beta_t, \eta_t$	(Ch. 5–6, overloaded with the risk-aversion γ of Ch. 8) fast/intermediate/slow RL learning-rate sequences
θ_t, ϕ_t, ψ_t	Critic, actor, and meta-learner parameters (Ch. 5–6)
$\mathcal{L}_{\text{idio}}, \mathcal{L}_{\text{mf}}, \mathcal{L}_{\text{common}}$	SPDE operators for idiosyncratic diffusion, mean-field interaction, and common-noise transport (Ch. 6)

Bibliography

- Saran Ahuja, Weiluo Ren, and Tzu-Wei Yang. Forward-backward stochastic differential equations with monotone functionals and mean field games with common noise. *Stochastic Processes and their Applications*, 129(10):3859–3892, 2019.
- Dominique Bakry, Ivan Gentil, and Michel Ledoux. *Analysis and Geometry of Markov Diffusion Operators*. Springer, 2014.
- Jean-David Benamou and Yann Brenier. A computational fluid mechanics solution to the Monge–Kantorovich mass transfer problem. *Numerische Mathematik*, 84:375–393, 2000.
- Haïm Brezis. *Functional Analysis, Sobolev Spaces and Partial Differential Equations*. Springer, 2011.
- Pierre Cardaliaguet, François Delarue, Jean-Michel Lasry, and Pierre-Louis Lions. *The Master Equation and the Convergence Problem in Mean Field Games*. Annals of Mathematics Studies. Princeton University Press, 2019.
- Pierre Cardaliaguet, Benjamin Seeger, and Panagiotis E. Souganidis. Regularity for mean field games with common noise and degenerate idiosyncratic noise. *Preprint*, 2025.
- René Carmona and François Delarue. *Probabilistic Theory of Mean Field Games with Applications I-II*. Probability Theory and Stochastic Modelling. Springer, 2018.
- Marco Cuturi. Sinkhorn distances: lightspeed computation of optimal transport. *Advances in Neural Information Processing Systems*, 26, 2013.
- Nicolas Fournier and Arnaud Guillin. On the rate of convergence in wasserstein distance of the empirical measure. *Probability Theory and Related Fields*, 162(3-4):707–738, 2015.
- Richard Jordan, David Kinderlehrer, and Felix Otto. The variational formulation of the Fokker–Planck equation. *SIAM Journal on Mathematical Analysis*, 29(1):1–17, 1998.
- Jean-Michel Lasry and Pierre-Louis Lions. Mean field games. *Japanese Journal of Mathematics*, 2(1):229–260, 2007.
- Felix Otto. The geometry of dissipative evolution equations: The porous medium equation. *Communications in Partial Differential Equations*, 26(1–2):101–174, 2001.
- Amnon Pazy. *Semigroups of Linear Operators and Applications to Partial Differential Equations*. Springer, 1983.

Justin Sirignano and Konstantinos Spiliopoulos. DGM: A deep learning algorithm for solving partial differential equations. *Journal of Computational Physics*, 375:1339–1364, 2018.