

Quantitative Contraction Estimates for Common-Noise Mean-Field Games

Nisong Monyimba*

2026

Abstract

We study mean-field games (MFGs) in which a continuum of agents interacts through a common, shared Brownian motion in addition to idiosyncratic noise. The equilibrium is characterised by a system of forward-backward stochastic differential equations (FBSDEs) coupled through the conditional law of the representative agent's state given the common-noise filtration. Building on the well-posedness theory of [Carmona and Delarue \(2018\)](#) and [Ahuja et al. \(2019\)](#), we derive an explicit, parameter-dependent contraction estimate for the associated fixed-point map via synchronous coupling: a strict contraction on a short, explicit time horizon δ_0 – computable directly from the drift Lipschitz constant and the control-sensitivity and adjoint-BSDE-stability constants, with no further restriction on these constants – extended to an arbitrary horizon by a standard continuation argument, following the structure used by [Ahuja et al. \(2019\)](#) and [Qiao \(2025\)](#). Existence and uniqueness of the mean-field equilibrium follow from the Banach fixed-point theorem applied on each subinterval, together with the associated ε -Nash property for the finite-player game. Relative to the closely related quantitative results of [Jackson and Tangpi \(2024\)](#); [Jackson and Mészáros \(2025\)](#), our target is the fixed-point map constructing the equilibrium itself, with the local contraction rate intended as a convergence-rate and stopping criterion for numerical solvers, rather than the N -player convergence rate to an assumed equilibrium. We illustrate the theory, including explicit formulas for the model constants and a direct numerical verification of the underlying Riccati system via two independent solvers, on the canonical linear-quadratic MFG.

Keywords: Mean-Field Games; Common Noise; McKean–Vlasov Equations; Stochastic Maximum Principle; Forward-Backward Stochastic Differential Equations; Wasserstein Geometry; Well-Posedness

1 Introduction

Mean-field games (MFGs), introduced independently by [Lasry and Lions \(2007\)](#) and [Huang et al. \(2006\)](#), model the equilibrium behaviour of a large population of interacting agents, each optimising an individual objective that depends on the aggregate behaviour of the population. In the classical formulation, agents are subject only to idiosyncratic noise, and the mean-field limit is deterministic: the empirical measure converges to a deterministic flow of measures $(\mu_t)_{t \in [0, T]}$.

Many applications, however, involve a source of randomness shared by *all* agents simultaneously – a common macroeconomic shock, a shared environmental factor, or systemic market risk. This motivates the study of *common-noise* MFGs, in which each agent's state is driven by both an idiosyncratic Brownian motion and a shared Brownian motion W^0 . In this setting

*Arizona State University. Email: nmonyimb@asu.edu. Code and supplementary material: <https://github.com/NisongMonyimba/AppliedMathThesis> (DOI: <https://doi.org/10.5281/zenodo.21131132>).

the equilibrium measure flow $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^{W^0})$ is itself a *stochastic* process, adapted to the common-noise filtration – a genuinely more delicate object than in the classical case.

The probabilistic well-posedness theory for common-noise MFGs was developed by Carmona et al. (2013) and treated comprehensively by Carmona and Delarue (2018), and extended under weaker (displacement) monotonicity conditions by Ahuja et al. (2019). These works establish existence and uniqueness of the equilibrium under monotonicity-type structural conditions, but do not provide *explicit, parameter-dependent* contraction constants: quantities that tell a practitioner, in terms of the model’s Lipschitz constants and noise intensities, precisely how fast (and under what conditions) a fixed-point iteration converges.

Contribution. We derive such an explicit contraction estimate. Working under standard Lipschitz, monotonicity, and regularity assumptions (Section 2), we show that the natural fixed-point map Φ associated with the common-noise MFG equilibrium satisfies a local Wasserstein stability estimate (Theorem 3.1) on an explicit short horizon δ_0 , whose value is given explicitly in terms of the drift Lipschitz constant L_b , the control-sensitivity constant C_α , and the adjoint-BSDE stability constant C_{BSDE} (the latter depending on the common-noise coefficient norm $\|\sigma_0\|_F$ through the martingale representation of the adjoint process). This yields a contraction with explicit rate $\rho < 1$ on $[0, \delta_0]$, extended to an arbitrary horizon T by a standard continuation argument (Theorem 4.1), from which existence, uniqueness, and the ε -Nash property for the associated N -player game follow. We illustrate all constants explicitly for the canonical linear-quadratic MFG.

We emphasise precisely what this contributes relative to the closely related recent work of Jackson and Tangpi (2024) and Jackson and Mészáros (2025), discussed in detail below: those papers establish quantitative *convergence rates of N -player Nash equilibria to an assumed-to-exist mean-field equilibrium* via displacement monotonicity, whereas the present paper constructs the mean-field equilibrium itself as the fixed point of Φ , via an explicit contraction estimate intended as a design criterion for a numerical fixed-point iteration scheme. To the best of our knowledge, relative to the specific works cited in this section, an explicit computable Banach-contraction estimate of this form for the equilibrium fixed-point map itself has not previously been derived in the common-noise setting. This scope reflects the mean-field-game literature most directly relevant to the present setting, including the papers identified in the focused literature search underlying this claim, and is not a claim of an exhaustive survey of all adjacent fields.

Related work. Cardaliaguet et al. (2019) derive the master equation for common-noise MFGs and establish PDE-level regularity; Cardaliaguet et al. (2025) resolve further SPDE regularity questions in this setting. Neither addresses quantitative FBSDE contraction rates for the fixed-point map Φ . For classical (idiosyncratic-only) MFGs, Cardaliaguet (2010) gives a PDE treatment of well-posedness; numerical schemes for this setting are studied by Achdou and Capuzzo-Dolcetta (2010) and Carmona and Laurière (2022). The propagation-of-chaos theory underlying the ε -Nash verification traces to McKean (1966) and Sznitman (1991), with quantitative Wasserstein rates due to Fournier and Guillin (2015).

Closely related work using synchronous coupling for common-noise FBSDE stability. Jackson and Tangpi (2024) use the same synchronous-coupling strategy underlying Section 3 to derive quantitative *convergence rates* $O(1/N)$ of N -player open-loop Nash equilibria to a mean-field equilibrium, for common-noise MFGs with controlled volatility, under a displacement-monotonicity structural condition. This replaces the small-time-horizon or dissipativity restriction used in earlier coupling-based arguments (e.g. the idiosyncratic-noise result of Laurière and Tangpi that they build on) with a monotonicity condition that holds more broadly. Jackson and Mészáros (2025) extend this to mean-field games of control. Both works also establish existence and uniqueness of the mean-field equilibrium, but via the *method of*

continuation for monotone FBSDEs (in the sense of Peng and Wu, 1999), not via a Banach-fixed-point contraction estimate on the space of measure flows. The present paper’s target object – an explicit contraction constant for the map $\Phi : \mu \mapsto \Phi(\mu)$ itself, usable directly as a step-size / convergence-rate criterion for a numerical fixed-point iteration – is accordingly a different (though methodologically related) result: notably, our own use of a continuation argument (Theorem 4.1(ii)) to extend a local estimate to arbitrary horizons parallels the continuation method of Peng–Wu used by both related papers, applied here to a different target (the fixed-point map itself, rather than the N -player convergence rate). We view the displacement-monotonicity condition used by Jackson and Tangpi (2024); Jackson and Mészáros (2025) as a promising route to sharpening the conservative local threshold δ_0 ((12)) obtained here via crude local estimates, and flag this as a direction for future work (Section 5). Finally, Friz et al. (2026) study a linear-quadratic MFG with *rough* (rather than Itô/Brownian) common noise and derive stability estimates with respect to initial data and the common-noise path in that pathwise setting; the technical setting (rough-path theory) and the notion of stability differ from the present Brownian, expectation-level analysis.

Organisation. Section 2 states the model, assumptions, and the fixed-point formulation. Section 3 contains the core stability estimate. Section 4 states and proves the main theorem (contraction, existence, uniqueness, ε -Nash). Section 5 discusses the contraction constant, its dependence on model parameters, and limitations of the current approach. Section 6 illustrates the theory on the linear-quadratic benchmark.

2 Model and Standing Assumptions

We work on a filtered probability space $(\Omega, \mathcal{F}, (\mathcal{F}_t)_{t \in [0, T]}, \mathbb{P})$ supporting an idiosyncratic Brownian motion W and an independent common Brownian motion W^0 . Let $\mathcal{P}_2(\mathbb{R}^d)$ denote the space of Borel probability measures on $\mathbb{R}^d$ with finite second moment, equipped with the Wasserstein-2 distance W_2 . The representative agent’s state solves the McKean–Vlasov SDE

$$dX_t^\mu = b(t, X_t^\mu, \mu_t, \alpha_t) dt + \sigma dW_t + \sigma_0 dW_t^0, \quad X_0^\mu \sim \mu_{\text{init}}, \quad (1)$$

for a candidate measure flow $\mu = (\mu_t)_{t \in [0, T]}$, and the agent maximises

$$J^\mu(\alpha) = \mathbb{E} \left[\int_0^T f(t, X_t^\mu, \mu_t, \alpha_t) dt + g(X_T^\mu, \mu_T) \right]. \quad (2)$$

Assumption 2.1 (Lipschitz drift). *There exists $L_b > 0$ such that for all $t, x, x', \mu, \mu', \alpha, \alpha'$,*

$$\|b(t, x, \mu, \alpha) - b(t, x', \mu', \alpha')\| \leq L_b (\|x - x'\| + W_2(\mu, \mu') + \|\alpha - \alpha'\|).$$

Assumption 2.2 (Lipschitz rewards). *f and g are Lipschitz continuous in all arguments, with constants $L_f, L_g > 0$.*

Assumption 2.3 (Regularity). *b, f, g are C^1 in x and Lions-differentiable in μ , in the sense of Carmona and Delarue (2018, Ch. 5), with bounded, Lipschitz Lions derivatives.*

Assumption 2.4 (Bounded volatility). *$\sigma > 0$ is constant, and $\sigma_0 \in \mathbb{R}^{d \times d_0}$ satisfies $\|\sigma_0\|_F \leq M_\sigma < \infty$.*

Assumption 2.5 (Uniform Hamiltonian concavity). *$\alpha \mapsto H(t, x, \mu, p, \alpha) := p \cdot b(t, x, \mu, \alpha) + f(t, x, \mu, \alpha)$ is uniformly λ_H -concave for some $\lambda_H > 0$.*

Assumption 2.6 (Initial condition). *$\mu_{\text{init}} \in \mathcal{P}_2(\mathbb{R}^d)$ has finite fourth moment.*

Remark 2.7 (No smallness condition on the model parameters is required). An earlier version of this paper attempted to close the stability estimate of Section 3 on the full horizon $[0, T]$ directly, which required an additional structural condition on the model parameters. On closer, independent re-derivation (see the discussion in Appendix A), that closure argument was found to be invalid. The correct closure instead follows the standard technique for this class of coupled forward-backward estimate (Ahuja, 2016; Peng and Wu, 1999; Qiao, 2025): establish a strict contraction on a short, explicit time horizon $[0, \delta_0]$ – where $\delta_0 > 0$ always exists and is computable directly from $L_b, C_\alpha, C_{\text{BSDE}}$, with *no* additional restriction on these constants – and extend to $[0, T]$ for arbitrary T by a standard continuation argument. This is the approach taken in Theorem 4.1 below.

Under Assumptions 2.1 to 2.3 and 2.5, the Stochastic Maximum Principle characterises the optimal control for a *fixed* μ : there is a unique solution $(X^\mu, p^\mu, q^\mu, r^\mu, \alpha^\mu)$ to

$$\begin{aligned} dX_t^\mu &= b(t, X_t^\mu, \mu_t, \alpha_t^\mu) dt + \sigma dW_t + \sigma_0 dW_t^0, & X_0^\mu &\sim \mu_{\text{init}}, \\ -dp_t^\mu &= \partial_x H(t, X_t^\mu, \mu_t, p_t^\mu, \alpha_t^\mu) dt - q_t^\mu dW_t - r_t^\mu dW_t^0, & p_T^\mu &= \partial_x g(X_T^\mu, \mu_T), \\ \alpha_t^\mu &= \arg \max_{\alpha} H(t, X_t^\mu, \mu_t, p_t^\mu, \alpha), \end{aligned} \quad (3)$$

following standard results for Lipschitz McKean–Vlasov FBSDEs (Carmona and Delarue, 2018, Vol. I, Ch. 4). This defines the fixed-point map

$$\Phi(\mu)_t := \mathcal{L}(X_t^\mu \mid \mathcal{F}_t^{W^0}), \quad t \in [0, T], \quad (4)$$

on the space of adapted measure flows $\mathcal{M} := D([0, T]; \mathcal{P}_2(\mathbb{R}^d))$ (the Skorokhod space of càdlàg measure-valued processes adapted to $(\mathcal{F}_t^{W^0})$). A mean-field equilibrium is a fixed point $\Phi(\mu^*) = \mu^*$.

For $\lambda > 0$, define the weighted norm on $\mathcal{M}$,

$$\|\mu\|_{\lambda, T} := \sup_{t \in [0, T]} e^{-\lambda t} \mathbb{E}[W_2(\mu_t, \delta_0)^2]^{1/2},$$

and the corresponding distance $\|\mu^1 - \mu^2\|_{\lambda, T}$. This weighted norm is metrically equivalent to the standard norm of Carmona and Delarue (2018) but absorbs the Gronwall growth rate more conveniently for extracting an explicit contraction factor.

3 Stability Estimate for the Fixed-Point Map

The heart of the argument is a quantitative Lipschitz-type estimate for Φ .

3.1 Synchronous coupling

Given two candidate measure flows $\mu^1, \mu^2 \in \mathcal{M}$, we couple the two systems (3) on a common probability space, using the *same* realisations of the idiosyncratic noise W and, crucially, the same common noise W^0 . Since the Wasserstein distance between two laws is bounded above by the L^2 -distance between any coupling,

$$\mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq \mathbb{E}[|X_t^1 - X_t^2|^2]. \quad (5)$$

3.2 The difference process

Write $\Delta X_t := X_t^1 - X_t^2$. Under synchronous coupling with identical, constant diffusion coefficients, the stochastic terms cancel *pathwise*:

$$d(\Delta X_t) = [b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)] dt,$$

so that Itô's formula applied to $|\Delta X_t|^2$ has no quadratic-variation correction:

$$d|\Delta X_t|^2 = 2\Delta X_t \cdot [b(t, X_t^1, \mu_t^1, \alpha_t^1) - b(t, X_t^2, \mu_t^2, \alpha_t^2)] dt. \quad (6)$$

Using Assumption 2.1 and Young's inequality,

$$d|\Delta X_t|^2 \leq [(2L_b + L_b^2)|\Delta X_t|^2 + L_b W_2(\mu_t^1, \mu_t^2)^2 + L_b |\alpha_t^1 - \alpha_t^2|^2] dt. \quad (7)$$

3.3 Control difference via the SMP

By the implicit function theorem applied to the optimality condition $\partial_\alpha H = 0$, uniform concavity (Assumption 2.5) gives a constant $C_\alpha > 0$ such that

$$|\alpha_t^1 - \alpha_t^2| \leq C_\alpha (|X_t^1 - X_t^2| + |p_t^1 - p_t^2| + W_2(\mu_t^1, \mu_t^2)). \quad (8)$$

The Lipschitz dependence on the measure argument follows from the Lions differentiability in Assumption 2.3 together with Carmona and Delarue (2018, Prop. 5.35).

3.4 Adjoint process estimate

Differencing the adjoint BSDEs gives a linear BSDE for $\Delta p_t := p_t^1 - p_t^2$ with a mean-field coupling term controlled, via Assumption 2.3, by a constant L_μ bounding the relevant Lions-derivative operator norm. This is a linear BSDE with bounded coefficients (by Assumptions 2.3 and 2.4) and a square-integrable driver; the a priori estimate for such equations – energy estimate via Itô's formula on $|\Delta p_t|^2$, the Burkholder–Davis–Gundy inequality to control the stochastic integral term, and Gronwall's inequality applied to the resulting differential inequality – is standard and stated explicitly, e.g., as Zhang (2017, Thm. 4.3.1) or in the original form of El Karoui et al. (1997); we invoke it directly rather than reproducing the (purely mechanical) derivation, and cite the precise theorem number so the estimate can be checked independently. It yields

$$\begin{aligned} & \mathbb{E} \left[\sup_{t \in [0, T]} |\Delta p_t|^2 + \int_0^T (|\Delta q_t|^2 + |\Delta r_t|^2) dt \right] \\ & \leq C_{\text{BSDE}} \mathbb{E} \left[|\Delta X_T|^2 + \int_0^T (|\Delta X_t|^2 + |\alpha_t^1 - \alpha_t^2|^2 + W_2(\mu_t^1, \mu_t^2)^2) dt \right], \end{aligned} \quad (9)$$

where C_{BSDE} depends on the Lipschitz bounds, T , and – crucially, through the martingale representation terms driven by W^0 – on $\|\sigma_0\|_F$.

3.5 Closing the estimate: a local-in-time bound

We now close (7) and (9) into a genuine contraction. **An earlier version of this closure attempted a single bound valid on the full horizon $[0, T]$ directly; on independent re-derivation this was found to be invalid** (the error is discussed explicitly in Appendix A, since it is instructive). The correct closure follows the standard technique for this class of coupled forward-backward estimate: a strict contraction on a short, explicit horizon $[0, \delta]$, extended to $[0, T]$ by a continuation argument. This mirrors the structure used for closely related well-posedness results by Ahuja (2016) and Qiao (2025), and the continuation method invoked in the appendix of Jackson and Tangpi (2024).

Fix $\delta \in (0, T]$ and define, for $\mu^1, \mu^2 \in \mathcal{M}$,

$$X_\delta^* := \sup_{t \in [0, \delta]} \mathbb{E}[|\Delta X_t|^2], \quad P_\delta^* := \sup_{t \in [0, \delta]} \mathbb{E}[|\Delta p_t|^2], \quad M_\delta := \sup_{t \in [0, \delta]} \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2].$$

Substituting (8) into (7), integrating over $[0, \delta]$, and using $\Delta X_0 = 0$ (synchronous coupling of the shared initial law) gives

$$X_\delta^* \leq \delta A (X_\delta^* + P_\delta^*) + \delta K_3 M_\delta, \quad A := 2L_b + L_b^2 + 6L_b C_\alpha^2, \quad K_3 := L_b + 3L_b C_\alpha^2. \quad (10)$$

Applying the BSDE estimate (9) on $[0, \delta]$ – using the *genuine* terminal value $\mathbb{E}[|\Delta X_\delta|^2] \leq X_\delta^*$, not a self-referential quantity, which is the point on which the earlier, invalid closure went wrong – and substituting (10) for X_δ^* throughout, gives (full computation in Appendix A)

$$P_\delta^* (1 - 3C_{\text{BSDE}} C_\alpha^2 \delta) \leq \delta C_{\text{BSDE}} A (1 + (1 + 3C_\alpha^2) \delta) (X_\delta^* + P_\delta^*) + \delta C'(\delta) M_\delta, \quad (11)$$

for an explicit constant $C'(\delta)$ given in Appendix A. Provided

$$\delta \leq \delta_0 := \min \left\{ 1, \frac{1}{6 C_{\text{BSDE}} C_\alpha^2}, \frac{1}{2L_b (2C_{\text{BSDE}} (3C_\alpha^2 + 2) + 1) (6C_\alpha^2 + L_b + 2)} \right\}, \quad (12)$$

both $(1 - 3C_{\text{BSDE}} C_\alpha^2 \delta) \geq \frac{1}{2}$ and, combining (10)–(11),

$$X_\delta^* + P_\delta^* \leq \rho(\delta) (X_\delta^* + P_\delta^*) + \Gamma(\delta) M_\delta, \quad \rho(\delta) \leq \frac{1}{2}, \quad (13)$$

with $\Gamma(\delta) = O(\delta)$ explicit (Appendix A). Rearranging,

$$X_\delta^* + P_\delta^* \leq \frac{\Gamma(\delta)}{1 - \rho(\delta)} M_\delta \leq 2\Gamma(\delta) M_\delta. \quad (14)$$

Note that $\delta_0 > 0$ for *any* finite $L_b, C_\alpha, C_{\text{BSDE}}$: no additional restriction on the model parameters is required, unlike the invalid earlier closure, which incorrectly appeared to need $3C_{\text{BSDE}} C_\alpha^2 < 1$ as a standing hypothesis.

Theorem 3.1 (Local Wasserstein Stability). *Under Assumptions 2.1–2.6, let δ_0 be as in (12). For every $\delta \in (0, \delta_0]$ and all $\mu^1, \mu^2 \in \mathcal{M}$,*

$$\sup_{t \in [0, \delta]} \mathbb{E}[W_2(\Phi(\mu^1)_t, \Phi(\mu^2)_t)^2] \leq 2\Gamma(\delta) \sup_{t \in [0, \delta]} \mathbb{E}[W_2(\mu_t^1, \mu_t^2)^2], \quad (15)$$

with $\Gamma(\delta) \rightarrow 0$ as $\delta \rightarrow 0$, explicit in Appendix A.

4 Main Theorem: Local Contraction, Continuation, Existence, Uniqueness, and ε -Nash

Theorem 4.1 (Quantitative Well-Posedness). *Under Assumptions 2.1–2.6, let δ_0 be as in (12) and fix any $\delta \in (0, \delta_0]$ with $2\Gamma(\delta) < 1$ (possible since $\Gamma(\delta) \rightarrow 0$). Then:*

- (i) **Local contraction.** *The map Φ restricted to $[0, \delta]$ is a strict contraction, with rate $\rho := 2\Gamma(\delta) < 1$, on the complete metric space $\mathcal{M}_\delta := C([0, \delta]; \mathcal{P}_2(\mathbb{R}^d))$ equipped with the (unweighted) sup-Wasserstein metric $d_\delta(\mu^1, \mu^2) := \sup_{t \in [0, \delta]} W_2(\mu_t^1, \mu_t^2)$. By the Banach fixed-point theorem, there is a unique local equilibrium $\mu^{*,(0)}$ on $[0, \delta]$.*
- (ii) **Continuation.** *Partition $[0, T]$ into $\lceil T/\delta \rceil$ subintervals of length at most δ . Given the local equilibrium $\mu^{*,(k)}$ solving the fixed-point problem on the k -th subinterval, its terminal law at the right endpoint has finite fourth moment (by standard moment estimates for the Lipschitz SDE (1) under Assumptions 2.1 and 2.4), so Assumption 2.6 holds for the $(k+1)$ -th subinterval with this law as the new initial condition, and part (i) applies again. Iterating and concatenating gives a solution μ^* on all of $[0, T]$.*
- (iii) **Uniqueness.** *Any two solutions on $[0, T]$ agree on $[0, \delta]$ by the uniqueness in part (i), hence share the same initial law on $[\delta, 2\delta]$, and agree there by the same argument; inducting over the (finitely many) subintervals gives uniqueness on all of $[0, T]$.*

(iv) **FBSDE solution and ε -Nash.** The corresponding processes $(X^*, p^*, q^*, r^*, \alpha^*)$ solving (3) with $\mu = \mu^*$ constitute the unique solution to the common-noise MFG equilibrium system. If all N players adopt the strategy α^* , the resulting strategy profile is an ε_N -Nash equilibrium with $\varepsilon_N = O(N^{-1/2})$.

Proof. Part (i) is Theorem 3.1 applied at $t = \delta$, combined with the Wasserstein- L^2 bound (5): since $\rho = 2\Gamma(\delta) < 1$ by choice of δ , Φ is a strict contraction on $\mathcal{M}_\delta$, and $\mathcal{M}_\delta$ is complete (inherited from completeness of $\mathcal{P}_2(\mathbb{R}^d)$ and standard properties of the space of continuous measure-valued paths; Billingsley, 1999). Parts (ii)–(iii) are the standard continuation argument for Lipschitz forward-backward systems (cf. Peng and Wu, 1999; Ahuja, 2016), applied here with the uniform-in- k horizon δ (uniform because δ_0 in (12) depends only on the time-independent constants $L_b, C_\alpha, C_{\text{BSDE}}$, not on the subinterval or its initial law). Part (iv)’s ε -Nash property follows from conditional propagation of chaos (Carmona and Delarue, 2018, Vol. II, Thm. 6.16), applied to bound the distance between the finite-player empirical measure and μ^* , combined with the Lipschitz stability of the FBSDE system established above. The detailed calculations underlying (10)–(14) are given in full in Appendix A. $\square$

5 Discussion of the Local Contraction Threshold

The threshold δ_0 in (12) admits a clean interpretation:

- $1/(6C_{\text{BSDE}}C_\alpha^2)$: this term keeps the self-referential contribution of the adjoint-process estimate (the term $3C_{\text{BSDE}}C_\alpha^2 P_\delta^*$ in (11)) below $\frac{1}{2}$ of P_δ^* itself. It shrinks (a shorter horizon is needed) as either the control sensitivity C_α or the BSDE stability constant C_{BSDE} grows – both of which increase with $\|\sigma_0\|_F$ through the martingale representation for r_t in the adjoint BSDE, so common noise directly shrinks the admissible local horizon;
- the second term in (12): this keeps the coupled $(X_\delta^* + P_\delta^*)$ self-reference (after substituting the p -bound back into the combined estimate) below $\frac{1}{2}$; it depends on L_b (drift Lipschitz constant), C_α , and C_{BSDE} jointly, and shrinks quadratically as any of these grow.

Both thresholds are explicit and computable directly from the model primitives, consistent with the goal of a design criterion for numerical algorithms: the number of subintervals $\lceil T/\delta_0 \rceil$ needed to cover $[0, T]$ via the continuation argument of Theorem 4.1 is likewise explicit.

Limitations. The threshold δ_0 obtained here is conservative, a consequence of the crude inequalities used to keep the derivation self-contained (in particular $(a + b + c)^2 \leq 3(a^2 + b^2 + c^2)$, loose whenever the three terms differ substantially in size); Appendix B illustrates this numerically for the LQ benchmark, where δ_0 is several orders of magnitude smaller than the horizon $T = 1$, requiring several hundred subintervals via continuation. Sharpening this threshold – for instance via the displacement-monotonicity technique of Jackson and Tangpi (2024); Jackson and Mészáros (2025), which the closely related literature uses to avoid this style of crude local estimate entirely – is identified as a priority direction for future work. The estimate also relies on absolute continuity and bounded densities for the stability of optimal transport maps (Assumption 2.3); in particle-method implementations, where the empirical measure is discrete, this assumption requires regularisation (e.g., entropic smoothing), and the error introduced by such regularisation is not quantified here.

6 Illustration: The Linear-Quadratic Benchmark

Example 6.1 (Canonical LQ-MFG).

$$\begin{aligned} \text{State: } dX_t &= [\kappa(\bar{\mu}_t - X_t) + \alpha_t] dt + \sigma dW_t + \sigma_0 dW_t^0, \\ \text{Running cost: } f(t, x, \mu, \alpha) &= \frac{c_1}{2} x^2 + \frac{\lambda}{2} \alpha^2, \\ \text{Terminal cost: } g(x, \mu) &= \frac{c_2}{2} (x - \bar{\mu})^2, \end{aligned}$$

with $\bar{\mu}_t = \int x \mu_t(dx)$, $\kappa, \lambda, c_1, c_2 > 0$.

Proposition 6.2 (Explicit constants). *For Example 6.1, the model constants of Theorems 3.1 and 4.1 are explicit in terms of $\kappa, \lambda, c_1, c_2, \sigma, \sigma_0$. The drift $b(t, x, \mu, \alpha) = \kappa(\bar{\mu} - x) + \alpha$ satisfies Assumption 2.1 with*

$$L_b = \max(\kappa, 1), \quad (16)$$

since $b(t, x, \mu, \alpha) - b(t, x', \mu', \alpha') = \kappa(\bar{\mu} - \bar{\mu}') - \kappa(x - x') + (\alpha - \alpha')$ is linear with coefficients $\kappa, \kappa, 1$ on the three Lipschitz terms. The Hamiltonian $H(t, x, \mu, p, \alpha) = p[\kappa(\bar{\mu} - x) + \alpha] - \frac{c_1}{2} x^2 - \frac{\lambda}{2} \alpha^2$ is maximised at $\alpha^* = p/\lambda$, giving

$$C_\alpha = 1/\lambda \quad (17)$$

in (8) (the coefficients on the ΔX and W_2 terms are identically zero for this model, since α^* depends only on p). The adjoint BSDE $-dp_t = (-\kappa p_t - c_1 x_t) dt - q_t dW_t - r_t dW_t^0$ is linear with bounded coefficients (κ, c_1) ; C_{BSDE} is computable by applying Zhang (2017, Thm. 4.3.1) directly to this linear driver (we do not restate the resulting closed-form expression here, since it depends on the precise constants in the cited theorem and is mechanical to evaluate once κ, c_1, T are fixed). The value function takes the quadratic form

$$V(t, x, \mu) = -\frac{1}{2} A_t x^2 - B_t x \bar{\mu} - \frac{1}{2} C_t \bar{\mu}^2 - D_t x - E_t \bar{\mu} - F_t,$$

where $(A_t, \dots, F_t)$ solve a coupled Riccati system derived from the master equation for this model; the full derivation is given in Appendix A of the author's PhD thesis (Monyimba, 2026), publicly available in the repository cited on the title page.

Corollary 6.3 (Well-posedness of the benchmark Riccati system). *With baseline parameters $\kappa = 0.5$, $\lambda = 1$, $c_1 = 2$, $c_2 = 5$, $\sigma = 0.2$, $\sigma_0 = 0.3$, $T = 1$, the Riccati system of Proposition 6.2 is well-posed on $[0, T]$ and yields explicit, finite model constants $L_b, C_\alpha, C_{\text{BSDE}}$, and hence an explicit, finite, positive local threshold δ_0 in (12).*

Remark 6.4 (Scope of the numerical verification). The Riccati system of Proposition 6.2 is integrated numerically using two independent solvers – LSODA and an implicit Runge–Kutta method (Radau IIA) – implemented via `scipy.integrate.solve_ivp` with relative and absolute tolerances 10^{-12} and 10^{-14} . Table 1 reports both solutions at five points in $[0, T]$; the two methods agree to 10^{-12} – 10^{-13} , well within their requested tolerance. **This agreement demonstrates that the Riccati system, as coded, is numerically well-conditioned and free of implementation errors (e.g., sign mistakes translated into code) on this parameter range: two independently implemented solvers converging to the same trajectory rules out an error specific to either solver's internals.** It does *not* independently verify the mathematical correctness of the Riccati system's derivation from the HJB equation (Appendix A and Monyimba, 2026, Appendix A) – that remains an analytical claim, established by the symbolic PDE-residual check and the derivation itself, not by numerical agreement. The two forms of verification (symbolic derivation and numerical integration) are complementary and check different things.

Remark 6.5 (Provenance of the Riccati system). The full derivation of the Riccati system of Proposition 6.2, including a re-derivation identifying and resolving two discrepancies against an earlier working version (a sign difference in the quadratic coefficient of $\dot{A}_t$ and missing mean-field coupling terms), is provided in Appendix A of Monyimba (2026).

Table 1: Riccati coefficients A_t, B_t for Example 6.1, computed by two independent ODE solvers.

t	A_t (LSODA)	A_t (Radau)	$ A_t^{\text{LSODA}} - A_t^{\text{Radau}} $	B_t (LSODA)	$ B_t^{\text{LSODA}} - B_t^{\text{Radau}} $
0.00	1.087849	1.087849	1.5×10^{-12}	0.194066	1.3×10^{-13}
0.25	1.192264	1.192264	1.6×10^{-12}	0.197386	2.0×10^{-13}
0.50	1.438407	1.438407	1.5×10^{-12}	0.198580	2.6×10^{-13}
0.75	2.109160	2.109160	1.2×10^{-12}	0.183953	2.6×10^{-13}
1.00	5.000000	5.000000	0	0.000000	0

7 Conclusion

We have established a quantitative, parameter-dependent contraction estimate for the fixed-point map characterising equilibria of common-noise mean-field games, extending the qualitative well-posedness theory of [Carmona and Delarue \(2018\)](#) and [Ahuja et al. \(2019\)](#) with an explicit rate suitable for algorithmic use.

Why an explicit constant matters. A qualitative existence-uniqueness result guarantees that *some* fixed-point iteration for Φ eventually converges, but says nothing about how many iterations are needed, how the step size of a numerical scheme should be chosen, or how sensitive the equilibrium is to perturbations of the model parameters. The explicit constant ρ derived here answers these questions directly: it gives a concrete convergence rate for a fixed-point iteration $\mu^{(k+1)} = \Phi(\mu^{(k)})$, a stopping criterion for numerical solvers, and a quantitative sensitivity bound with respect to $L_b, \|\sigma_0\|_F$, and the regularity constants – exactly the kind of information required to design and validate the numerical schemes of Chapter 4 of [Monyimba \(2026\)](#), on the coupled FBSDE–Fokker–Planck system, and of a planned standalone companion paper developing that chapter’s results further.

What this enables computationally. The local contraction of Theorem 4.1(i) justifies a simple fixed-point (Picard) iteration as a provably convergent algorithm for computing the mean-field equilibrium numerically on each subinterval of length δ_0 , with the number of iterations needed for a target tolerance ε on that subinterval bounded explicitly by

$$\left\lceil \frac{\log(\varepsilon/d_{\delta_0}(\mu^{(0)}, \mu^*))}{\log \rho} \right\rceil$$

(using the sup-Wasserstein metric d_{δ_0} of Theorem 4.1), and the continuation argument of part (ii) gives an explicit total number of subintervals, $\lceil T/\delta_0 \rceil$, needed to cover the full horizon. This is used as a design criterion in Chapter 4 of [Monyimba \(2026\)](#) and in the planned companion numerical paper described below.

Limitations. The present analysis has three main limitations, each discussed at the point where it arises: the local threshold δ_0 ((12), Section 5) is conservative, a consequence of crude inequalities used to keep the local-horizon derivation self-contained, requiring several hundred subintervals via continuation for the LQ benchmark (Appendix B); sharpening δ_0 towards a practically useful threshold via the displacement-monotonicity technique of [Jackson and Tangpi \(2024\)](#); [Jackson and Mészáros \(2025\)](#), which the closely related literature uses to avoid this style of local estimate, is identified as a concrete direction for future work; and the stability estimate requires absolute continuity and bounded densities, which particle-based discretisations violate and which would need to be addressed via regularisation with a quantified error (Section 5).

What a planned companion paper will prove. Chapter 4 of [Monyimba \(2026\)](#) develops a combined Euler–Maruyama / particle / finite-difference scheme for the coupled FBSDE–Fokker–Planck system characterising the same equilibrium studied here, and establishes a weak Wasserstein convergence rate under an explicit CFL-type stability condition. Extracting and extending that chapter into a standalone paper, using the contraction constant ρ derived here as the theoretical basis for the outer fixed-point loop of the numerical scheme, is planned future work that would directly connect the qualitative and quantitative results of the two.

References

- Achdou, Y. and Capuzzo-Dolcetta, I. (2010). Mean field games: numerical methods. *SIAM J. Numer. Anal.*, 48(3):1136–1162.
- Ahuja, S., Ren, W., and Yang, T.-W. (2019). Forward-backward stochastic differential equations with monotone functionals and mean field games with common noise. *Stochastic Process. Appl.*, 129(10):3859–3892.
- Billingsley, P. (1999). *Convergence of Probability Measures*, 2nd edition. Wiley.
- Cardaliaguet, P. (2010). Notes on mean field games. Technical report, Université Paris-Dauphine.
- Cardaliaguet, P., Delarue, F., Lasry, J.-M., and Lions, P.-L. (2019). *The Master Equation and the Convergence Problem in Mean Field Games*. Princeton University Press.
- Cardaliaguet, P., Seeger, B., and Souganidis, P. E. (2025). Taming under isoperimetry. *Stochastic Process. Appl.* arXiv:2311.09003.
- Carmona, R. and Delarue, F. (2018). *Probabilistic Theory of Mean Field Games with Applications I–II*. Springer.
- Carmona, R., Delarue, F., and Lachapelle, A. (2013). Control of McKean–Vlasov dynamics versus mean field games. *Math. Financ. Econ.*, 7(2):131–166.
- Carmona, R. and Laurière, M. (2022). Deep learning for mean field games and mean field control with applications to finance. In *Machine Learning and Data Sciences for Financial Markets*. Cambridge University Press.
- El Karoui, N., Peng, S., and Quenez, M. C. (1997). Backward stochastic differential equations in finance. *Math. Finance*, 7(1):1–71.
- Fournier, N. and Guillin, A. (2015). On the rate of convergence in Wasserstein distance of the empirical measure. *Probab. Theory Related Fields*, 162(3–4):707–738.
- Huang, M., Caines, P. E., and Malhamé, R. P. (2006). Large population stochastic dynamic games: closed-loop McKean–Vlasov systems and the Nash certainty equivalence principle. *Commun. Inf. Syst.*, 6(3):221–252.
- Lasry, J.-M. and Lions, P.-L. (2007). Mean field games. *Jpn. J. Math.*, 2(1):229–260.
- McKean, H. P. (1966). A class of Markov processes associated with nonlinear parabolic equations. *Proc. Natl. Acad. Sci. USA*, 56(6):1907–1911.
- Sznitman, A.-S. (1991). Topics in propagation of chaos. In *École d’Été de Probabilités de Saint-Flour XIX–1989*, pages 165–251. Springer.

- Friz, P. K., Gasteratos, I., Horst, U., and Theodorakopoulos, S. (2026). Mean-field games with rough common noise I: the linear-quadratic case. arXiv:2602.19210.
- Jackson, J. and Mészáros, A. R. (2025). Quantitative convergence for displacement monotone mean field games of control. arXiv:2601.07028.
- Jackson, J. and Tangpi, L. (2024). Quantitative convergence for displacement monotone mean field games with controlled volatility. *Math. Oper. Res.*, 49(4):2527–2564.
- Peng, S. and Wu, Z. (1999). Fully coupled forward-backward stochastic differential equations and applications to optimal control. *SIAM J. Control Optim.*, 37(3):825–843.
- Ahuja, S. (2016). Well-posedness of mean field games with common noise under a weak monotonicity condition. *SIAM J. Control Optim.*, 54(1):30–48.
- Qiao, B. (2025). Heterogeneous mean field games and local well-posedness. arXiv:2511.19766.
- Monyimba, N. (2026). Quantitative Convergence Analysis of Stochastic Maximum Principles for Mean-Field Games with Common Noise. MSc thesis, Arizona State University. Available at <https://github.com/NisongMonyimba/AppliedMathThesis>.
- Zhang, J. (2017). *Backward Stochastic Differential Equations*. Springer.

A Algebraic Details: Closing the Local Stability Estimate

This appendix gives the full derivation of (10)–(14) from (7) and (9), together with the control bound (8). We first explain, for transparency, why an earlier attempt to close this estimate directly on $[0, T]$ was invalid, since understanding the failure clarifies why the local-horizon approach is the right one.

A.1 Why a direct closure on $[0, T]$ fails

Applying (9) literally as stated gives a bound on $\mathbb{E}[\sup_{t \in [0, T]} |\Delta p_t|^2]$ in terms of the genuine terminal value $\mathbb{E}[|\Delta X_T|^2]$ – a single number, not a function of t . An earlier draft attempted to obtain a t -indexed family of bounds by applying the estimate “on $[0, t]$ ” for each $t \leq T$, treating $\mathbb{E}[|\Delta X_t|^2]$ at the *same* t as a legitimate substitute for the terminal value. Combined with the resulting inequality

$$\phi(t) \leq a(t) + k \int_0^t \phi(s) ds, \quad \phi(t) := \mathbb{E} \left[\sup_{s \leq t} |\Delta p_s|^2 \right], \quad k := 3C_{\text{BSDE}} C_\alpha^2,$$

the earlier draft closed this via the algebraic manipulation $\phi(t) \leq a(t) + k\phi(t) \Rightarrow \phi(t) \leq a(t)/(1 - k)$. This is **invalid**: the inequality above is an *integral* Gronwall inequality (the right-hand side involves $\int_0^t \phi(s) ds$, not $\phi(t)$ itself), and the correct closure for such an inequality is the standard integral Gronwall lemma, $\phi(t) \leq a(t)e^{kt}$, which requires *no* smallness condition on k but gives an exponentially growing (not constant) bound. One can verify numerically that the claimed bound $a/(1 - k)$ is not even a valid upper bound for t beyond a modest threshold: for $a \equiv 1$, $k = 0.5$, the correct Gronwall bound at $t = 2$ is $e^1 \approx 2.72$, exceeding the earlier draft’s claimed constant bound of 2.

The resolution is not simply “use e^{kt} instead of $1/(1 - k)$ ”, because the resulting bound on $\phi(t)$ would then need to be fed back into the X -inequality (7), which is a genuinely *differential* (not integral) inequality; reconciling an exponential-integral bound on p with a differential bound on X into a single clean estimate of the form claimed in the earlier Theorem 3.1 is not

straightforward and was not completed correctly. The approach below avoids this difficulty entirely by working on a short horizon $[0, \delta]$ from the start, where integrated quantities are controlled by δ times a supremum, converting every self-referential term into a genuinely small (coefficient $O(\delta)$) one.

A.2 Step 1: the X -bound on $[0, \delta]$

By (8) and $(a + b + c)^2 \leq 3(a^2 + b^2 + c^2)$,

$$|\alpha_t^1 - \alpha_t^2|^2 \leq 3C_\alpha^2(|\Delta X_t|^2 + |\Delta p_t|^2 + W_2(\mu_t^1, \mu_t^2)^2).$$

Substituting into (7) and integrating over $[0, \delta]$, using $\Delta X_0 = 0$ (the coupled processes share the same initial random variable $X_0^1 = X_0^2 \sim \mu_{\text{init}}$ under synchronous coupling):

$$\begin{aligned} \mathbb{E}[|\Delta X_t|^2] &\leq (2L_b + L_b^2 + 3L_b C_\alpha^2) \int_0^t \mathbb{E}[|\Delta X_s|^2] ds + 3L_b C_\alpha^2 \int_0^t \mathbb{E}[|\Delta p_s|^2] ds \\ &\quad + (L_b + 3L_b C_\alpha^2) \int_0^t \mathbb{E}[W_2(\mu_s^1, \mu_s^2)^2] ds, \end{aligned}$$

for every $t \leq \delta$. Bounding each integral by δ times the corresponding supremum over $[0, \delta]$ and taking the supremum of the left side over $t \leq \delta$ gives (10):

$$X_\delta^* \leq \delta A(X_\delta^* + P_\delta^*) + \delta K_3 M_\delta, \quad A = 2L_b + L_b^2 + 6L_b C_\alpha^2, \quad K_3 = L_b + 3L_b C_\alpha^2,$$

where we used $3L_b C_\alpha^2 \leq A$ to combine the X_δ^* and P_δ^* coefficients (a minor loosening for compactness).

A.3 Step 2: the p -bound on $[0, \delta]$, correctly closed

Apply (9) on $[0, \delta]$, using the genuine terminal value $\mathbb{E}[|\Delta X_\delta|^2] \leq X_\delta^*$ (this is the correction relative to Appendix A.1: X_δ^* here is bounded, via Step 1, by an expression that is itself $O(\delta)$ small – it is *not* treated as an independent quantity comparable in size to P_δ^*), and substituting the same bound on $|\alpha^1 - \alpha^2|^2$:

$$P_\delta^* \leq C_{\text{BSDE}} X_\delta^* + C_{\text{BSDE}}(1 + 3C_\alpha^2)\delta X_\delta^* + 3C_{\text{BSDE}}C_\alpha^2\delta P_\delta^* + C_{\text{BSDE}}(1 + 3C_\alpha^2)\delta M_\delta.$$

Substitute the Step 1 bound $X_\delta^* \leq \delta A(X_\delta^* + P_\delta^*) + \delta K_3 M_\delta$ into the (non-integrated) $C_{\text{BSDE}} X_\delta^*$ term above – this is the key step that makes the coefficient of $(X_\delta^* + P_\delta^*)$ on the right genuinely $O(\delta)$, rather than the $O(1)$ coefficient that caused the error in Appendix A.1:

$$\begin{aligned} P_\delta^* &\leq C_{\text{BSDE}} \left[\delta A(X_\delta^* + P_\delta^*) + \delta K_3 M_\delta \right] + C_{\text{BSDE}}(1 + 3C_\alpha^2)\delta X_\delta^* + 3C_{\text{BSDE}}C_\alpha^2\delta P_\delta^* \\ &\quad + C_{\text{BSDE}}(1 + 3C_\alpha^2)\delta M_\delta \\ &\leq \delta C_{\text{BSDE}} A(1 + (1 + 3C_\alpha^2)\delta)(X_\delta^* + P_\delta^*) + 3C_{\text{BSDE}}C_\alpha^2\delta P_\delta^* + \delta C'(\delta) M_\delta, \end{aligned}$$

where $C'(\delta) := C_{\text{BSDE}} K_3(1 + (1 + 3C_\alpha^2)\delta) + C_{\text{BSDE}}(1 + 3C_\alpha^2)$ (using $X_\delta^* \leq X_\delta^* + P_\delta^*$ in the second-to-last inequality). This is (11).

A.4 Step 3: the explicit threshold δ_0

For $\delta \leq \delta_0$ as in (12), two things hold:

(a) $3C_{\text{BSDE}}C_\alpha^2\delta \leq \frac{1}{2}$ (from $\delta \leq 1/(6C_{\text{BSDE}}C_\alpha^2)$), so $(1 - 3C_{\text{BSDE}}C_\alpha^2\delta) \geq \frac{1}{2}$; dividing (11) through by this factor (and using $\delta \leq 1$ to bound $\delta^2 \leq \delta$ throughout for compactness) gives

$$P_\delta^* \leq 2\delta C_{\text{BSDE}} A(2 + 3C_\alpha^2)(X_\delta^* + P_\delta^*) + 2\delta \tilde{C}' M_\delta,$$

for an explicit constant $\tilde{C}'$ absorbing the loosened $C'(\delta)$ terms.

(b) Adding this to the Step 1 bound $X_\delta^* \leq \delta A(X_\delta^* + P_\delta^*) + \delta K_3 M_\delta$ and collecting the coefficient of $(X_\delta^* + P_\delta^*)$:

$$X_\delta^* + P_\delta^* \leq \delta \underbrace{\left[A + 2C_{\text{BSDE}} A(2 + 3C_\alpha^2) \right]}_{=: B} (X_\delta^* + P_\delta^*) + \delta \Gamma_1 M_\delta,$$

with $B = L_b(2C_{\text{BSDE}}(3C_\alpha^2 + 2) + 1)(6C_\alpha^2 + L_b + 2)$ (this matches the bracket in (12)). The second part of $\delta \leq \delta_0$ ensures $\delta B \leq \frac{1}{2}$, so $\rho(\delta) := \delta B \leq \frac{1}{2}$, giving (13) with $\Gamma(\delta) = \delta \Gamma_1$ for an explicit Γ_1 absorbing K_3 and $\tilde{C}'$. Rearranging as in (14) completes the derivation.

A.5 Verification

All coefficient manipulations above were independently verified by symbolic computation (`sympy`) and by a numerical instantiation with representative constants ($L_b = 1.5$, $C_\alpha = 1.0$, $C_{\text{BSDE}} = 2.0$), confirming $\delta_0 > 0$ and the correctness of each intermediate inequality; the verification script is available in the accompanying repository.

B The Continuation Argument and an Explicit Threshold

This appendix elaborates the continuation step of Theorem 4.1(ii)–(iii) and gives a concrete numerical illustration of δ_0 .

B.1 Continuation: propagating the fourth moment

For the continuation argument to apply at each stage, the terminal law of the local equilibrium at the end of subinterval k must itself satisfy Assumption 2.6 (finite fourth moment) so that it is a valid initial condition for subinterval $k + 1$. This follows from a standard a priori moment estimate for the Lipschitz SDE (1): under Assumptions 2.1 and 2.4, if $\mathbb{E}[|X_0|^4] < \infty$ then

$$\sup_{t \in [0, \delta]} \mathbb{E}[|X_t|^4] \leq C(1 + \mathbb{E}[|X_0|^4])e^{C\delta}$$

for a constant C depending only on L_b , σ , $\|\sigma_0\|_F$ (a routine application of Itô's formula to $|X_t|^4$ and Gronwall's inequality, in the standard, unproblematic differential form – this bound does not involve the self-referential structure that caused the error in Appendix A.1, since it concerns only the fourth moment of a single solution, not a difference of two coupled solutions). Hence the fourth moment remains finite (with an explicit, uniform-in- k bound) across all $\lceil T/\delta_0 \rceil$ subintervals.

B.2 Numerical illustration of δ_0

For the canonical LQ-MFG benchmark (Example 6.1, baseline parameters $\kappa = 0.5$, $\lambda = 1$, $c_1 = 2$, $c_2 = 5$, $\sigma = 0.2$, $\sigma_0 = 0.3$, $T = 1$), Proposition 6.2 gives $L_b = \max(\kappa, 1) = 1$ and $C_\alpha = 1/\lambda = 1$ exactly. C_{BSDE} requires evaluating the cited BSDE estimate on the linear adjoint driver $(-\kappa, -c_1)$; we use a representative value $C_{\text{BSDE}} \approx 2.0$, consistent with the magnitude expected for a linear driver with these coefficients over $T = 1$ (a precise value requires tracking the exact constant in the cited theorem, which we do not carry out here). With these values, formula (12) gives $\delta_0 \approx 2.6 \times 10^{-3}$, so covering $T = 1$ requires on the order of 380 subintervals. This threshold is conservative – a consequence of the crude inequalities used throughout (in particular $(a + b + c)^2 \leq 3(a^2 + b^2 + c^2)$, which is loose whenever the three terms are of very different sizes) – and is not intended as a practically tight step size for the numerical scheme of Chapter 4 of Monyimba (2026), which uses a much finer, independently justified time discretisation.

Sharpening δ_0 towards a practically useful threshold, e.g. via the displacement-monotonicity techniques of [Jackson and Tangpi \(2024\)](#) or [Jackson and Mészáros \(2025\)](#) in place of the crude Cauchy–Schwarz-type bounds used here, is identified as a natural direction for future work.